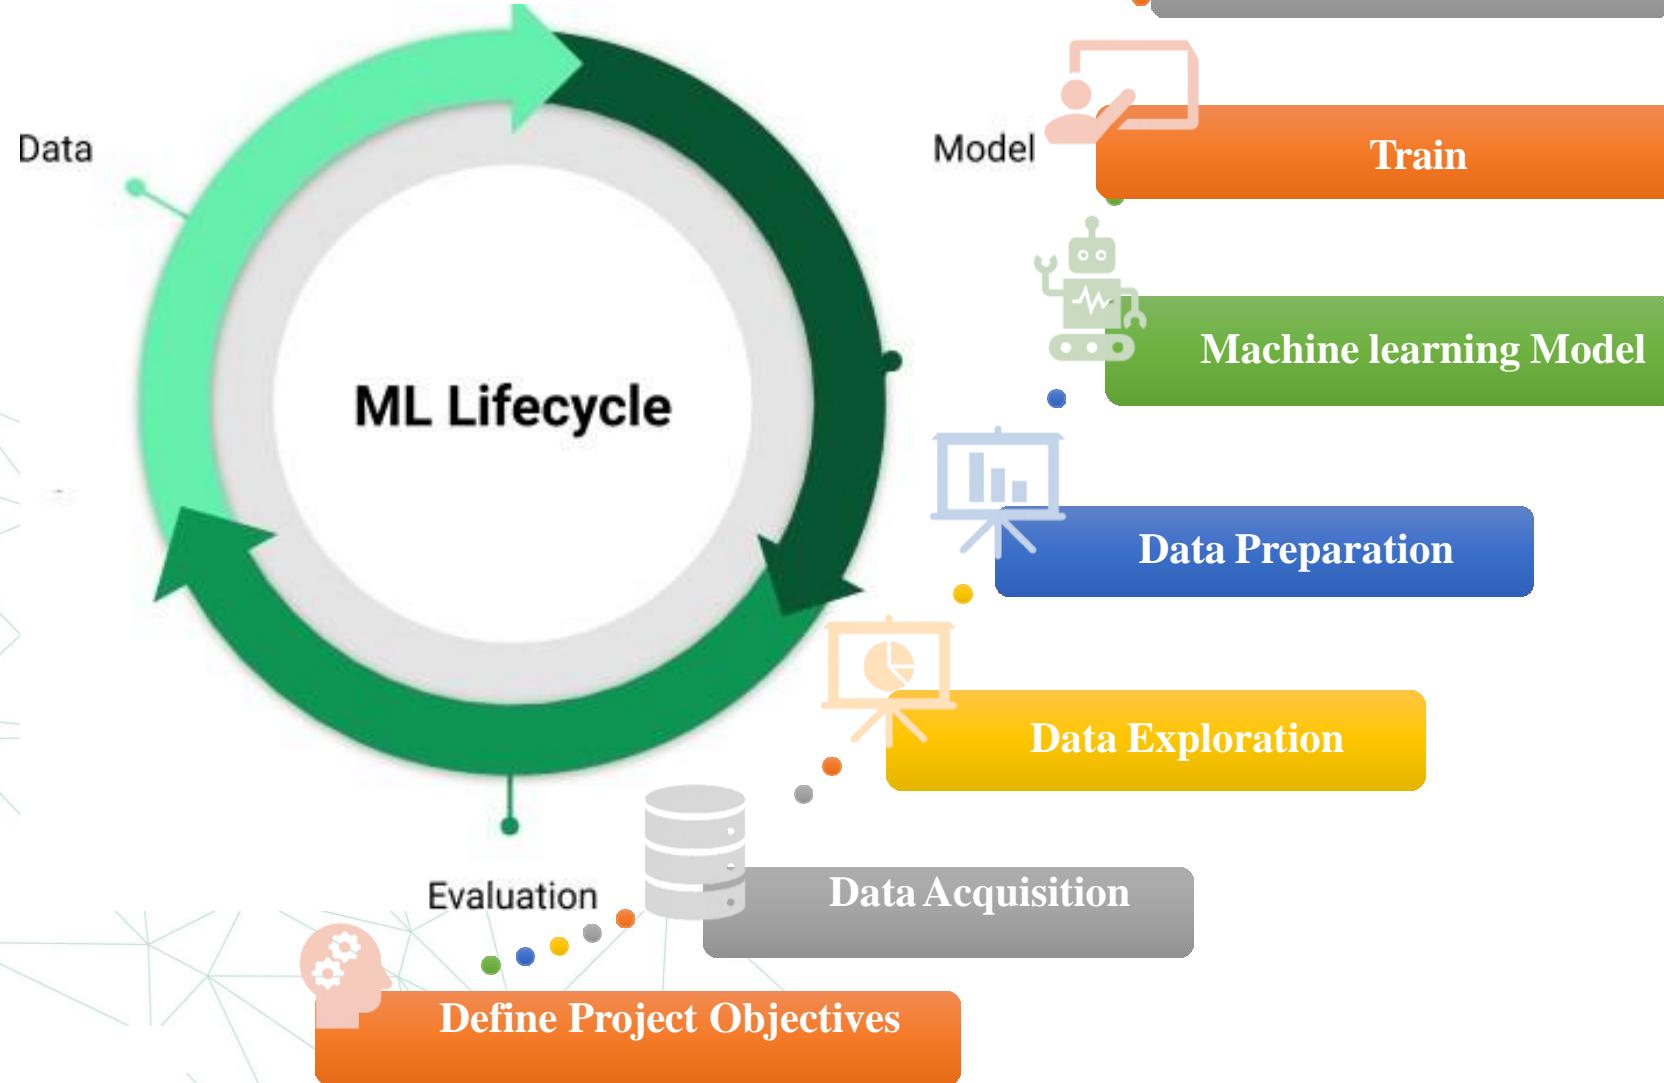


Machine Learning



Supervised Learning

Supervised Learning



Regression



What will be the temperature tomorrow?

84°

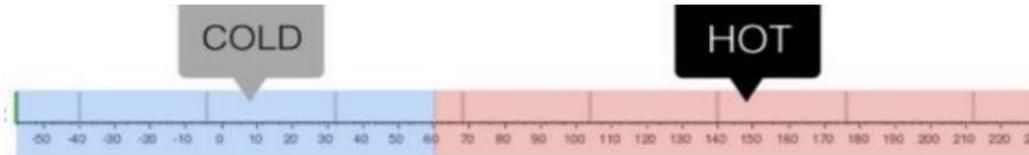


Fahrenheit

Classification



Will it be hot or cold tomorrow?



Fahrenheit

Model

A Mathematical representation of a real world process in the form of input-output relationship.

$$\text{slices of pizza I'll eat} = 2 * (\text{hours since last meal}) + 1$$

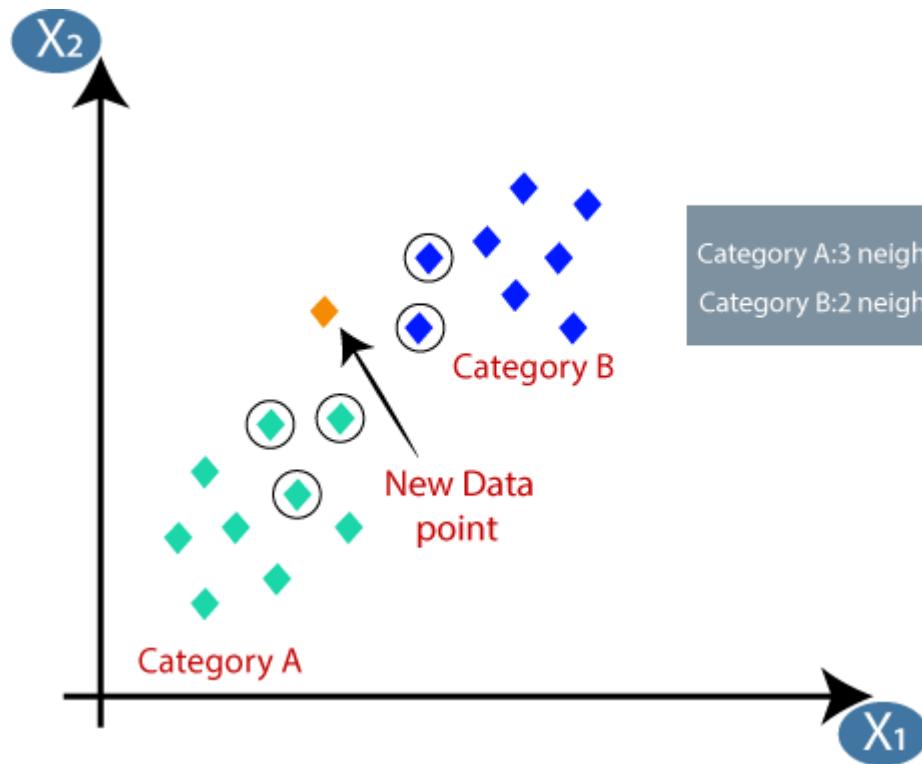
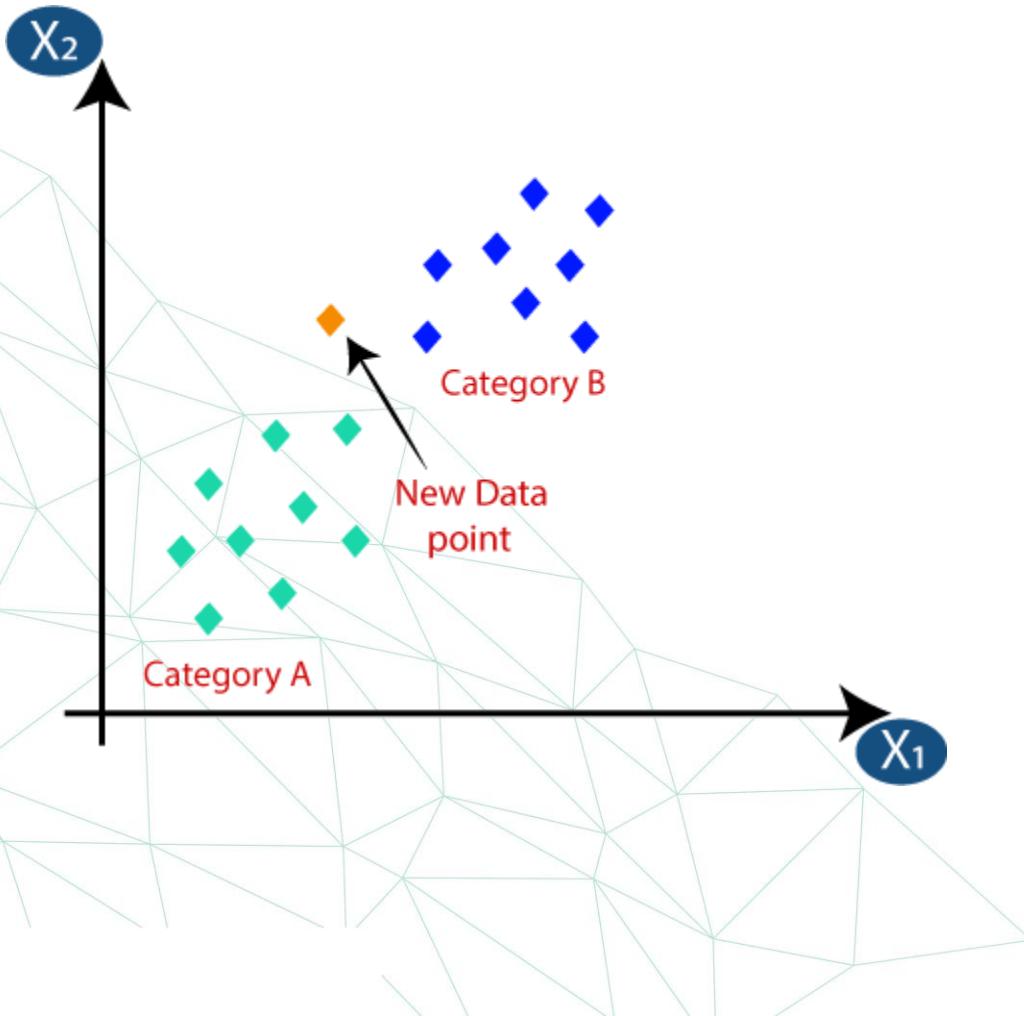
KNN



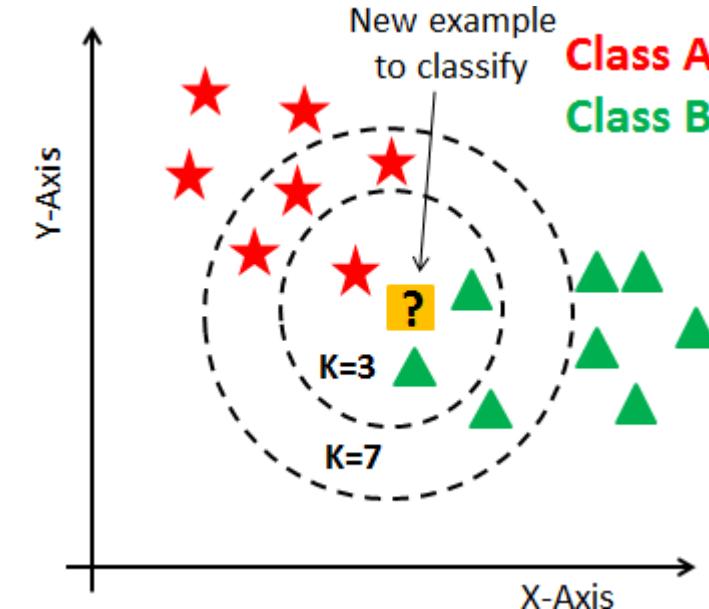
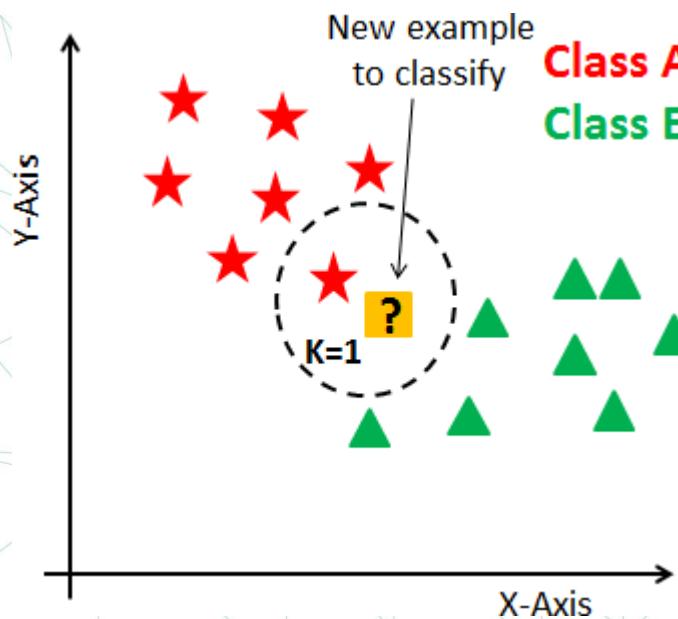
How does K-NN work?

- The K-NN working can be explained on the basis of the below algorithm:
 - **Step-1:** Select the number **K** of the **neighbors**
 - **Step-2:** Calculate the **Euclidean distance** of **K number of neighbors**
 - **Step-3:** Take the **K nearest neighbors** as per the calculated Euclidean distance.
 - **Step-4:** Among these k neighbors, **count** the number of the data points in **each category**.
 - **Step-5:** **Assign** the **new data points** to that category for which the number of the neighbor is **maximum**.

KNN - lazy learner algorithm



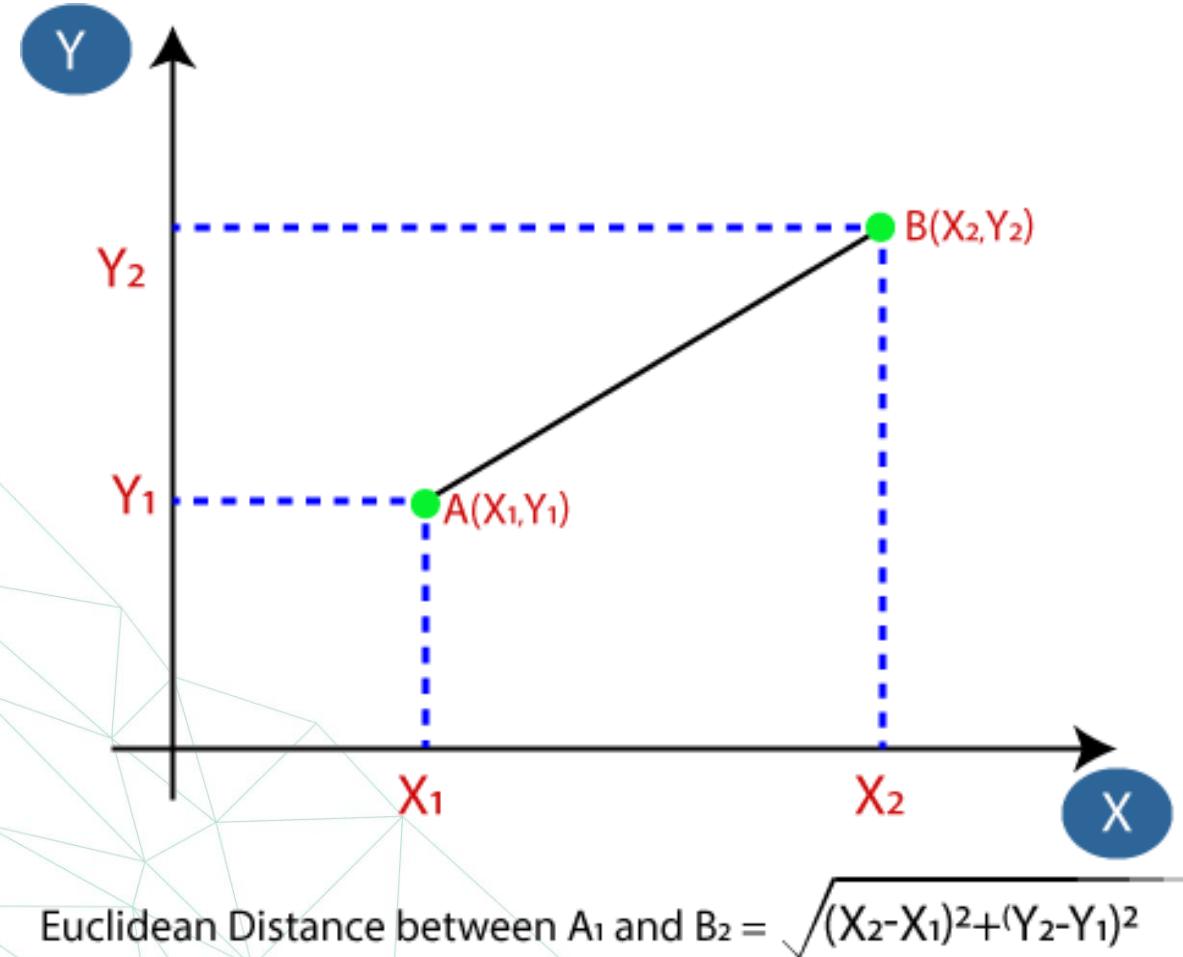
How do you Decide the Number of Neighbors in KNN?



How do you Decide the Number of Neighbors in KNN?

- Defining **k** can be a balancing act as different values can lead to **overfitting** or **underfitting**.
- **Lower** values of k can have high variance, but low bias[**overfitting**].
- **larger** values of k may lead to high bias and lower variance[**underfitting**].
- The choice of k will largely depend on the input data as data with more **outliers** or **noise** will likely perform better with **higher** values of k.
- Overall, it is recommended to have an **odd** number for k to avoid ties in classification, and **cross-validation** tactics can help you choose the optimal k for your dataset.

Euclidean Distance



Advantages

- **Easy to implement**
- **Adapts easily:** As **new training samples are added**, the algorithm **adjusts** to account for any new data since all training data is stored into memory.
- **Few hyperparameters:** KNN only requires a **k** value and a **distance metric**.

Disadvantages

- **Does not scale well:** Since KNN is a lazy algorithm, it takes up **more memory** and data storage compared to other classifiers. This can be **costly** from both a **time** and **money** perspective.
- **Curse of dimensionality:** The KNN algorithm tends to fall victim to the curse of dimensionality, which means that it **doesn't perform well** with **high-dimensional** data inputs.
- **Prone to overfitting:** Due to the “curse of dimensionality”, KNN is also more prone to **overfitting**. While **feature selection** and **dimensionality reduction** techniques are leveraged to prevent this from occurring, the value of **k** can also impact the model’s behavior.

Preparing Data For KNN

- **Rescale Data:** KNN performs much better if all of the data has the same scale. Normalizing your data to the range between 0 and 1 is a good idea. It may also be a good idea to standardize your data if it has a Gaussian distribution.
- **Address Missing Data:** Missing data will mean that the distance between samples cannot be calculated. These samples could be excluded or the missing values could be imputed.
- **Lower Dimensionality:** KNN is suited for lower dimensional data. You can try it on high dimensional data (hundreds or thousands of input variables) but be aware that it may not perform as well as other techniques. KNN can benefit from feature selection that reduces the dimensionality of the input feature space.

KNN Lab

Classification Metrics



What is Confusion Matrix?

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive (TP)	False Negative (FN) Type II Error
	Negative	False Positive (FP) Type I Error	True Negative (TN)

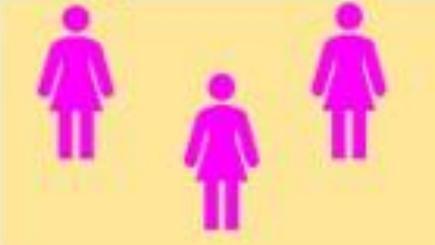
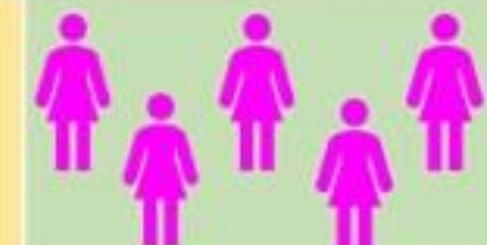
What is Confusion Matrix?

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive (TP)	False Negative (FN) Type II Error
	Negative	False Positive (FP) Type I Error	True Negative (TN)

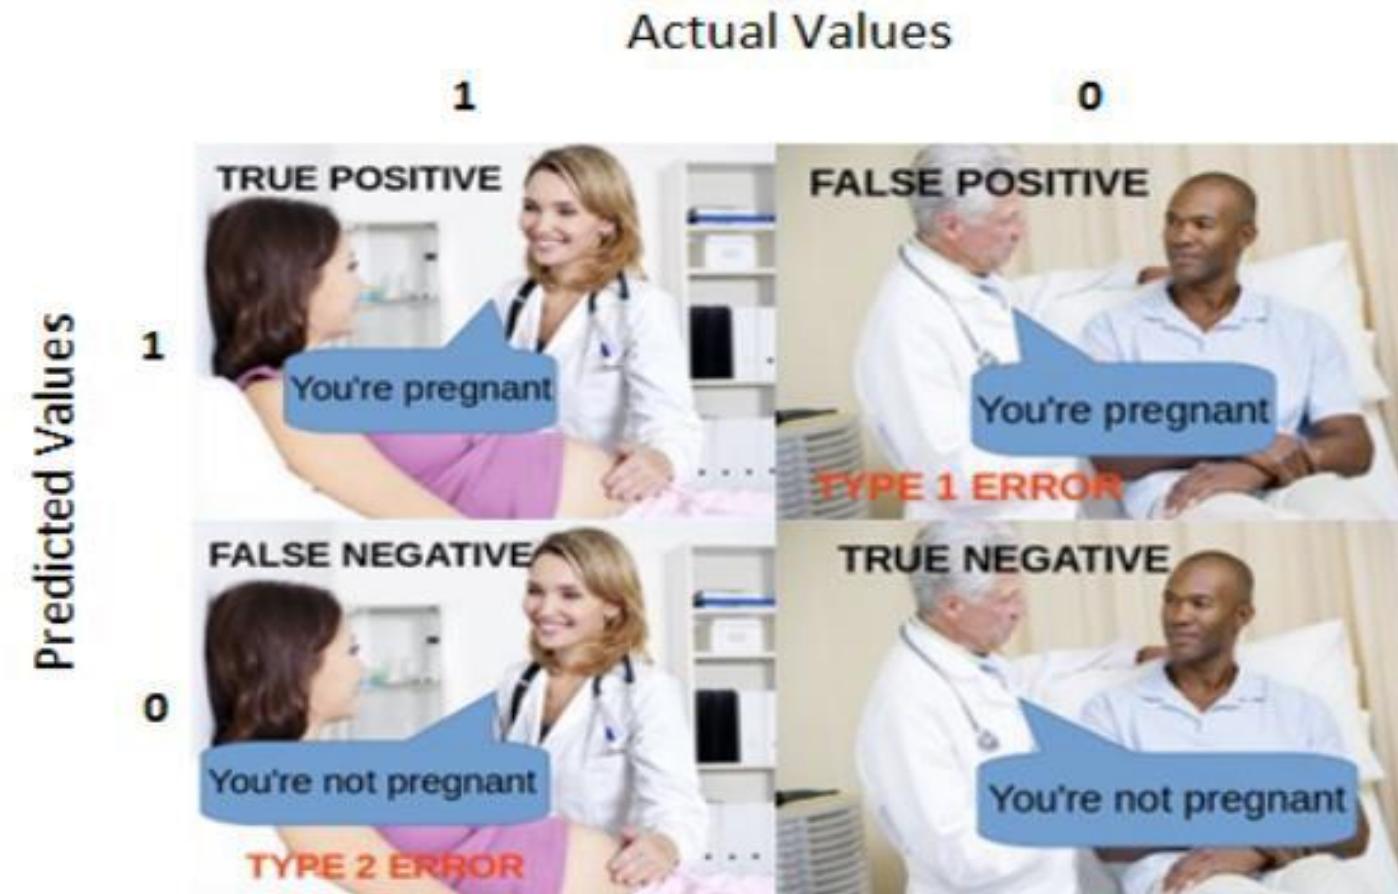
- **True Positive (TP)** — When the model says that the **patient has cancer** and the patient **actually has it**
- **False Positive (FP)** — When the model says that the **patient has cancer** but the patient **doesn't have it**
- **True Negative (TN)** — When the model says that the **patient does not have cancer** and the patient **actually doesn't have it**
- **False Negative (FN)** — When the model says that the **patient doesn't have cancer** but the patient **actually has it**. *We don't want this, do we?*



What is Confusion Matrix?

		PREDICTED VALUES	
		Pregnant	Not Pregnant
ACTUAL VALUES	Pregnant		
	Not Pregnant		

What is Confusion Matrix?



Performance evaluation Measures

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive (TP)	False Negative (FN) Type II Error
	Negative	False Positive (FP) Type I Error	True Negative (TN)

$$\text{Accuracy} = (TP+TN) / (TP+FP+FN+TN)$$

Exercise

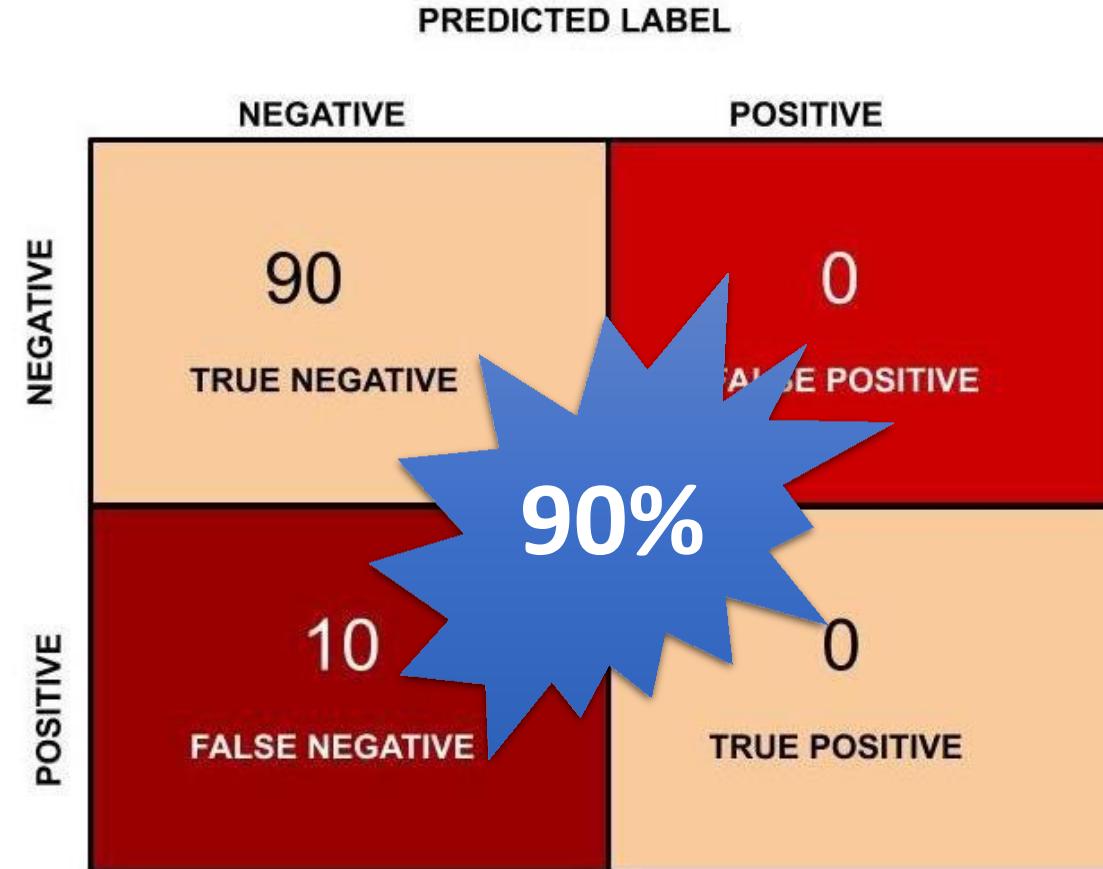
- We select 100 people which includes pregnant women, not pregnant women and men with fat belly. Let us assume out of this 100 people 40 are pregnant and the remaining 60 people include not pregnant women and men with fat belly. We now use a machine learning algorithm to predict the outcome.
- Out of 40 pregnant women 30 pregnant women are classified correctly and the remaining 10 pregnant women are classified as not pregnant by the machine learning algorithm.
- On the other hand, out of 60 people in the not pregnant category, 55 are classified as not pregnant and the remaining 5 are classified as pregnant.

Example

		PREDICTED LABEL	
		NEGATIVE	POSITIVE
TRUE LABEL	NEGATIVE	55 TRUE NEGATIVE	5 FALSE POSITIVE
	POSITIVE	10 FALSE NEGATIVE	30 TRUE POSITIVE



Is accuracy the best measure?



$$\text{Accuracy} = (\text{TP} + \text{TN}) / (\text{TP} + \text{FP} + \text{FN} + \text{TN})$$

Performance evaluation Measures

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positive (TP)	False Negative (FN) <i>Type II Error</i>
	Negative	False Positive (FP) <i>Type I Error</i>	True Negative (TN)

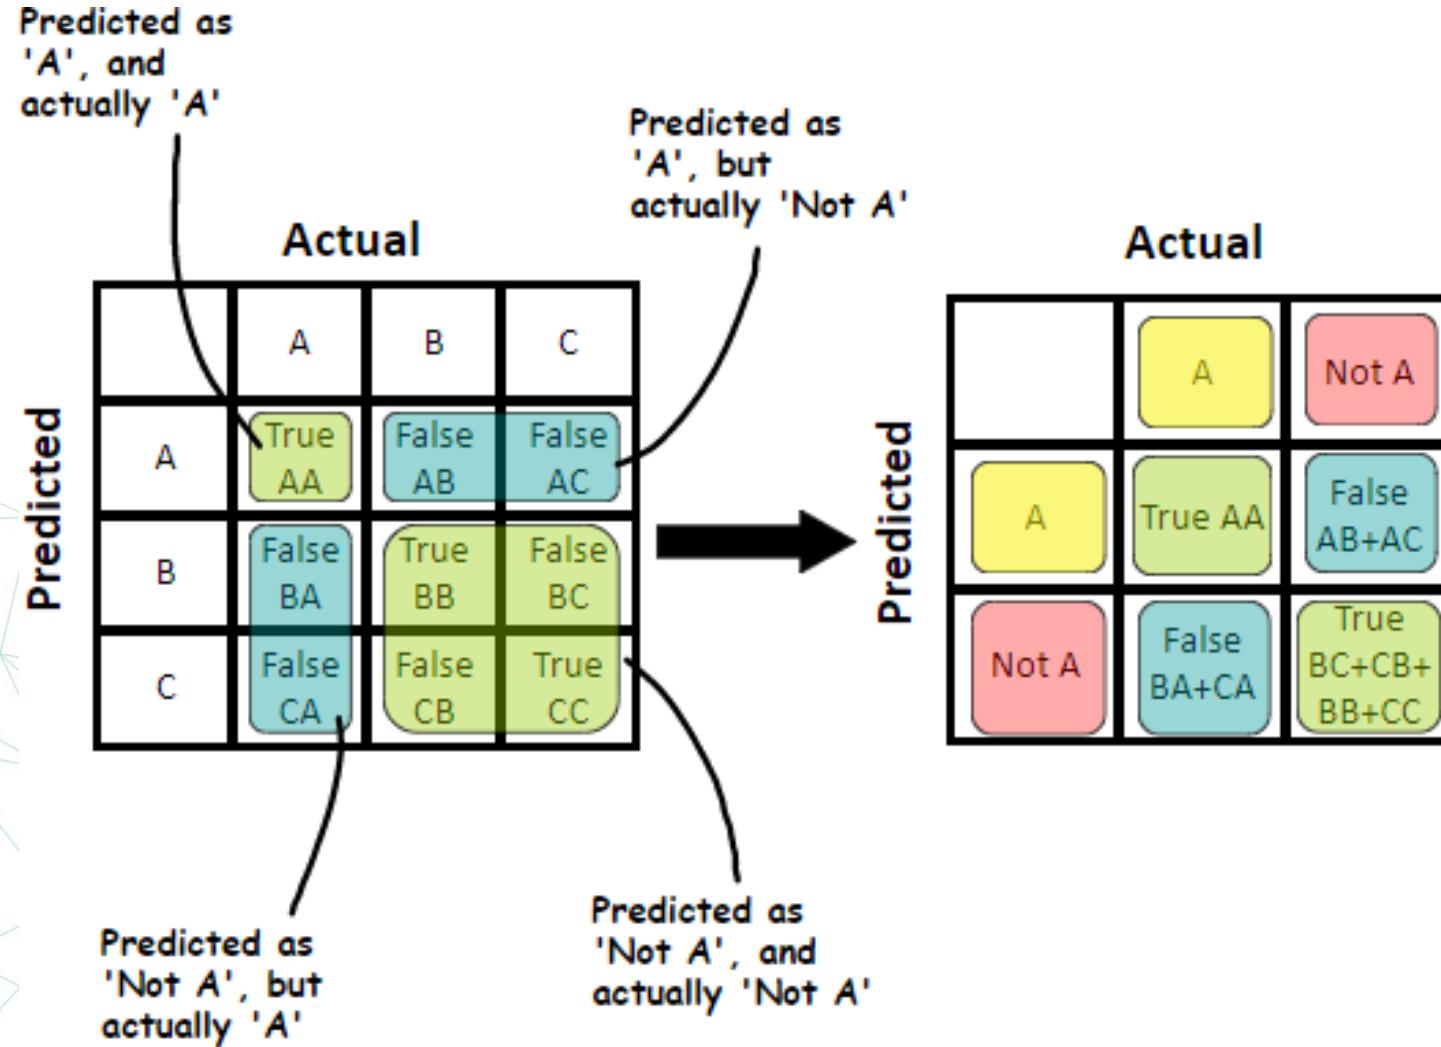
$$\text{Accuracy} = (TP+TN) / (TP+FP+FN+TN)$$

$$\text{Precision} = TP / (TP+FP)$$

$$\text{Recall} = TP / (TP+FN)$$

$$\text{F1 Score} = 2 * (\text{Recall} * \text{Precision}) / (\text{Recall} + \text{Precision})$$

MultiClass Classification



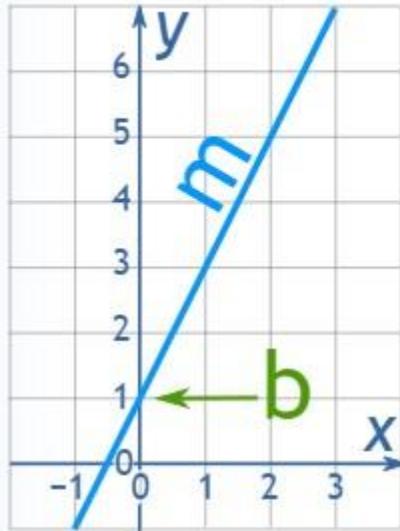
- In case of stock market crash **or not** prediction, false negative will be more important **as if stock is predicted as not** crash so you didn't withdraw money but actual it is crashed as result you lost money so here you are in loss(false negative),
- Second case :- false positive where you predict stock market crash so you withdraw money but actual it **is not** crash so result you were **not** able to earn profit becoz you withdraw money
- Conclusion:- false negative **is** more important **in** case of stock market crash prediction.
- **False negative is riskier**;because **if** the person has infected **with** disease but predicted **as negative** it will be dangerous because it **is** contagious.
- As classifier assumes it **as** Spam **and** moved to trash, **while** it **is not** spam, then the user may misses the email **and** may causes a lot of trouble. While **if** it **is** spam but classifier failed to detect, then user will see the spam then it **is not** a big deal, use

- (1) In case of stock market crash or not prediction, false negative will be more important as if stock is predicted as not crash so you didn't withdraw money but actual it is crashed as result you lost money so here you are in loss(false negative), Second case :- false positive where you predict stock market crash so you withdraw money but actual it is not crash so result you were not able to earn profit becoz you withdraw money . Conclusion:- false negative is more important in case of stock market crash prediction.
- (2) False negative is riskier;because if the person has infected with disease but predicted as negative it will be dangerous because it is contagious.
- (3) Credit Card Fraud
- (4)As classifier assumes it as Spam and moved to trash, while it is not spam, then the user may misses the email and may causes a lot of trouble. While if it is spam but classifier failed to detect, then user will see the spam then it is not a big deal, user will delete it by one click.

Lines and Gradient



Equation of a Straight Line



$$y = mx + b$$

Slope or Gradient

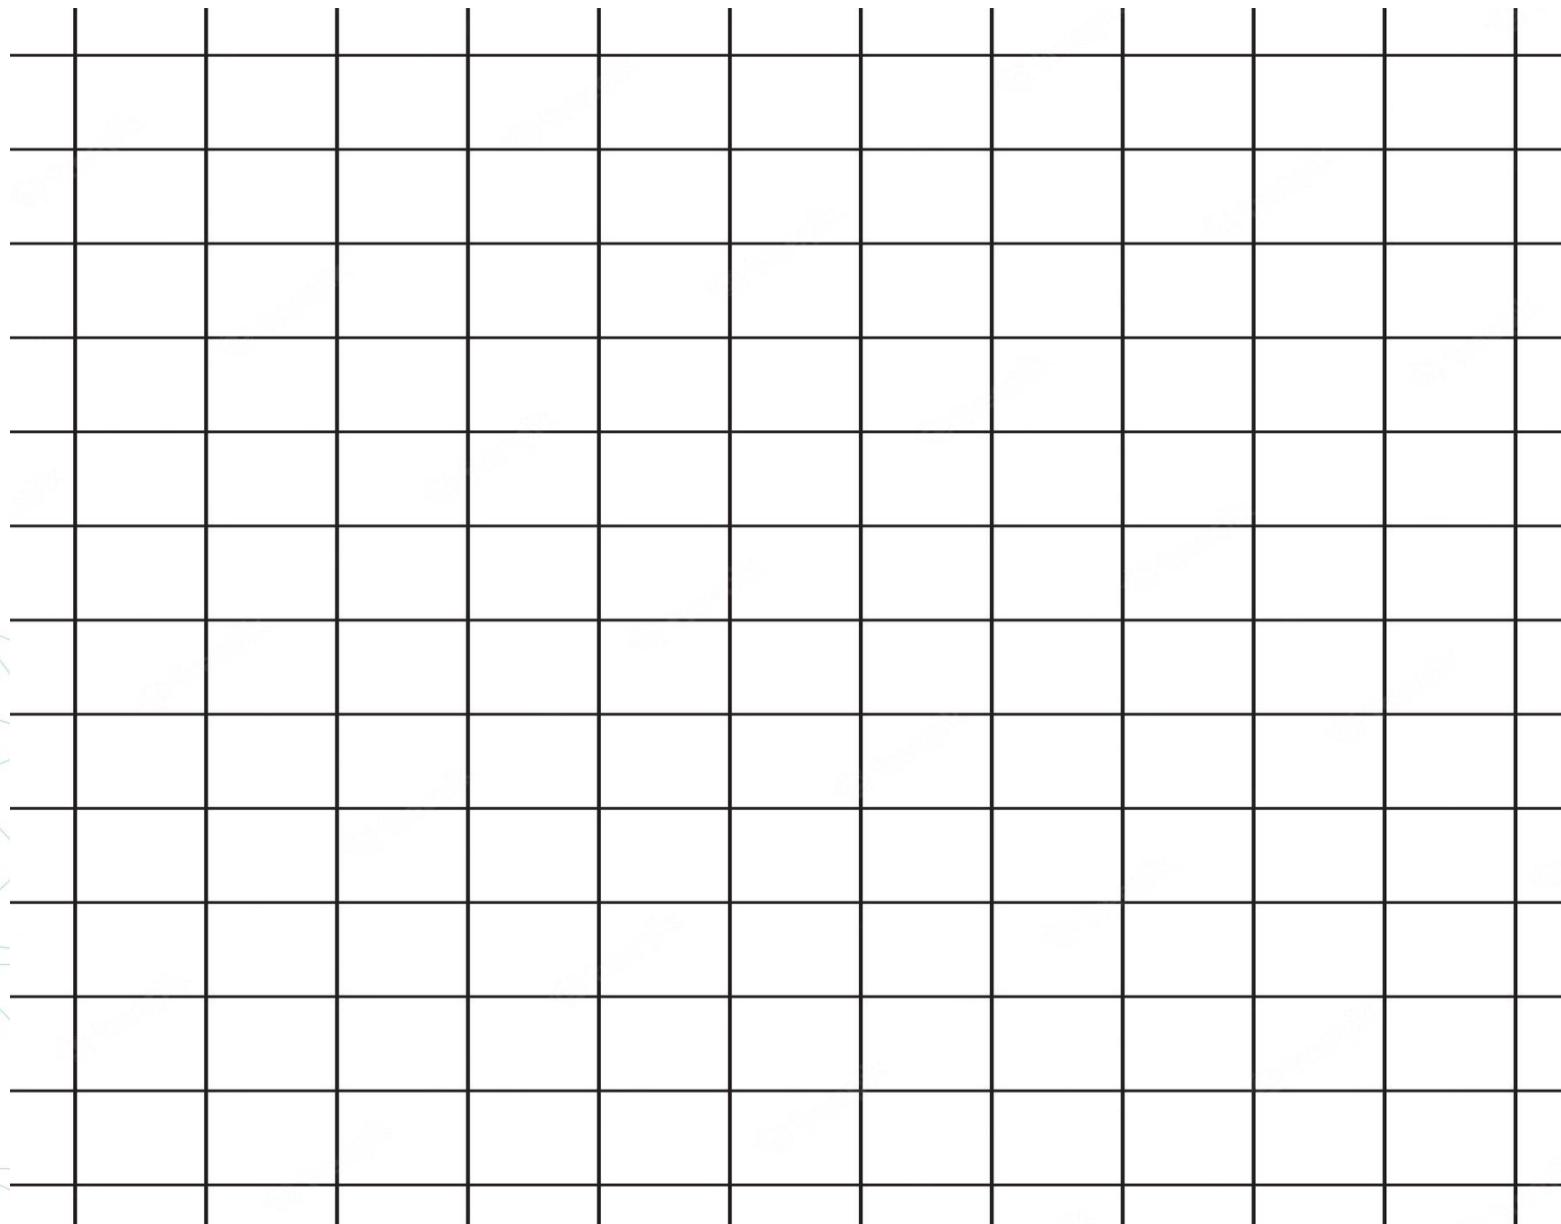
y value when $x=0$
(see Y Intercept)

y = how far up

x = how far along

m = Slope or Gradient (how steep the line is)

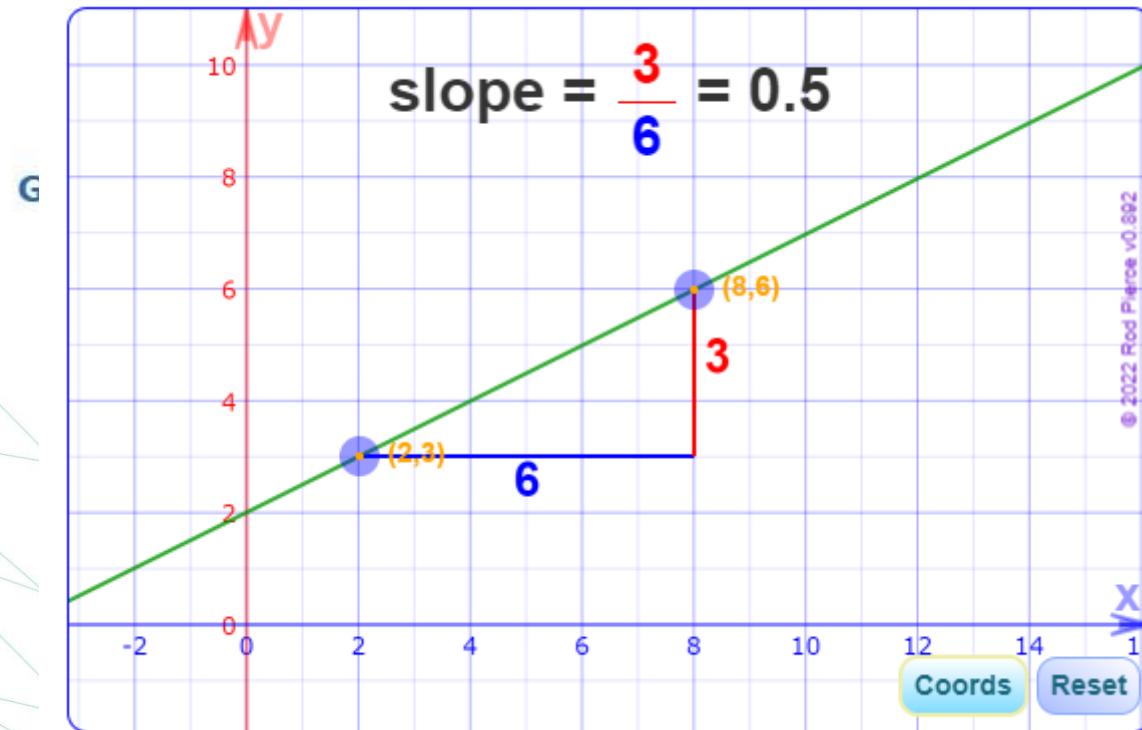
b = value of y when $x=0$



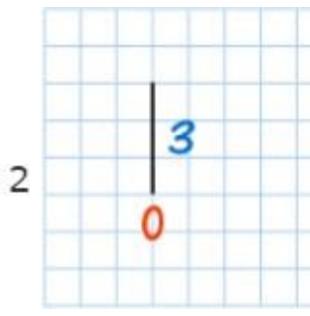
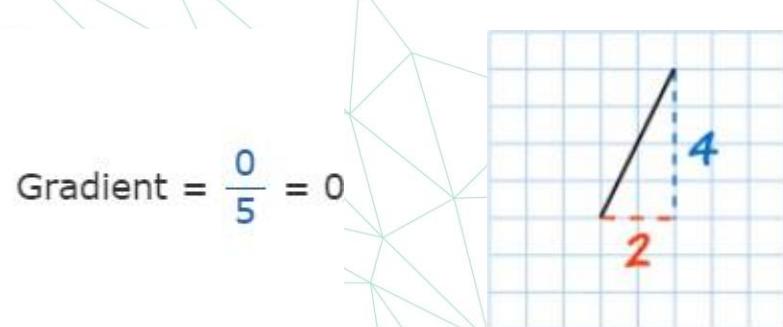
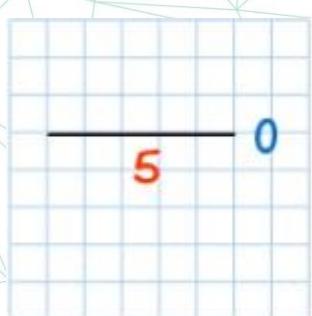
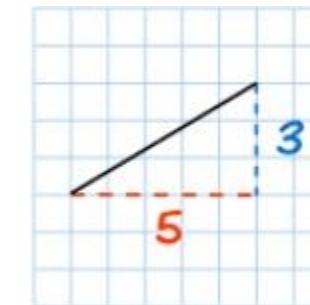
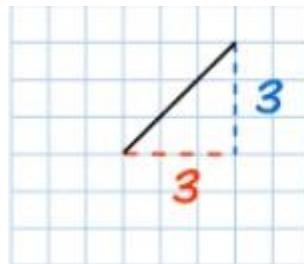
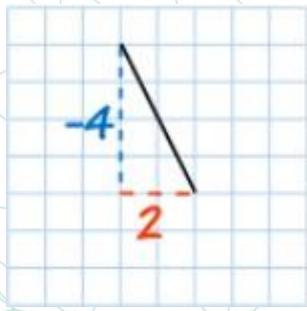


Gradient (Slope) of a Straight Line

Divide the **change in height** by the **change in horizontal distance**



Gradient (Slope) of a Straight Line



Derivatives

But how do we find the slope **at a point**?

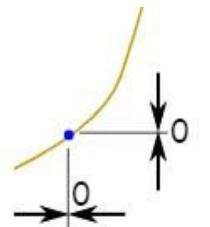
There is nothing to measure!

Sometimes we can't work something out **directly** ... but we **can** see what it should be as we get **closer and closer**!

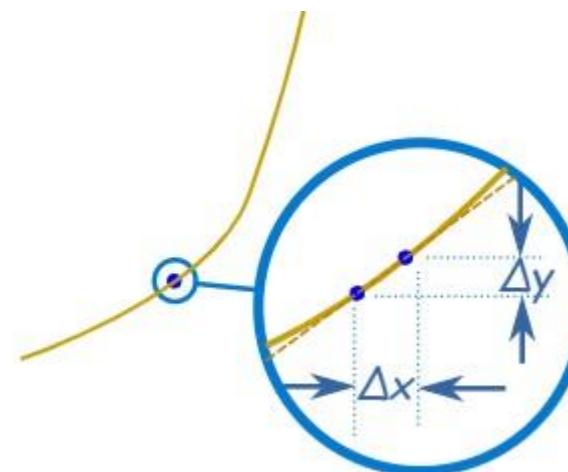
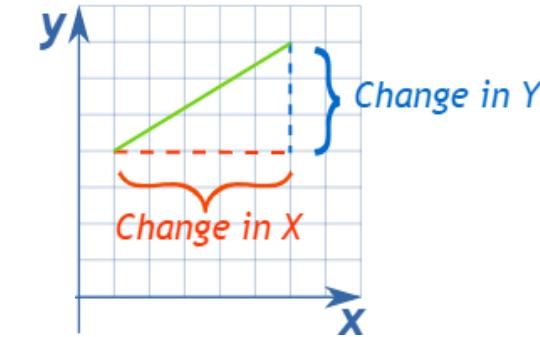
But with derivatives we use a small difference ...

... then have it **shrink towards zero**.

$$\text{Slope} = \frac{\text{Change in } Y}{\text{Change in } X}$$



$$\text{slope} = \frac{0}{0} = ???$$



Derivatives

To find the derivative of a function $y = f(x)$ we use the slope formula:

$$\text{Slope} = \frac{\text{Change in Y}}{\text{Change in X}} = \frac{\Delta y}{\Delta x}$$

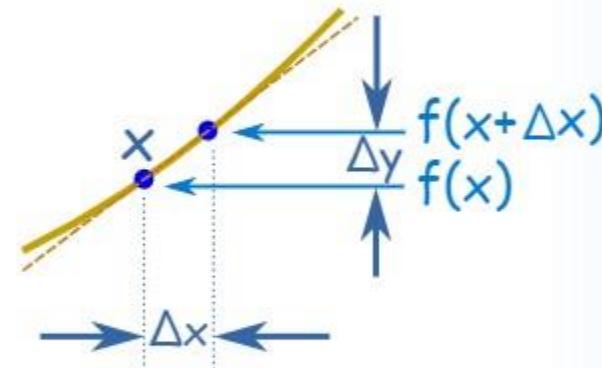
And (from the diagram) we see that:

x changes from x to $x + \Delta x$

y changes from $f(x)$ to $f(x + \Delta x)$

Now follow these steps:

- Fill in this slope formula: $\frac{\Delta y}{\Delta x} = \frac{f(x + \Delta x) - f(x)}{\Delta x}$
- Simplify it as best we can
- Then make Δx shrink towards zero.





Derivatives

We know $f(x) = x^2$, and we can calculate $f(x+\Delta x)$:

Start with: $f(x+\Delta x) = (x+\Delta x)^2$

Expand $(x + \Delta x)^2$: $f(x+\Delta x) = x^2 + 2x \Delta x + (\Delta x)^2$

The slope formula is: $\frac{f(x+\Delta x) - f(x)}{\Delta x}$

Put in $f(x+\Delta x)$ and $f(x)$: $\frac{x^2 + 2x \Delta x + (\Delta x)^2 - x^2}{\Delta x}$

Simplify (x^2 and $-x^2$ cancel): $\frac{2x \Delta x + (\Delta x)^2}{\Delta x}$

Simplify more (divide through by Δx): $= 2x + \Delta x$

Then, as Δx heads towards 0 we get: $= 2x$

Derivatives

We write **dx** instead of " **Δx heads towards 0**".

And "the derivative of" is commonly written $\frac{d}{dx}$ like this:

$$\frac{d}{dx}x^2 = 2x$$

"The derivative of x^2 equals $2x$ "
or simply "d dx of x^2 equals $2x$ "



Derivatives

So what does $\frac{d}{dx} x^2 = 2x$ mean?

It means that, for the function x^2 , the slope or "rate of change" at any point is $2x$.

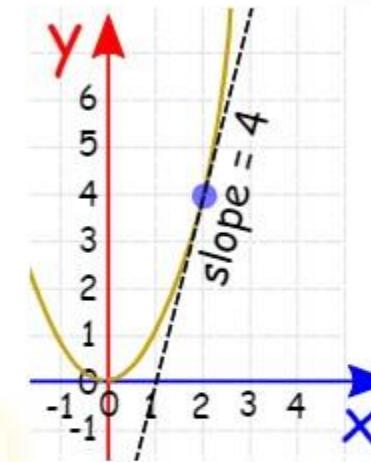
So when $x=2$ the slope is $2x = 4$, as shown here:

Or when $x=5$ the slope is $2x = 10$, and so on.

Note: $f'(x)$ can also be used for "the derivative of":

$$f'(x) = 2x$$

"The derivative of $f(x)$ equals $2x$ "
or simply "f-dash of x equals $2x$ "



Derivatives

Example: What is $\frac{d}{dx}x^3$?

We know $f(x) = x^3$, and can calculate $f(x+\Delta x)$:

Start with: $f(x+\Delta x) = (x+\Delta x)^3$

Expand $(x + \Delta x)^3$: $f(x+\Delta x) = x^3 + 3x^2 \Delta x + 3x (\Delta x)^2 + (\Delta x)^3$

The slope formula: $\frac{f(x+\Delta x) - f(x)}{\Delta x}$

Put in $f(x+\Delta x)$ and $f(x)$: $\frac{x^3 + 3x^2 \Delta x + 3x (\Delta x)^2 + (\Delta x)^3 - x^3}{\Delta x}$

Simplify (x^3 and $-x^3$ cancel): $\frac{3x^2 \Delta x + 3x (\Delta x)^2 + (\Delta x)^3}{\Delta x}$

Simplify more (divide through by Δx): $3x^2 + 3x \Delta x + (\Delta x)^2$

Then, as Δx heads towards 0 we get: $3x^2$

Result: the derivative of x^3 is $3x^2$



Derivatives

"Shrink towards zero" is actually written as a **limit** like this:

$$f'(x) = \lim_{\Delta x \rightarrow 0} \frac{f(x+\Delta x) - f(x)}{\Delta x}$$

"The derivative of f equals
the limit as Δx goes to zero of $f(x+\Delta x) - f(x)$ over Δx "

Or sometimes the derivative is written like this (explained on [Derivatives as dy/dx](#)):

$$\frac{dy}{dx} = \frac{f(x+dx) - f(x)}{dx}$$

The process of finding a derivative is called "differentiation".



You **do** differentiation ... to **get** a derivative.

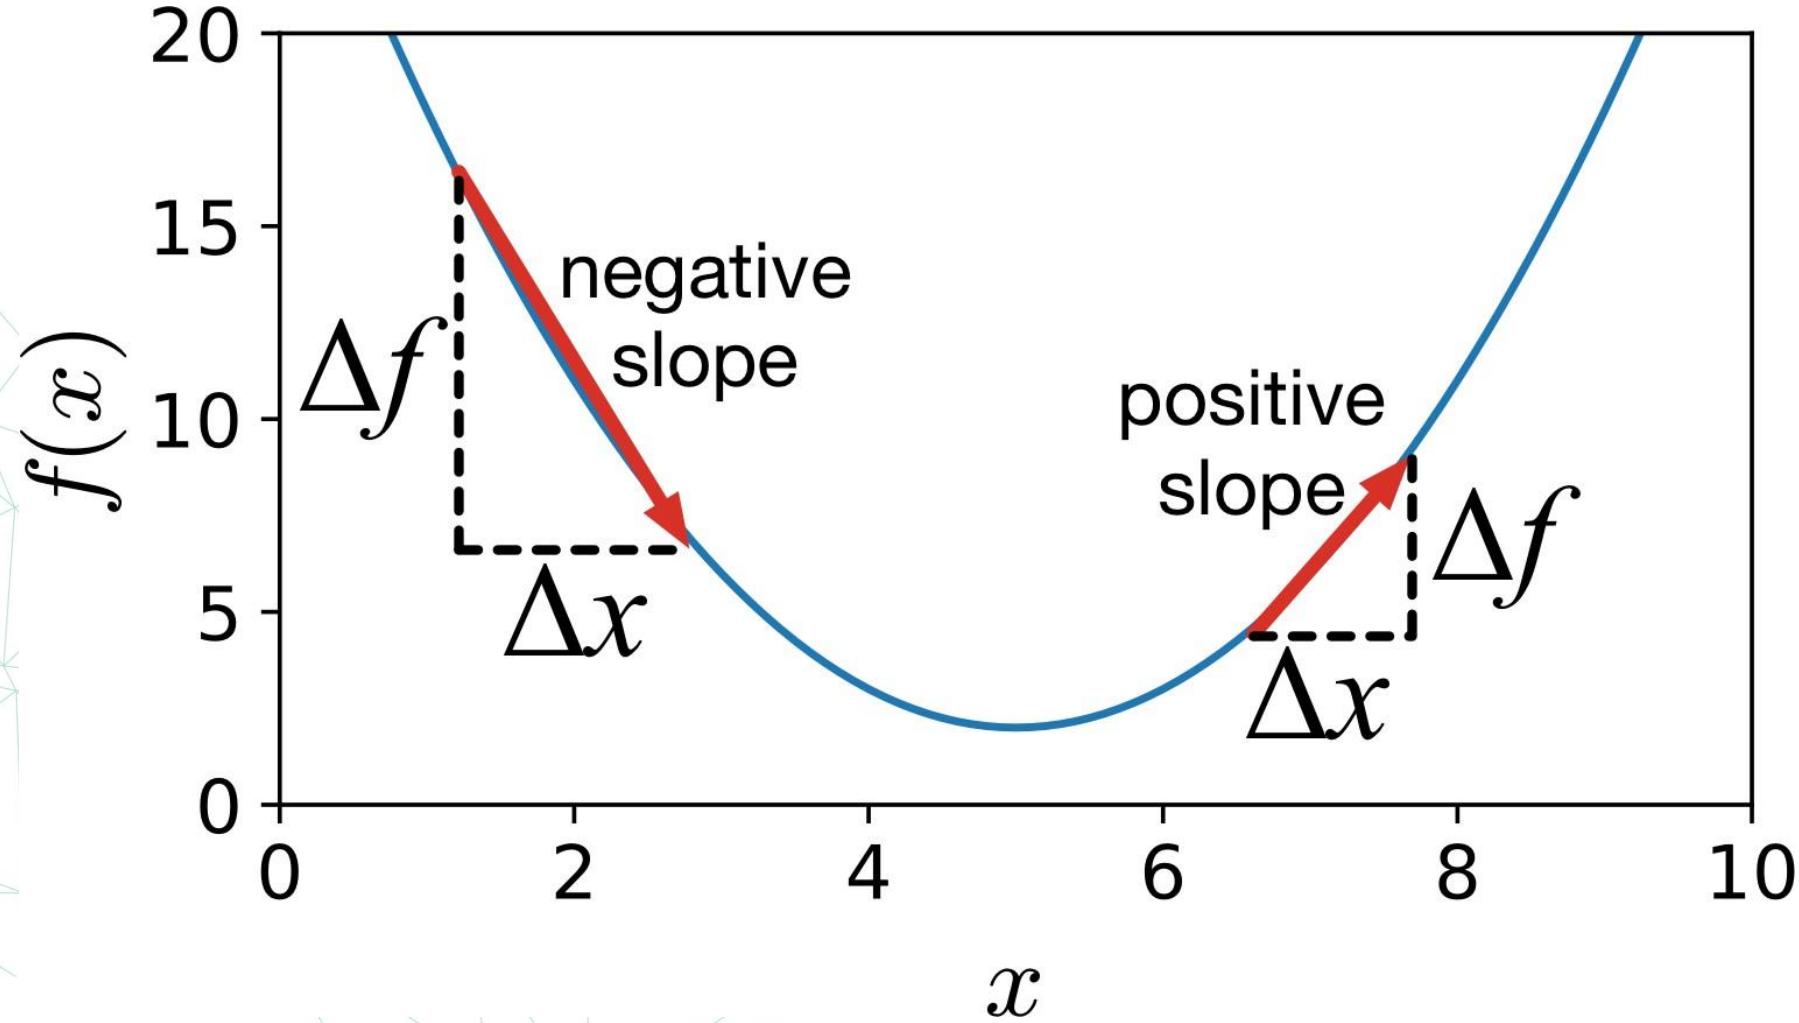
Derivatives

- <https://www.mathsisfun.com/calculus/derivative-plotter.html>
- <https://www.mathsisfun.com/calculus/slope-function-point.html>
- <https://www.mathsisfun.com/calculus/derivatives-rules.html>

Derivative Rules

Common Functions	Function	Derivative
Constant	c	0
Line	x	1
Square	ax	a
Square Root	x^2	$2x$
Exponential	\sqrt{x}	$(\frac{1}{2})x^{-\frac{1}{2}}$
Logarithms	e^x	e^x
	a^x	$\ln(a) a^x$
	$\ln(x)$	$1/x$

Derivative

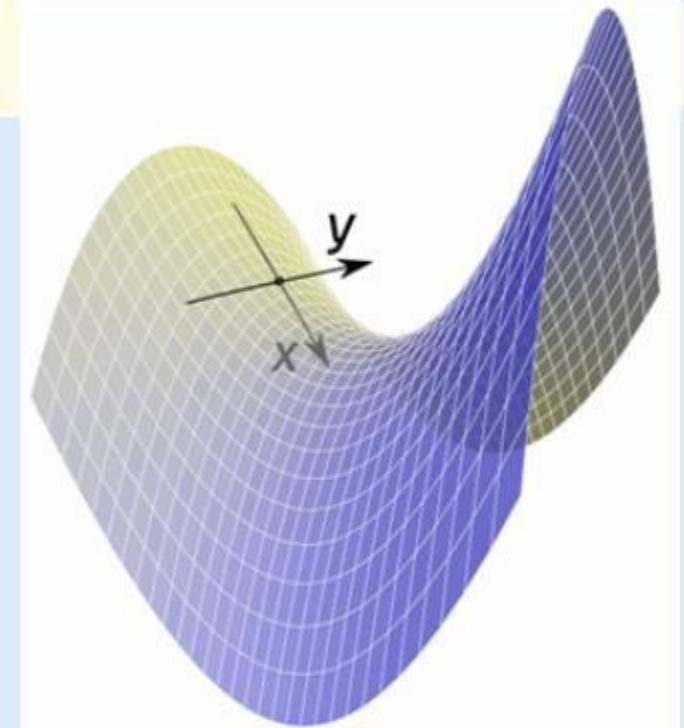


Partial Derivatives

Example: a function for a surface that depends on two variables x and y

When we find the slope in the x direction (while keeping y fixed) we have found a partial derivative.

Or we can find the slope in the y direction (while keeping x fixed).



Partial Derivatives

But what about a function of **two variables** (x and y):

$$f(x, y) = x^2 + y^3$$

We can find its **partial derivative with respect to x** when we treat **y as a constant** (imagine y is a number like 7 or something):

$$f'_x = 2x + 0 = 2x$$

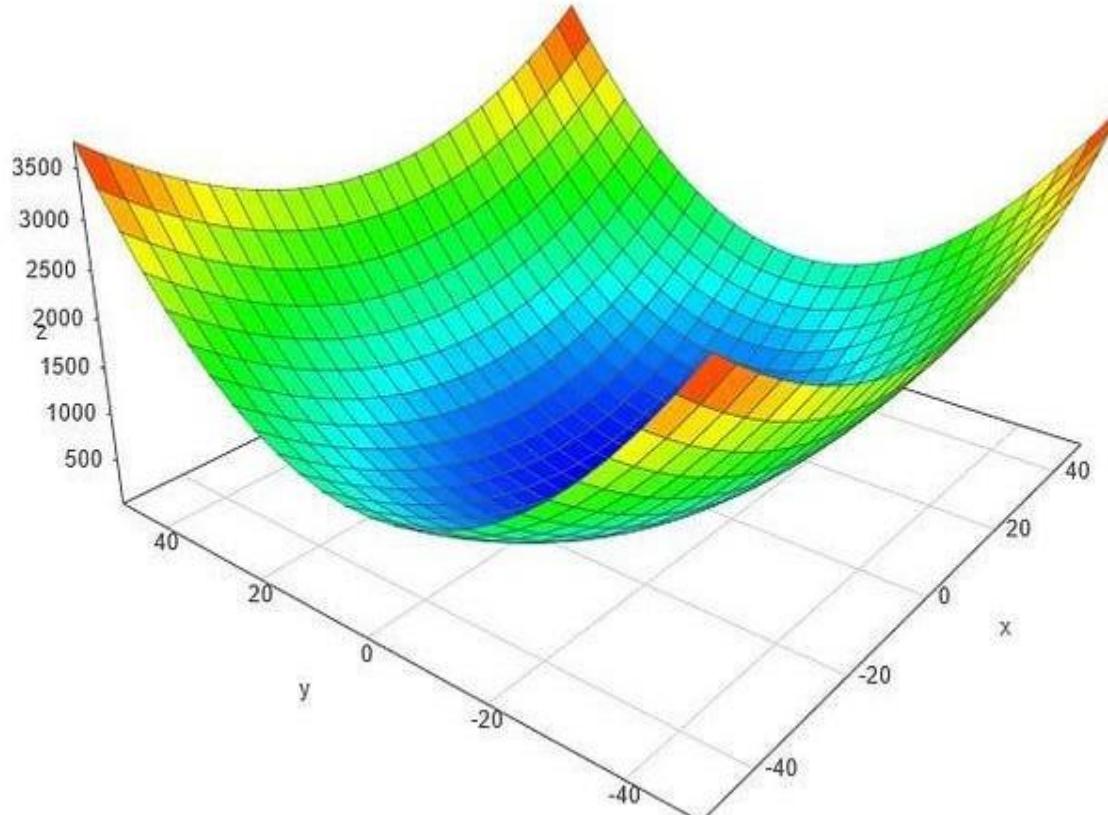
Explanation:

- *the derivative of x^2 (with respect to x) is $2x$*
- *we treat y as a constant, so y^3 is also a constant (imagine $y=7$, then $7^3=343$ is also a constant), and the derivative of a constant is 0*

To find the partial derivative **with respect to y** , we treat **x as a constant**:

$$f'_y = 0 + 3y^2 = 3y^2$$

$$f(x) = 0.5x^2 + y^2$$



Let's assume we are interested in a gradient at point p(10,10):

$$\frac{\partial f(x, y)}{\partial x} = x, \quad \frac{\partial f(x, y)}{\partial y} = 2y$$

so consequently:

$$\nabla f(x, y) = \begin{bmatrix} x \\ 2y \end{bmatrix}$$

$$\nabla f(10, 10) = \begin{bmatrix} 10 \\ 20 \end{bmatrix}$$

By looking at these values we conclude that the slope is twice steeper along the y axis.



Finding Maxima and Minima using Derivatives

- Where is a function at a high or low point?
 - Calculus can help!



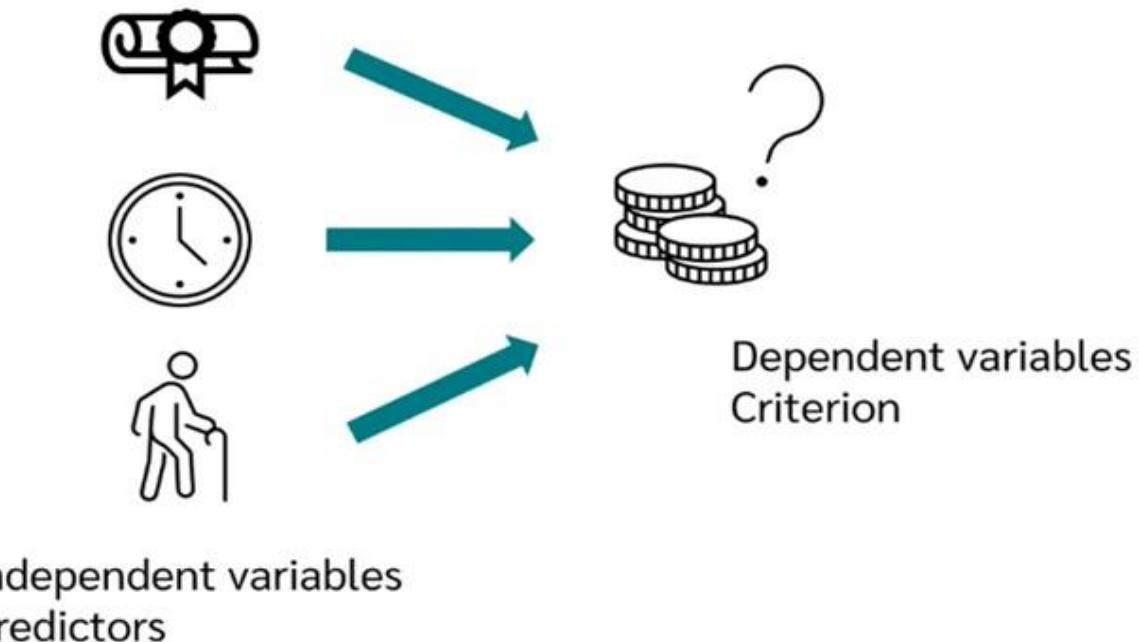
um is always where the function

Where does it flatten out? Where the **slope is zero**.
Where is the slope zero? The **Derivative** tells us!

Linear Regression

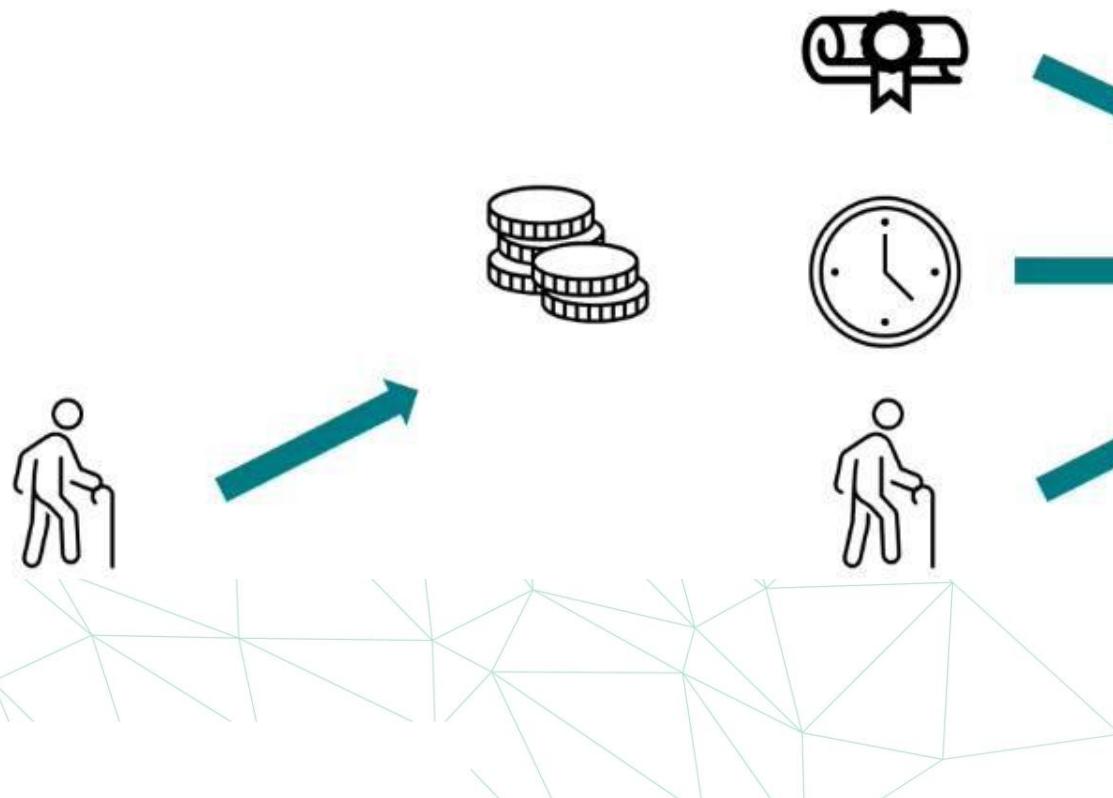
Regression Analysis

A **regression analysis** makes it possible to infer or predict another variable on the basis of one or more variables.

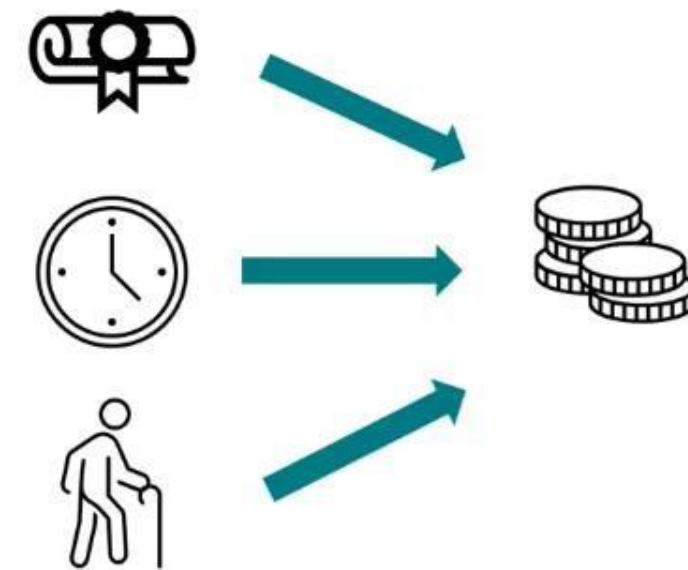


Forms of Regression Analysis

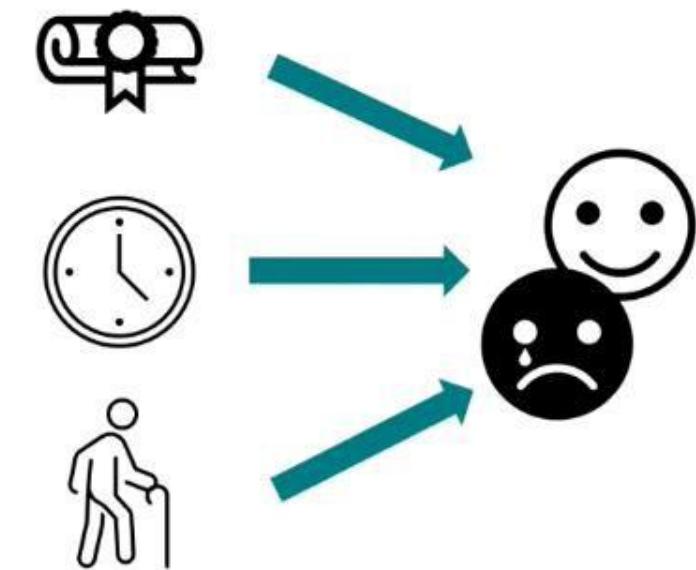
Simple linear Regression



Multiple linear Regression

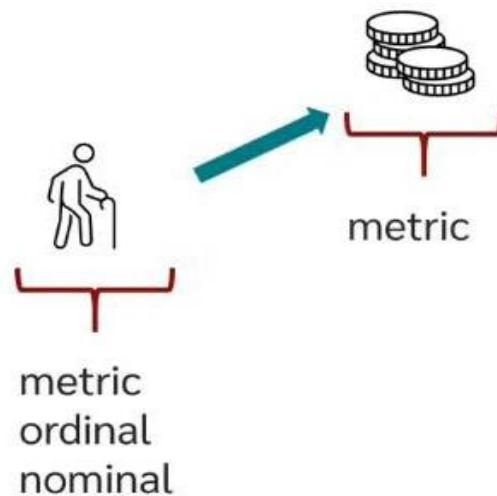


Logistic Regression

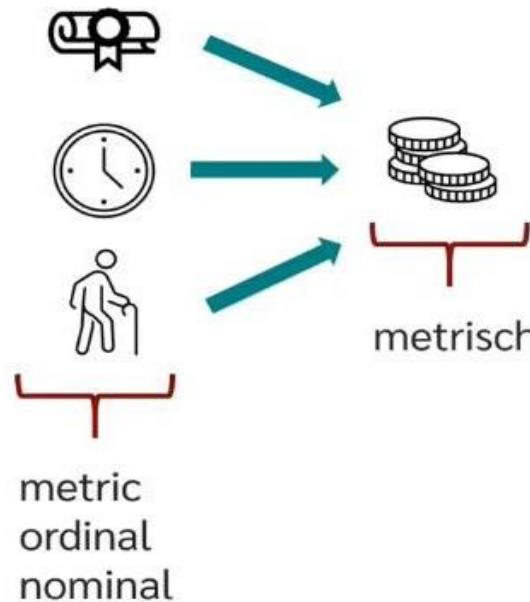


Forms of Regression Analysis

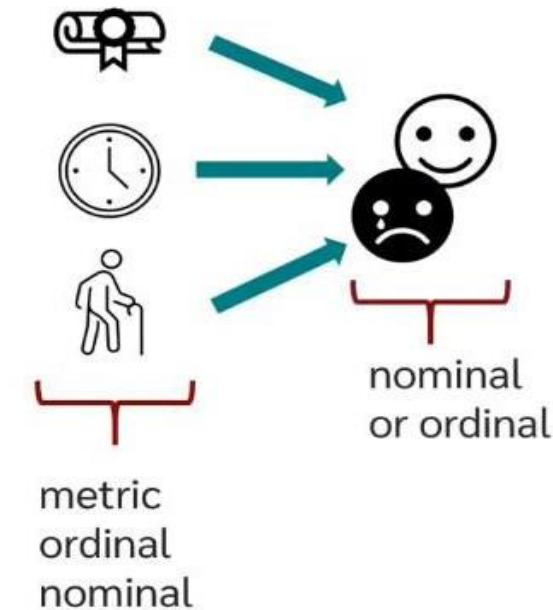
Simple linear Regression



Multiple linear Regression



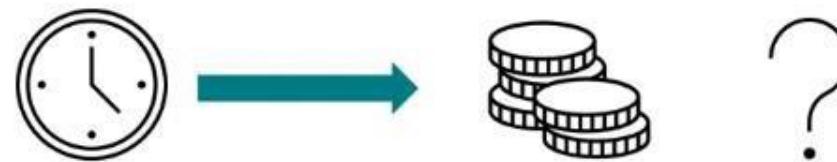
Logistic Regression



Forms of Regression Analysis

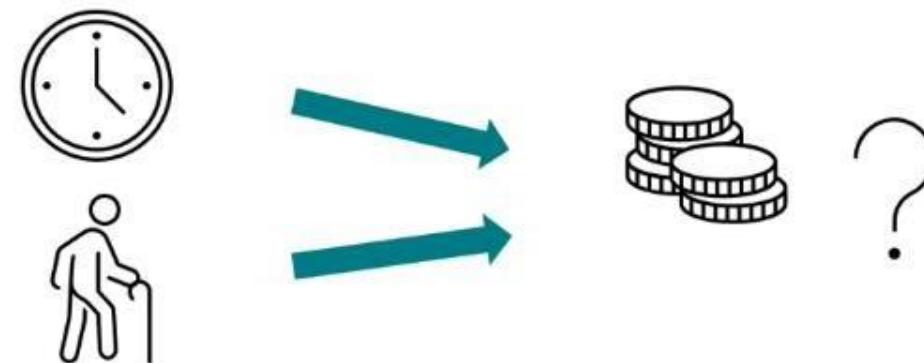
Simple linear regression

Does the **weekly working time** have an influence on the **hourly salary** of employees?



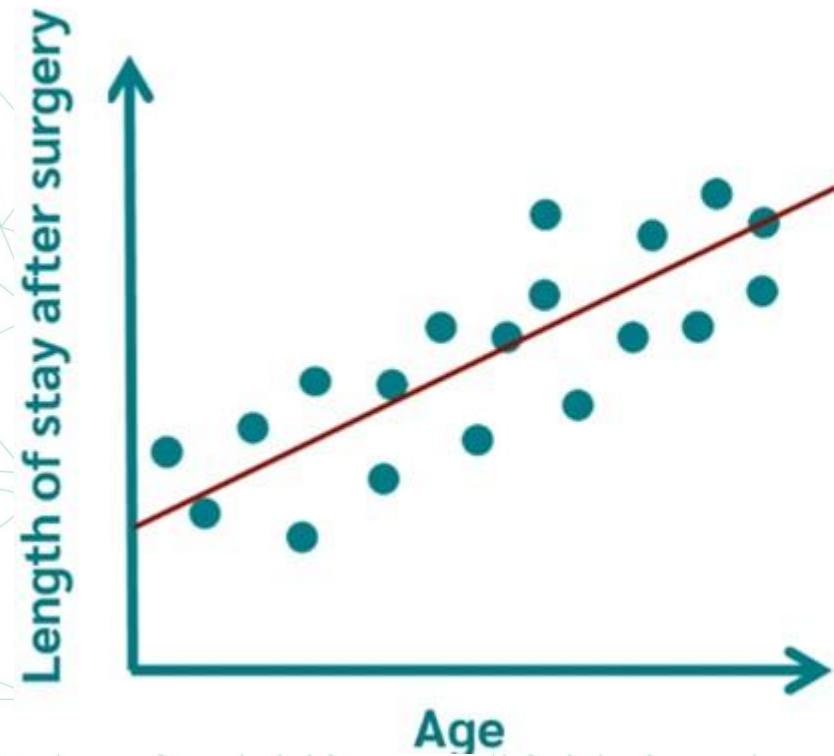
Multiple linear regression

Do the **weekly working hours** and the **age** of employees have an influence on their **hourly salary**?





Simple Linear Regression



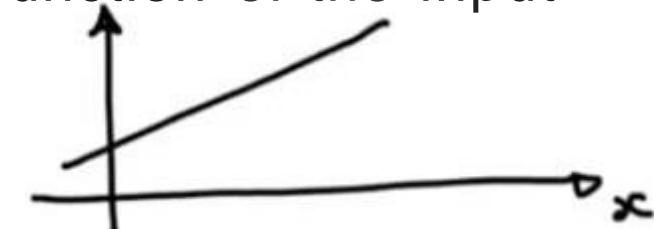
Estimated length of stay

$$\hat{y} = b \cdot x + a$$
$$\hat{y} = 0.14 \cdot x + 1.2$$
$$5.82 = 0.14 \cdot 33 + 1.2$$

The model

A simple linear regression model predicts the output as a linear function of the input feature x :

$$f(x; w_0, w_1) = w_0 + w_1 x$$



We refer to w_0 and w_1 as the *parameters* of the model.

To choose w_0 and w_1 , we are given a data set of previous input-output measurements:

$$\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(N)}, y^{(N)})\}$$

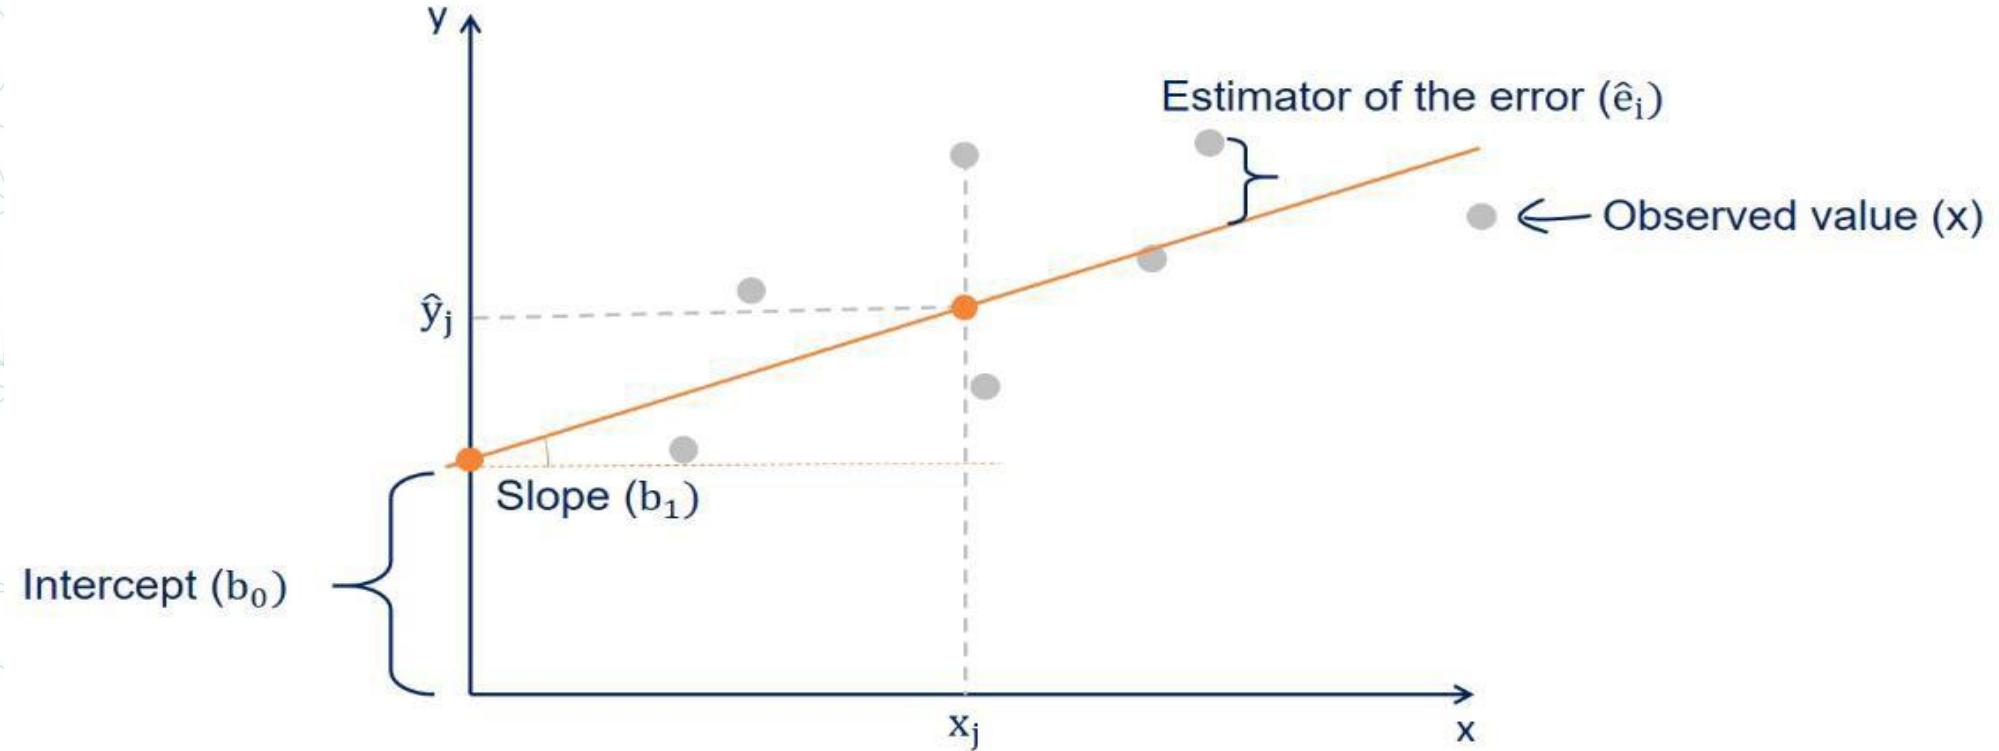
I will sometimes just write this as:

$$\{(x^{(n)}, y^{(n)})\}_{n=1}^N$$

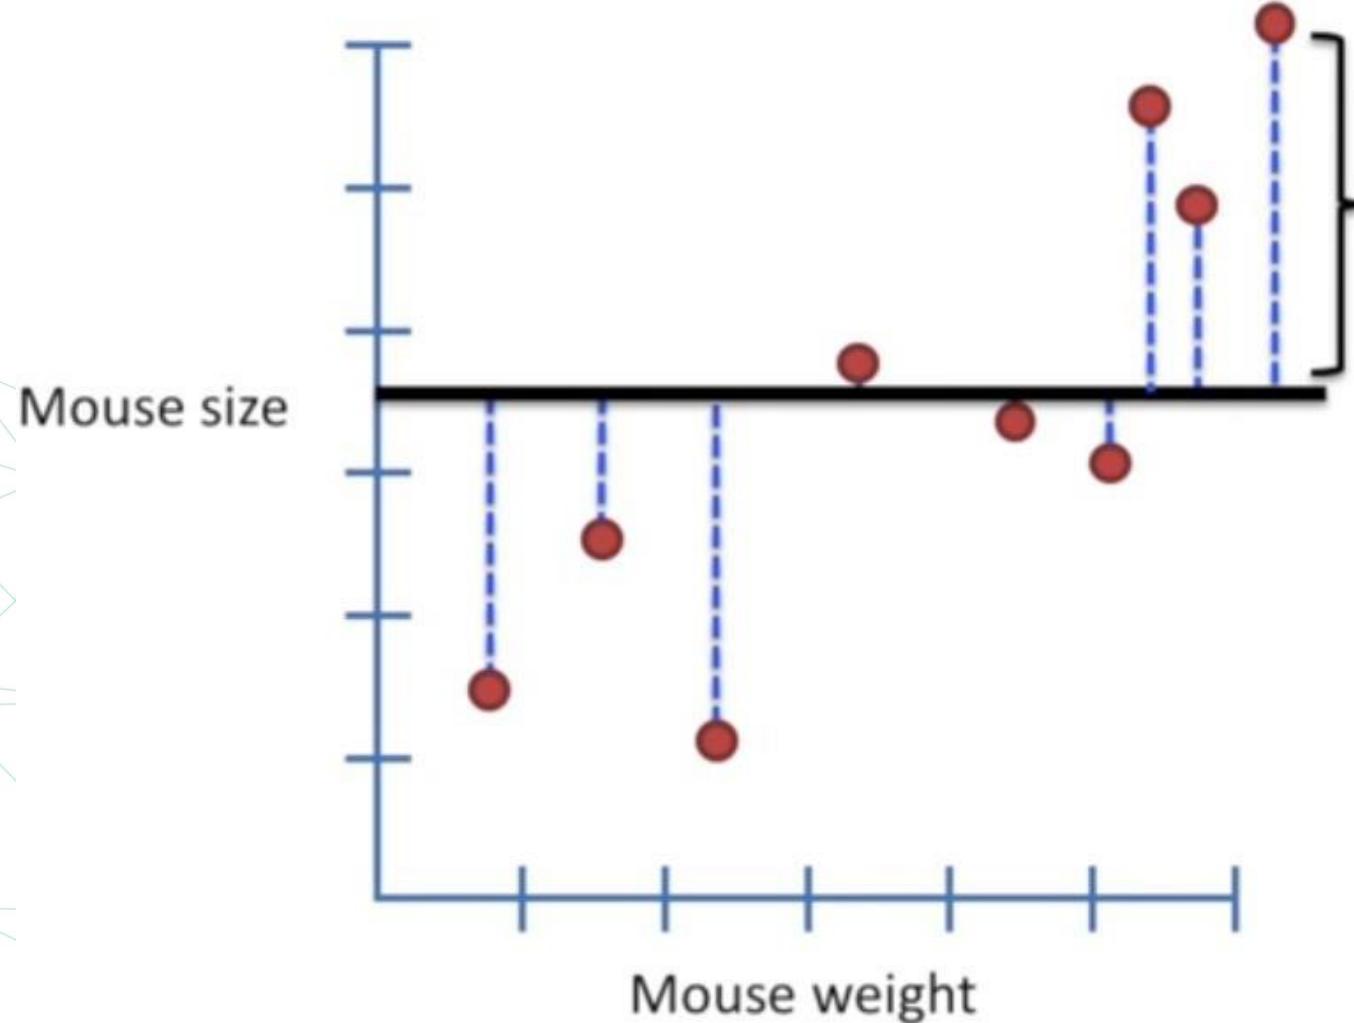
How do we choose w_0 and w_1 based on the data? We need some way to measure the "goodness" or "badness" of the parameters, given the data.

Simple Linear Regression

$$\hat{y}_i = b_0 + b_1 x_i$$



Simple Linear Regression

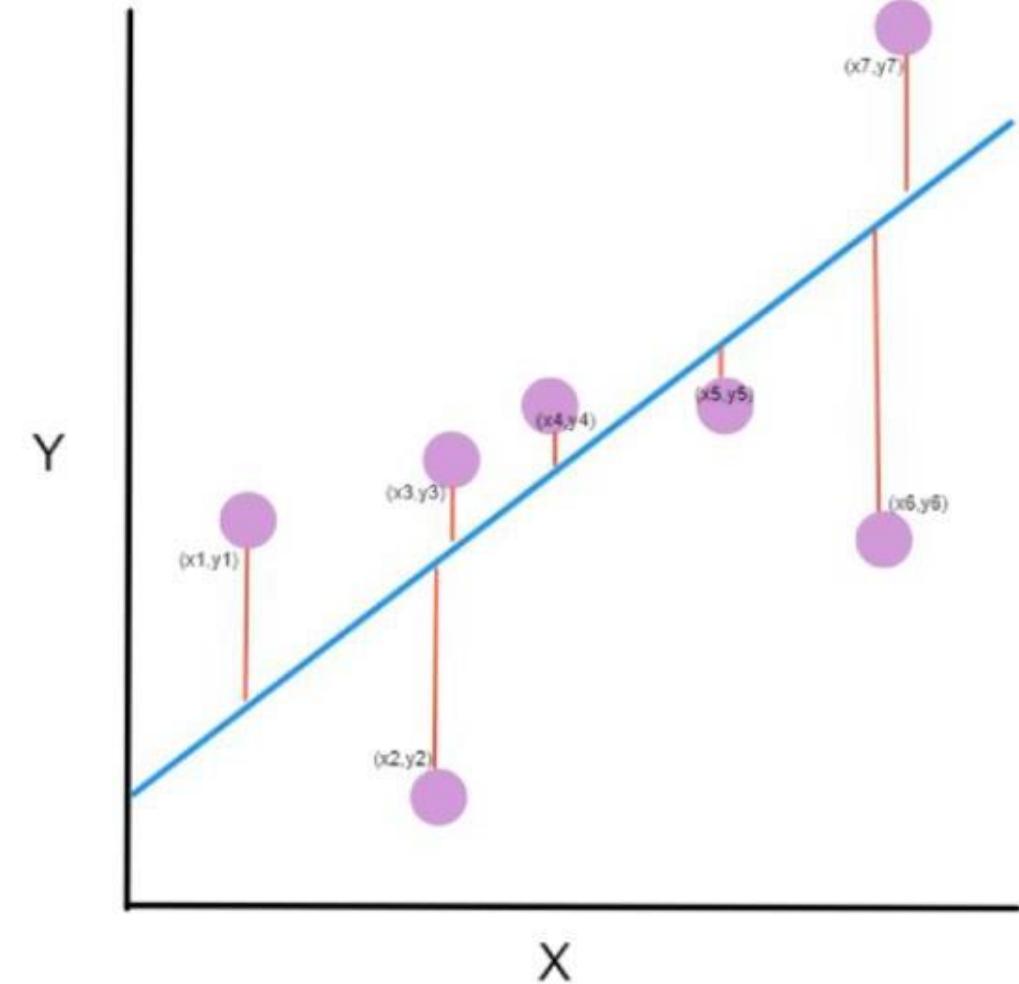


Second, measure the distance from the line to the data, square each distance, and then add them up.

Terminology alert!
The distance from a line to a data point is called a “**residual**”.

Error/Loss/Cost Function

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$





Multiple Linear Regression

- Unlike simple linear regression, multiple linear regression can include two or more independent variables.
- The goal is to estimate one variable based on several other variables.
- The variable to be estimated is called the **dependent variable (criterion)**. The variables that are used for prediction are called **independent variables (predictors)**.
- Multiple linear regression is often used in empirical social research as well as in market research. In both areas it is of interest to find out what influence different factors have on a variable.



Good models require
multiple regressions, in order
to address the higher
complexity of problems



$$\text{College GPA} = 0.275 + 0.0017 * \text{SAT}$$





MULTIPLE REGRESSION EQUATION

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 + \cdots + b_kx_k$$

↑ ↑ ↑
independent independent independent
variable variable variable

Multiple Linear Regression

Simple linear
Regression

$$\hat{y} = b \cdot x + a \quad \rightarrow$$

Multiple linear
Regression

$$\hat{y} = b_1 \cdot x_1 + b_2 \cdot x_2 + \dots + b_k \cdot x_k + a$$

- The coefficients are interpreted similarly to the linear regression equation.
- If all independent variables are 0, the value a is obtained.
- If an independent variable changes by one unit, the associated coefficient b indicates by how much the dependent variable changes.





Metrics

$$MAE = \frac{1}{n} \sum_{i=1}^n |\hat{y}_i - y_i|$$

$$MSE = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

less robust to outliers

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2}$$

less robust to outliers

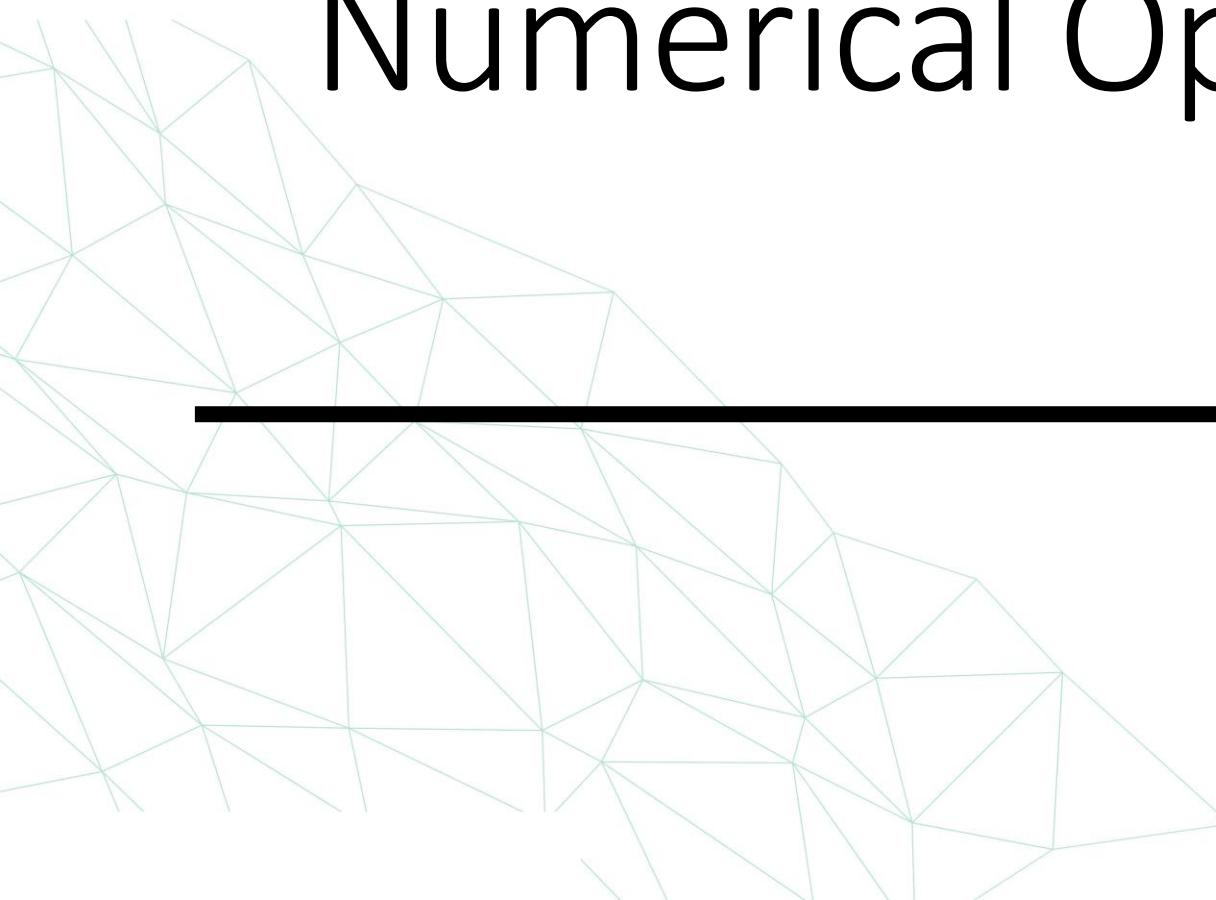
If you want to train a model which focuses on **reducing large outlier errors** then **MSE** is the better choice, whereas if this **isn't important** and you would prefer greater interpretability then **MAE** would be **better**.

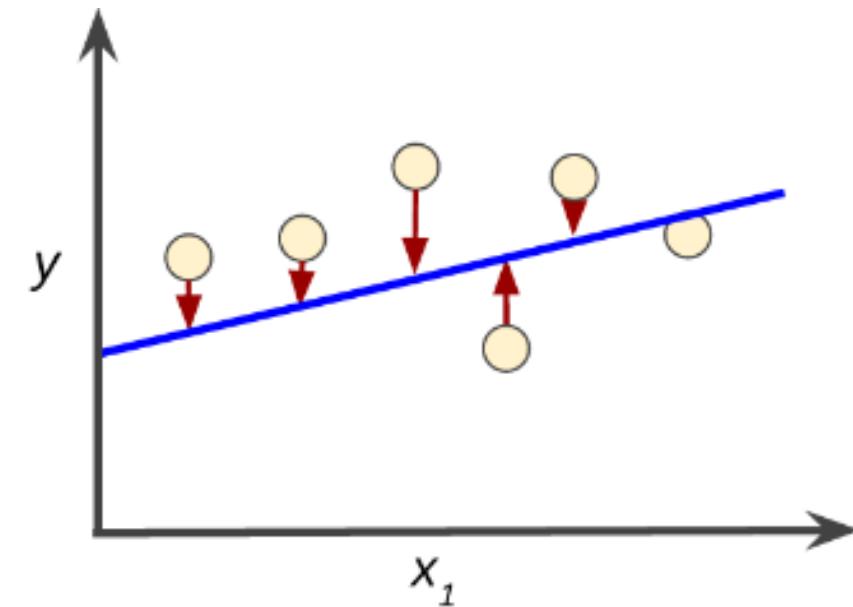
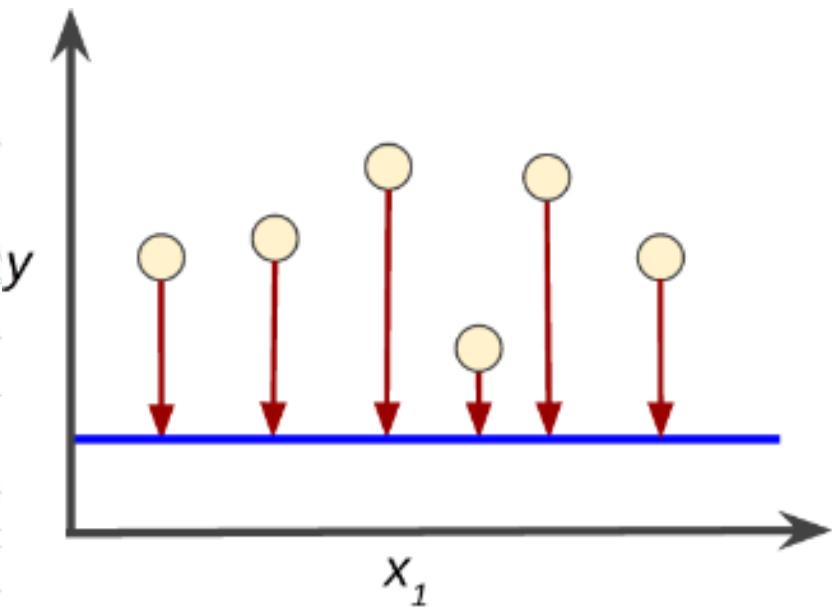
So a robust system or metric must be **less affected by outliers**. In this scenario it is easy to conclude that **MSE** may be less robust than **MAE**, since the **squaring** of the errors will **enforce a higher importance on outliers**.

Q&A

Questions and answers

Numerical Optimization





We have established the fact that all the straight lines are just different combination of model parameters a_0 and a_1

$$J(a) = \frac{1}{2m} \sum_{i=1}^m (a_0 + a_1 x_1^{(i)} - y_{i \text{ (act)}})^2$$

and **Cost function (J)** is a function of parameter space $a = (a_0, a_1)$.

Therefore, by changing a_0 and a_1 we can change the **cost function**.

We will keep changing a_0 and a_1 till we find a combination where cost function is **minimized**

Hypothesis: $h_{\theta}(x) = \theta_0 + \theta_1 x$

Parameters: θ_0, θ_1

Cost Function: $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

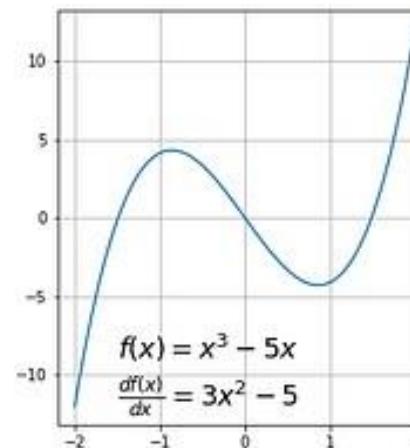
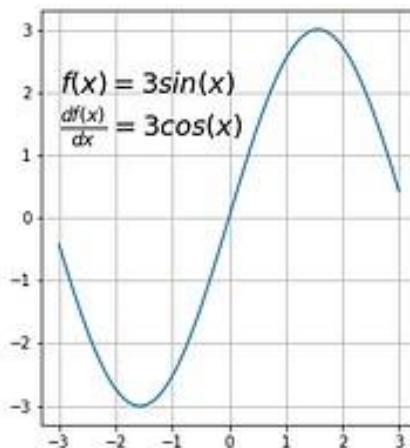
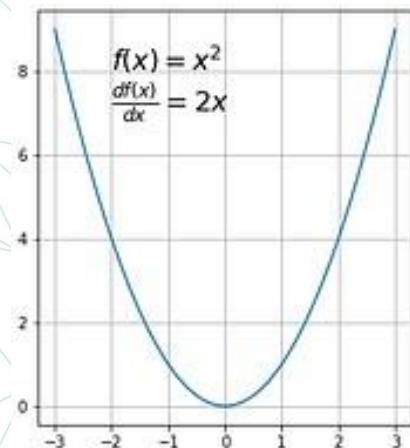
Goal: $\underset{\theta_0, \theta_1}{\text{minimize}} J(\theta_0, \theta_1)$

The Gradient Descent Algorithm

- It's a powerful tool used to **optimize the parameters** of a model and minimize its error or **cost** function.
- In simple terms, gradient descent is a process of **finding the minimum point of a function** by following the steepest descent direction.
- It's used in a wide range of **machine learning algorithms**, Such as
 - **Linear regression,**
 - **Logistic regression,**
 - **Neural networks.**

Gradient

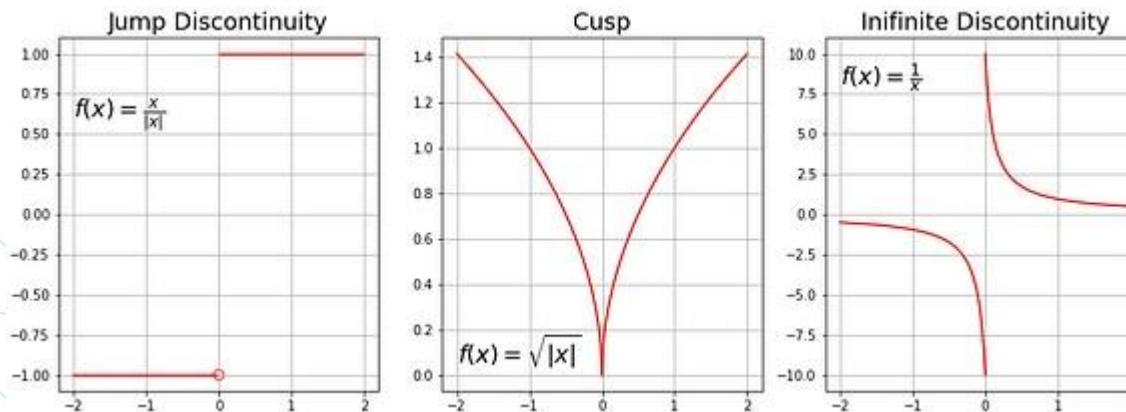
- Gradient descent algorithm does not work for all functions. There are two specific requirements. A function has to be:
 - **Differentiable** (has a derivative for each point in its domain)
 - **convex**



differentiable functions

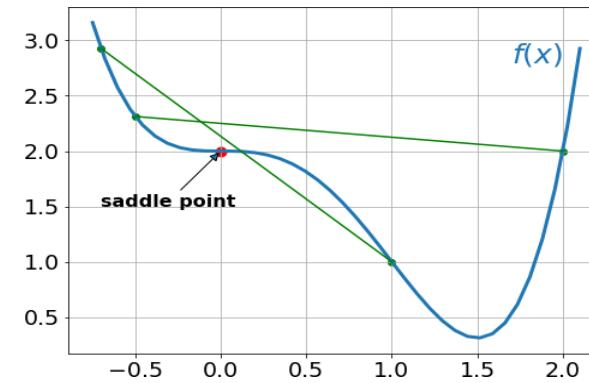
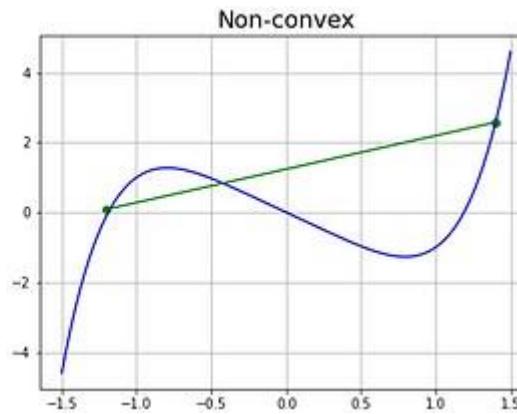
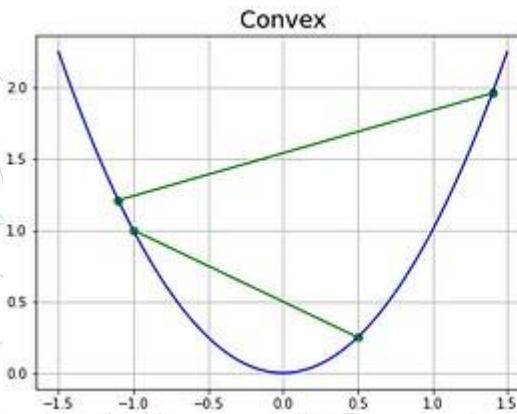
Gradient

- Typical non-differentiable functions have a step a cusp or a discontinuity:





Gradient

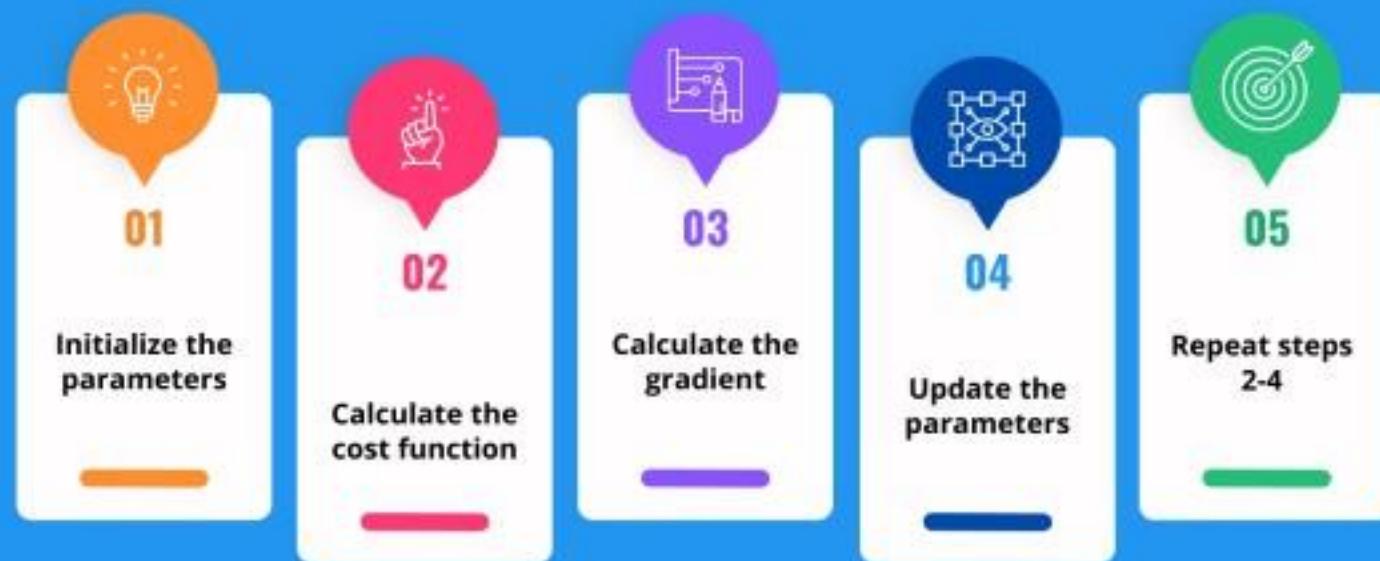


Gradient

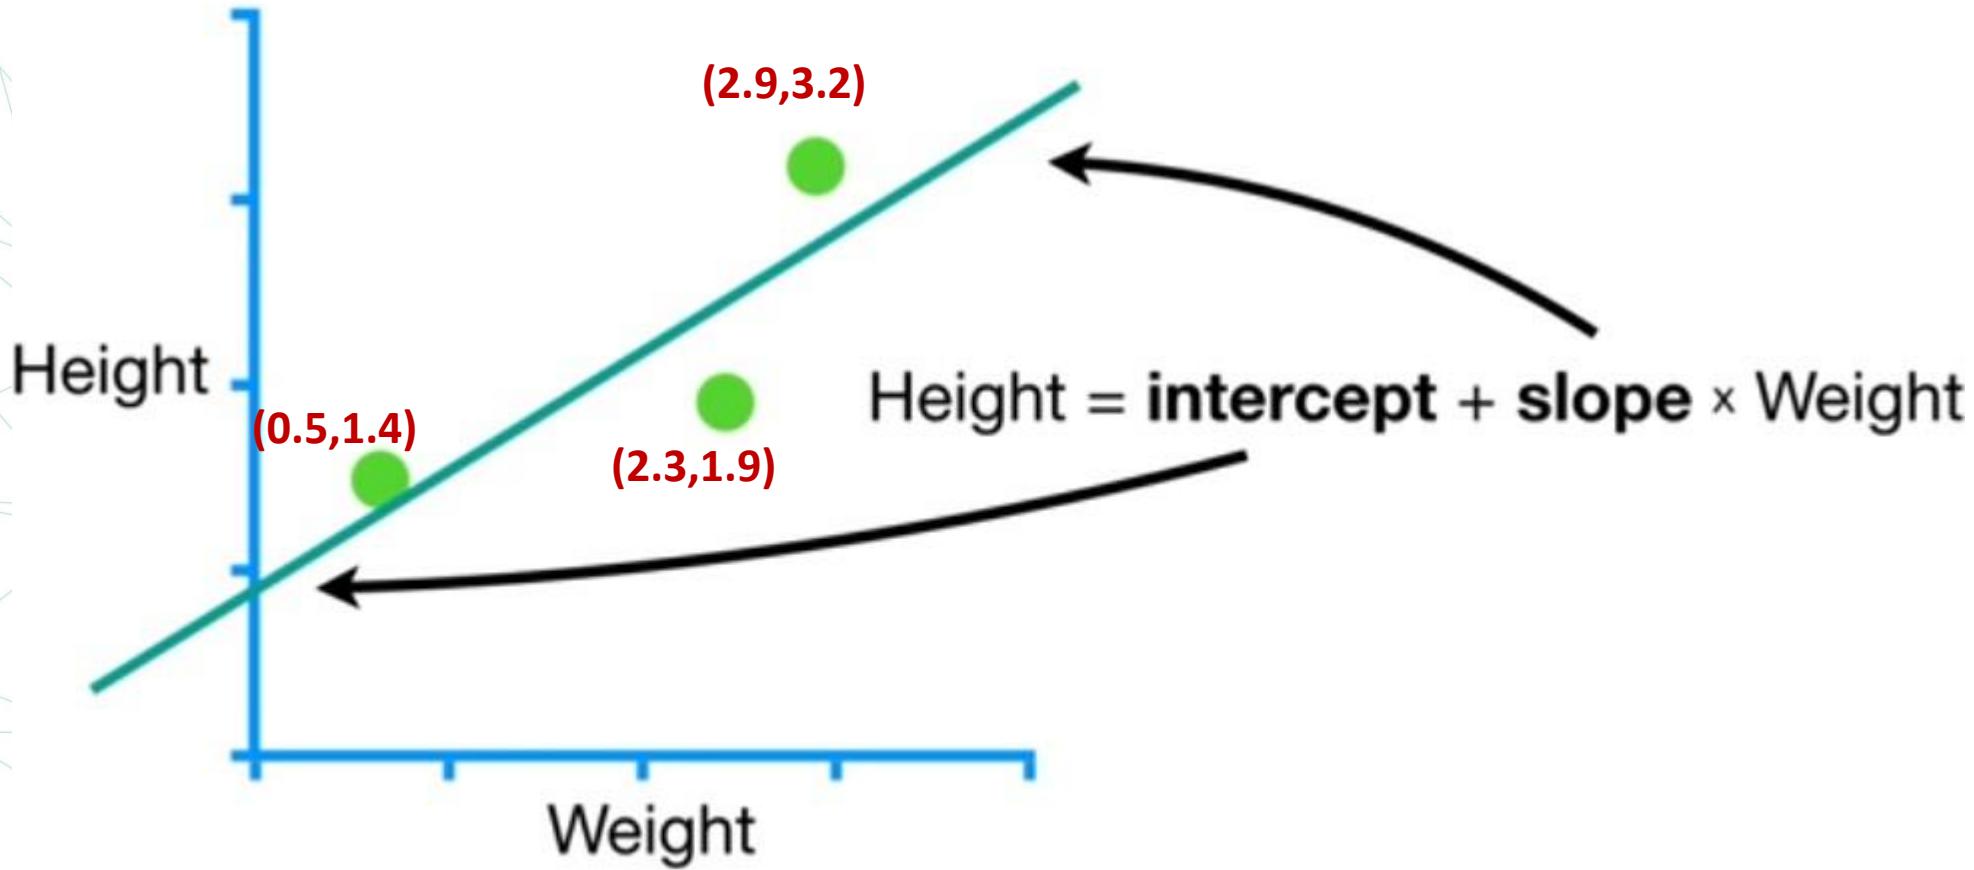
- Intuitively it is a **slope of a curve** at a given point in a specified **direction**.
- In the case of a **univariate function**, it is simply the **first derivative at a selected point**.
- In the case of a **multivariate function**, it is a **vector of derivatives** in each main direction (along variable axes). Because we are interested only in a slope along one axis and we don't care about others these derivatives are called **partial derivatives**.
- A gradient for an **n-dimensional function** $f(x)$ at a given point p is defined as follows:

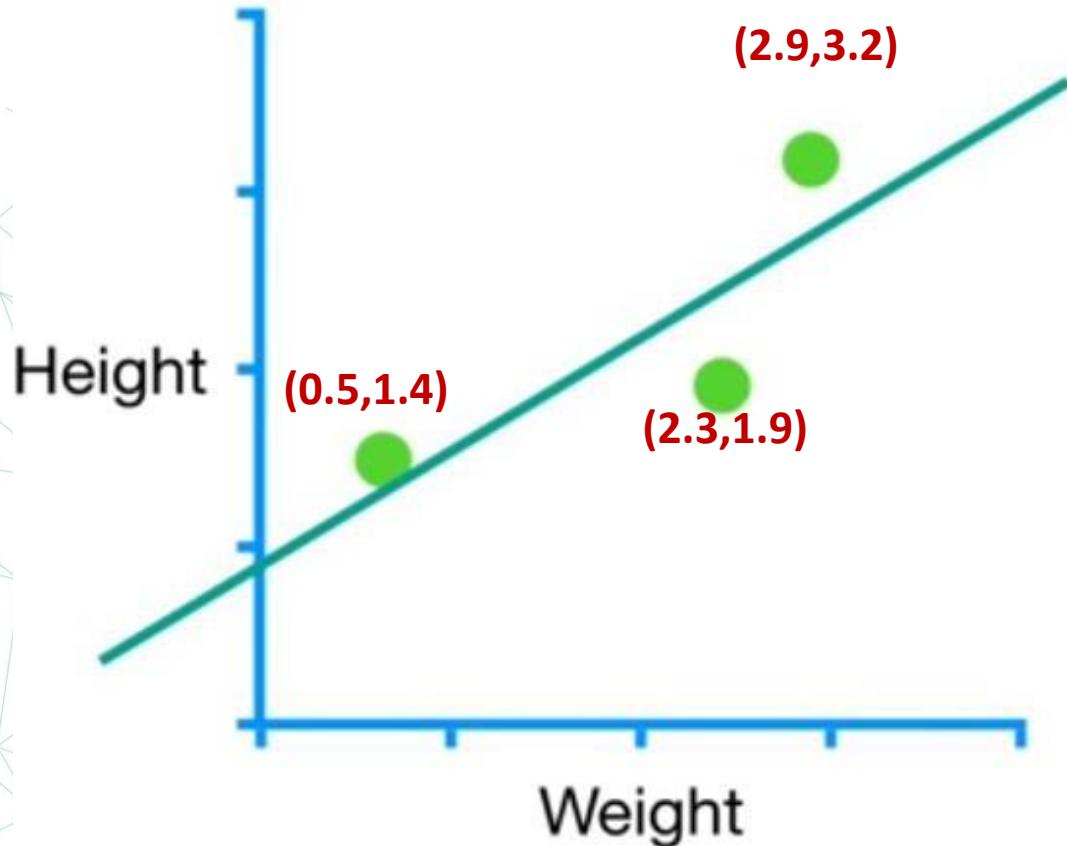
$$\nabla f(p) = \begin{bmatrix} \frac{\partial f}{\partial x_1}(p) \\ \vdots \\ \frac{\partial f}{\partial x_n}(p) \end{bmatrix}$$

The Step-by-step Process of Performing Gradient Descent



When we fit a line with **Linear Regression**, we optimize the **Intercept** and **Slope**.

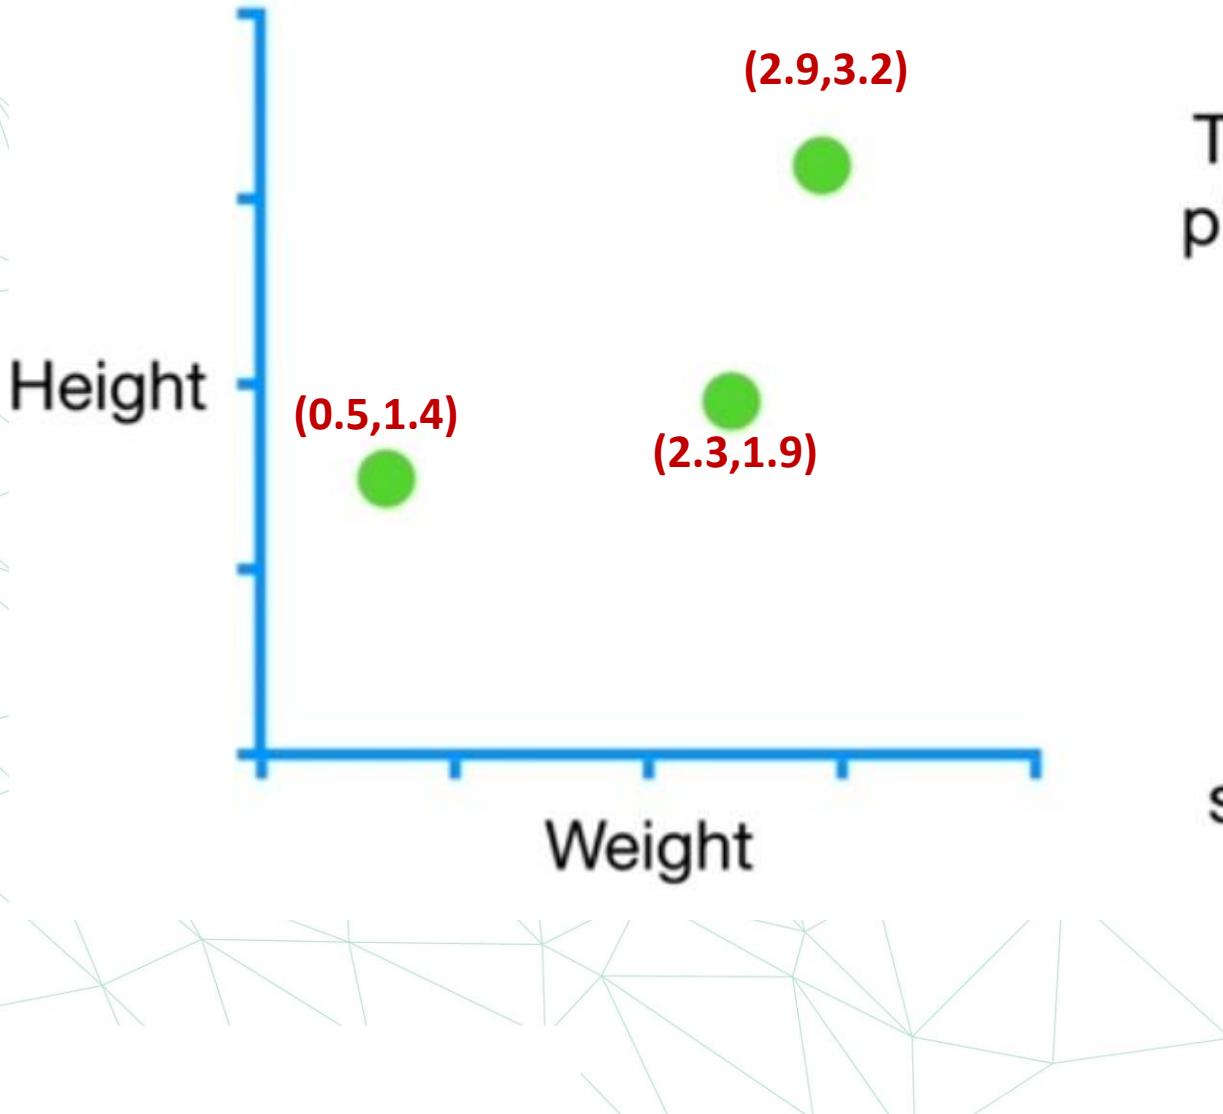




Predicted Height = intercept + slope × Weight

So for now, let's just plug in
the **Least Squares** estimate
for the **Slope, 0.64**.

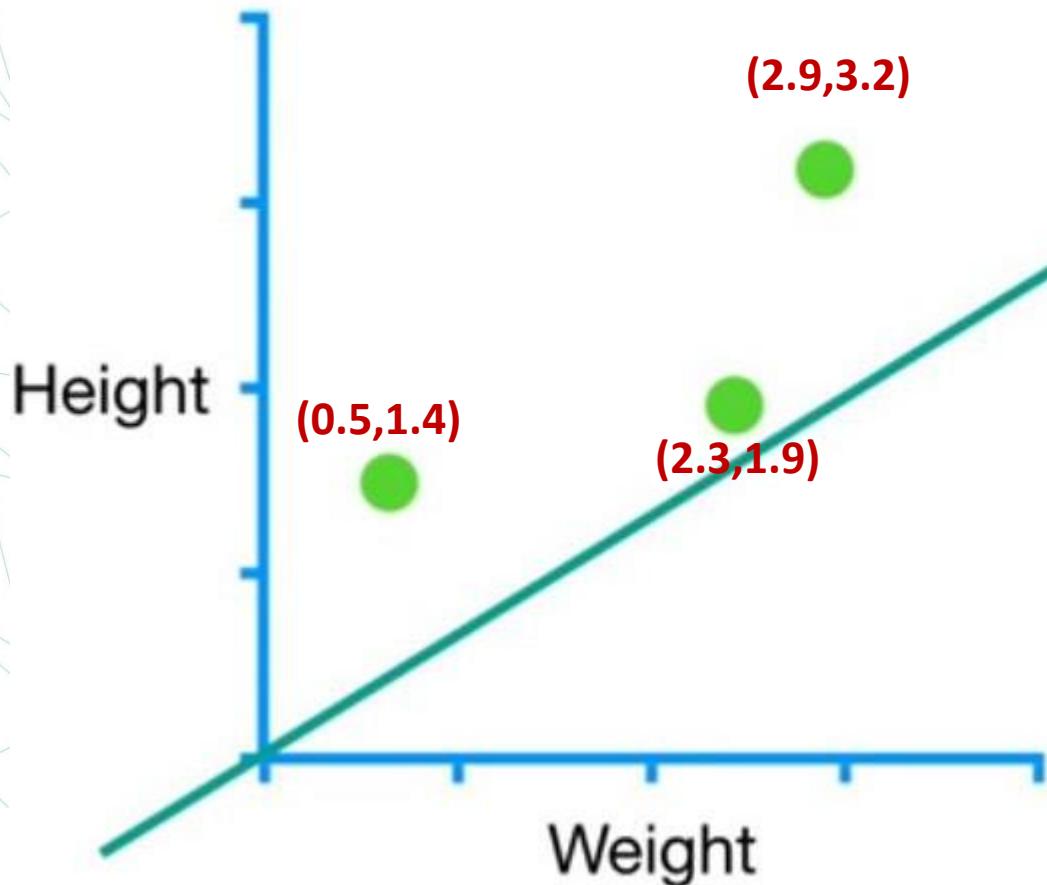




$$\text{Predicted Height} = \text{intercept} + 0.64 \times \text{Weight}$$

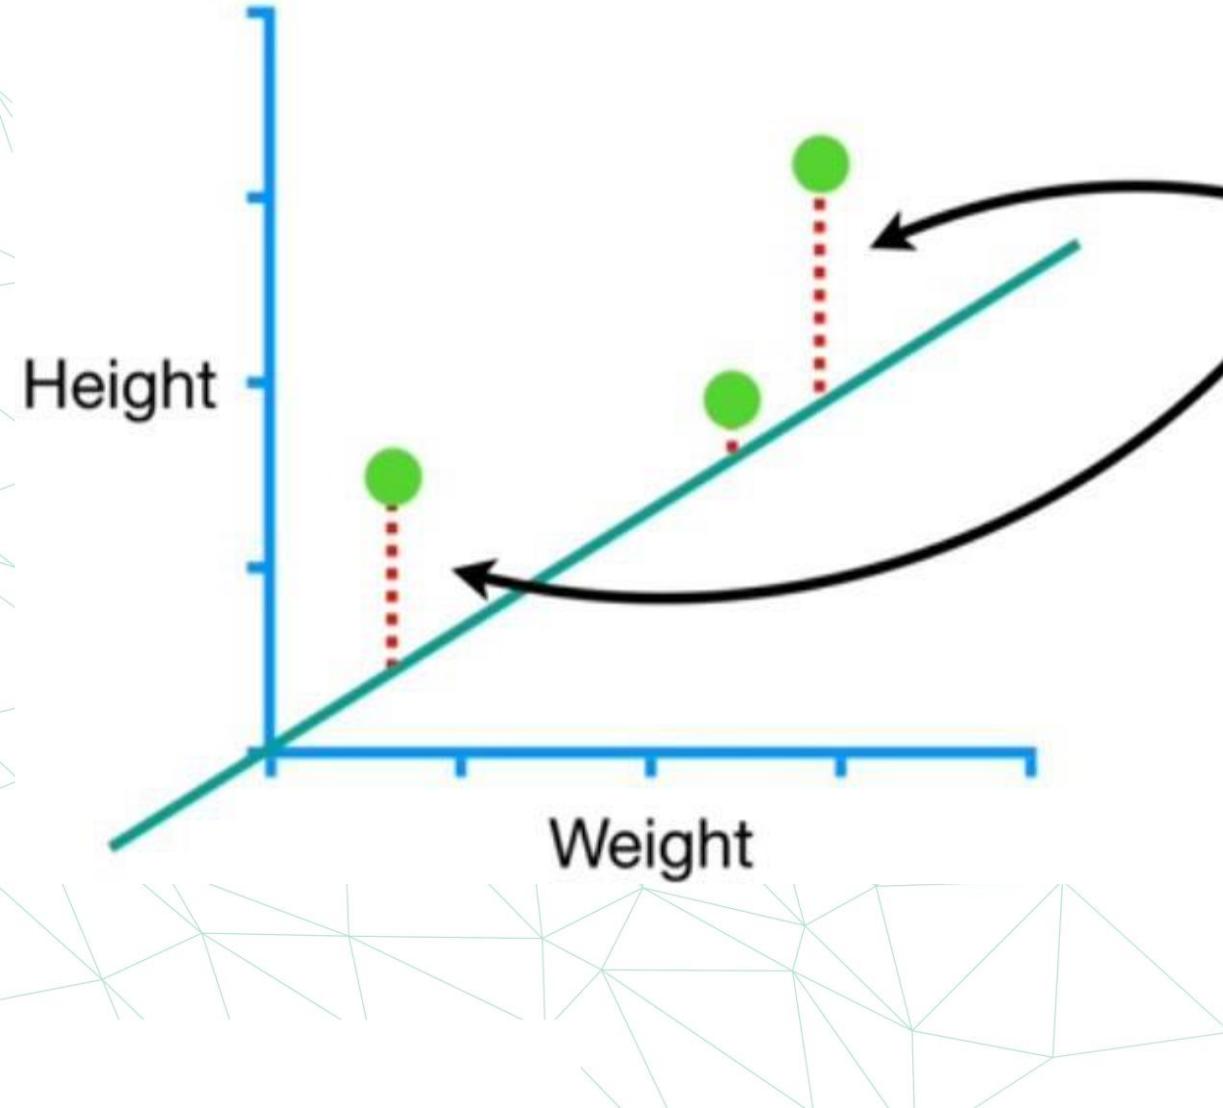
The first thing we do is pick a random value for the **Intercept**.

This is just an initial guess that gives **Gradient Descent** something to improve upon.



$$\text{Predicted Height} = \boxed{0} + 0.64 \times \text{Weight}$$

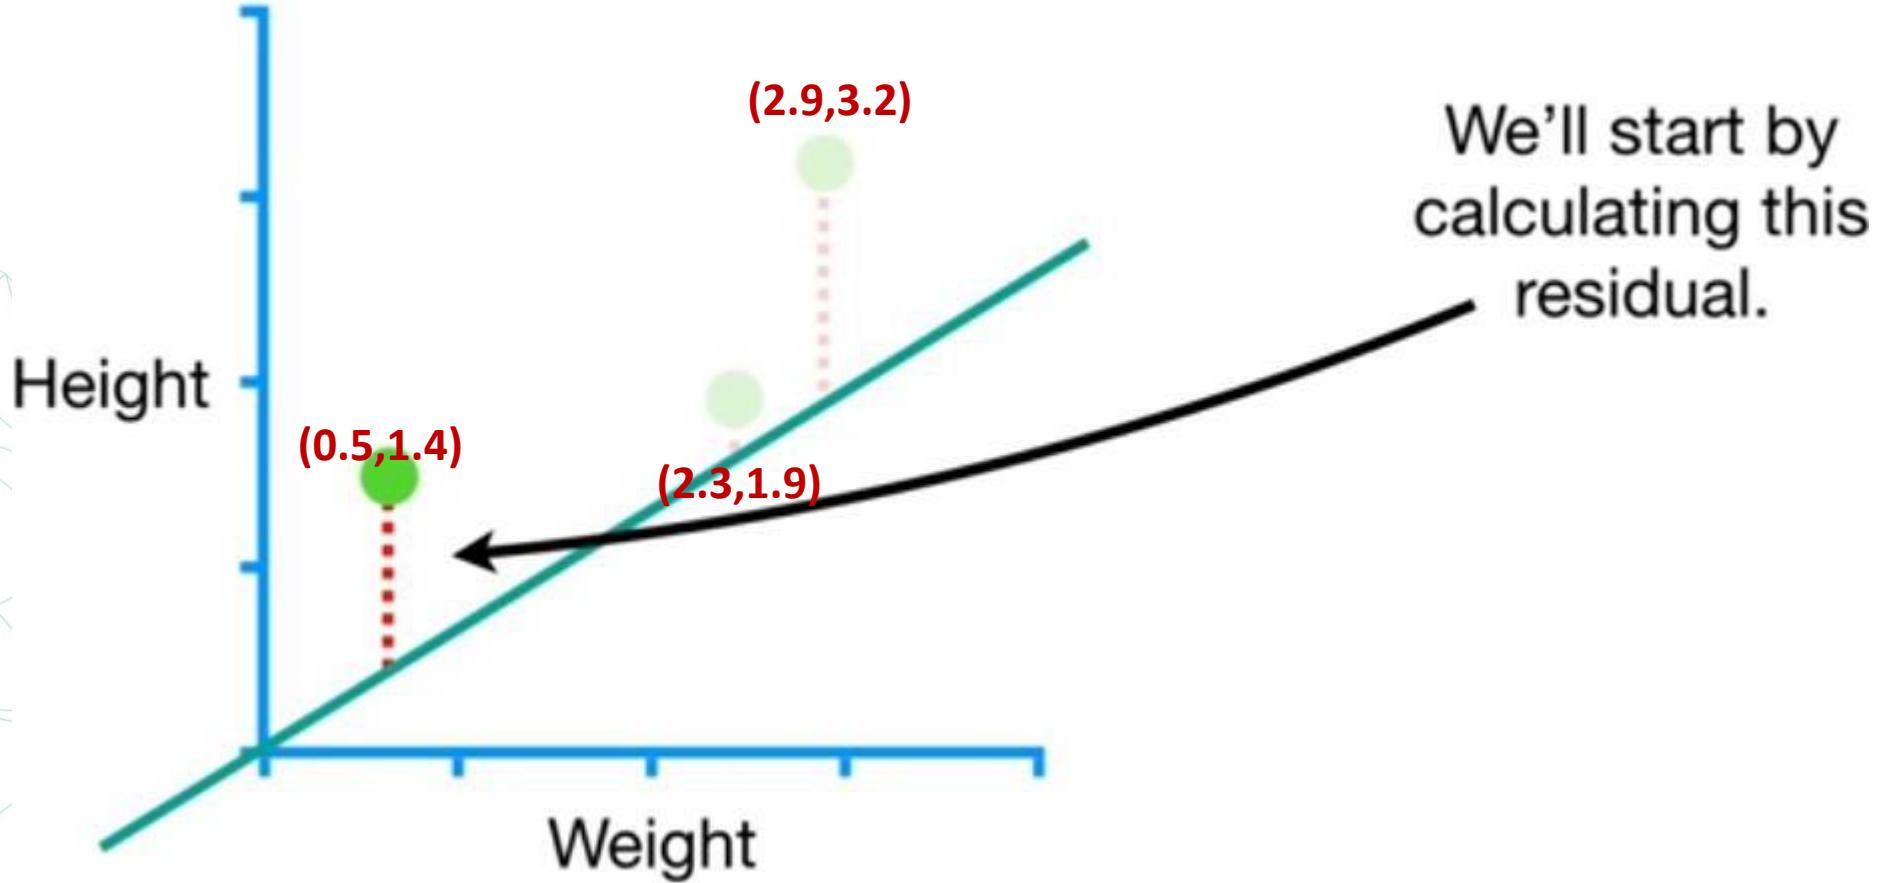
And that gives us the equation for this line.

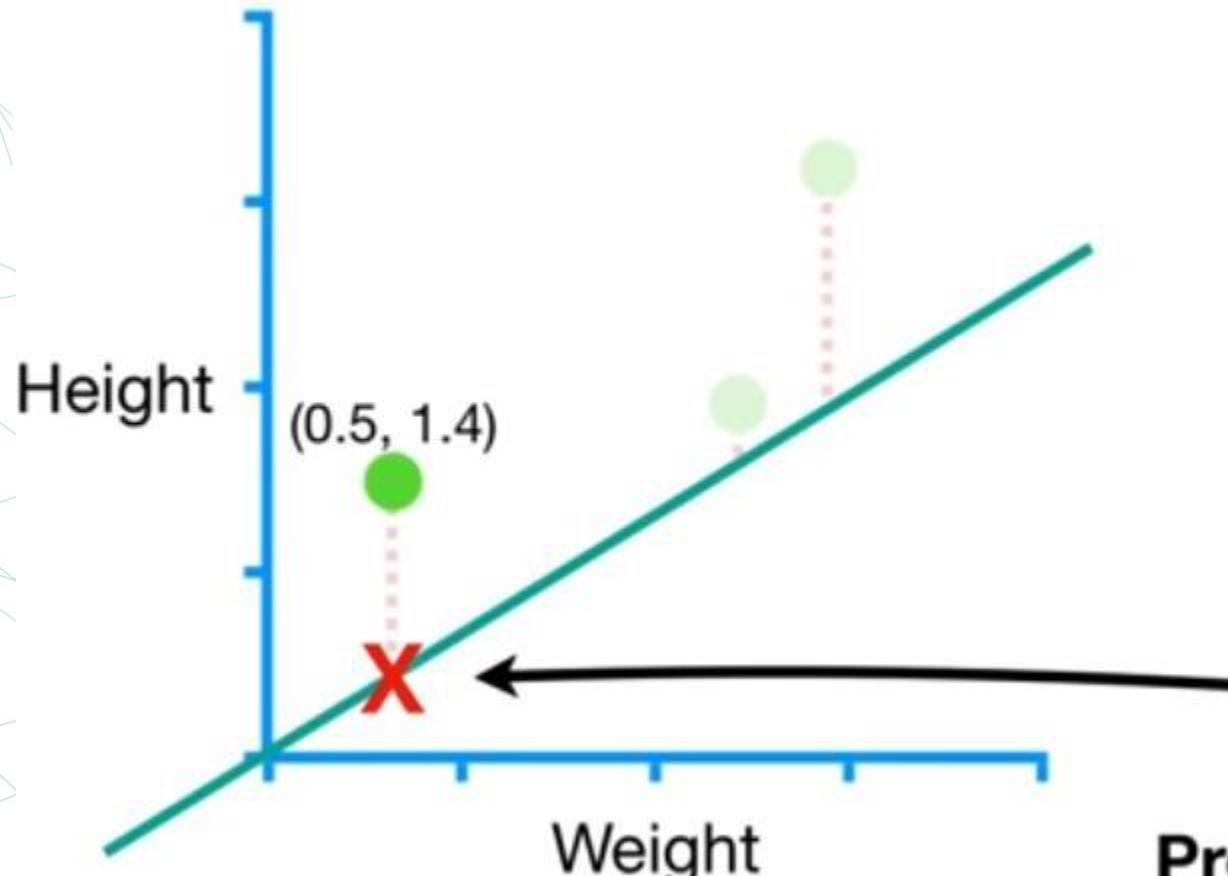


In this example, we will evaluate how well this line fits the data with the **Sum of the Squared Residuals.**

NOTE: In Machine Learning lingo, The Sum of the Squared Residuals is a type of **Loss Function.**

$$\text{residual sum of squares} = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

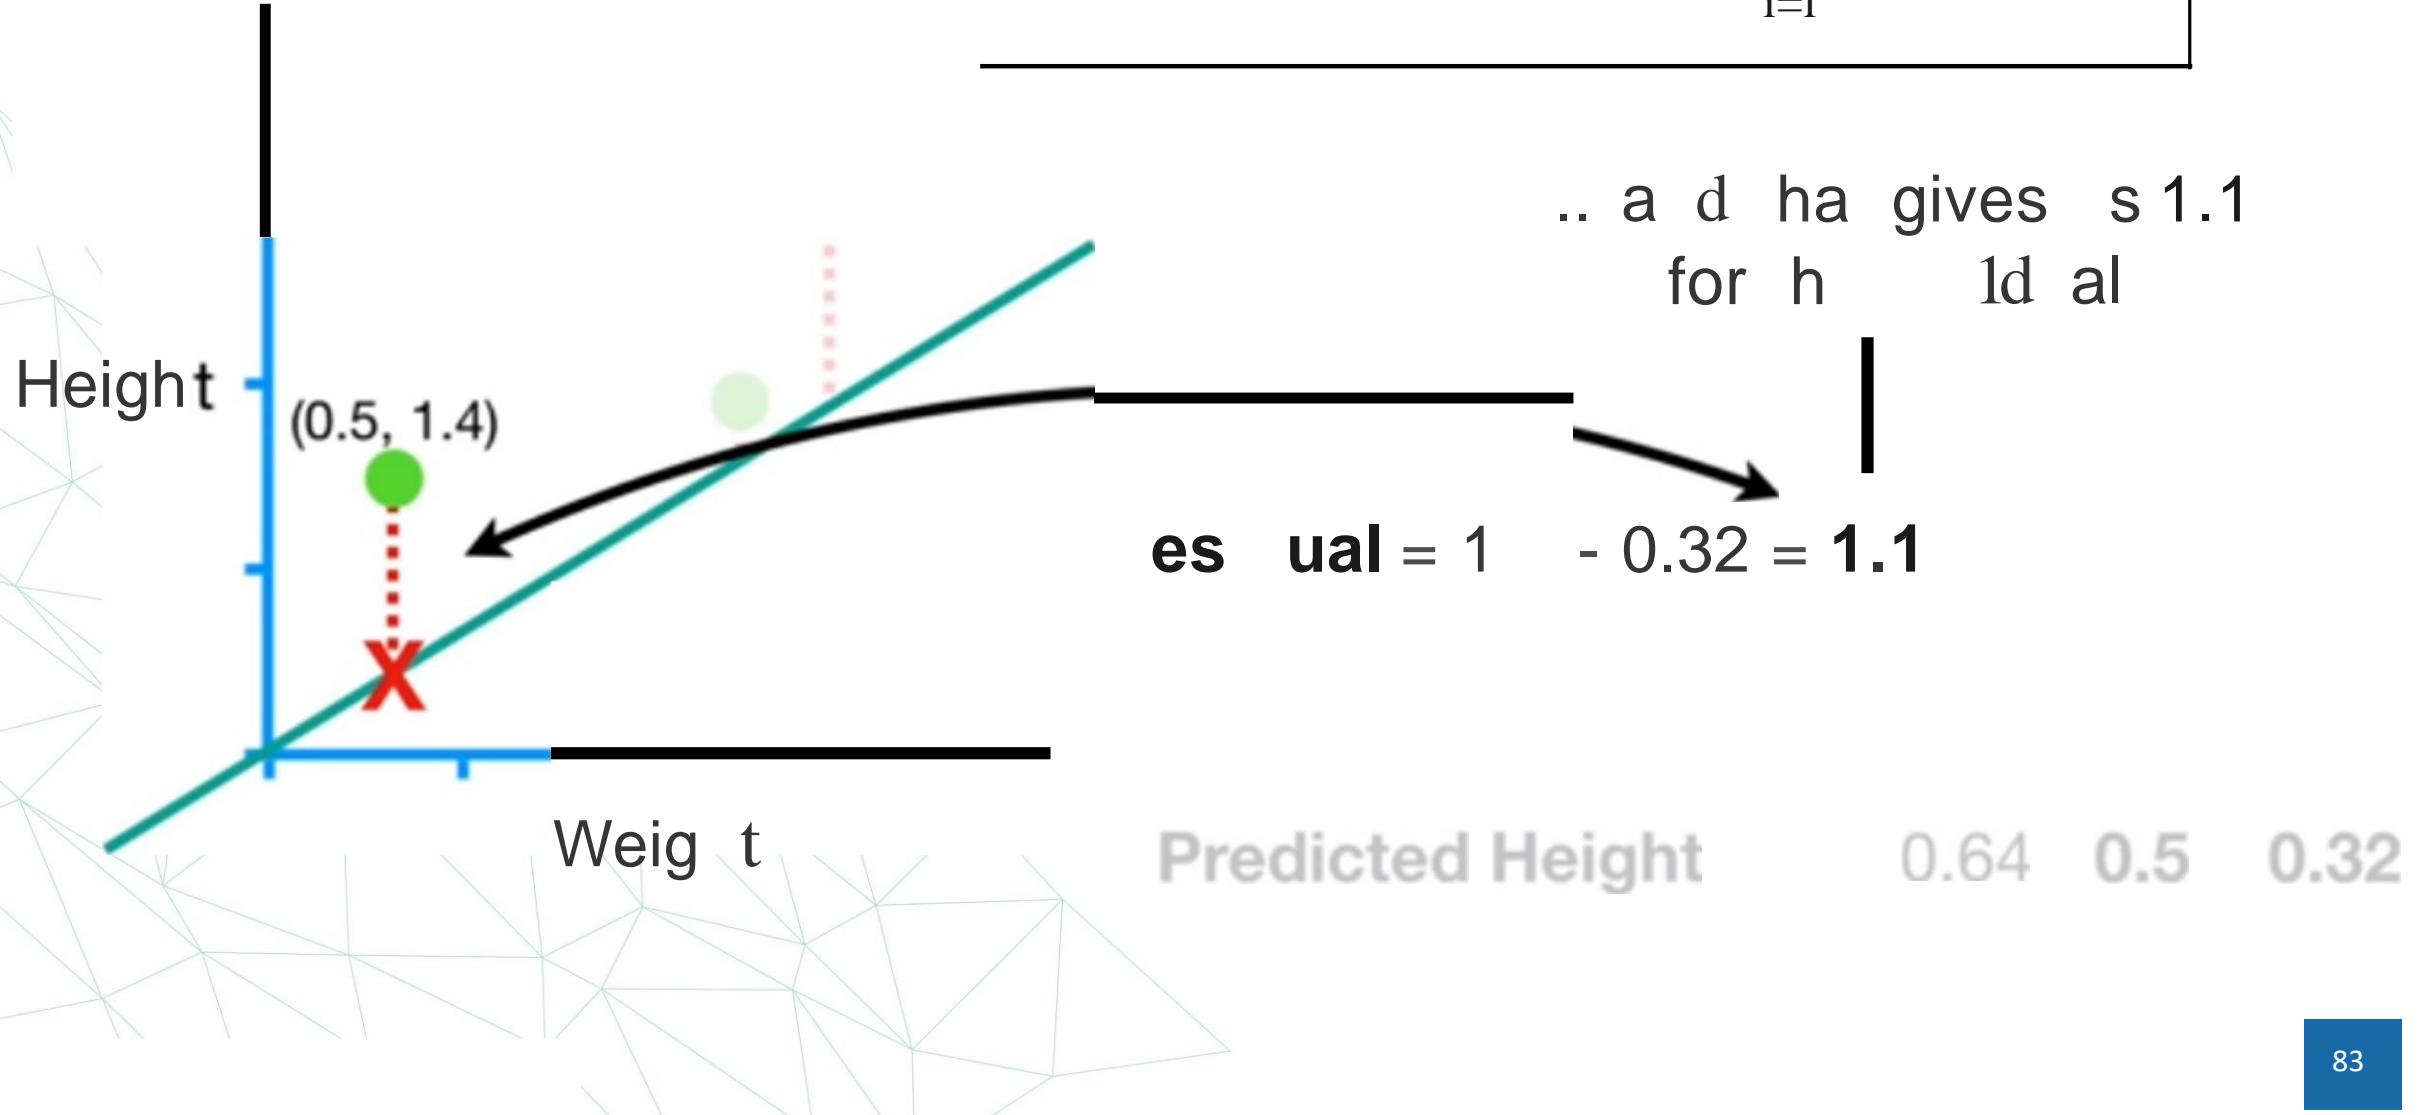




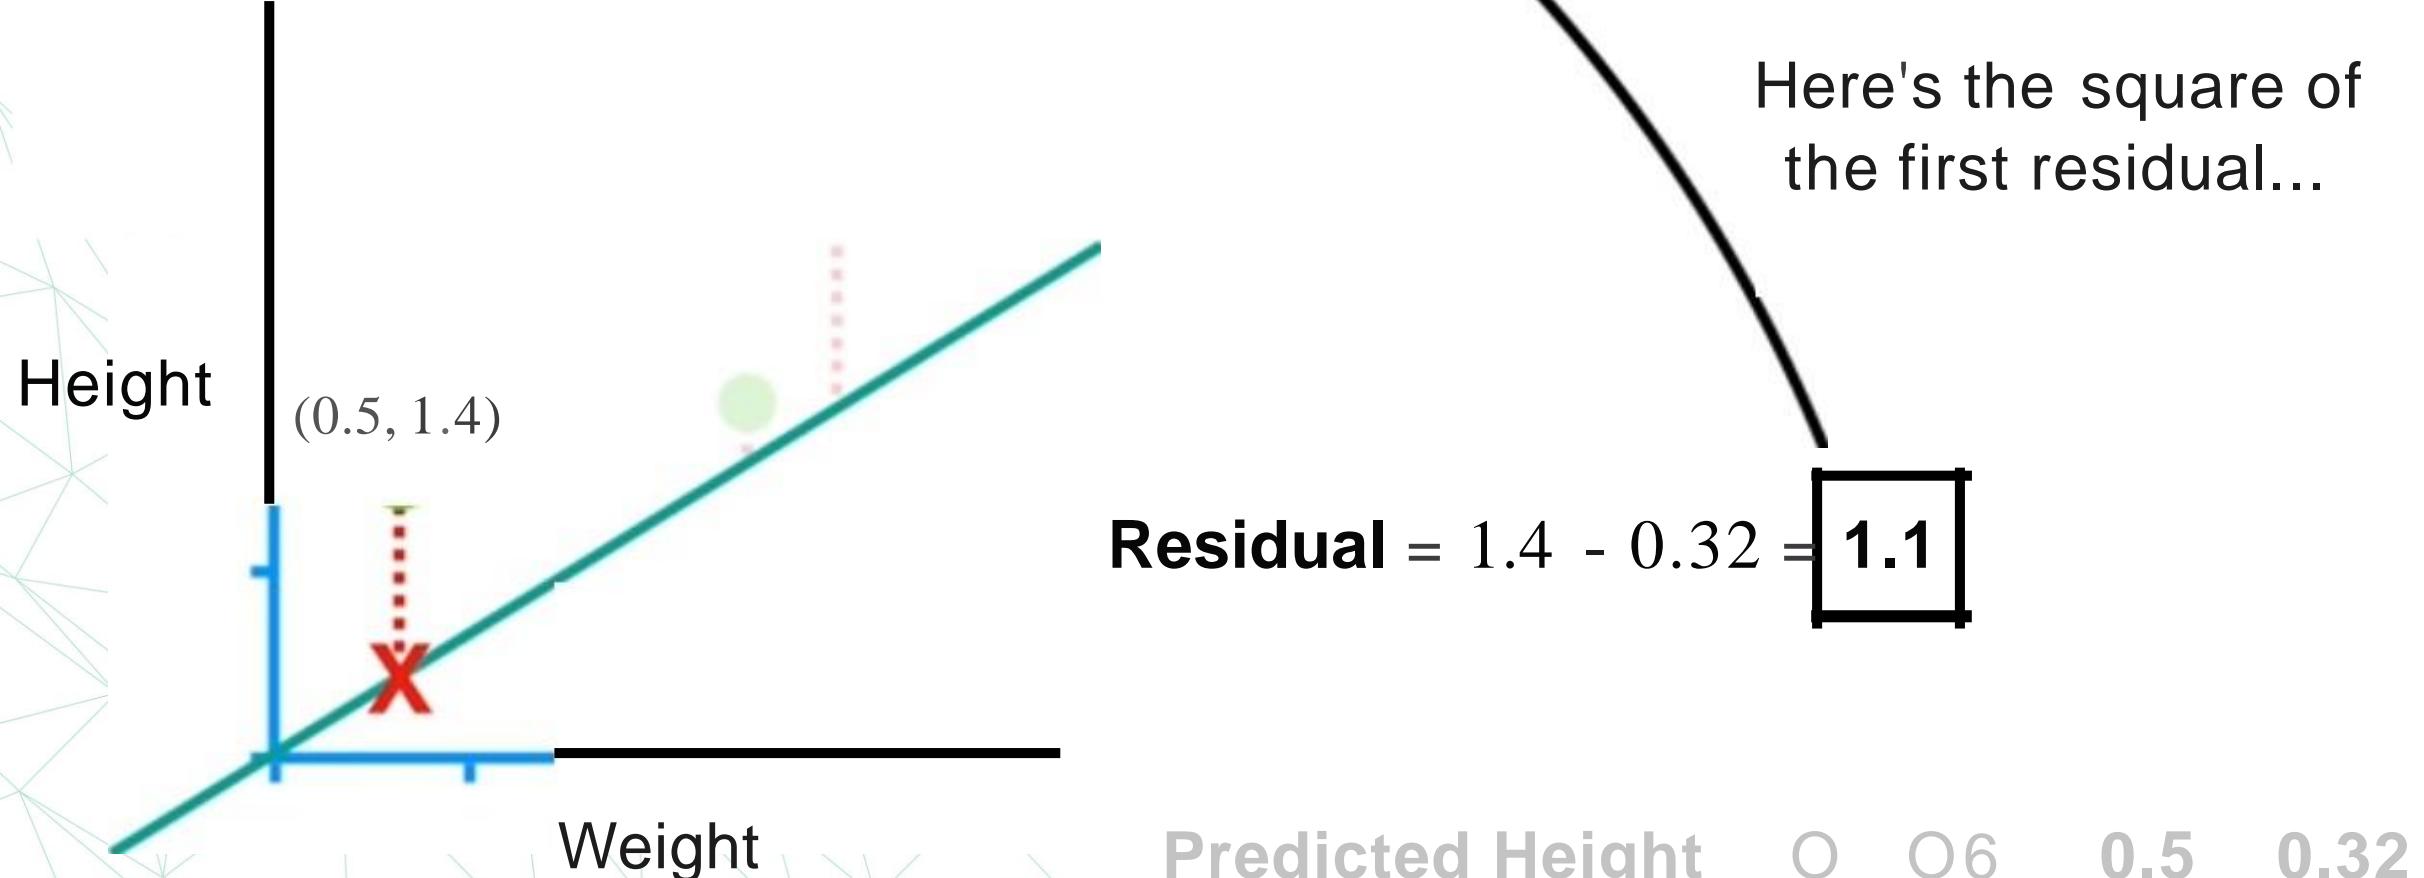
...and the **Predicted Height**
is **0.32**.

$$\text{Predicted Height} = 0 + 0.64 \times 0.5 = 0.32$$

$$\text{residual sum of squares} = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

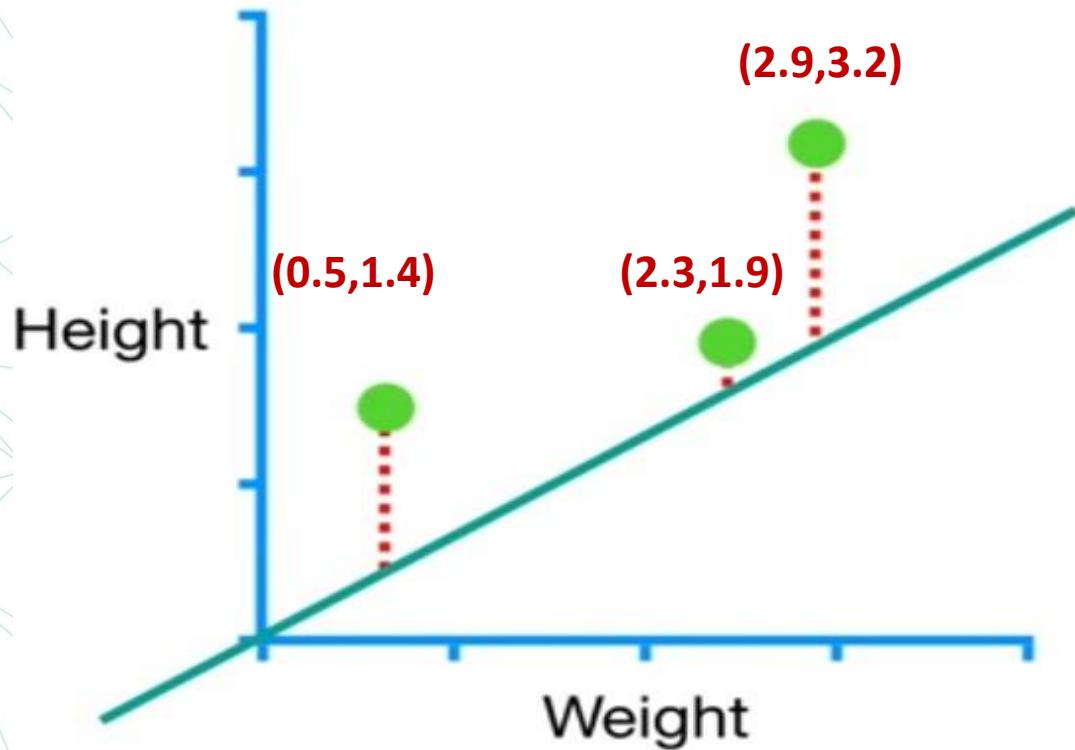


Sum of squared residuals = 1.1^2



$$\sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Sum of squared residuals = $1.1^2 + 0.4^2 + 1.3^2 = 3.1$



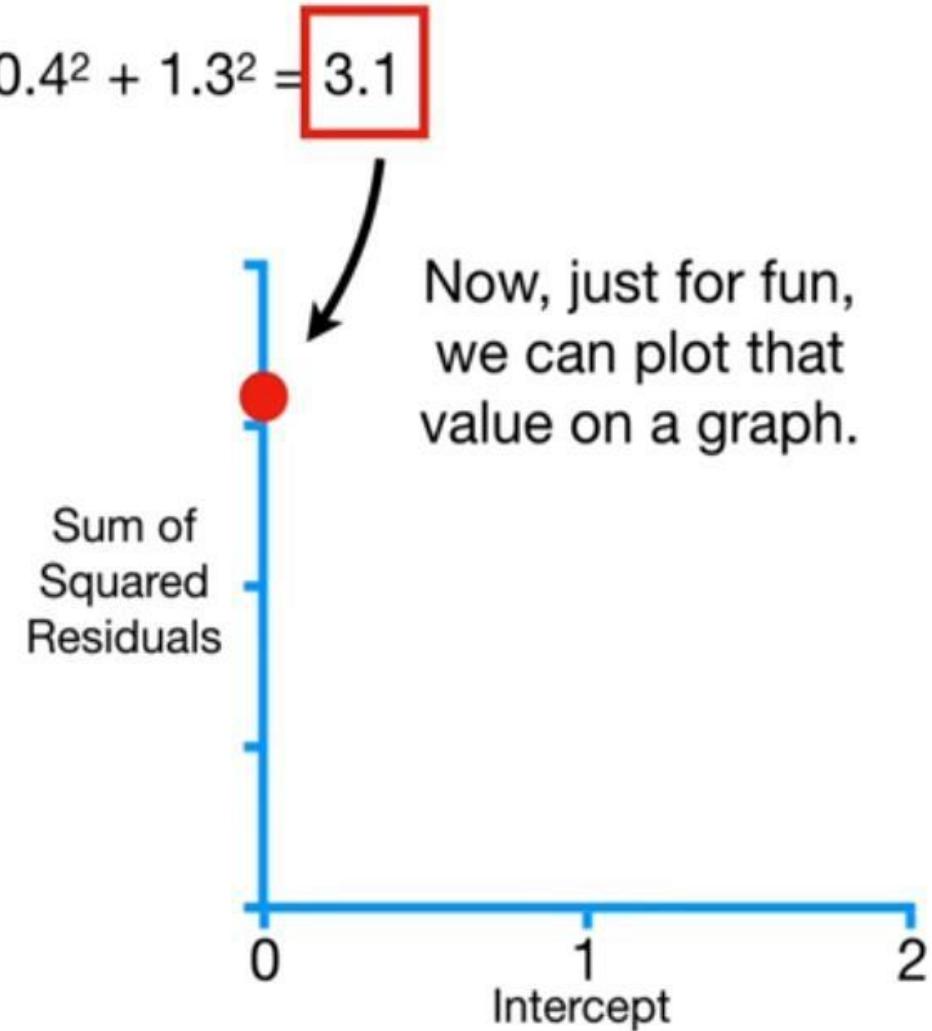
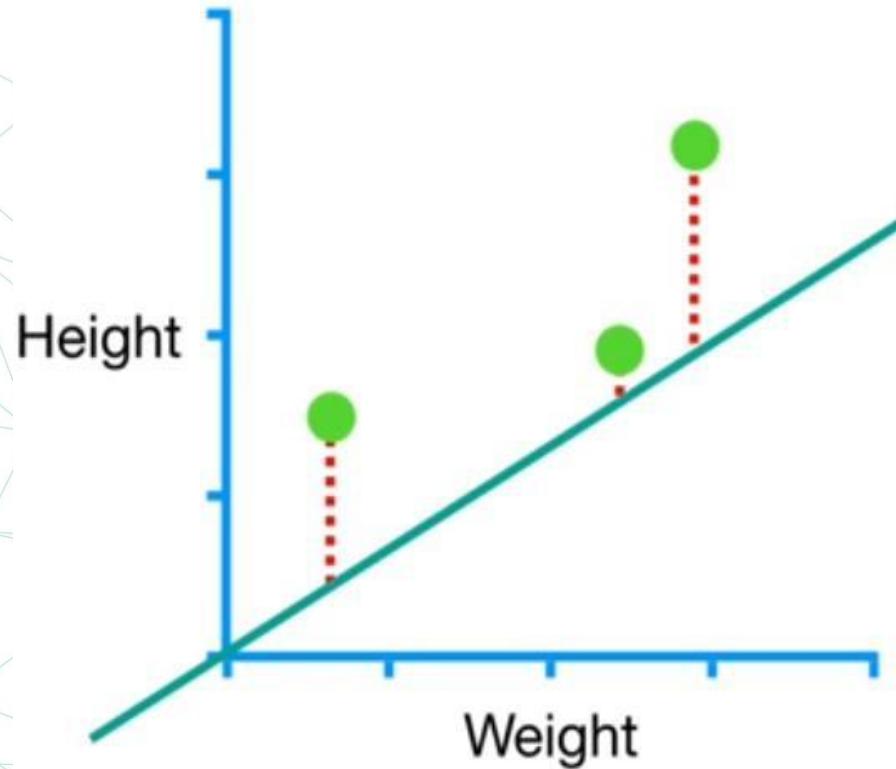
In the end, 3.1 is the Sum of the Squared Residuals.

3.1



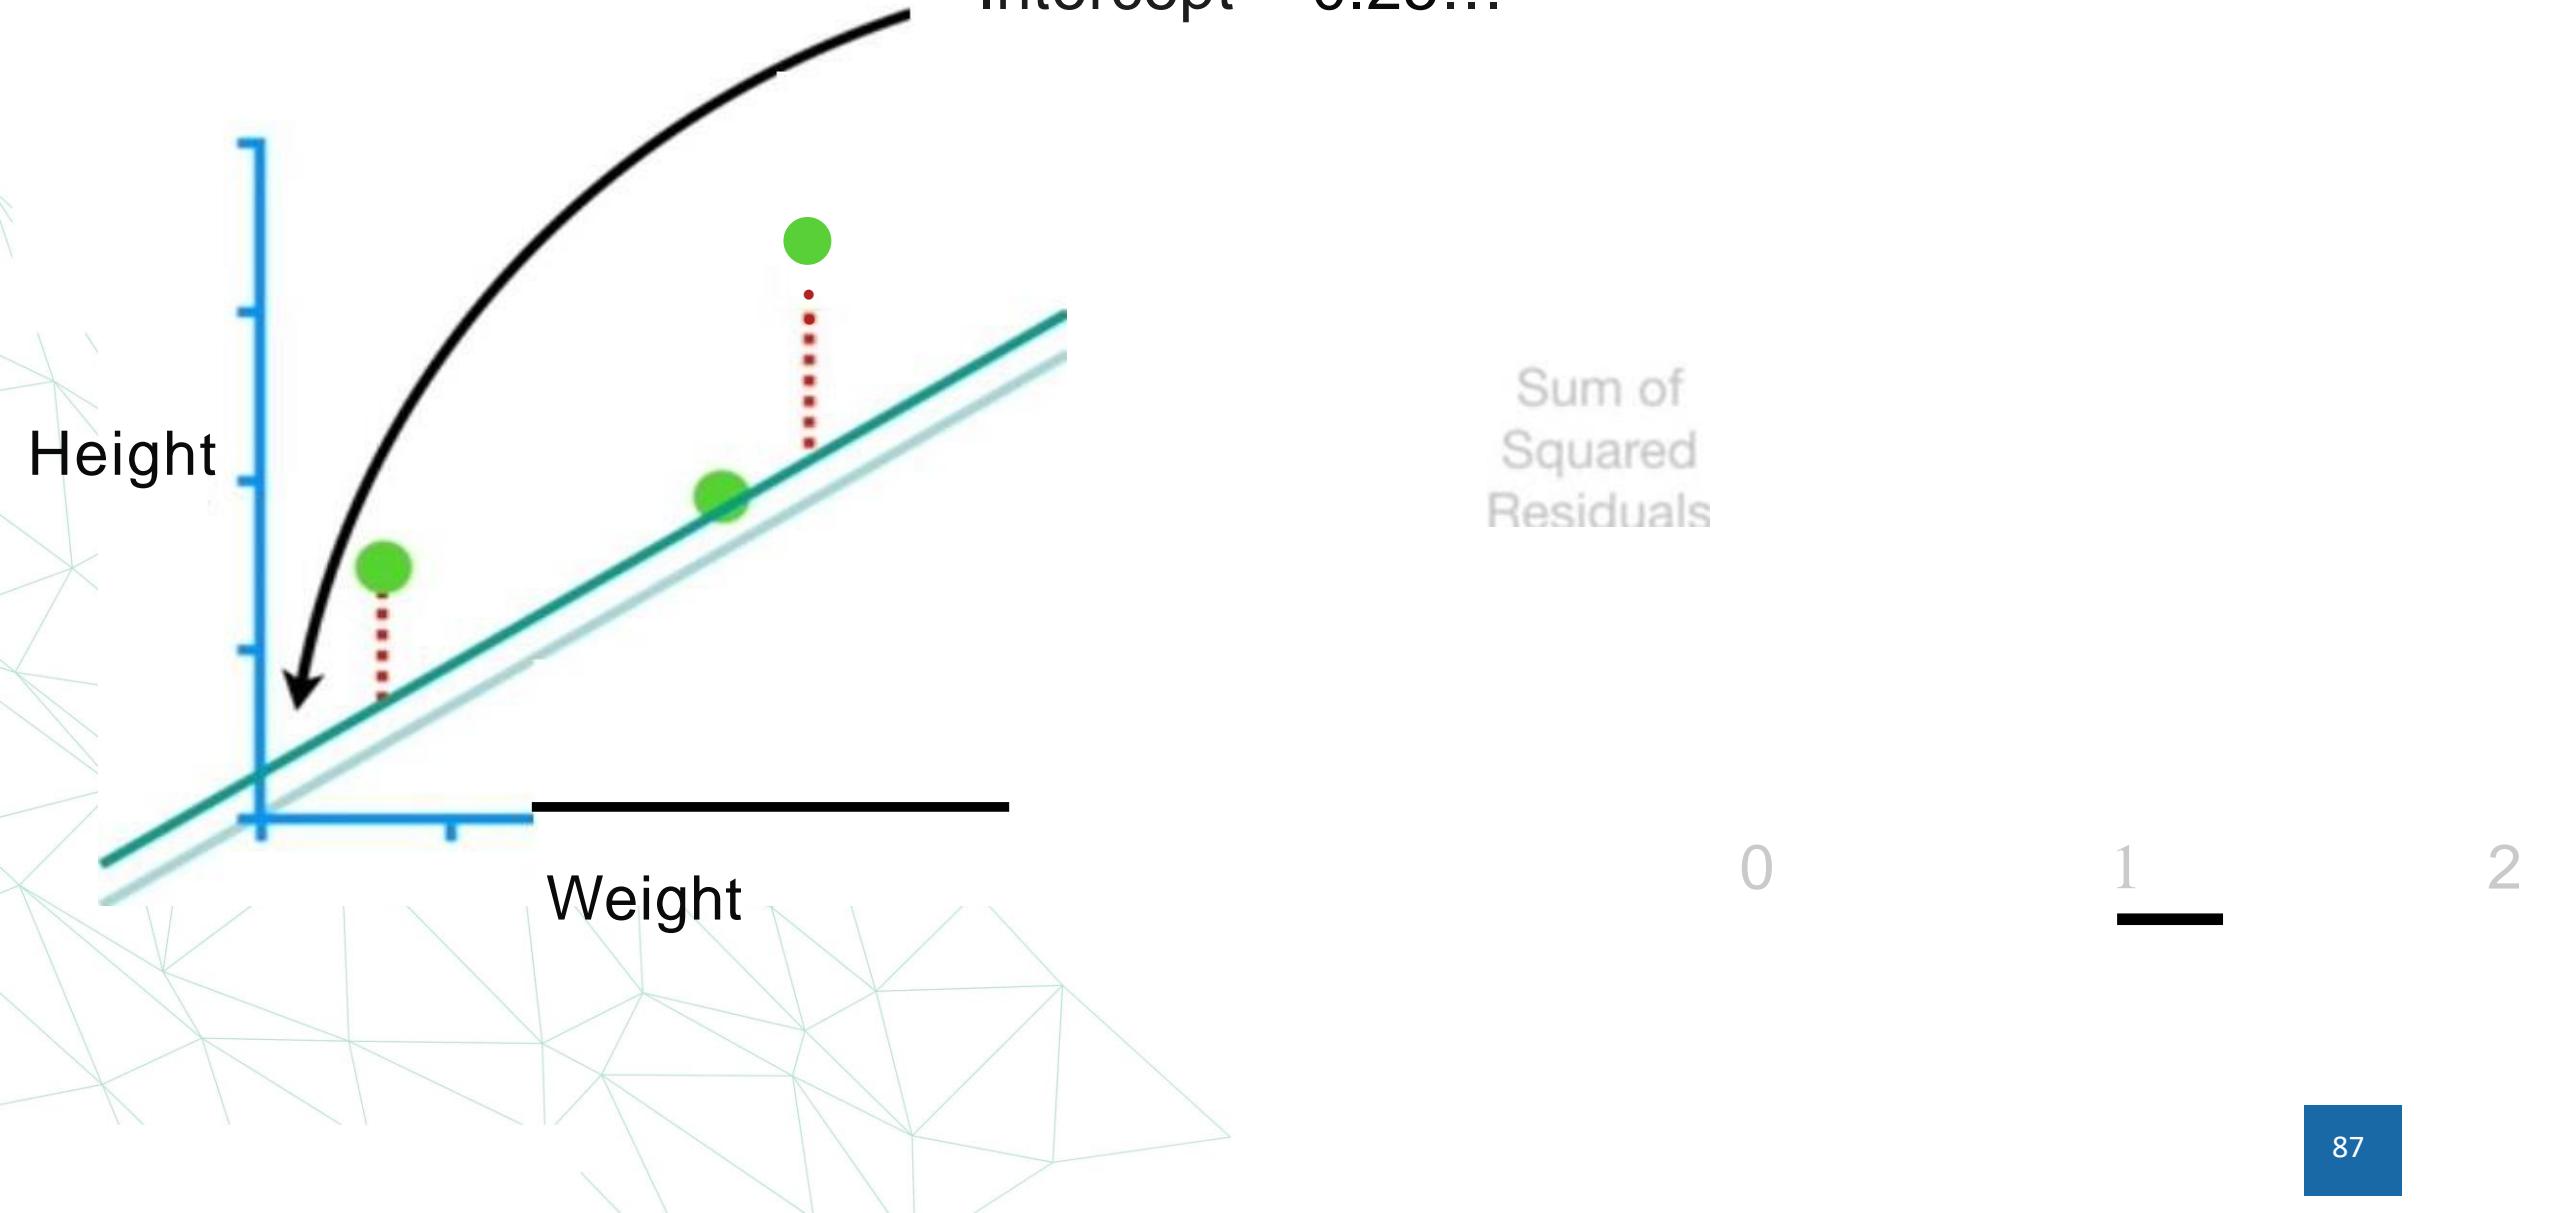


$$\text{Sum of squared residuals} = 1.1^2 + 0.4^2 + 1.3^2 = \boxed{3.1}$$

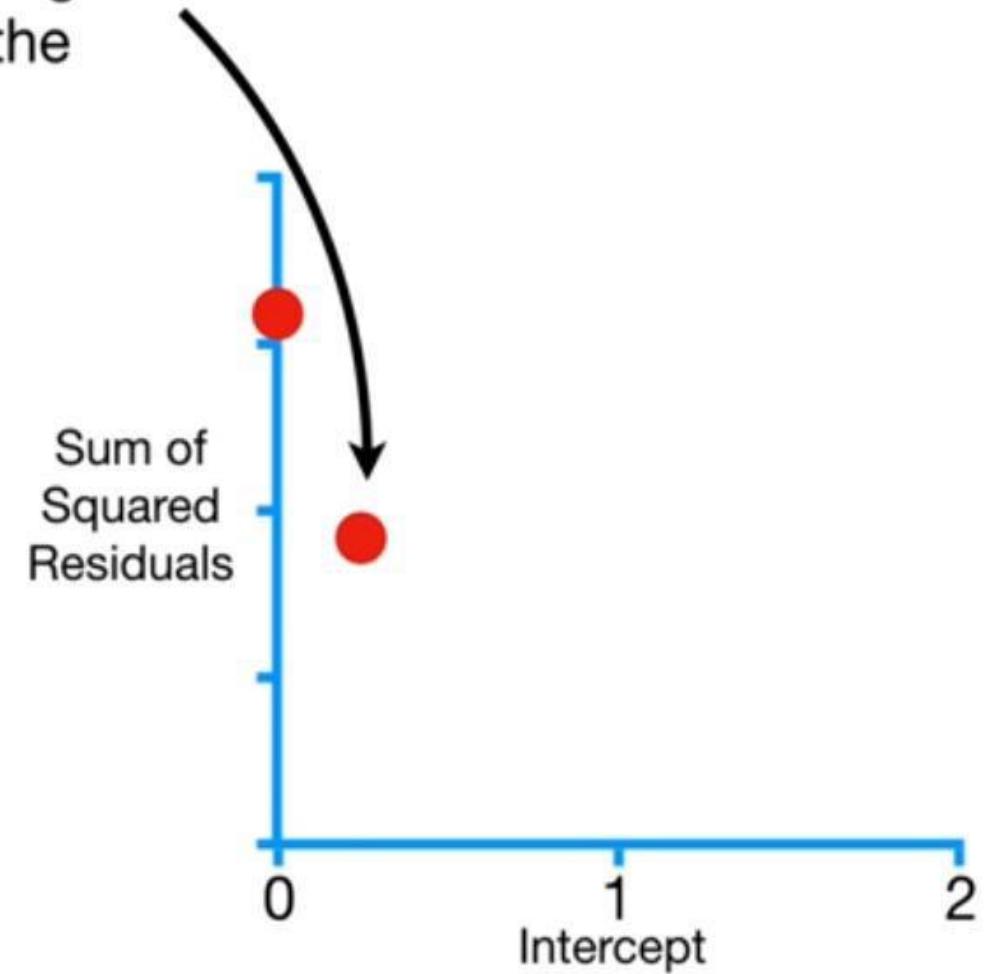
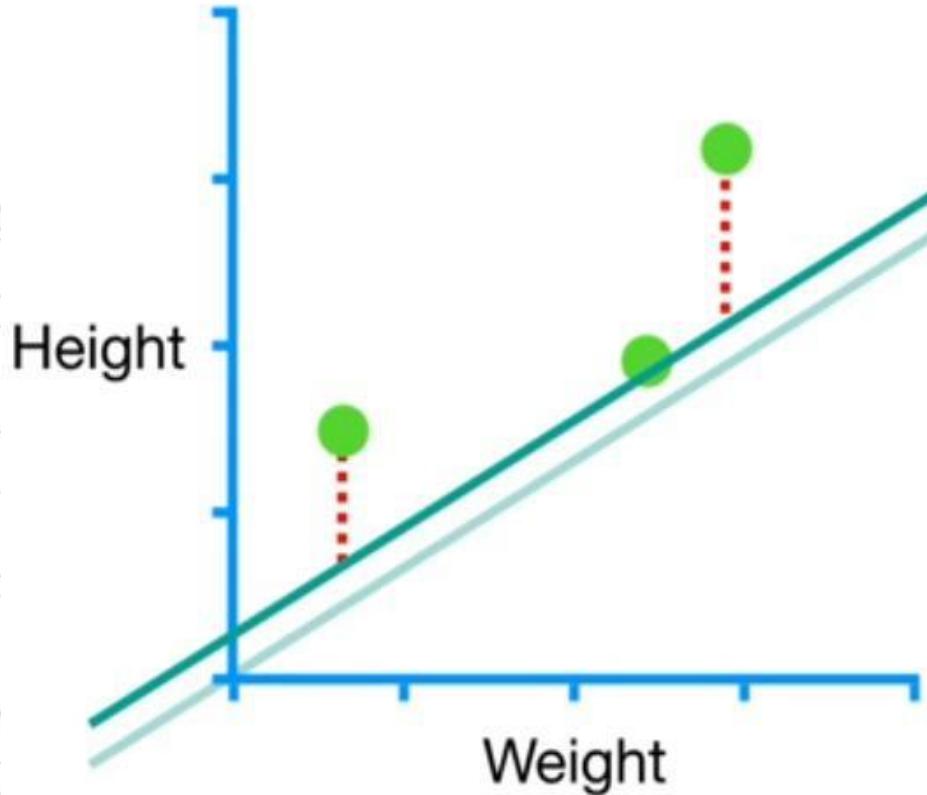


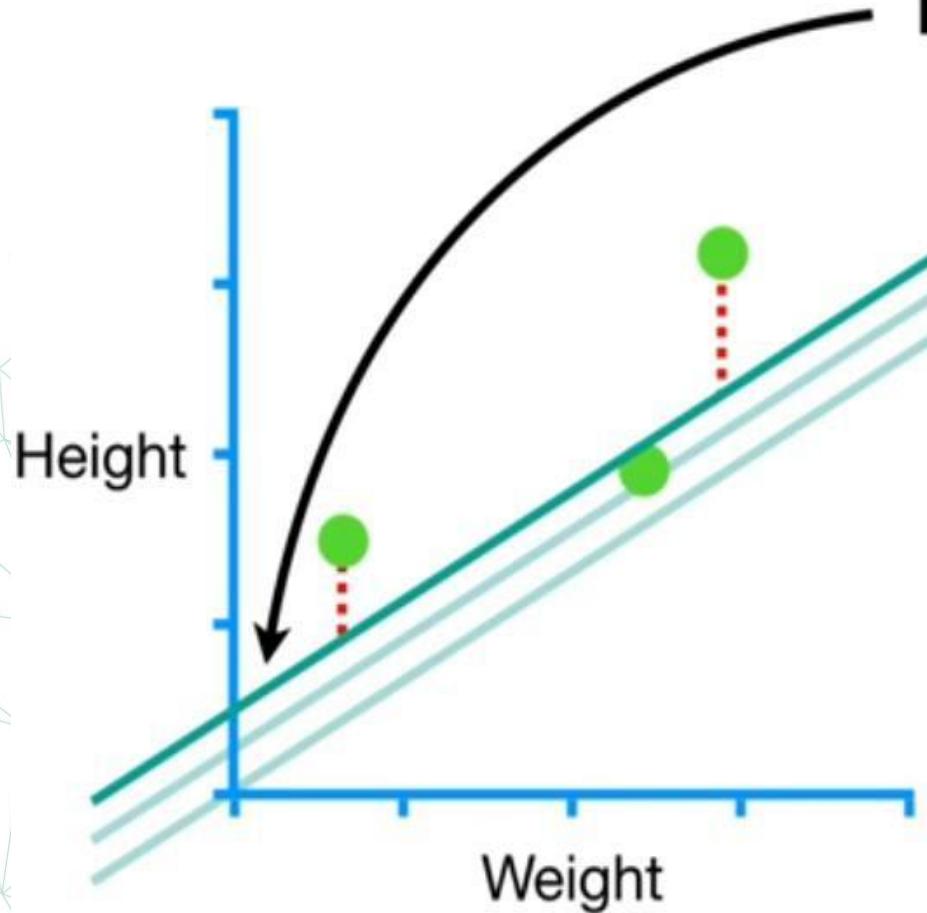


However, if the
Intercept = 0.25...

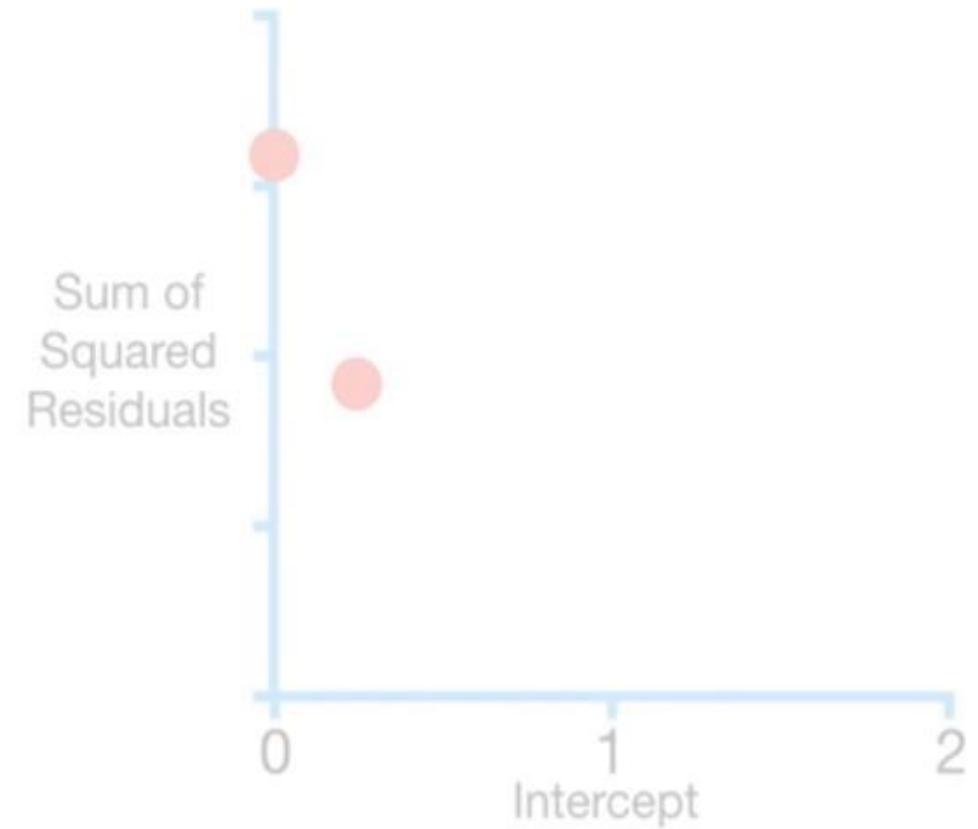


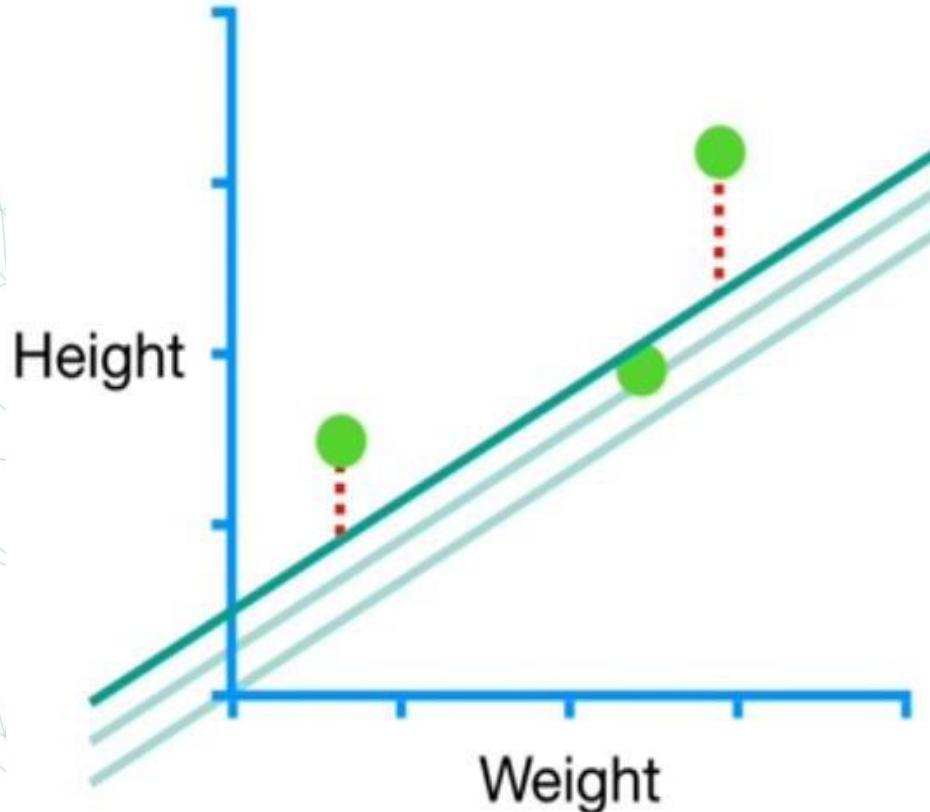
...then we would get
this point on the
graph.



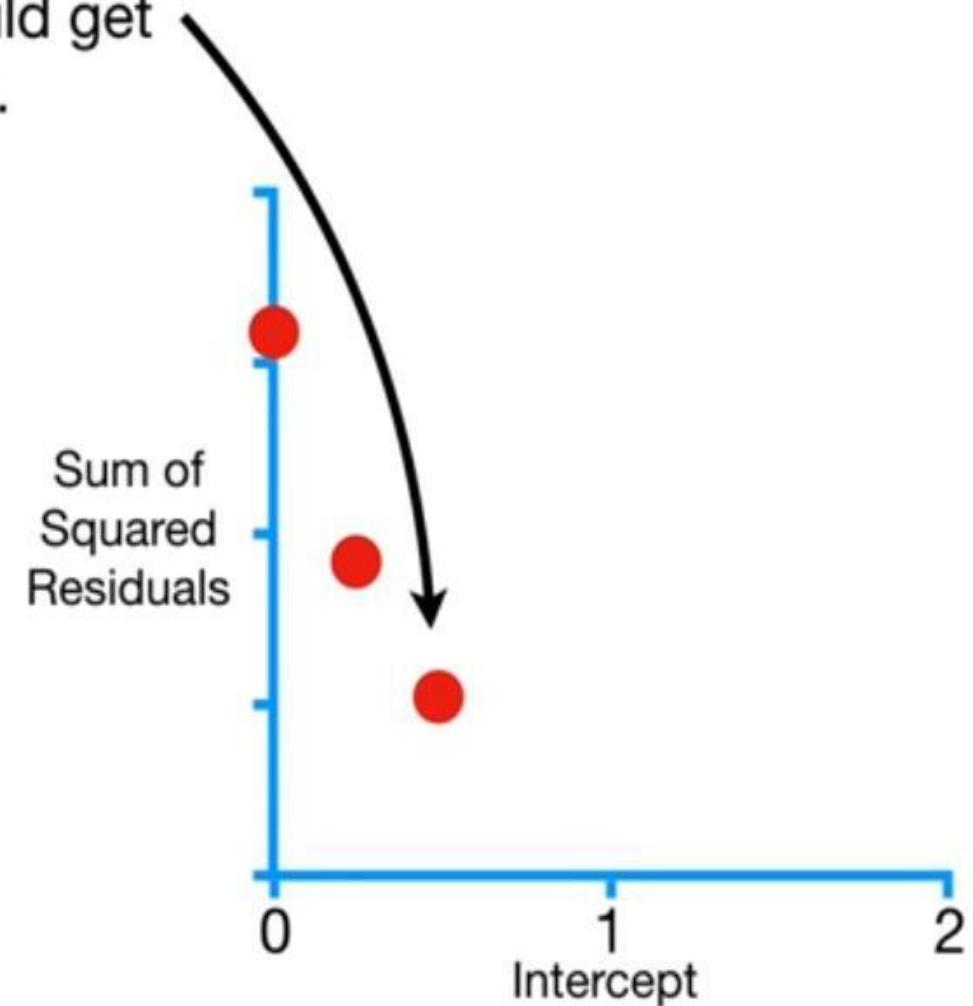


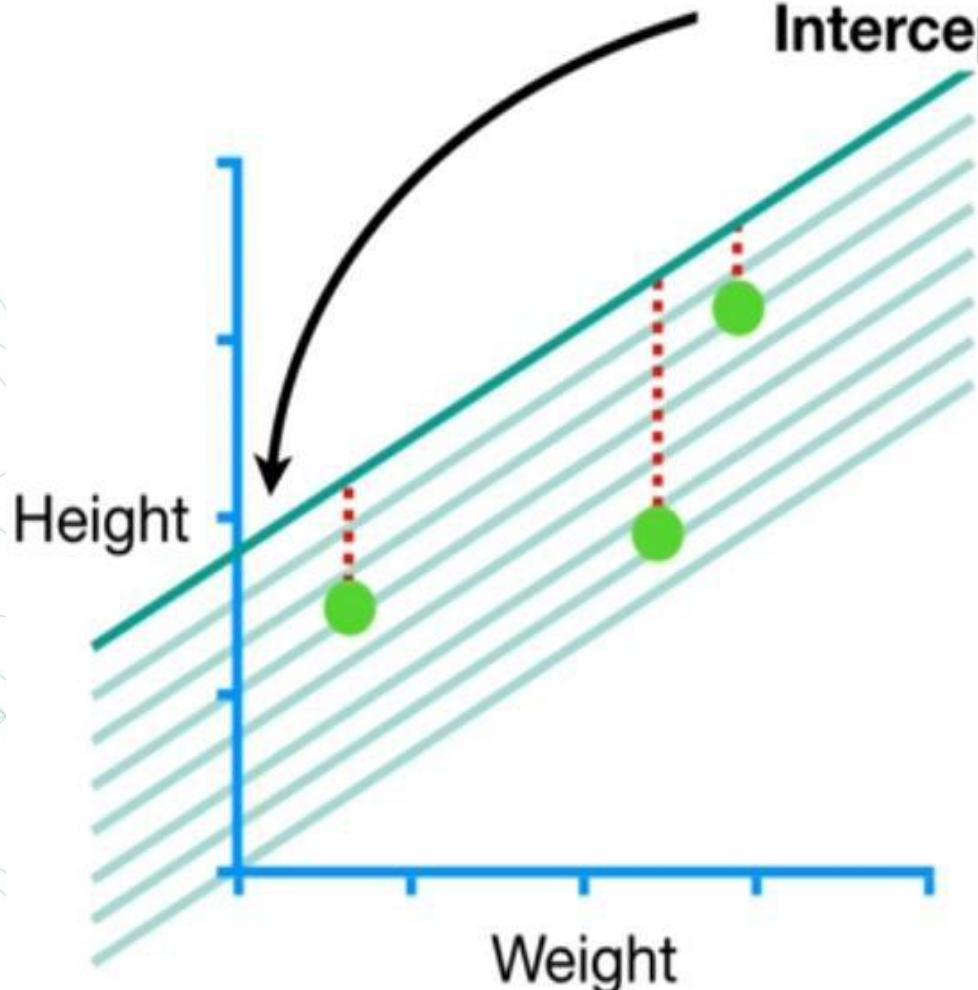
And if the
Intercept = 0.5...



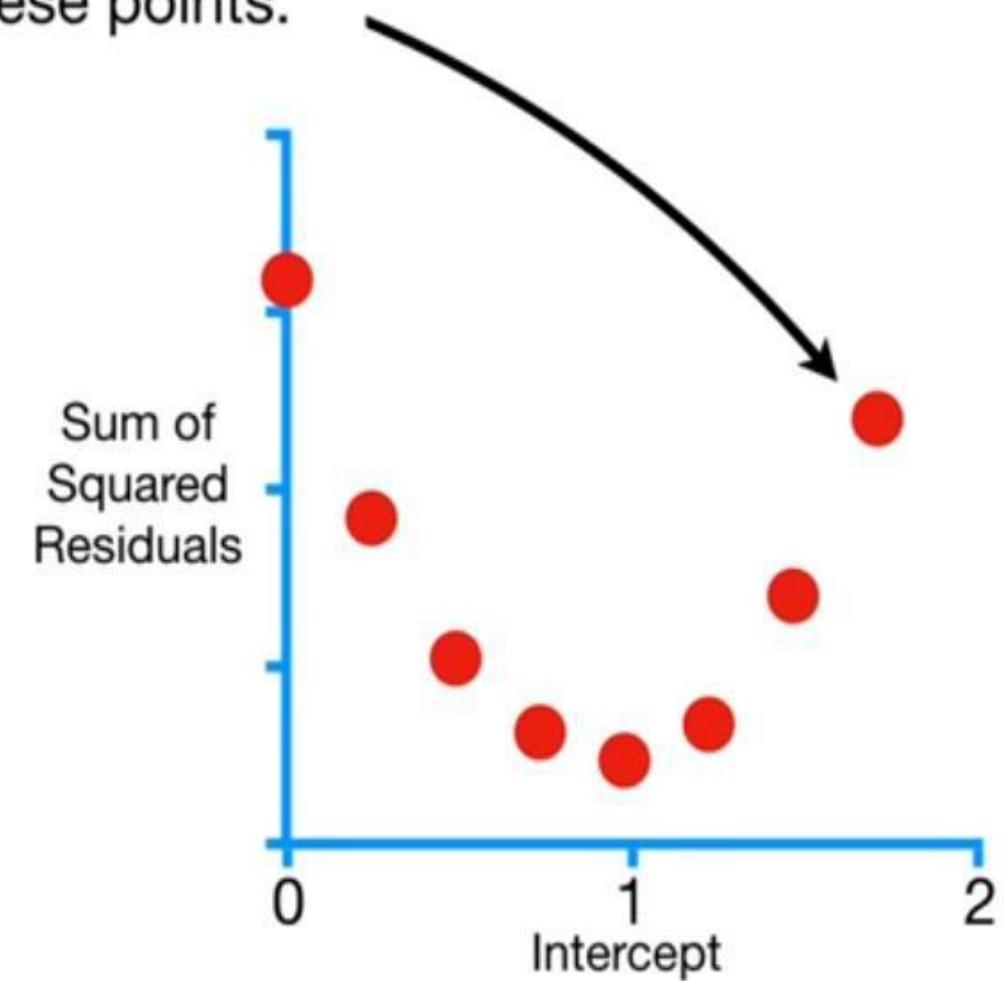


...then we would get
this point.



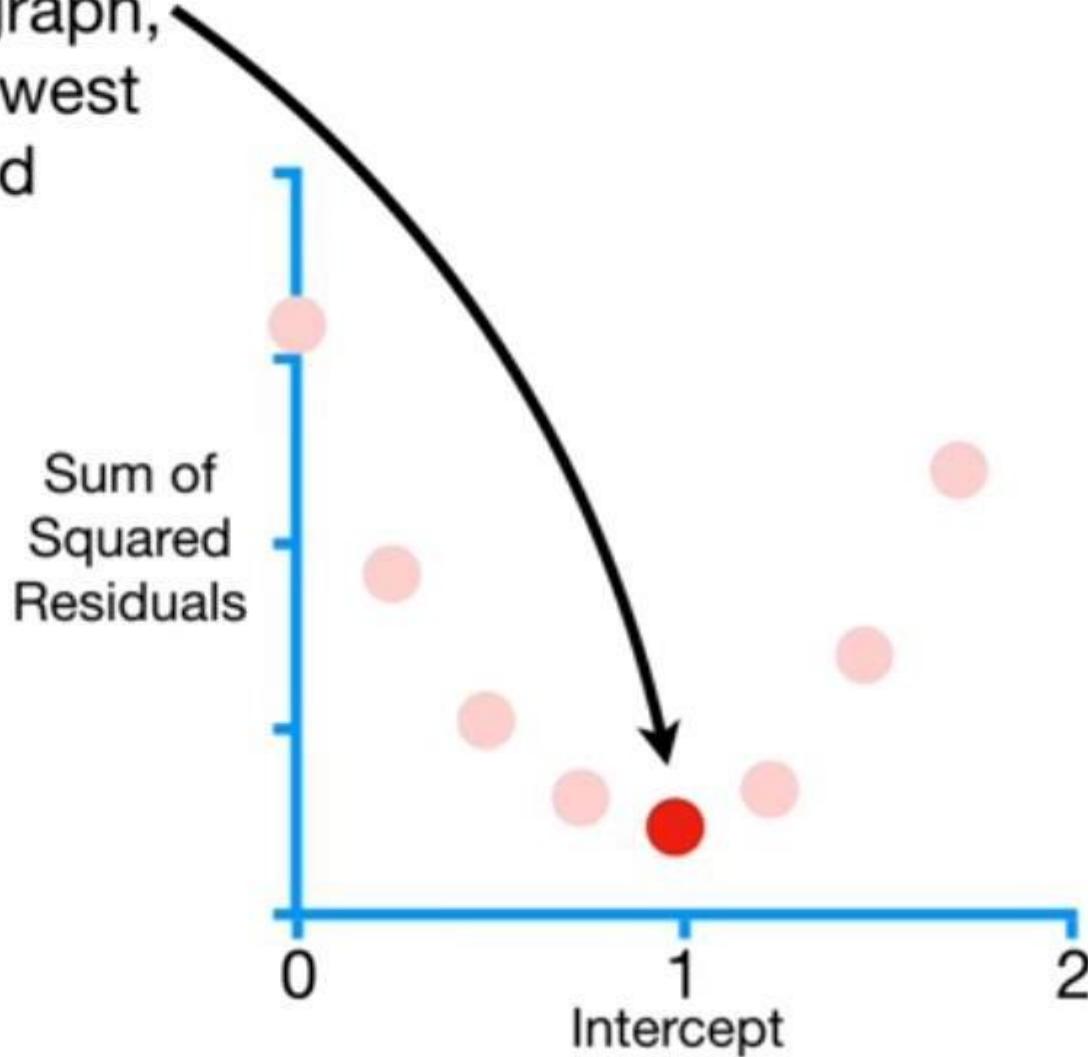


And for increasing values for the **Intercept**, we get these points.

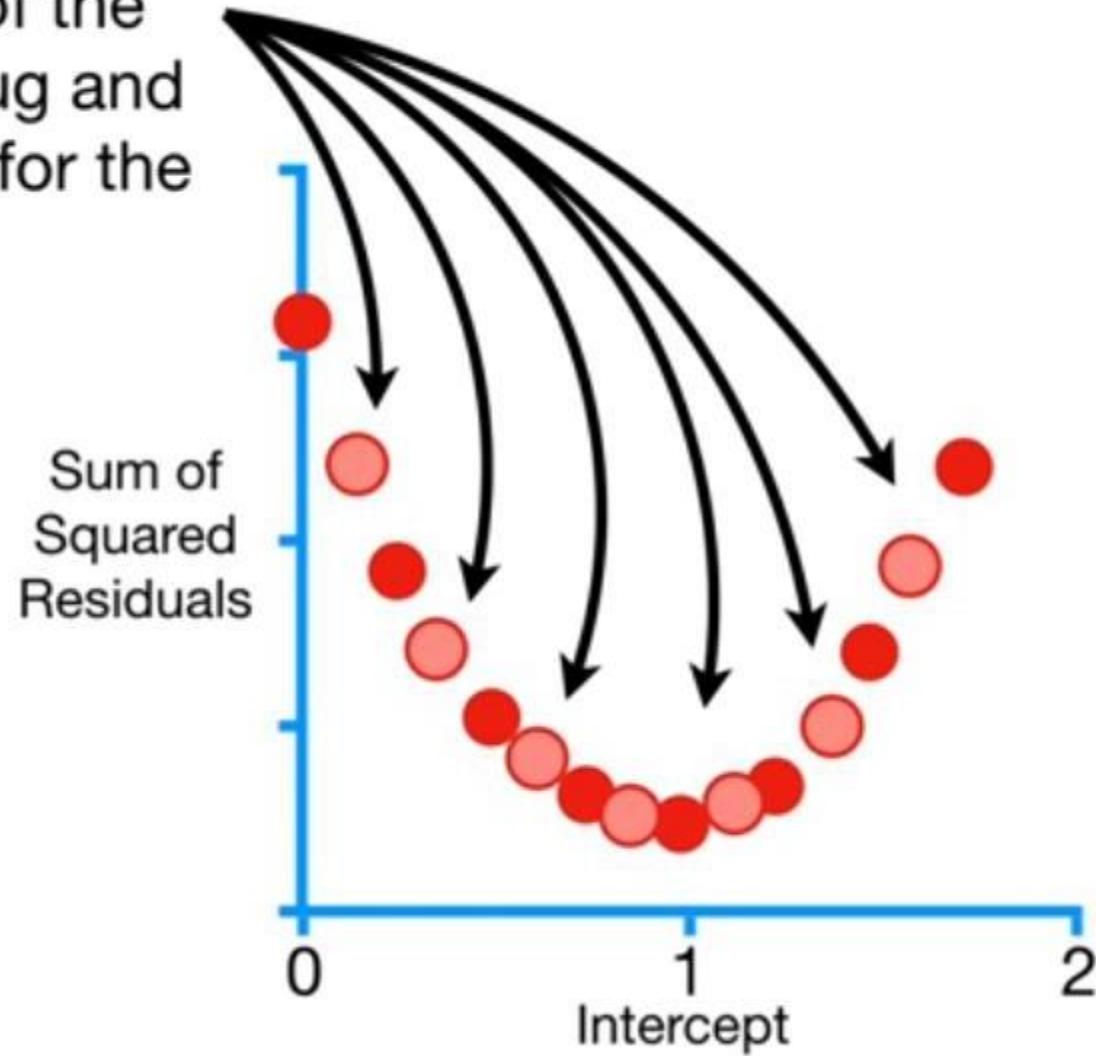




Of the points that we calculated for the graph, this one has the lowest Sum of Squared Residuals...

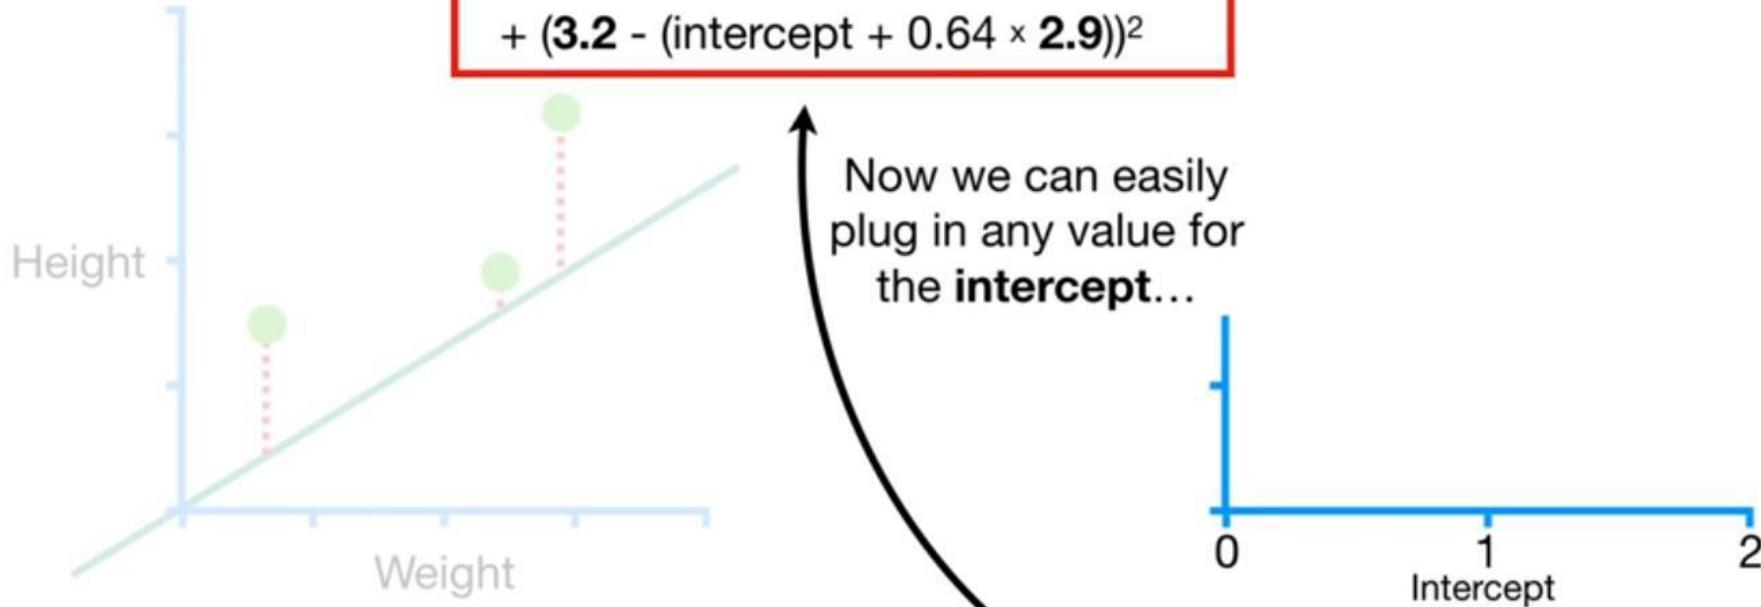


A slow and painful method for finding the minimal Sum of the Squared Residuals is to plug and chug a bunch more values for the **Intercept**.





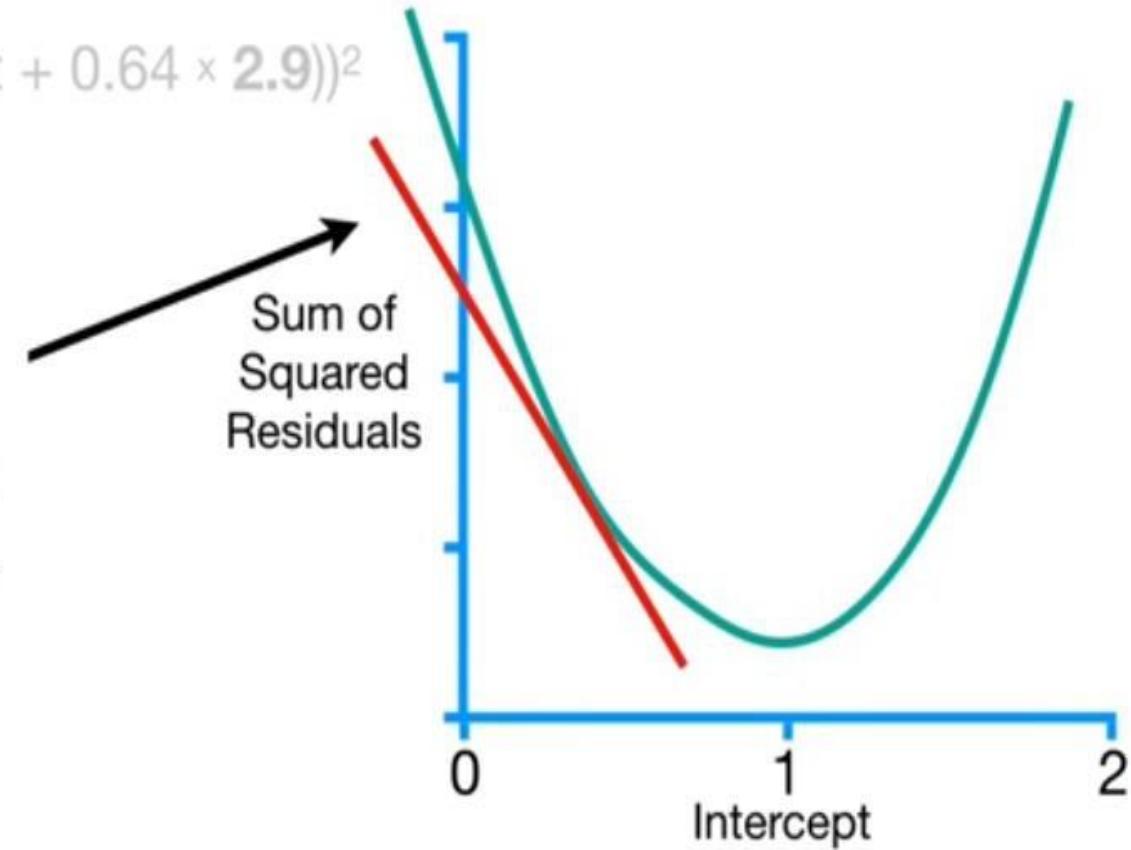
$$\begin{aligned}\text{Sum of squared residuals} = & (1.4 - (\text{intercept} + 0.64 \times 0.5))^2 \\ & + (1.9 - (\text{intercept} + 0.64 \times 2.3))^2 \\ & + (3.2 - (\text{intercept} + 0.64 \times 2.9))^2\end{aligned}$$

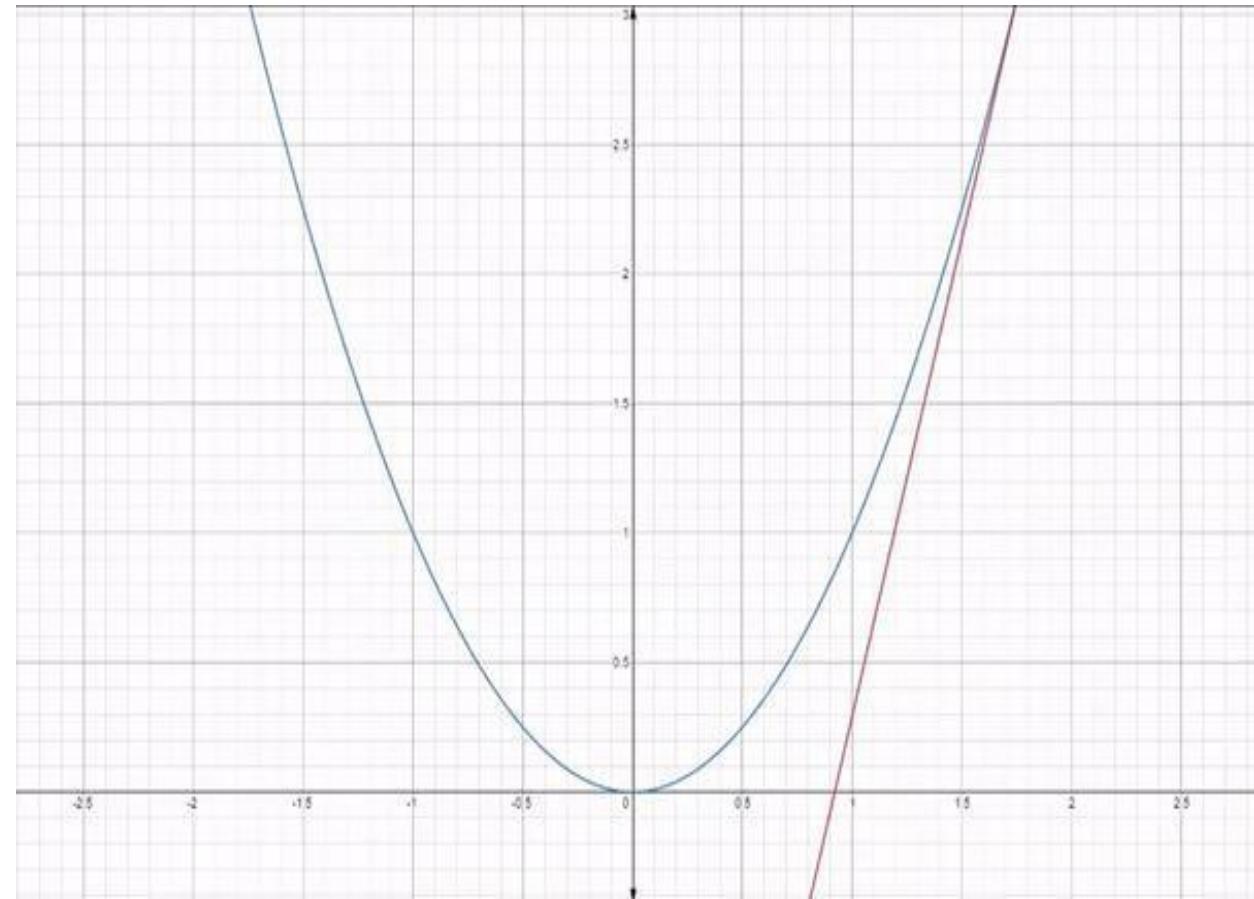
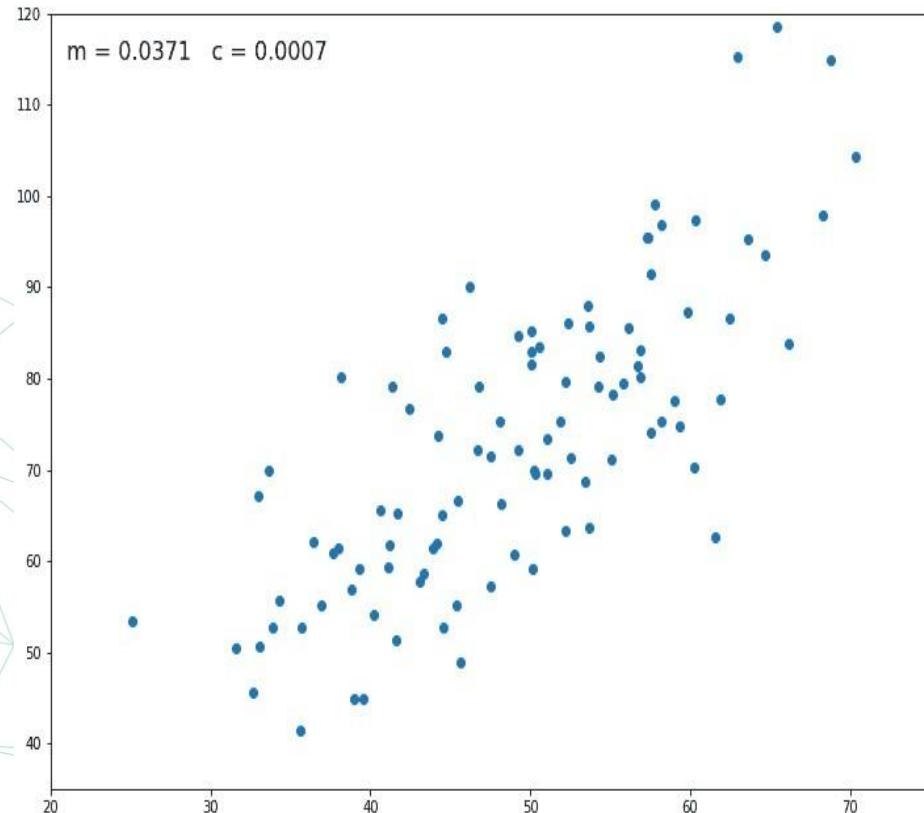




$$\begin{aligned}\text{Sum of squared residuals} = & (1.4 - (\text{intercept} + 0.64 \times 0.5))^2 \\ & + (1.9 - (\text{intercept} + 0.64 \times 2.3))^2 \\ & + (3.2 - (\text{intercept} + 0.64 \times 2.9))^2\end{aligned}$$

...and we can take the derivative of this function and determine the slope at any value for the **Intercept**.







$$\begin{aligned}\text{Sum of squared residuals} &= (1.4 - (\text{intercept} + 0.64 \times 0.5))^2 \\ &+ (1.9 - (\text{intercept} + 0.64 \times 2.3))^2 \\ &+ (3.2 - (\text{intercept} + 0.64 \times 2.9))^2\end{aligned}$$

Let's start by taking the derivative
of the first part.



$$\begin{aligned}\frac{d}{d \text{ intercept}} \text{ Sum of squared residuals} &= \frac{d}{d \text{ intercept}} (1.4 - (\text{intercept} + 0.64 \times 0.5))^2 \\ &+ \frac{d}{d \text{ intercept}} (1.9 - (\text{intercept} + 0.64 \times 2.3))^2 \\ &+ \frac{d}{d \text{ intercept}} (3.2 - (\text{intercept} + 0.64 \times 2.9))^2\end{aligned}$$

$$\frac{d}{d \text{ intercept}} (1.4 - (\text{intercept} + 0.64 \times 0.5))^2 = 2(1.4 - (\text{intercept} + 0.64 \times 0.5)) \times -1$$

$$\frac{d}{d \text{ intercept}} 1.4 - (\text{intercept} + 0.64 \times 0.5)$$

$$\frac{d}{d \text{ intercept}} \cancel{1.4} + (-1)\text{intercept} - 0.64 \cancel{\times 0.5} = -1$$

These parts don't contain a term for the **Intercept**, so they go away.

$$\frac{d}{d \text{ intercept}} (1.4 - (\text{intercept} + 0.64 \times 0.5))^2 = 2(1.4 - (\text{intercept} + 0.64 \times 0.5)) \times -1$$

$$= -2(1.4 - (\text{intercept} + 0.64 \times 0.5))$$

...and this...

...is the derivative
of the first part...

$$\frac{d}{d \text{ intercept}} \text{ Sum of squared residuals} = \frac{d}{d \text{ intercept}} (1.4 - (\text{intercept} + 0.64 \times 0.5))^2$$

$$+ \frac{d}{d \text{ intercept}} (1.9 - (\text{intercept} + 0.64 \times 2.3))^2$$

$$+ \frac{d}{d \text{ intercept}} (3.2 - (\text{intercept} + 0.64 \times 2.9))^2$$

I'll leave that as an exercise for
the viewer.

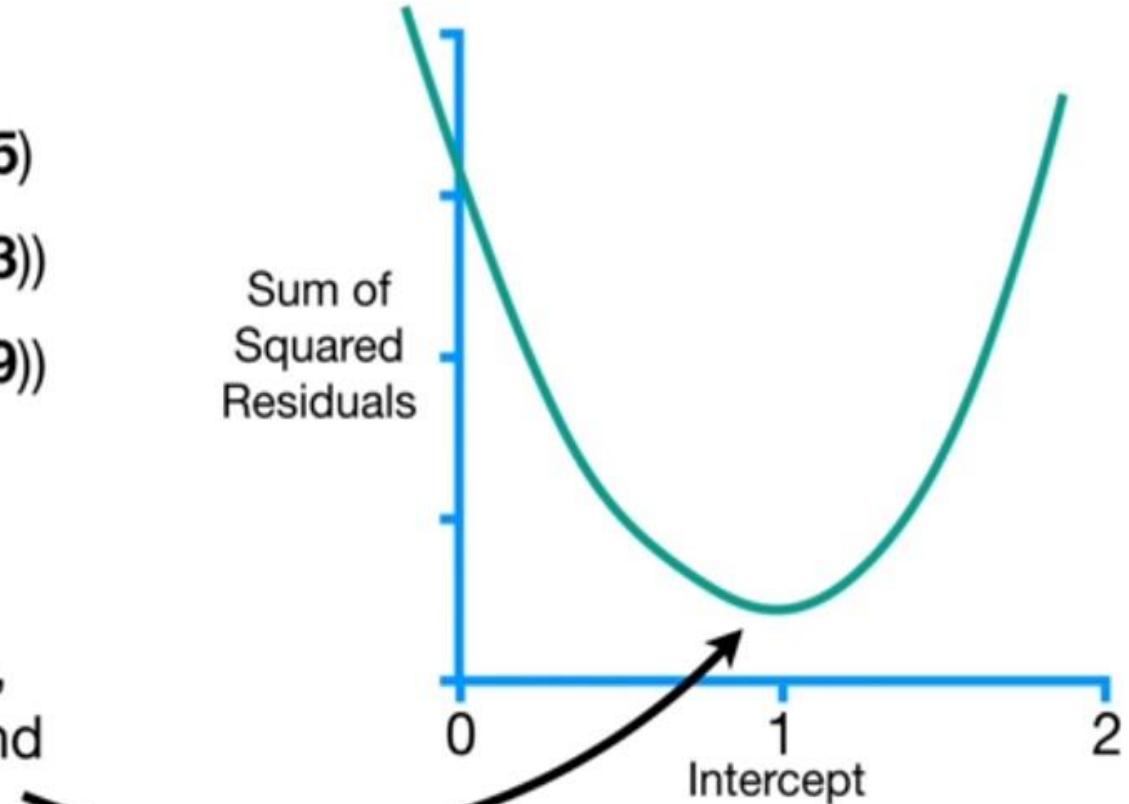
$$\frac{d}{d \text{ intercept}} \text{ Sum of squared residuals} = -2(1.4 - (\text{intercept} + 0.64 \times 0.5))$$

$$+ -2(1.9 - (\text{intercept} + 0.64 \times 2.3))$$

$$+ -2(3.2 - (\text{intercept} + 0.64 \times 2.9))$$

$$\frac{d}{d \text{ intercept}} \text{ Sum of squared residuals} =$$
$$-2(1.4 - (\text{intercept} + 0.64 \times 0.5))$$
$$+ -2(1.9 - (\text{intercept} + 0.64 \times 2.3))$$
$$+ -2(3.2 - (\text{intercept} + 0.64 \times 2.9))$$

Now that we have the derivative,
Gradient Descent will use it to find
where the Sum of Squared
Residuals is lowest.



$$\frac{d}{d \text{ intercept}} \text{Sum of squared residuals} =$$

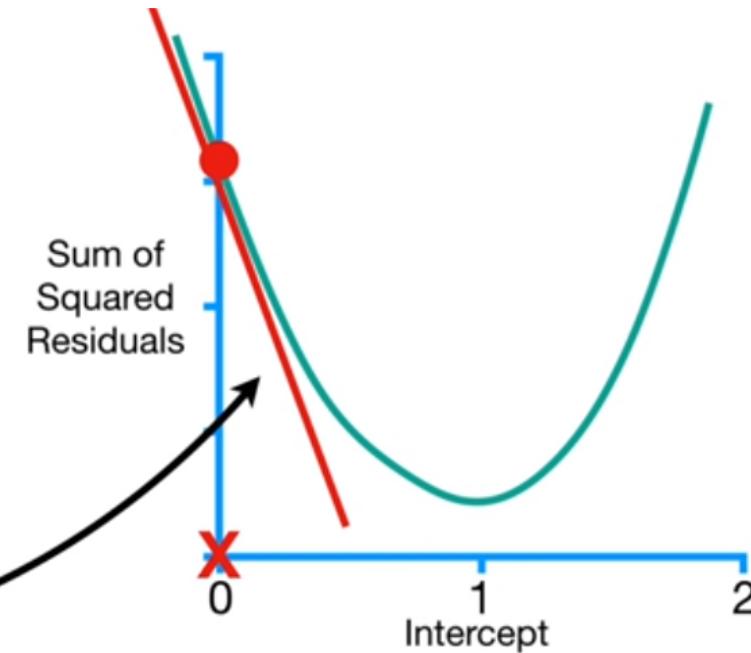
$$-2(1.4 - (0 + 0.64 \times 0.5))$$

$$+ -2(1.9 - (0 + 0.64 \times 2.3))$$

$$+ -2(3.2 - (0 + 0.64 \times 2.9))$$

$$= -5.7$$

So when the **Intercept = 0**,
the slope of the curve = **-5.7**.



$$\frac{d}{d \text{ intercept}} \text{ Sum of squared residuals} =$$

$$-2(1.4 - (0 + 0.64 \times 0.5))$$

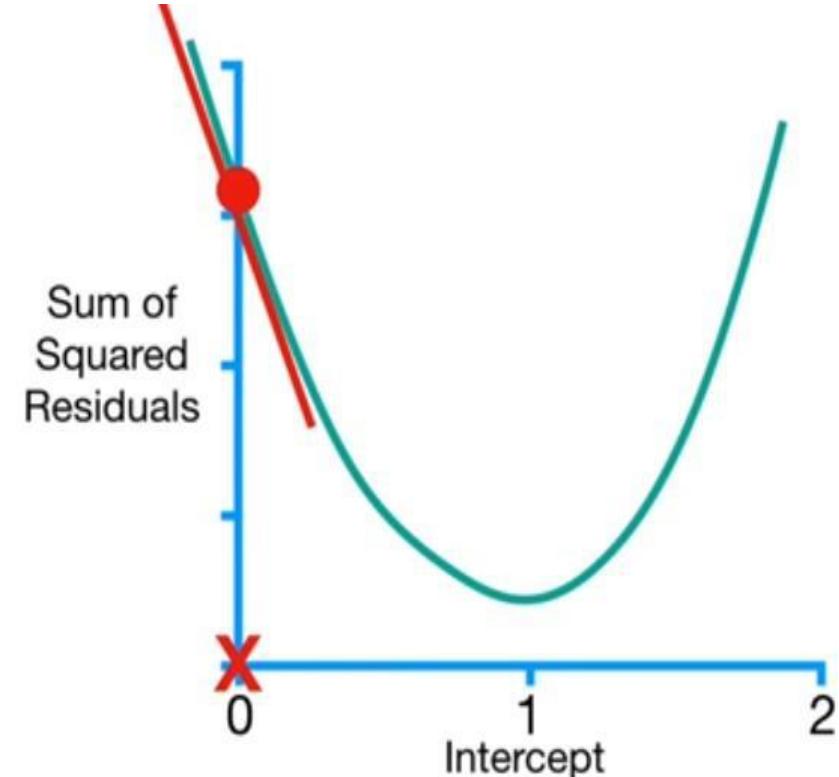
$$+ -2(1.9 - (0 + 0.64 \times 2.3))$$

$$+ -2(3.2 - (0 + 0.64 \times 2.9))$$

$$= -5.7$$

Step Size = -5.7×0.1

...by a small number called
The Learning Rate.



$$\frac{d}{d \text{ intercept}} \text{ Sum of squared residuals} =$$

$$-2(1.4 - (0 + 0.64 \times 0.5))$$

$$+ -2(1.9 - (0 + 0.64 \times 2.3))$$

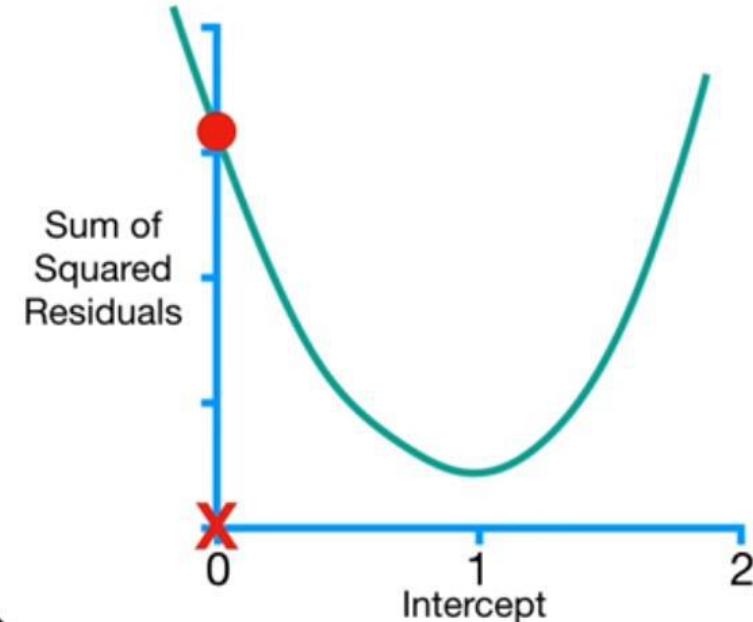
$$+ -2(3.2 - (0 + 0.64 \times 2.9))$$

$$= -5.7$$

Step Size = $-5.7 \times 0.1 = \boxed{-0.57}$

New Intercept = Old Intercept - Step Size

...minus the **Step Size**.



$$\frac{d}{d \text{ intercept}} \text{ Sum of squared residuals} =$$

$$-2(1.4 - (0 + 0.64 \times 0.5))$$

$$+ -2(1.9 - (0 + 0.64 \times 2.3))$$

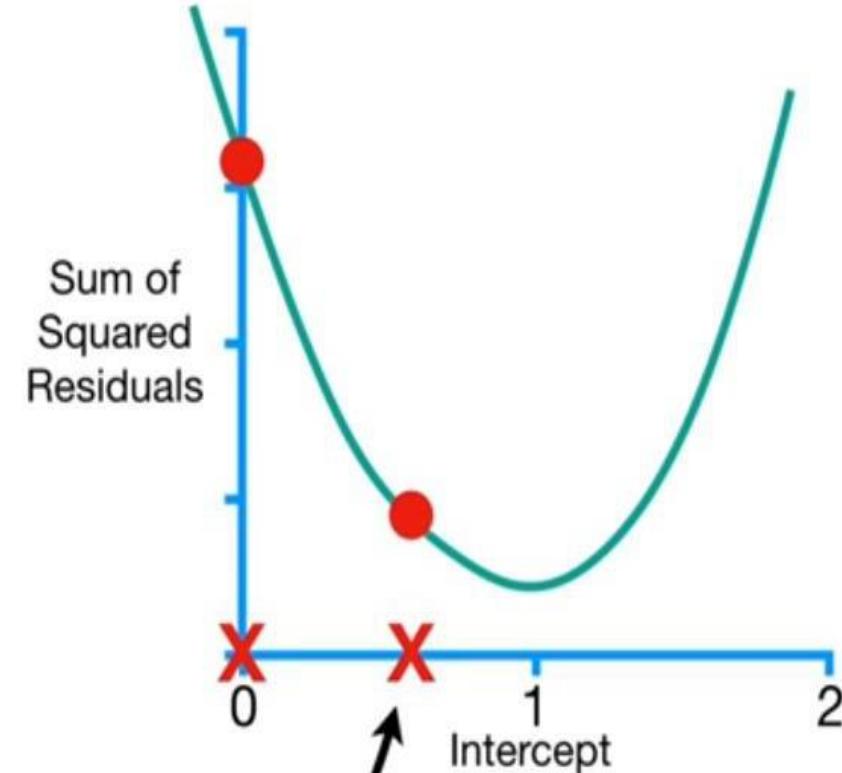
$$+ -2(3.2 - (0 + 0.64 \times 2.9))$$

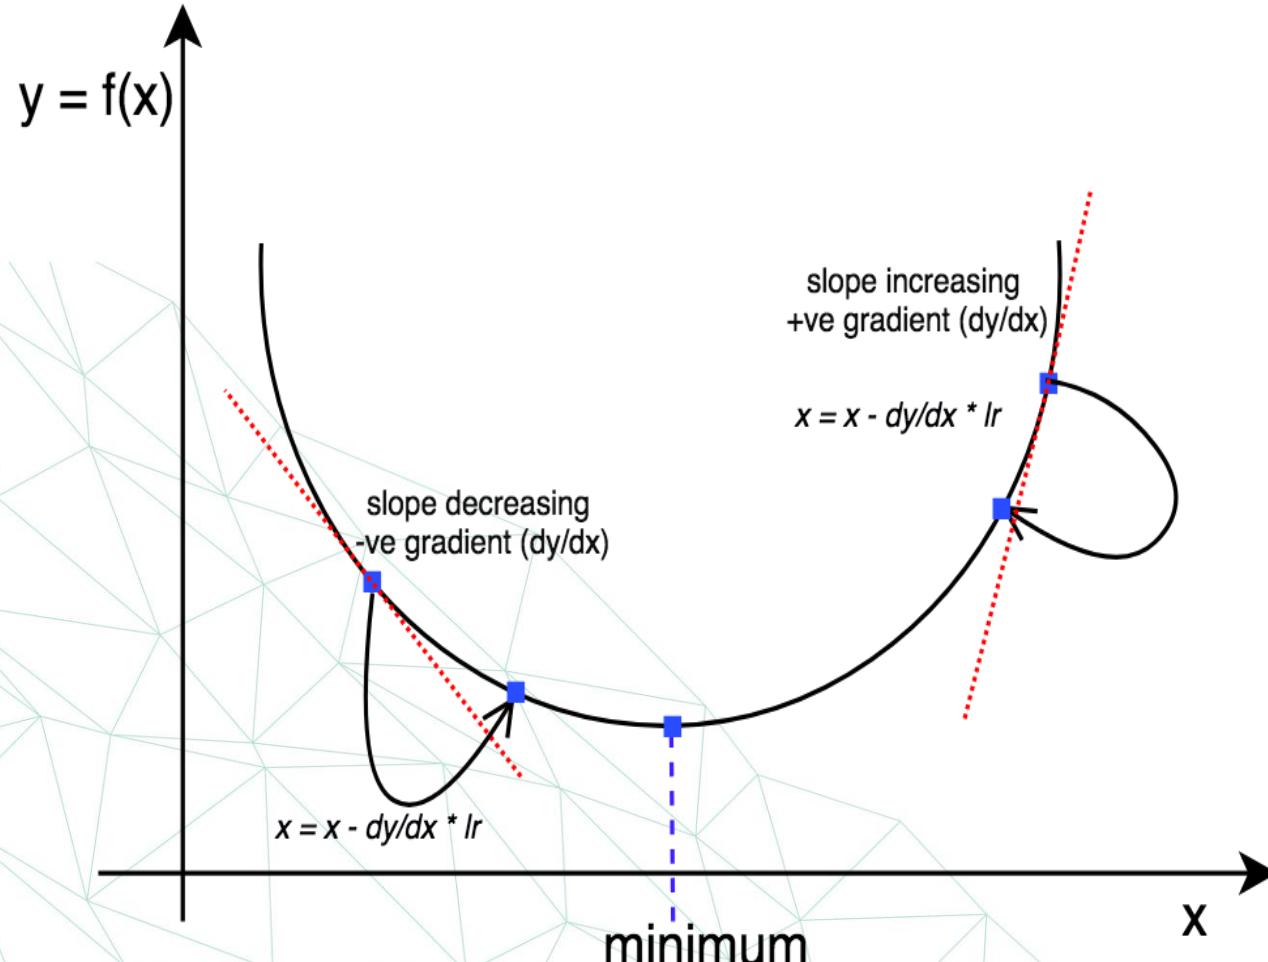
$$= -5.7$$

$$\text{Step Size} = -5.7 \times 0.1 = -0.57$$

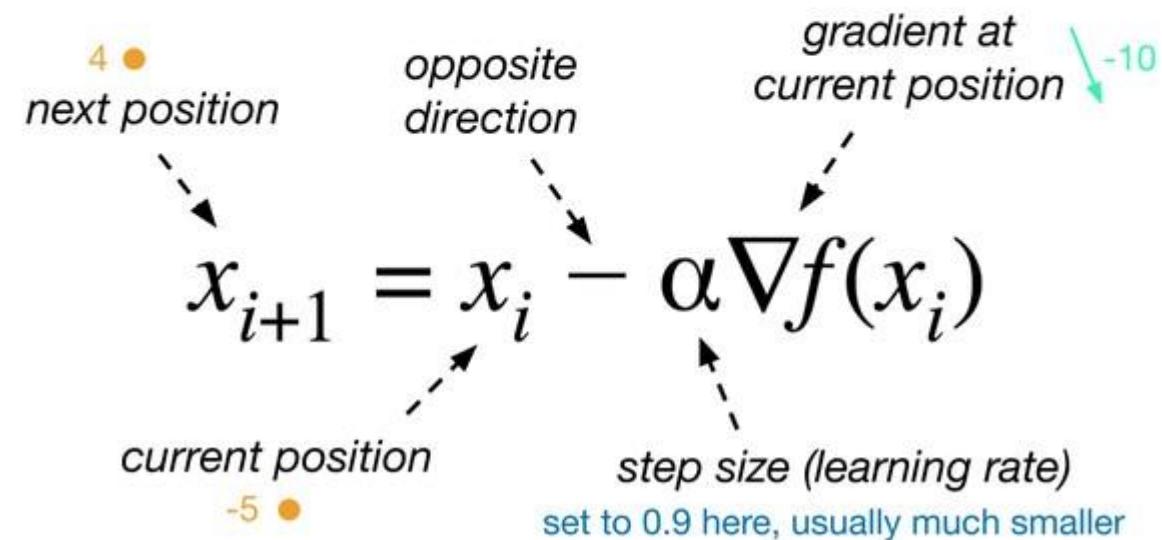
New Intercept = $0 - (-0.57) = 0.57$

...and the New Intercept = 0.57.





Gradient descent is an iterative optimization algorithm for finding the **minimum** of a function.



We will use the Mean Squared Error function to calculate the loss. There are three steps in this function:

- Find the difference between the actual y and predicted y value($y = mx + c$), for a given x .
- Square this difference.
- Find the mean of the squares for every value in X .

$$E = \frac{1}{n} \sum_{i=0}^n (y_i - \bar{y}_i)^2$$

Here y_i is the actual value and \bar{y}_i is the predicted value. Lets substitute the value of \bar{y}_i :

$$E = \frac{1}{n} \sum_{i=0}^n (y_i - (mx_i + c))^2$$



Calculating the partial derivative of the cost function with respect to slope (m), let the partial derivative of the cost function with respect to m be D_m

Let's try applying gradient descent to m and c and approach it step by step:

- Initially let $m = 0$ and $c = 0$. Let L be our learning rate. This controls how much the value of m changes with each step. L could be a small value like 0.0001 for good accuracy.
- Calculate the partial derivative of the loss function with respect to m, and plug in the current values of x, y, m and c in it to obtain the derivative value D.

$$D_m = \frac{1}{n} \sum_{i=0}^n 2(y_i - (mx_i + c))(-x_i)$$

$$D_m = \frac{-2}{n} \sum_{i=0}^n x_i(y_i - \bar{y}_i)$$

D_m is the value of the partial derivative with respect to m. Similarly lets find the partial derivative with respect to c, D_c :

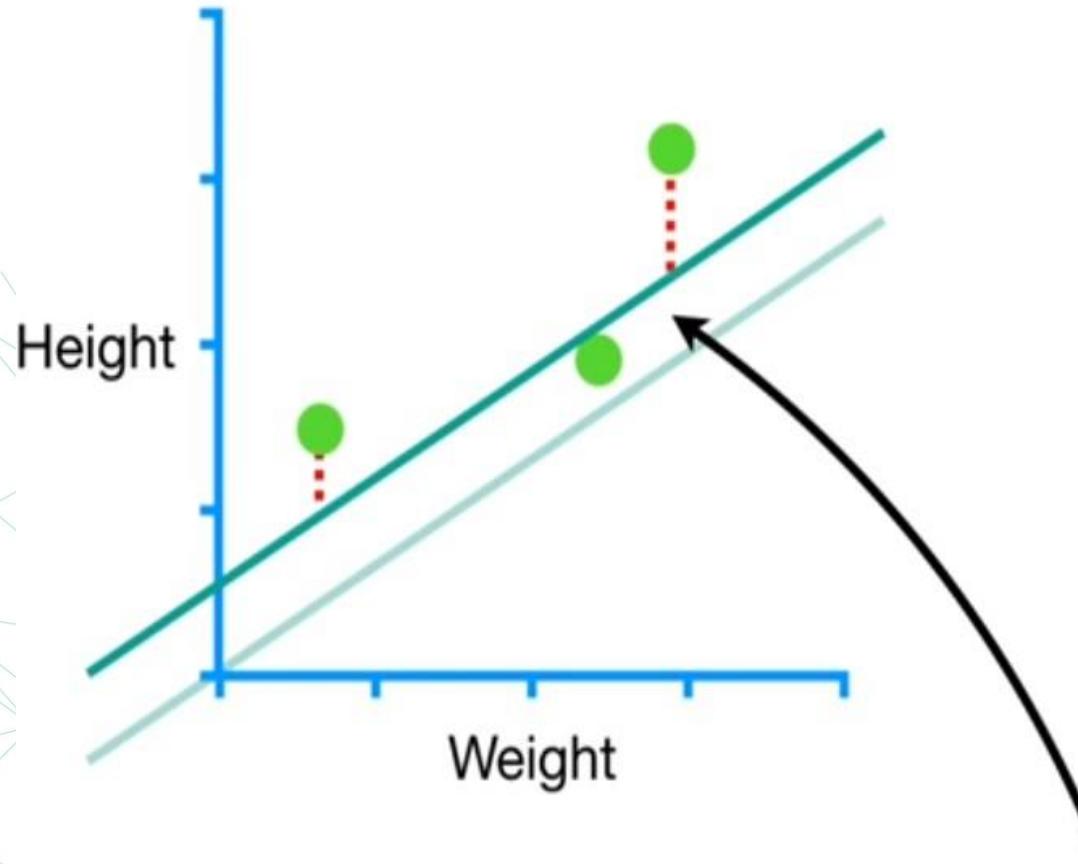
$$D_c = \frac{-2}{n} \sum_{i=0}^n (y_i - \bar{y}_i)$$

- Now we update the current value of m and c using the following equation:

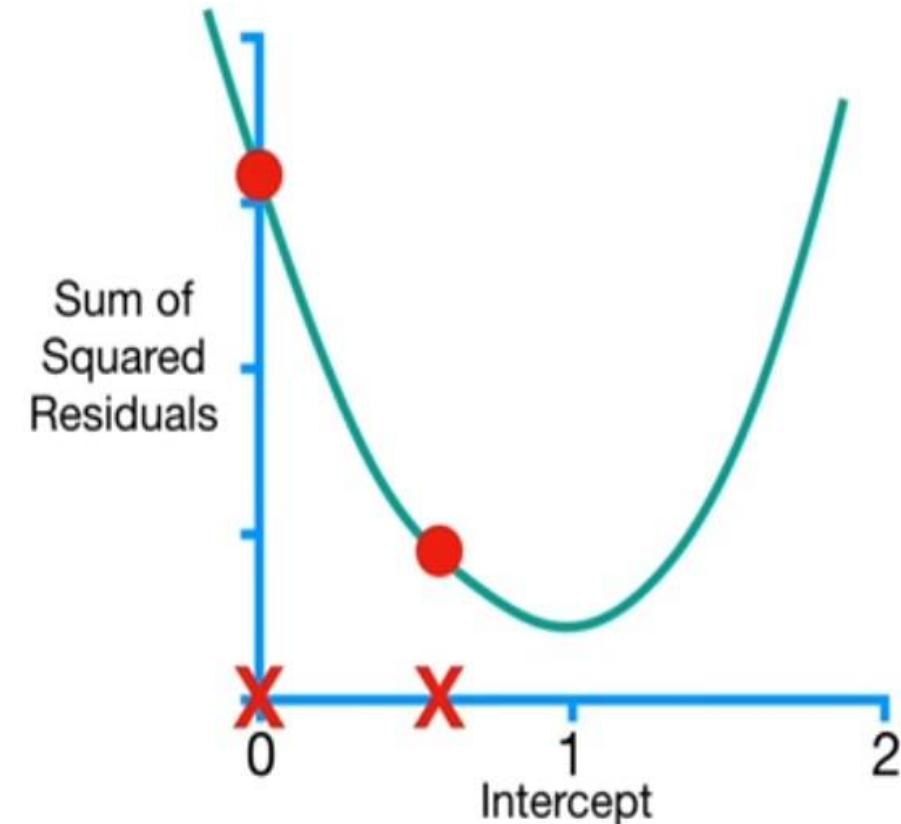
$$m = m - L \times D_m$$

$$c = c - L \times D_c$$

- We repeat this process until our loss function is a very small value or ideally 0 (which means 0 error or 100% accuracy). The value of m and c that we are left with now will be the optimum values.



...we can see how much the residuals shrink when the
Intercept = 0.57.

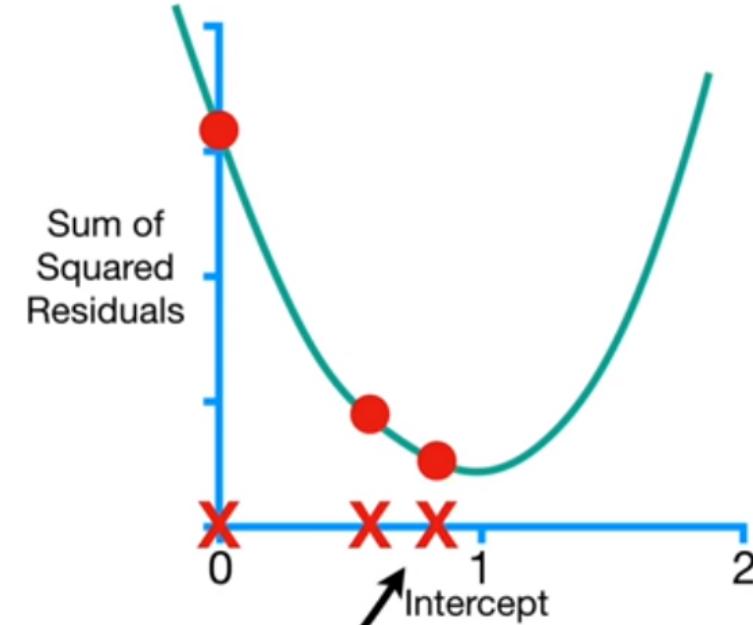


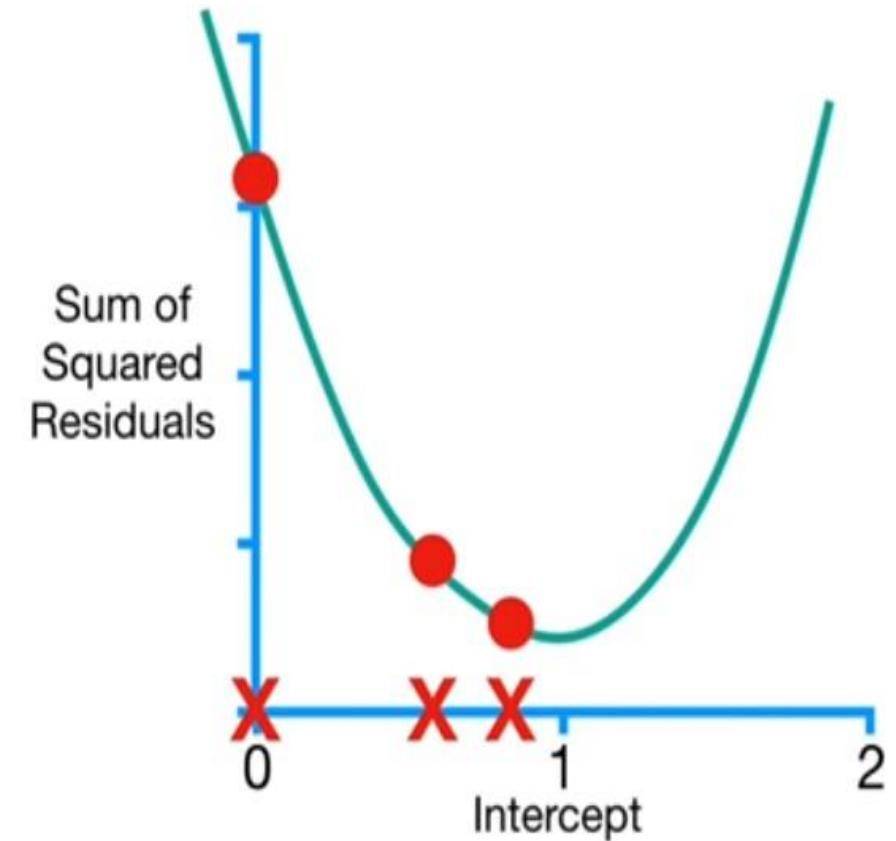
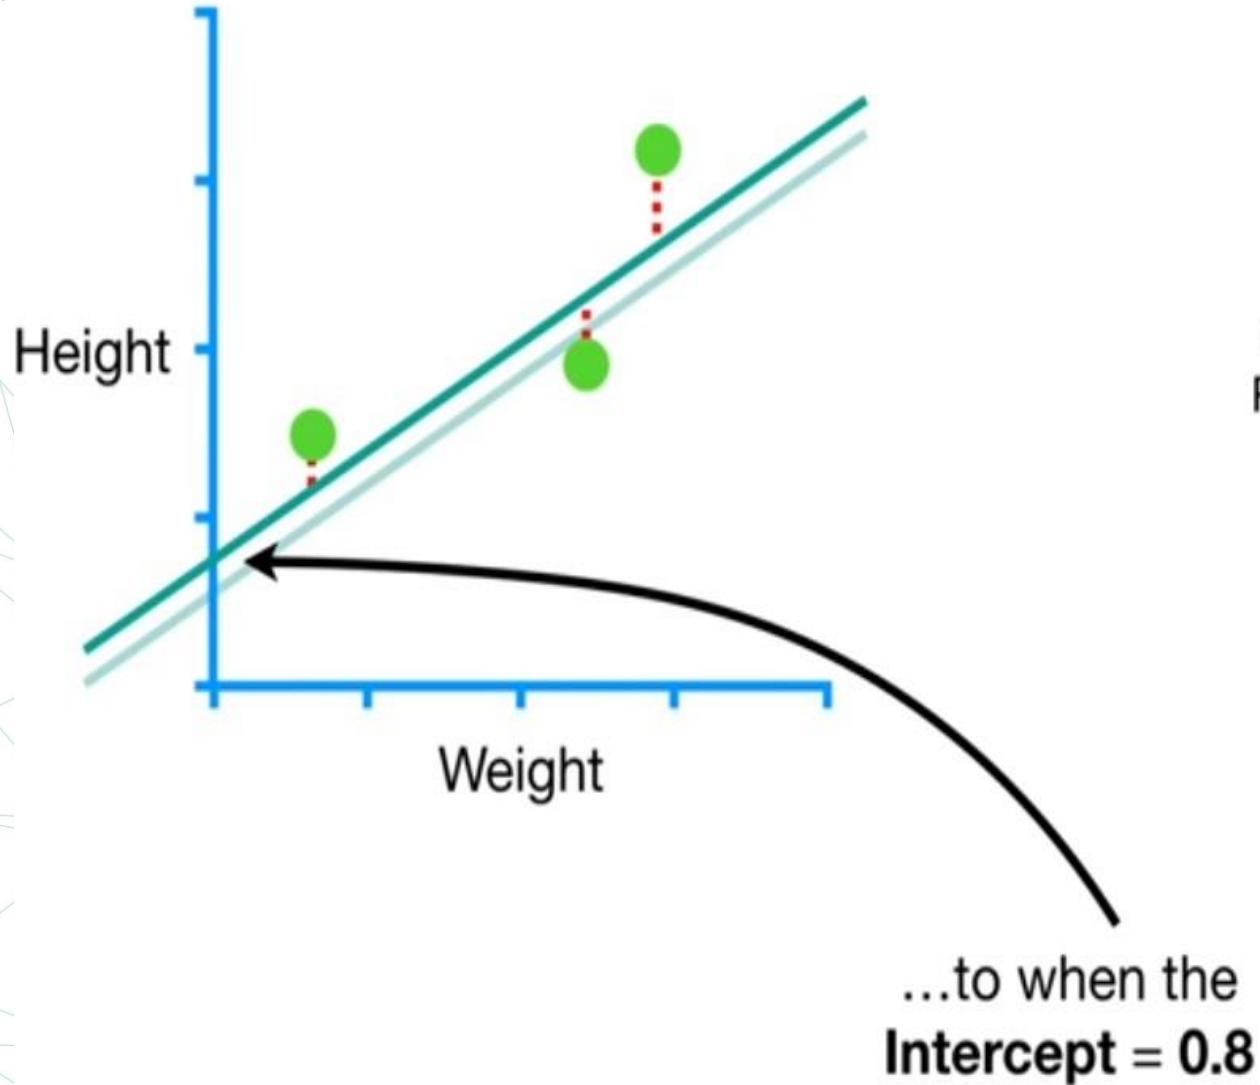
$$\frac{d}{d \text{ intercept}} \begin{aligned} \text{Sum of squared residuals} &= \\ &-2(1.4 - (0.57 + 0.64 \times 0.5)) \\ &+ -2(1.9 - (0.57 + 0.64 \times 2.3)) \\ &+ -2(3.2 - (0.57 + 0.64 \times 2.9)) \\ &= -2.3 \end{aligned}$$

$$\text{Step Size} = -2.3 \times 0.1 = -0.23$$

$$\text{New Intercept} = 0.57 - (-0.23) = \boxed{0.8}$$

...and the **New Intercept = 0.8**







$$\frac{d}{d \text{ intercept}} \text{Sum of squared residuals} =$$

$$-2(1.4 - (0.8 + 0.64 \times 0.5))$$

$$+ -2(1.9 - (0.8 + 0.64 \times 2.3))$$

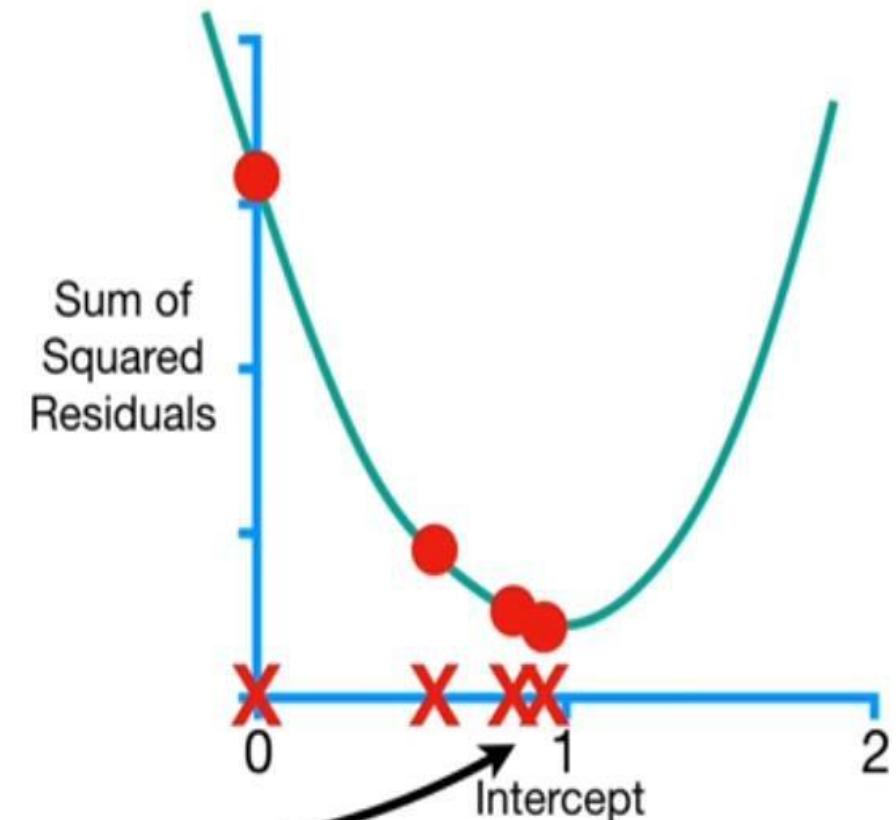
$$+ -2(3.2 - (0.8 + 0.64 \times 2.9))$$

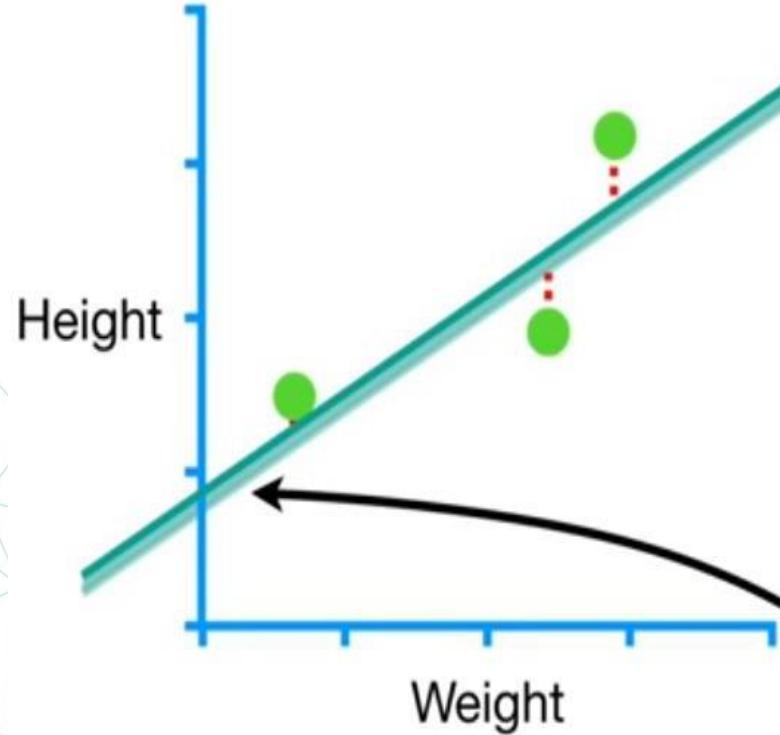
$$= -0.9$$

$$\text{Step Size} = -0.9 \times 0.1 = -0.09$$

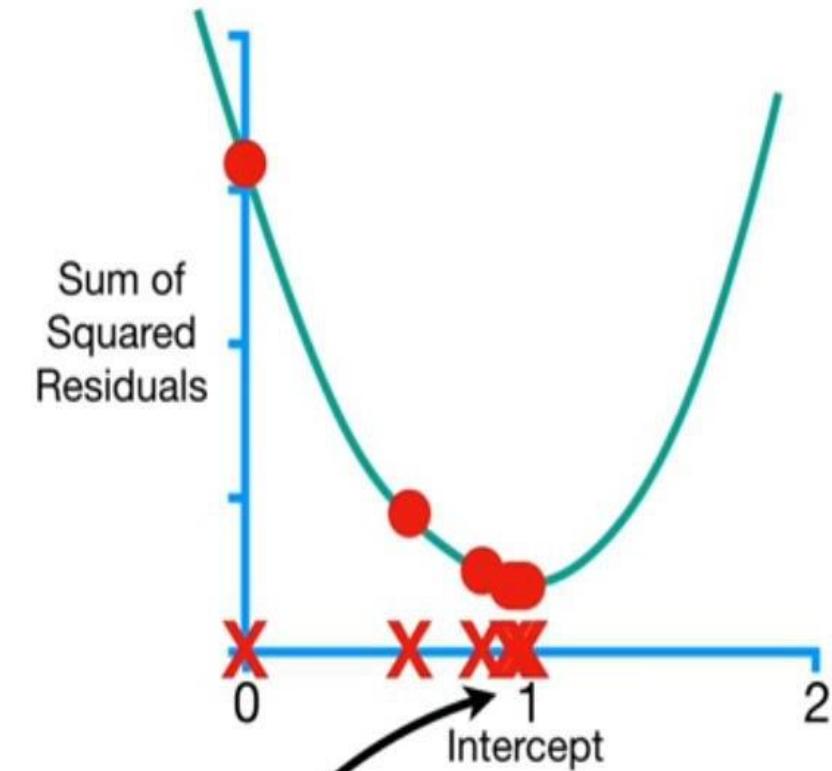
New Intercept = $0.8 - (-0.09) = 0.89$

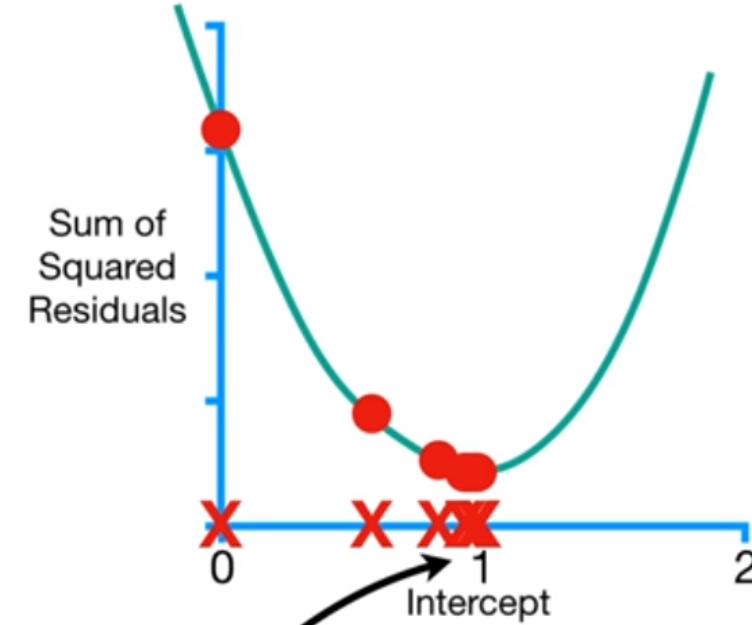
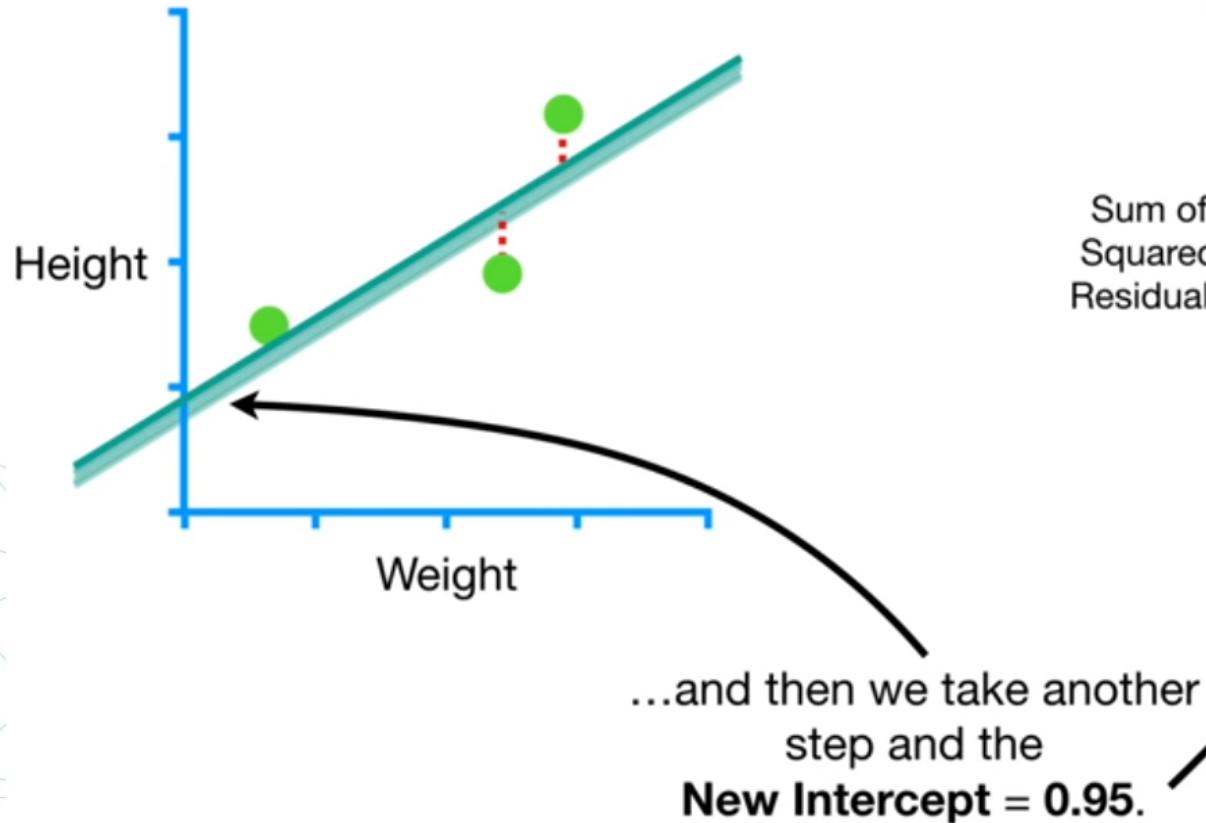
...and the **New Intercept** = 0.89

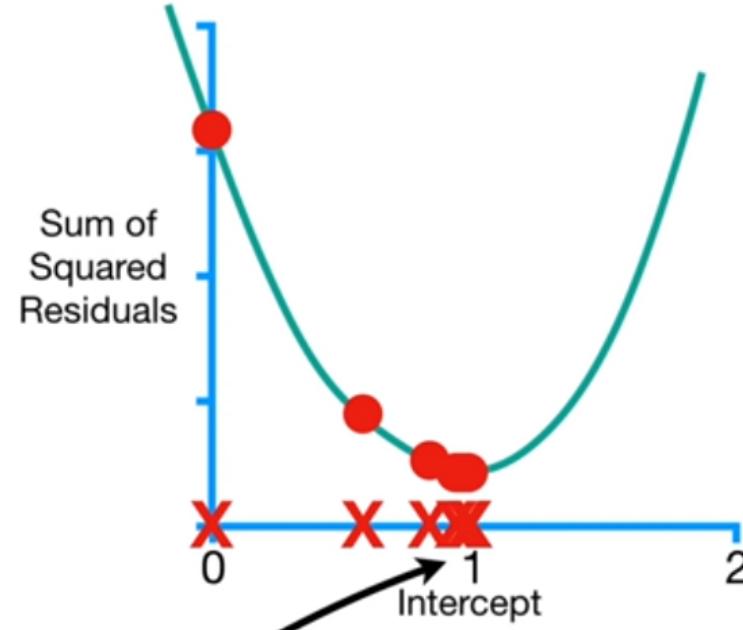




...and then we take another
step and the
New Intercept = 0.94...





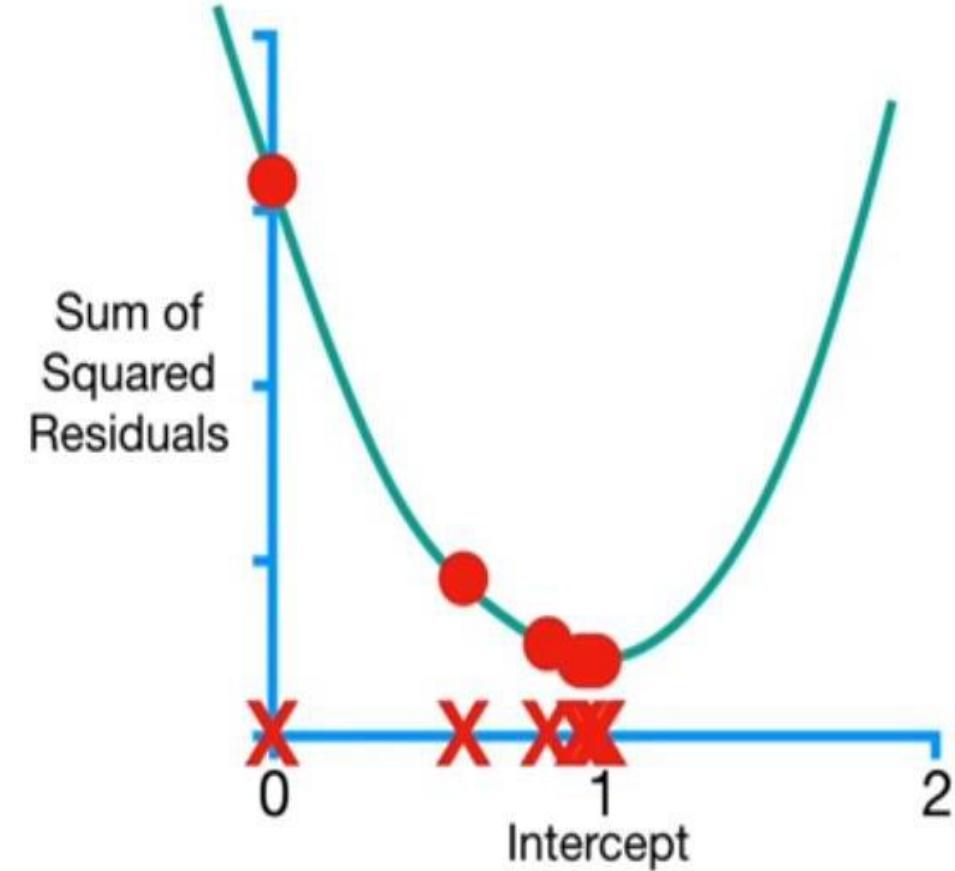


Notice how each step gets smaller and smaller the closer we get to the bottom of the curve.



Gradient Descent stops
when the **Step Size** is **Very**
Close To 0.

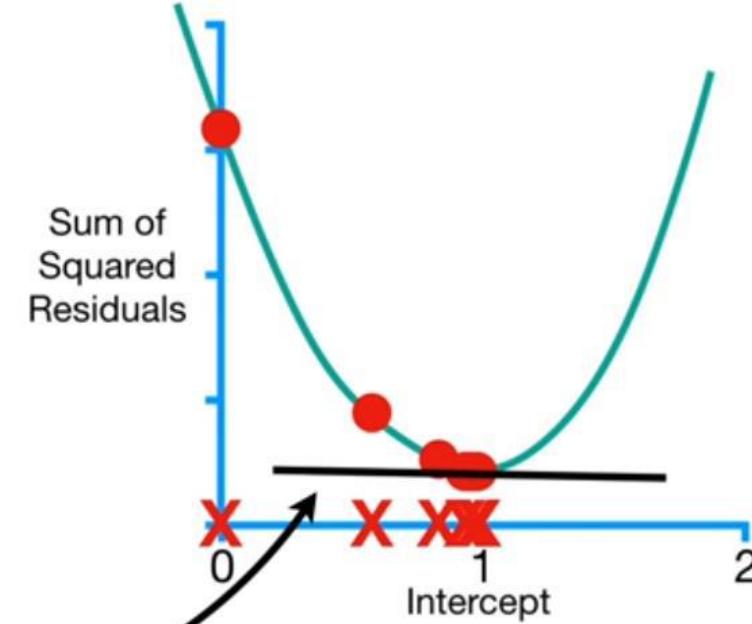
Step Size = Slope × Learning Rate





The **Step Size** will be **Very Close to 0** when the **Slope** is very close to **0**.

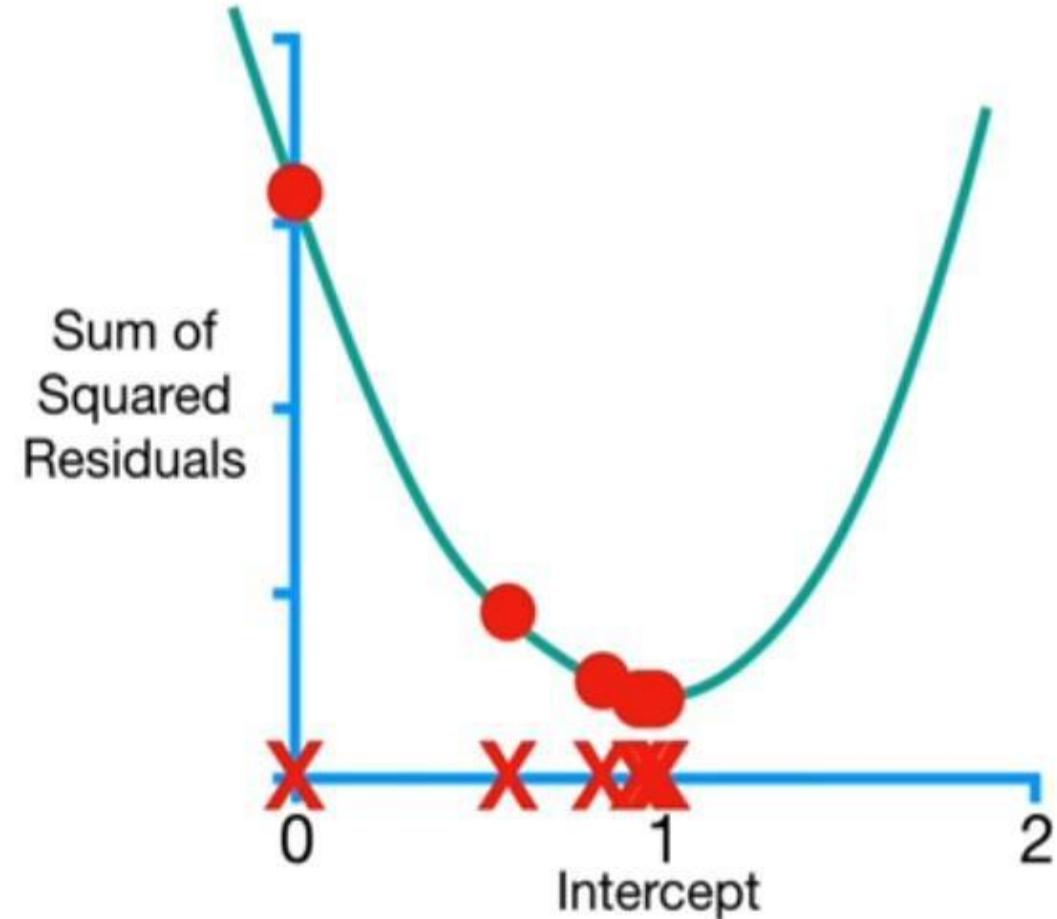
Step Size = **Slope** × Learning Rate





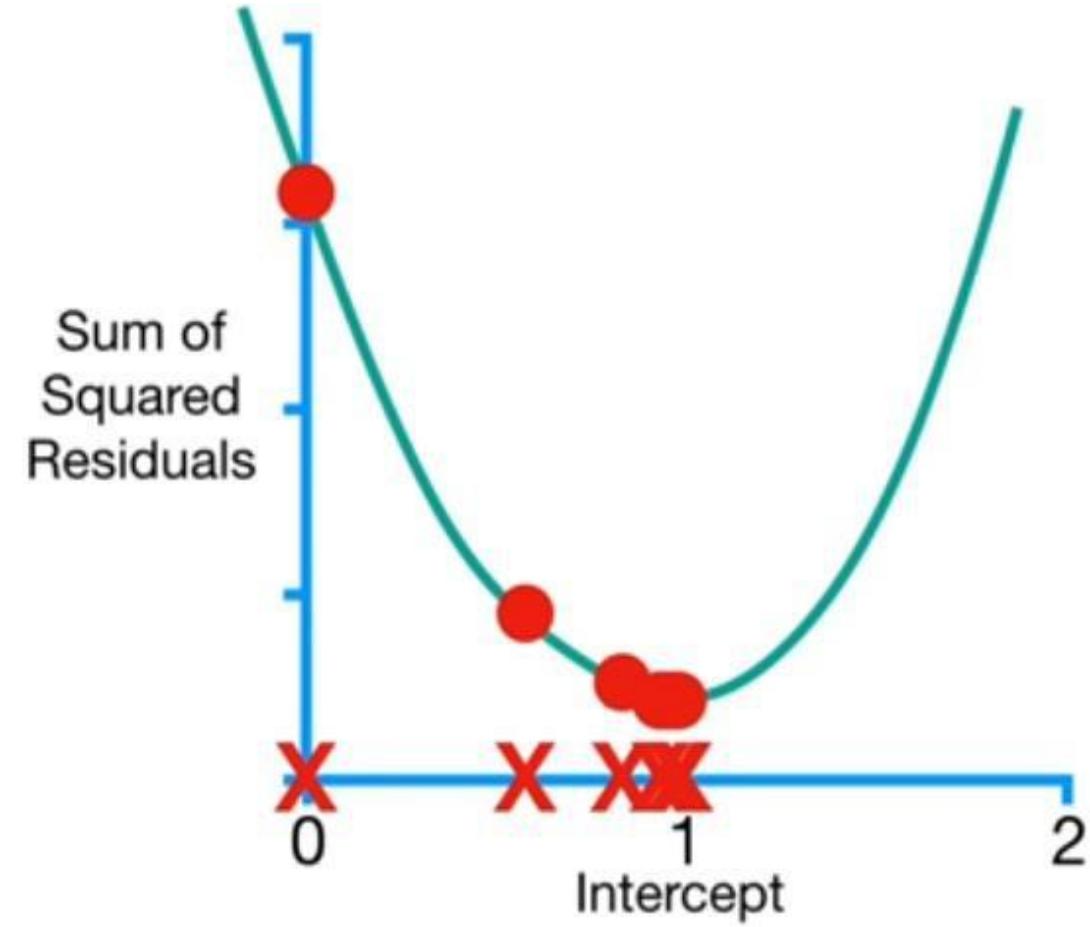
That said, **Gradient Descent** also includes a limit on the number of steps it will take before giving up.

In practice, the **Maximum Number of Steps = 1,000** or greater.

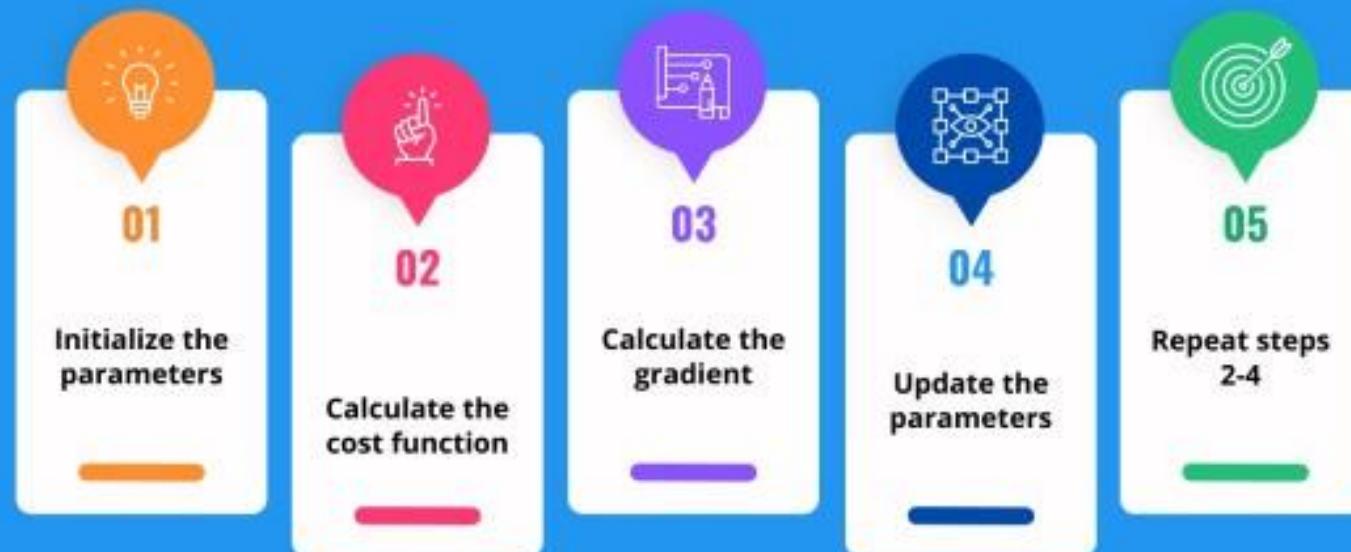




So, even if the **Step Size** is large, if there have been more than the **Maximum Number of Steps**, Gradient Descent will stop.

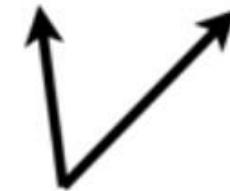


The Step-by-step Process of Performing Gradient Descent

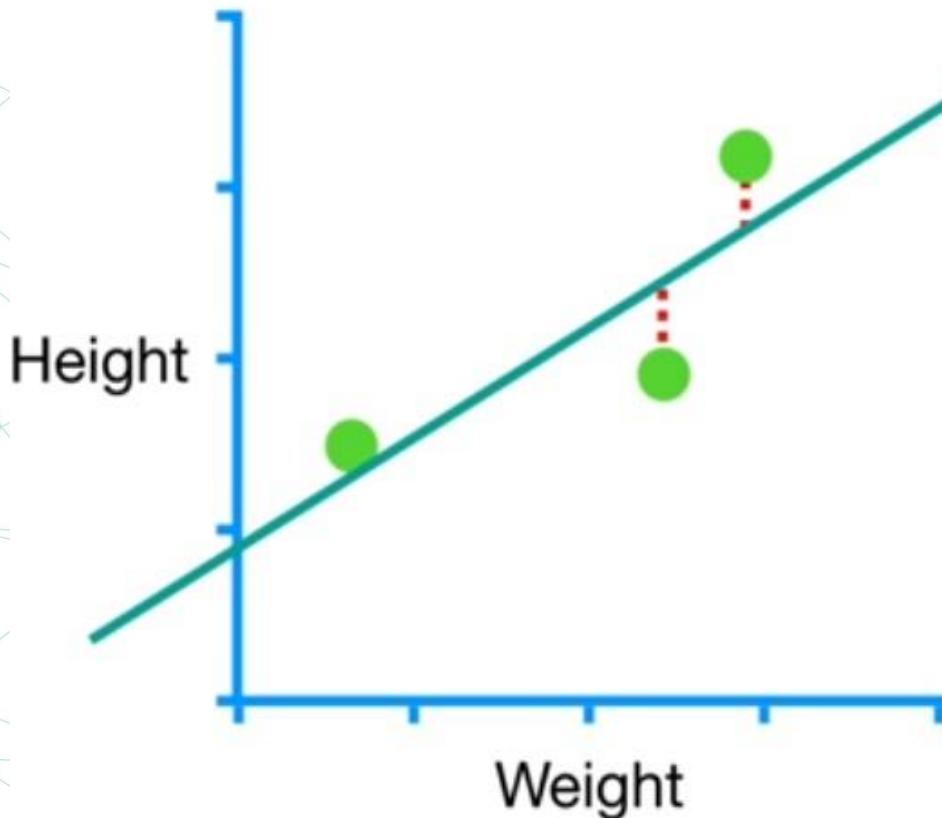




Predicted Height = intercept + slope × Weight



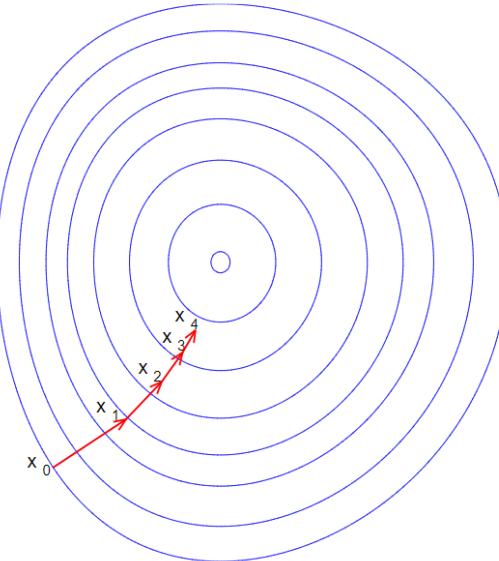
...let's talk about how to
estimate the **Intercept** and
the **Slope**.



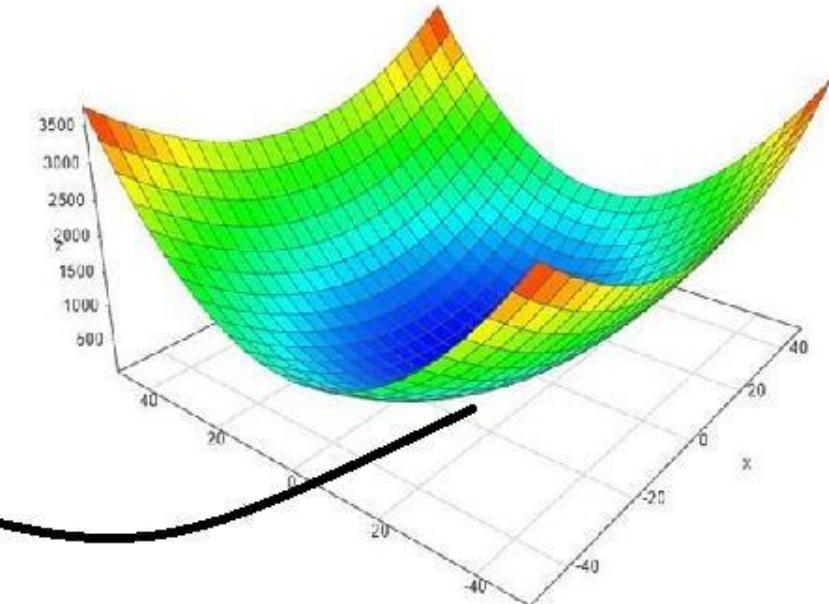


Sum of squared residuals = $(1.4 - (\text{intercept} + \text{slope} \times 0.5))^2$

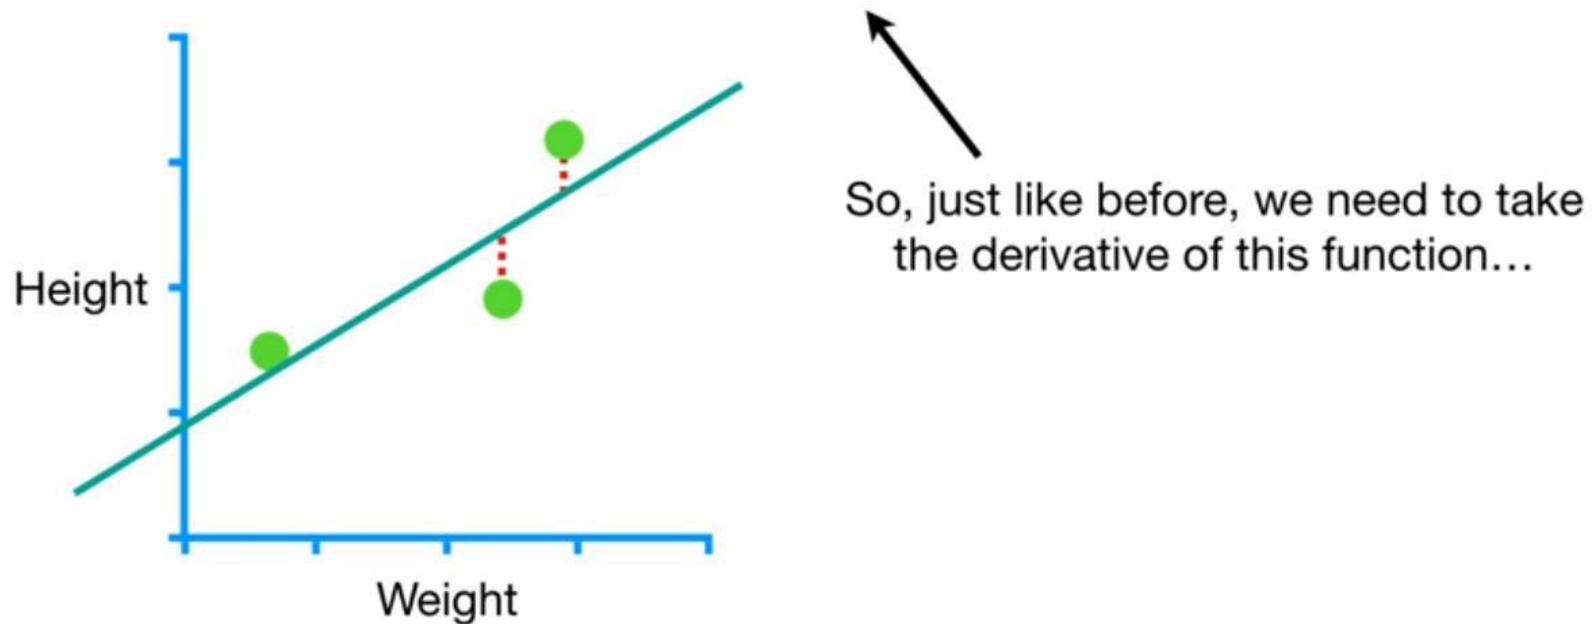
$$+ (1.9 - (\text{intercept} + \text{slope} \times 2.3))^2$$
$$+ (3.2 - (\text{intercept} + \text{slope} \times 2.9))^2$$



We want to find the values for the **Intercept** and **Slope** that give us the minimum Sum of the Squared Residuals.



Sum of squared residuals = $(1.4 - (\text{intercept} + \text{slope} \times 0.5))^2$
+ $(1.9 - (\text{intercept} + \text{slope} \times 2.3))^2$
+ $(3.2 - (\text{intercept} + \text{slope} \times 2.9))^2$



$$\begin{aligned}\text{Sum of squared residuals} = & (1.4 - (\text{intercept} + \text{slope} \times 0.5))^2 \\ & + (1.9 - (\text{intercept} + \text{slope} \times 2.3))^2 \\ & + (3.2 - (\text{intercept} + \text{slope} \times 2.9))^2\end{aligned}$$

We'll start by taking the derivative with respect to the intercept.

$$\frac{d}{d \text{ intercept}} \text{ Sum of squared residuals} =$$

$$\frac{d}{d \text{ intercept}} (1.4 - (\text{intercept} + \text{slope} \times 0.5))^2 = 2(1.4 - (\text{intercept} + \text{slope} \times 0.5)) \times -1$$

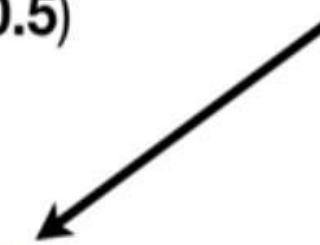


$$\frac{d}{d \text{ intercept}} 1.4 - (\text{intercept} + \text{slope} \times 0.5)$$

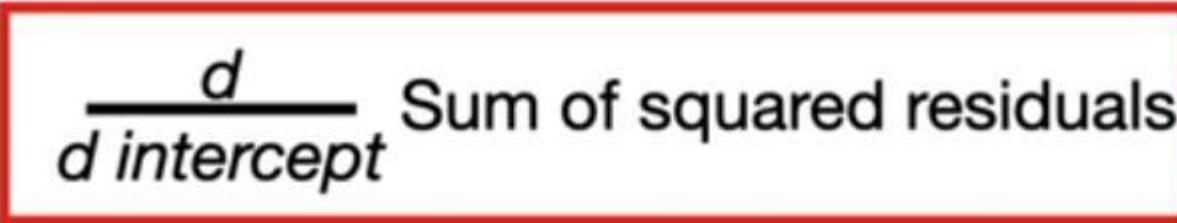


$$\frac{d}{d \text{ intercept}} 1.4 + (-1)\text{intercept} - \text{slope} \times 0.5$$

Since we are taking the derivative with respect to the **Intercept**, we treat the **Slope** like a constant, and the derivative of a constant is **0**.



....and this whole thing is the derivative of
the Sum of the Squared Residuals with
respect to the **Intercept**.


$$\frac{d}{d \text{ intercept}} \text{Sum of squared residuals} = -2(1.4 - (\text{intercept} + \text{slope} \times 0.5))$$

$$+ -2(1.9 - (\text{intercept} + \text{slope} \times 2.3))$$

$$+ -2(3.2 - (\text{intercept} + \text{slope} \times 2.9))$$

$$\begin{aligned}\text{Sum of squared residuals} &= (1.4 - (\text{intercept} + \text{slope} \times 0.5))^2 \\ &+ (1.9 - (\text{intercept} + \text{slope} \times 2.3))^2 \\ &+ (3.2 - (\text{intercept} + \text{slope} \times 2.9))^2\end{aligned}$$

Now let's take the derivative of the Sum of the Squared Residuals with respect to the **Slope**.


$$\frac{d}{d \text{ slope}} \text{Sum of squared residuals} =$$

$$\frac{d}{d \text{ slope}}$$

$$(1.4 - (\text{intercept} + \text{slope} \times 0.5))^2 = 2(1.4 - (\text{intercept} + \text{slope} \times 0.5)) \times -0.5$$



$$\frac{d}{d \text{ slope}}$$

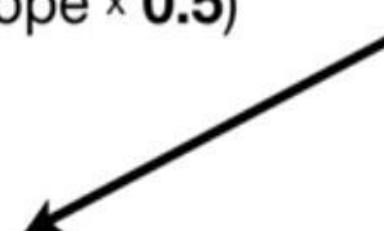
$$1.4 - (\text{intercept} + \text{slope} \times 0.5)$$

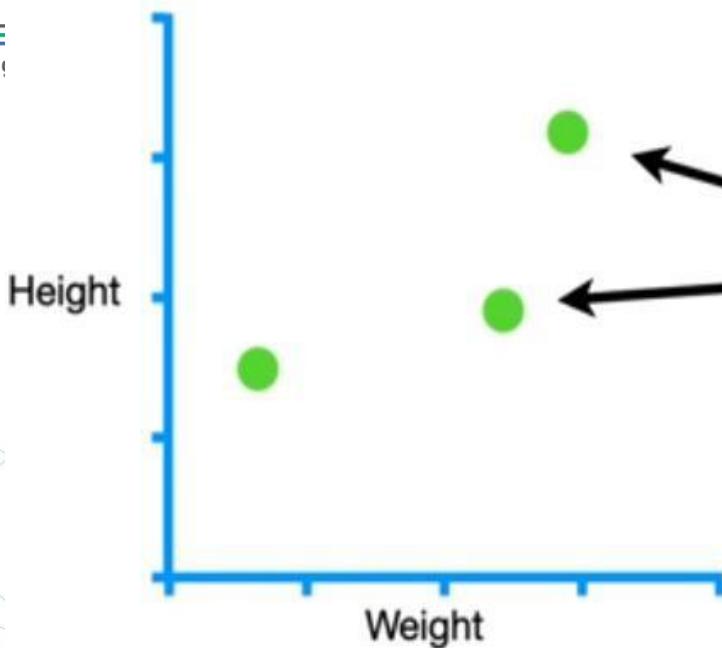


$$\frac{d}{d \text{ slope}}$$

$$1.4 + (-1)\text{intercept} - \text{slope} \times 0.5$$

Since we are taking the derivative with respect to the **Slope**, we treat the **Intercept** like a constant, and the derivative of a constant is **0**.





Again, **2.3** and **2.9** are in bold to remind us that they are the weights of the second and third samples.

$$\frac{d}{dslope} \text{ Sum of squared residuals} = -2 \times 0.5(1.4 - (\text{intercept} + \text{slope} \times 0.5))$$
$$+ -2 \times 2.3(1.9 - (\text{intercept} + \text{slope} \times 2.3))$$
$$+ -2 \times 2.9(3.2 - (\text{intercept} + \text{slope} \times 2.9))$$

$$\frac{d}{d \text{ intercept}} \text{ Sum of squared residuals} =$$

$$-2(1.4 - (\text{intercept} + \text{slope} \times 0.5))$$

$$+ -2(1.9 - (\text{intercept} + \text{slope} \times 2.3))$$

$$+ -2(3.2 - (\text{intercept} + \text{slope} \times 2.9))$$

NOTE: When you have two or more derivatives of the same function, they are called a **Gradient**.

$$\frac{d}{d \text{ slope}} \text{ Sum of squared residuals} =$$

$$-2 \times 0.5(1.4 - (\text{intercept} + \text{slope} \times 0.5))$$

$$+ -2 \times 2.9(3.2 - (\text{intercept} + \text{slope} \times 2.9))$$

$$+ -2 \times 2.3(1.9 - (\text{intercept} + \text{slope} \times 2.3))$$

$$\frac{d}{d \text{ intercept}} \text{ Sum of squared residuals} =$$
$$-2(1.4 - (\text{intercept} + \text{slope} \times 0.5))$$
$$+ -2(1.9 - (\text{intercept} + \text{slope} \times 2.3))$$
$$+ -2(3.2 - (\text{intercept} + \text{slope} \times 2.9))$$

Just like before, we will start by picking a random number for the **Intercept**. In this case we'll set the **Intercept = 0...**

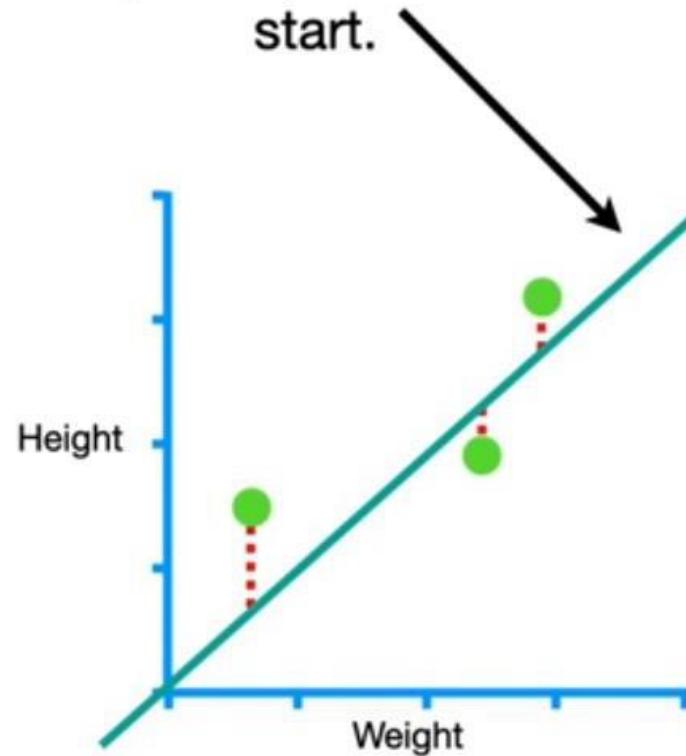
...and we'll pick a random number for the **Slope**. In this case we'll set the **Slope = 1.**

$$\frac{d}{d \text{ slope}} \text{ Sum of squared residuals} =$$
$$-2 \times 0.5(1.4 - (\text{intercept} + \text{slope} \times 0.5))$$
$$+ -2 \times 2.9(3.2 - (\text{intercept} + \text{slope} \times 2.9))$$
$$+ -2 \times 2.3(1.9 - (\text{intercept} + \text{slope} \times 2.3))$$

$$\frac{d}{d \text{ intercept}} \begin{aligned} & \text{Sum of squared residuals} = \\ & -2(1.4 - (\text{intercept} + \text{slope} \times 0.5)) \\ & + -2(1.9 - (\text{intercept} + \text{slope} \times 2.3)) \\ & + -2(3.2 - (\text{intercept} + \text{slope} \times 2.9)) \end{aligned}$$

$$\frac{d}{d \text{ slope}} \begin{aligned} & \text{Sum of squared residuals} = \\ & -2 \times 0.5(1.4 - (\text{intercept} + \text{slope} \times 0.5)) \\ & + -2 \times 2.9(3.2 - (\text{intercept} + \text{slope} \times 2.9)) \\ & + -2 \times 2.3(1.9 - (\text{intercept} + \text{slope} \times 2.3)) \end{aligned}$$

Thus, this line, with **Intercept = 0** and **Slope = 1**, is where we will start.



$$\frac{d}{d \text{ intercept}} \text{ Sum of squared residuals} =$$

$$-2(1.4 - (\text{intercept} + \text{slope} \times 0.5))$$

$$+ -2(1.9 - (\text{intercept} + \text{slope} \times 2.3))$$

$$+ -2(3.2 - (\text{intercept} + \text{slope} \times 2.9))$$

Now let's plug in **0** for the
Intercept and **1** for the **Slope**...

$$\frac{d}{d \text{ slope}} \text{ Sum of squared residuals} =$$

$$-2 \times 0.5(1.4 - (\text{intercept} + \text{slope} \times 0.5))$$

$$+ -2 \times 2.9(3.2 - (\text{intercept} + \text{slope} \times 2.9))$$

$$+ -2 \times 2.3(1.9 - (\text{intercept} + \text{slope} \times 2.3))$$

$\frac{d}{d \text{ intercept}}$ Sum of squared residuals =

$$-2(1.4 - (0 + 1 \times 0.5))$$

$$+ -2(1.9 - (0 + 1 \times 2.3))$$

$$+ -2(3.2 - (0 + 1 \times 2.9))$$

$$= -1.6$$

Step Size_{Intercept} = $-1.6 \times \text{Learning Rate}$



...now we plug the
Slopes into the **Step
Size** formulas...

$\frac{d}{d \text{ slope}}$ Sum of squared residuals =

$$-2 \times 0.5(1.4 - (0 + 1 \times 0.5))$$

$$+ -2 \times 2.9(3.2 - (0 + 1 \times 2.9))$$

$$+ -2 \times 2.3(1.9 - (0 + 1 \times 2.3))$$

Step Size_{Slope} = $-0.8 \times \text{Learning Rate}$





$$\frac{d}{d \text{ intercept}} \text{ Sum of squared residuals} =$$

$$-2(1.4 - (0 + 1 \times 0.5))$$

$$+ -2(1.9 - (0 + 1 \times 2.3))$$

$$+ -2(3.2 - (0 + 1 \times 2.9)) = -1.6$$

Step Size_{Intercept} = -1.6×0.01

NOTE: The larger **Learning Rate** that we used in the first example doesn't work this time. Even after a bunch of steps, **Gradient Descent** doesn't arrive at the correct answer.

$$\frac{d}{d \text{ slope}} \text{ Sum of squared residuals} =$$

$$-2 \times 0.5(1.4 - (0 + 1 \times 0.5))$$

$$+ -2 \times 2.9(3.2 - (0 + 1 \times 2.9))$$

$$+ -2 \times 2.3(1.9 - (0 + 1 \times 2.3)) = -0.8$$

Step Size_{Slope} = -0.8×0.01

$\frac{d}{d \text{ intercept}}$ Sum of squared residuals =

$$-2(1.4 - (0 + 1 \times 0.5))$$

$$+ -2(1.9 - (0 + 1 \times 2.3))$$

$$+ -2(3.2 - (0 + 1 \times 2.9)) = -1.6$$

Step Size_{Intercept} = $-1.6 \times 0.01 = -0.016$

$\frac{d}{d \text{ slope}}$ Sum of squared residuals =

$$-2 \times 0.5(1.4 - (0 + 1 \times 0.5))$$

$$+ -2 \times 2.9(3.2 - (0 + 1 \times 2.9))$$

$$+ -2 \times 2.3(1.9 - (0 + 1 \times 2.3)) = -0.8$$

Anyway, we do the math and get two **Step Sizes**.

Step Size_{Slope} = $-0.8 \times 0.01 = -0.008$

$\frac{d}{d \text{ intercept}}$ Sum of squared residuals =

$$-2(1.4 - (0 + 1 \times 0.5))$$

$$+ -2(1.9 - (0 + 1 \times 2.3))$$

$$+ -2(3.2 - (0 + 1 \times 2.9)) = -1.6$$

$$\text{Step Size}_{\text{Intercept}} = -1.6 \times 0.01 = -0.016$$

New Intercept = Old Intercept - Step Size

Now we calculate the
New Intercept and **New Slope** by plugging in the
Old Intercept and the
Old Slope...

$\frac{d}{d \text{ slope}}$ Sum of squared residuals =

$$-2 \times 0.5(1.4 - (0 + 1 \times 0.5))$$

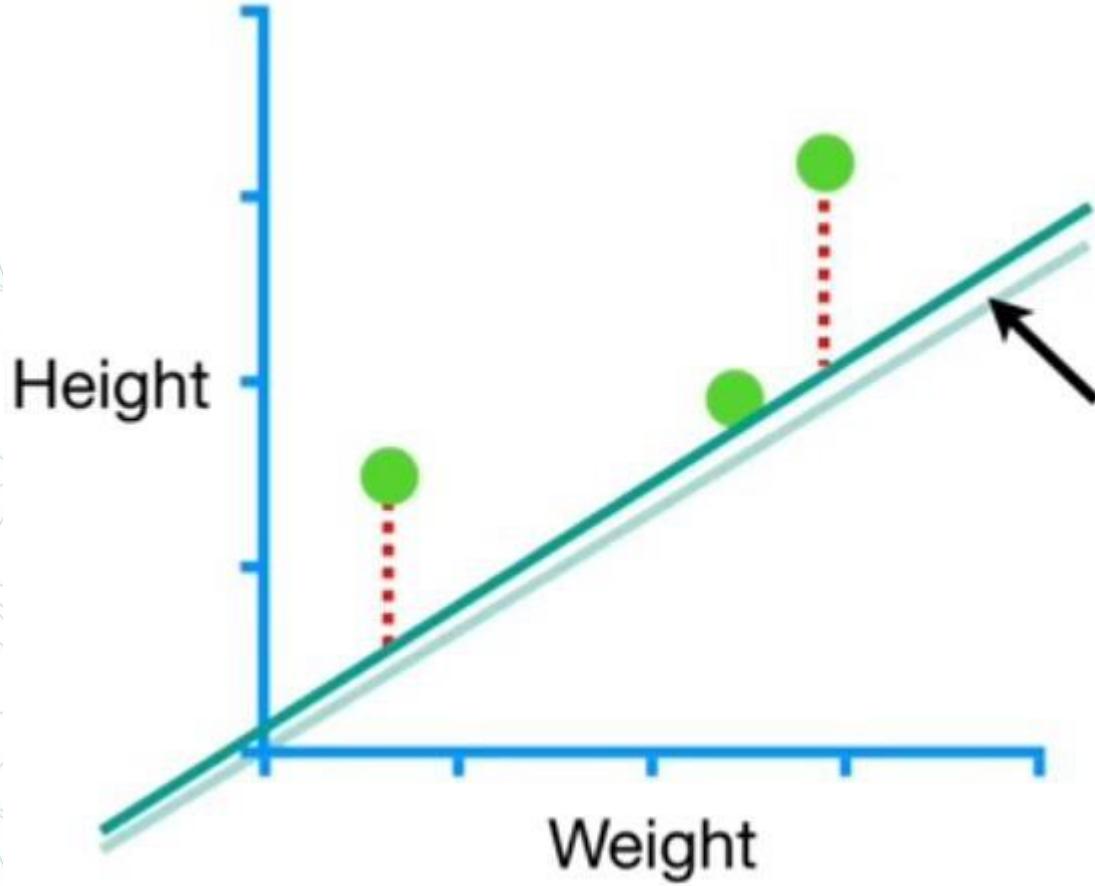
$$+ -2 \times 2.9(3.2 - (0 + 1 \times 2.9))$$

$$+ -2 \times 2.3(1.9 - (0 + 1 \times 2.3))$$

$$\text{Step Size}_{\text{Slope}} = -0.8 \times 0.01 = -0.008$$

New Slope = Old Slope - Step Size

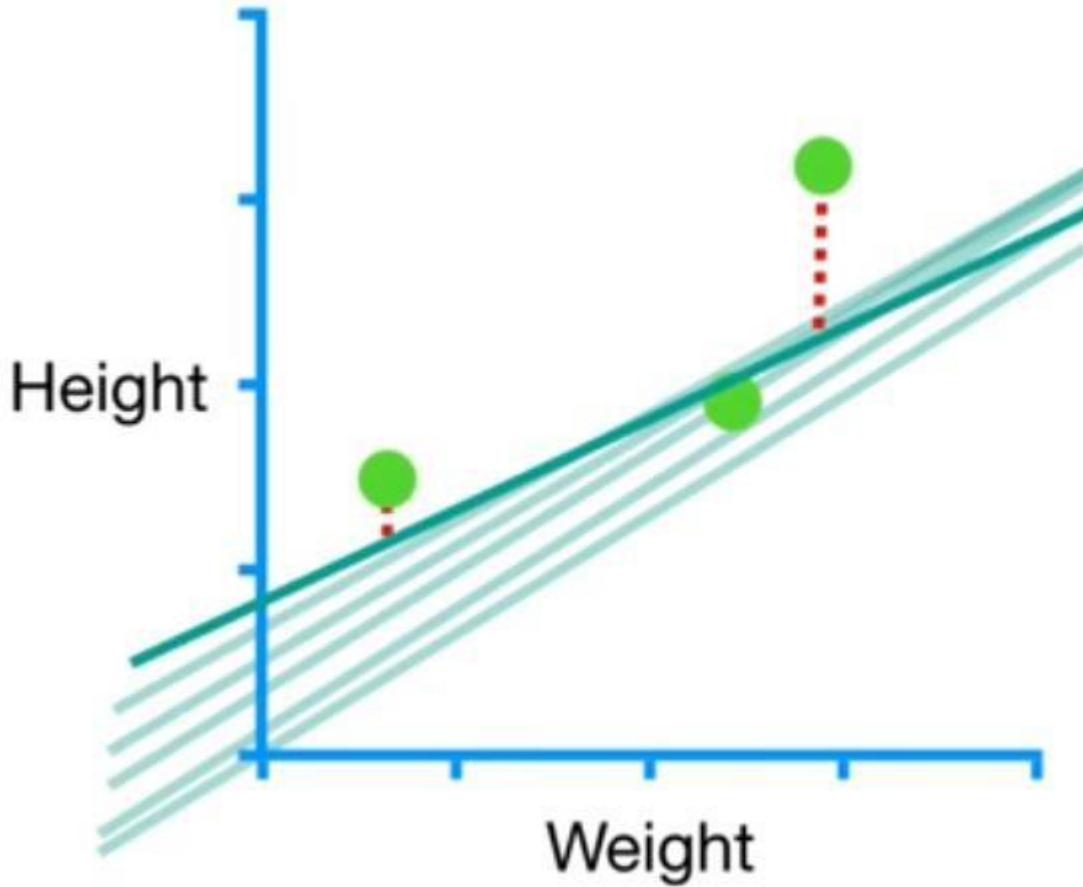
$$= -0.8$$



$$\text{New Intercept} = 0 - (-0.016) = 0.016$$

...and this is the new line
(with **Slope = 1.008** and
Intercept = 0.016) after
the first step.

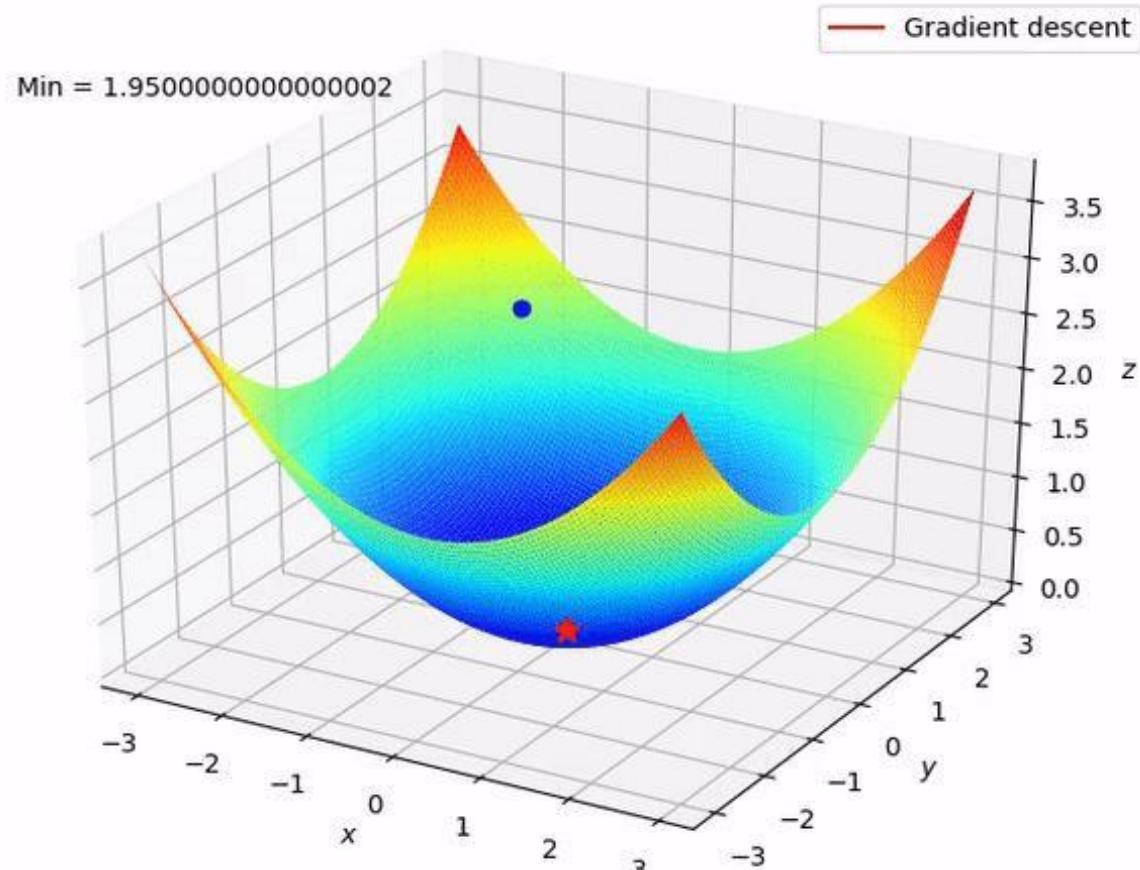
$$\text{New Slope} = 1 - (-0.008) = 1.008$$



Now we just repeat what we did until all of the **Steps Sizes** are very small or we reach the **Maximum Number of Steps**.



```
1 import numpy as np
2
3 def gradient_descent(start, gradient, learn_rate, max_iter, tol=0.01):
4     steps = [start] # history tracking
5     x = start
6
7     for _ in range(max_iter):
8         diff = learn_rate*gradient(x)
9         if np.abs(diff)<tol:
10             break
11         x = x - diff
12         steps.append(x) # history tracing
13
14 return steps, x
```



Cost Function

$$J(\Theta_0, \Theta_1) = \frac{1}{2m} \sum_{i=1}^m [h_\Theta(x_i) - y_i]^2$$

↑ Predicted Value
↑ True Value

Gradient Descent

$$\Theta_j := \Theta_j - \alpha \frac{\partial}{\partial \Theta_j} J(\Theta_0, \Theta_1)$$

↑ Learning Rate

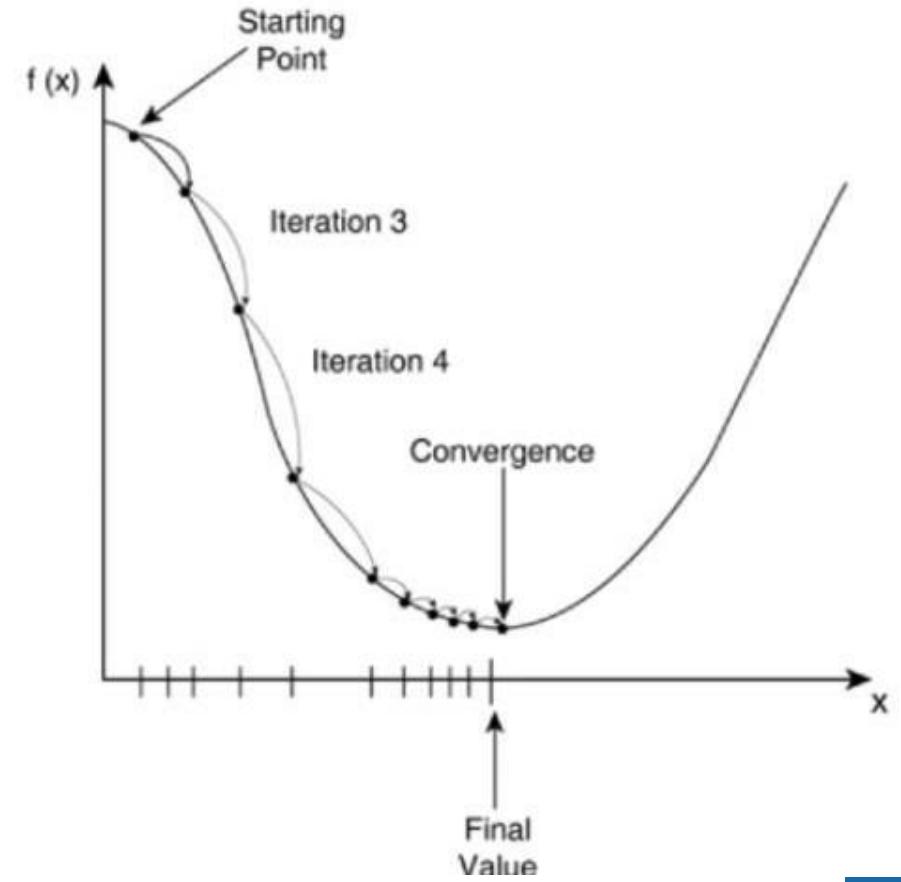
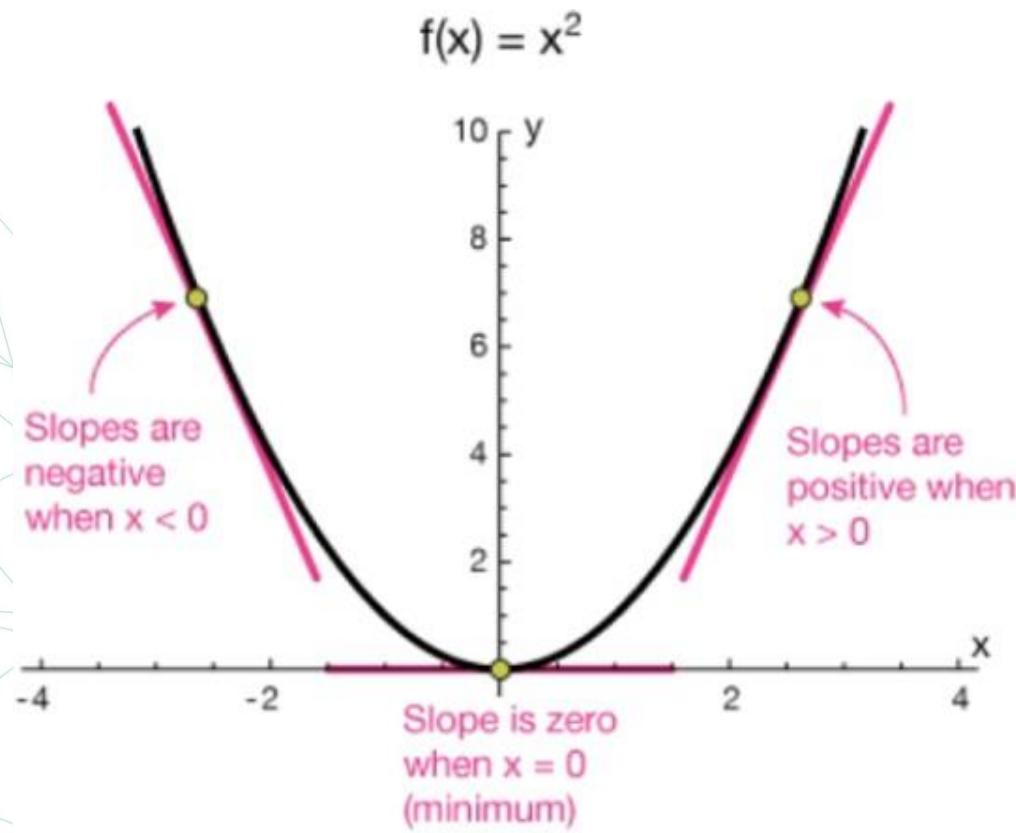
Now,

$$\begin{aligned}
 \frac{\partial}{\partial \Theta} J_\Theta &= \frac{\partial}{\partial \Theta} \frac{1}{2m} \sum_{i=1}^m [h_\Theta(x_i) - y_i]^2 \\
 &= \frac{1}{m} \sum_{i=1}^m (h_\Theta(x_i) - y_i) \frac{\partial}{\partial \Theta_j} (\Theta x_i - y) \\
 &= \frac{1}{m} (h_\Theta(x_i) - y) x_i
 \end{aligned}$$

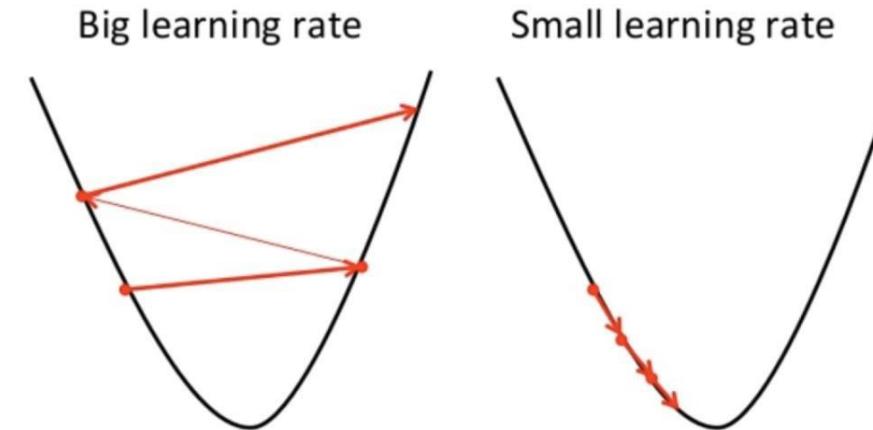
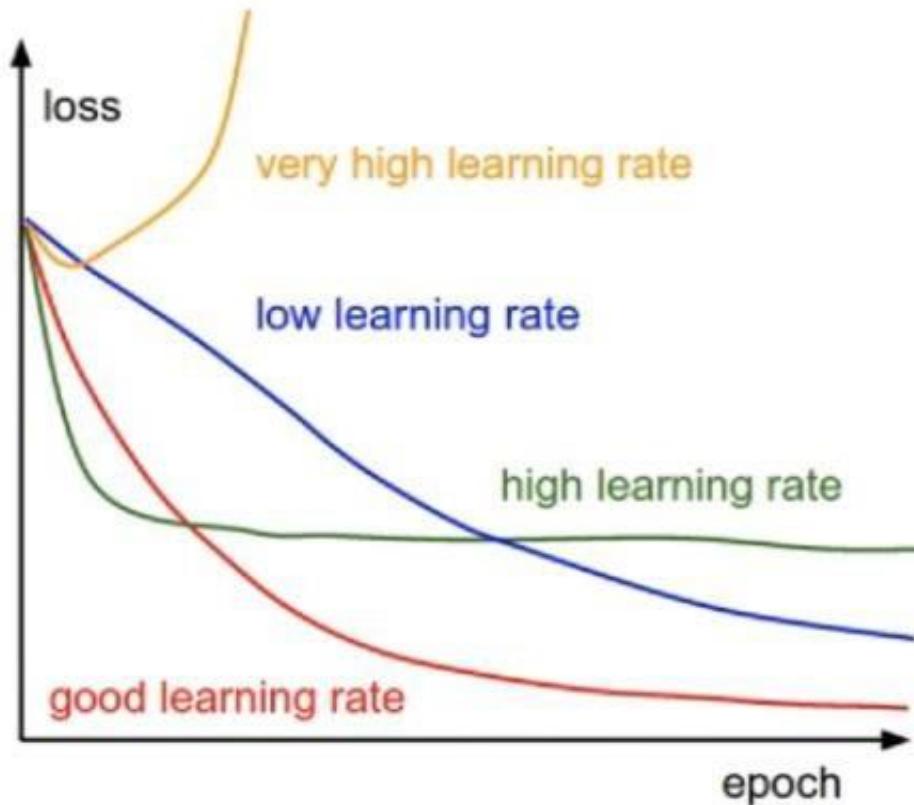
Therefore,

$$\Theta_j := \Theta_j - \frac{\alpha}{m} \sum_{i=1}^m [(h_\Theta(x_i) - y) x_i]$$

$$\begin{aligned}
 \theta_0 &:= \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \\
 \theta_1 &:= \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \cdot x^{(i)}
 \end{aligned}$$

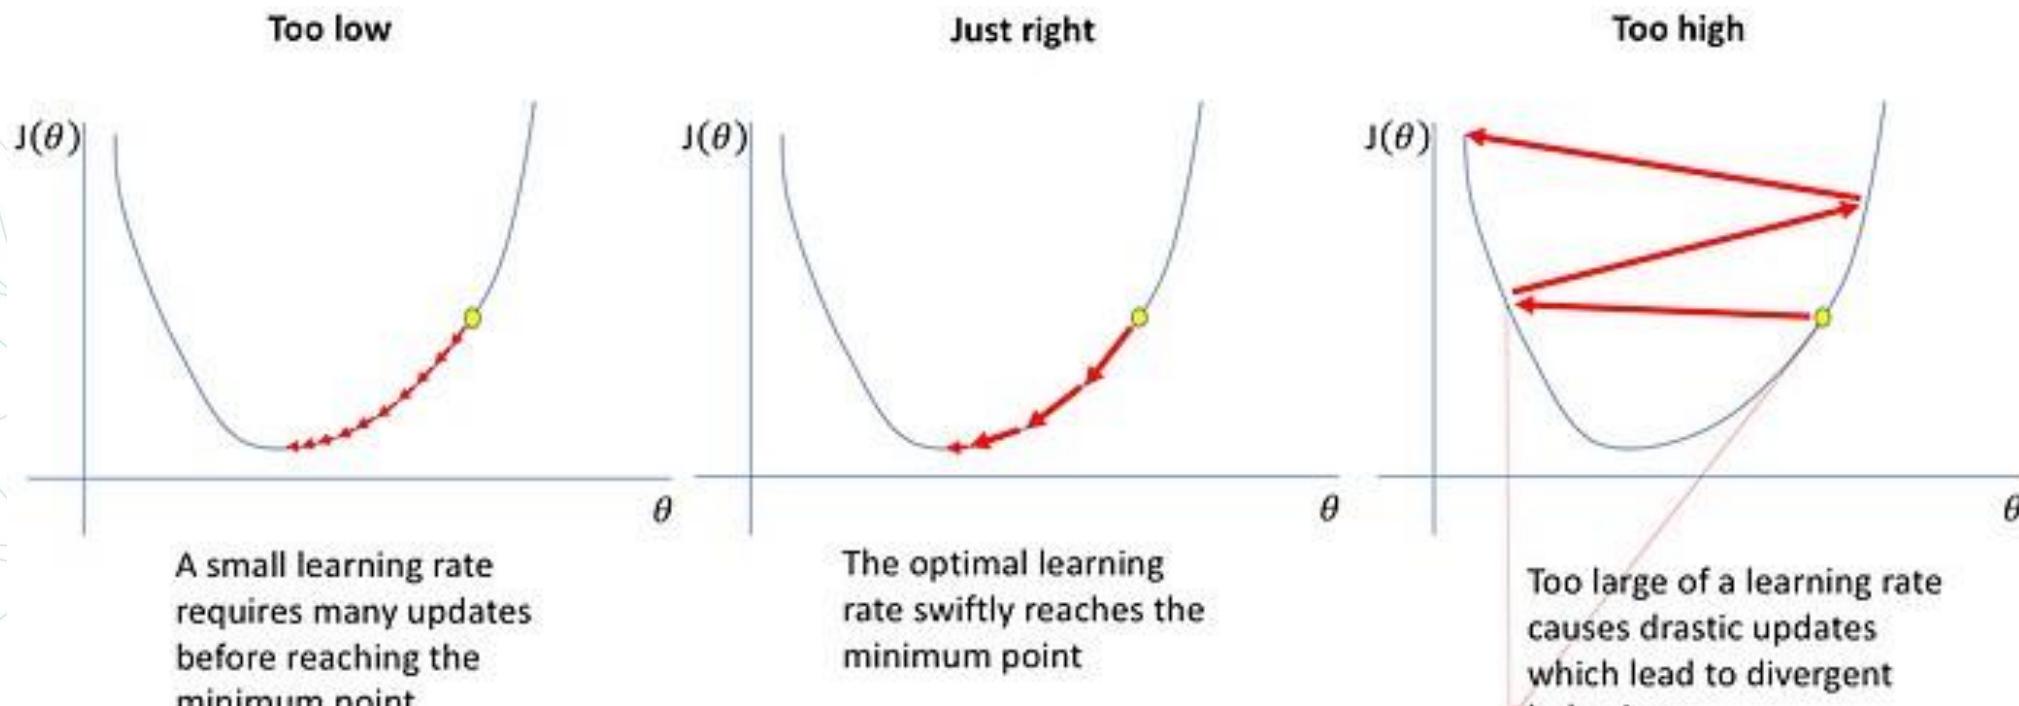


How to choose learning rate?

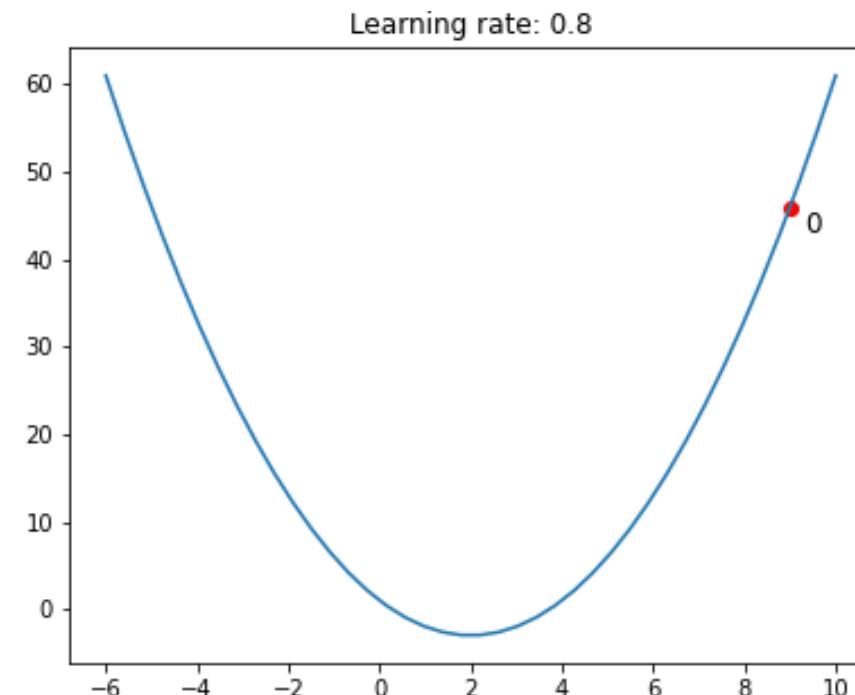
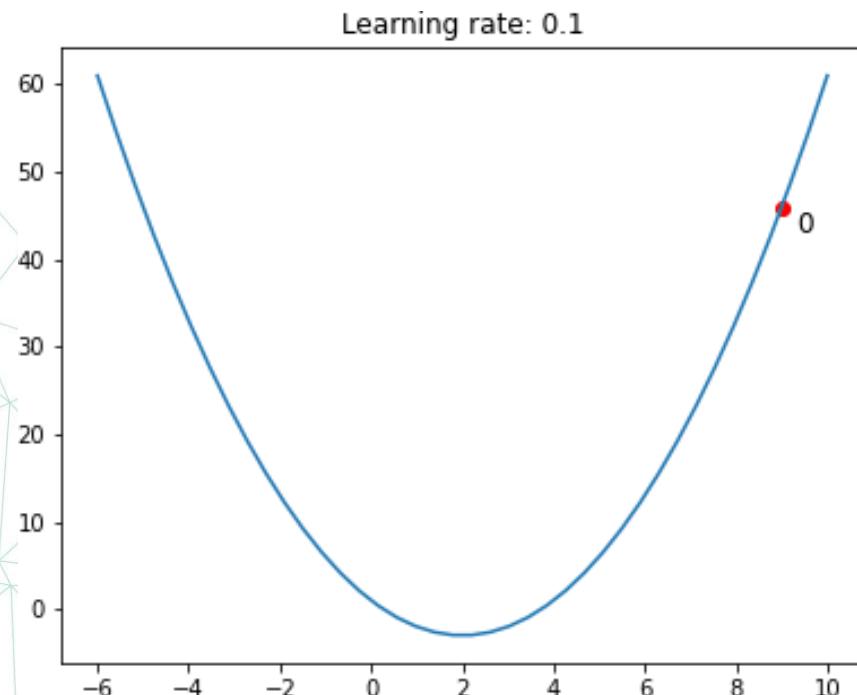




How to choose learning rate?

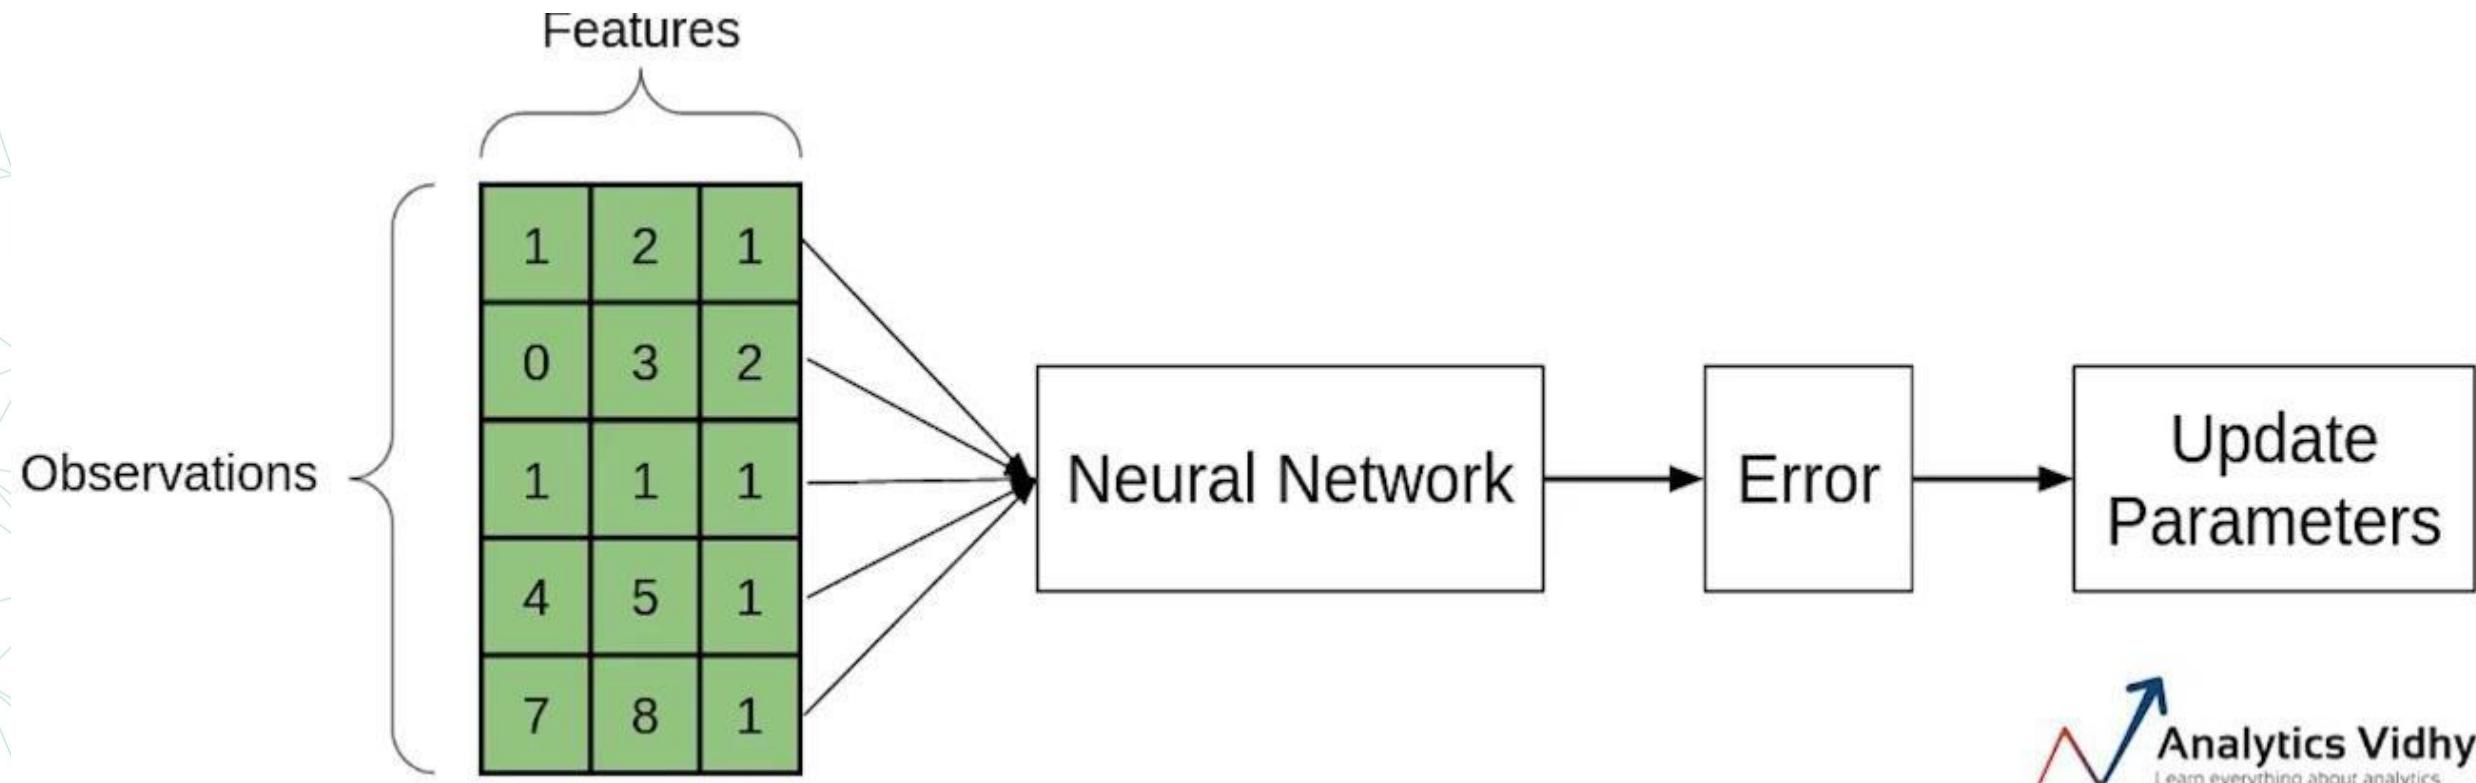


How to choose learning rate?



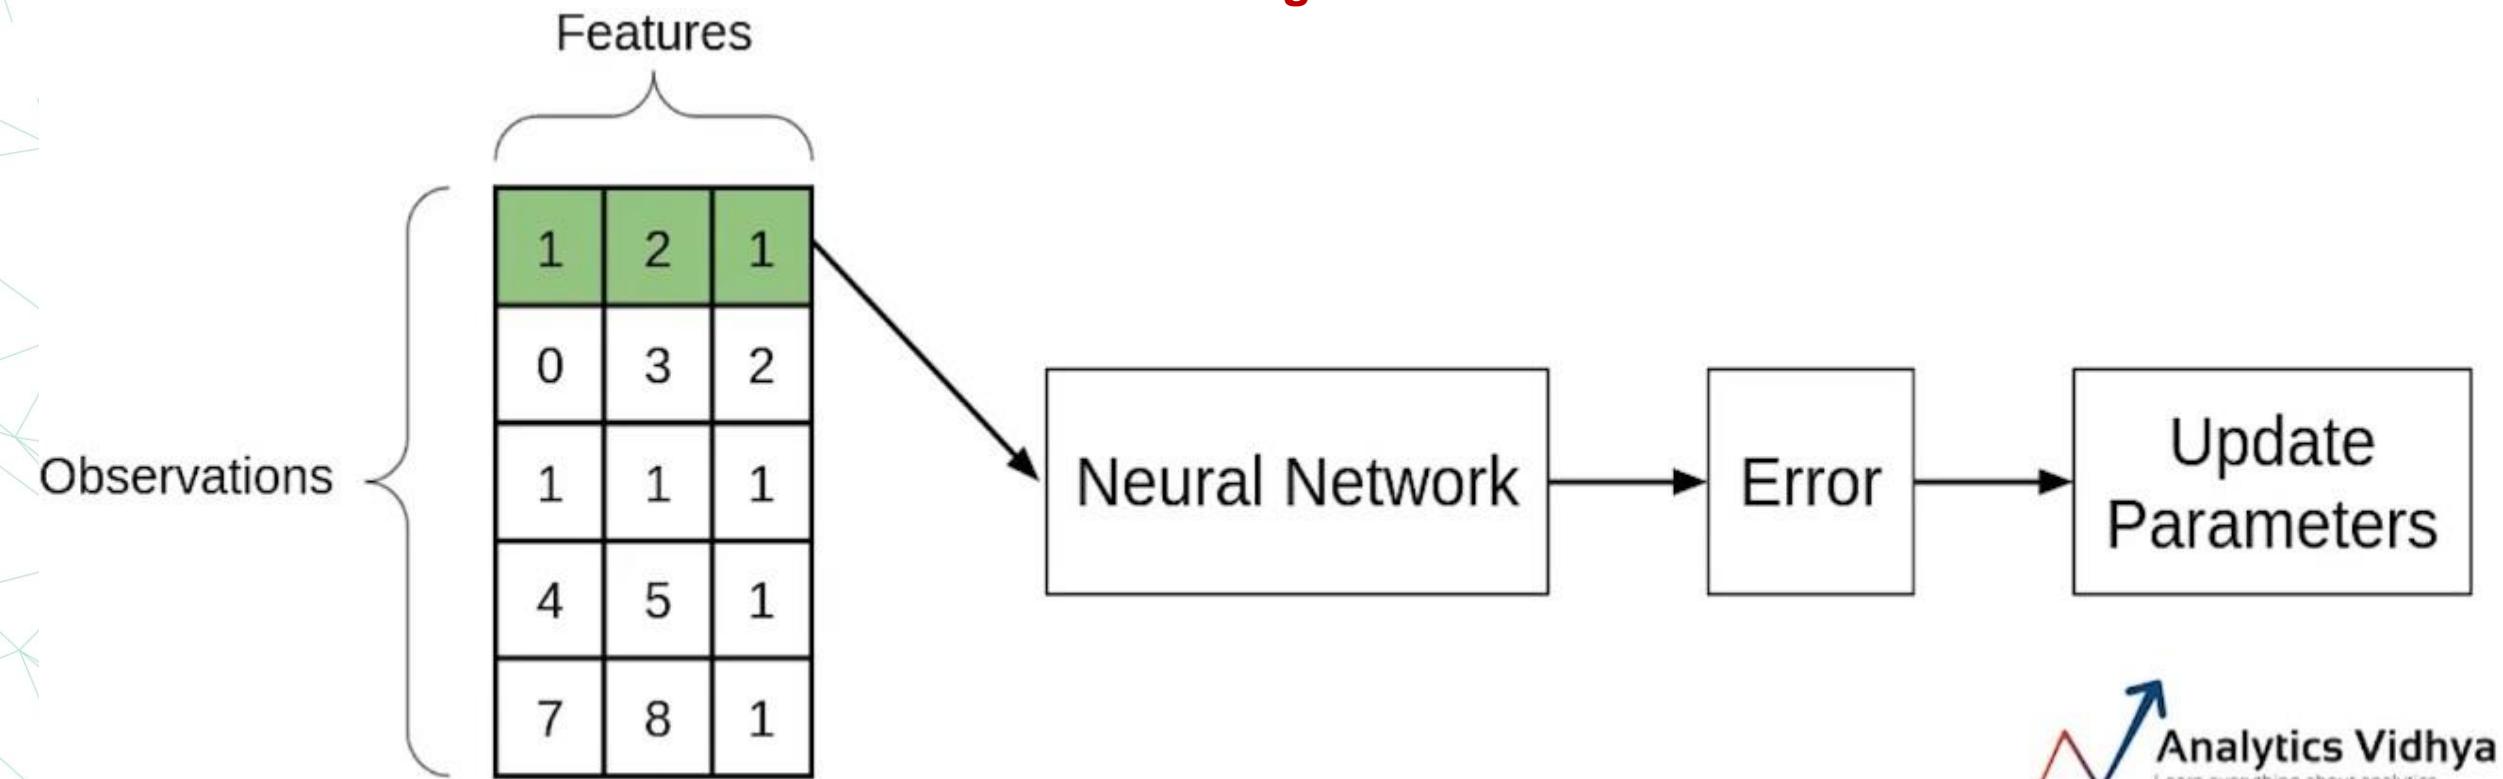
Variants of Gradient Descent Algorithm

Batch gradient descent

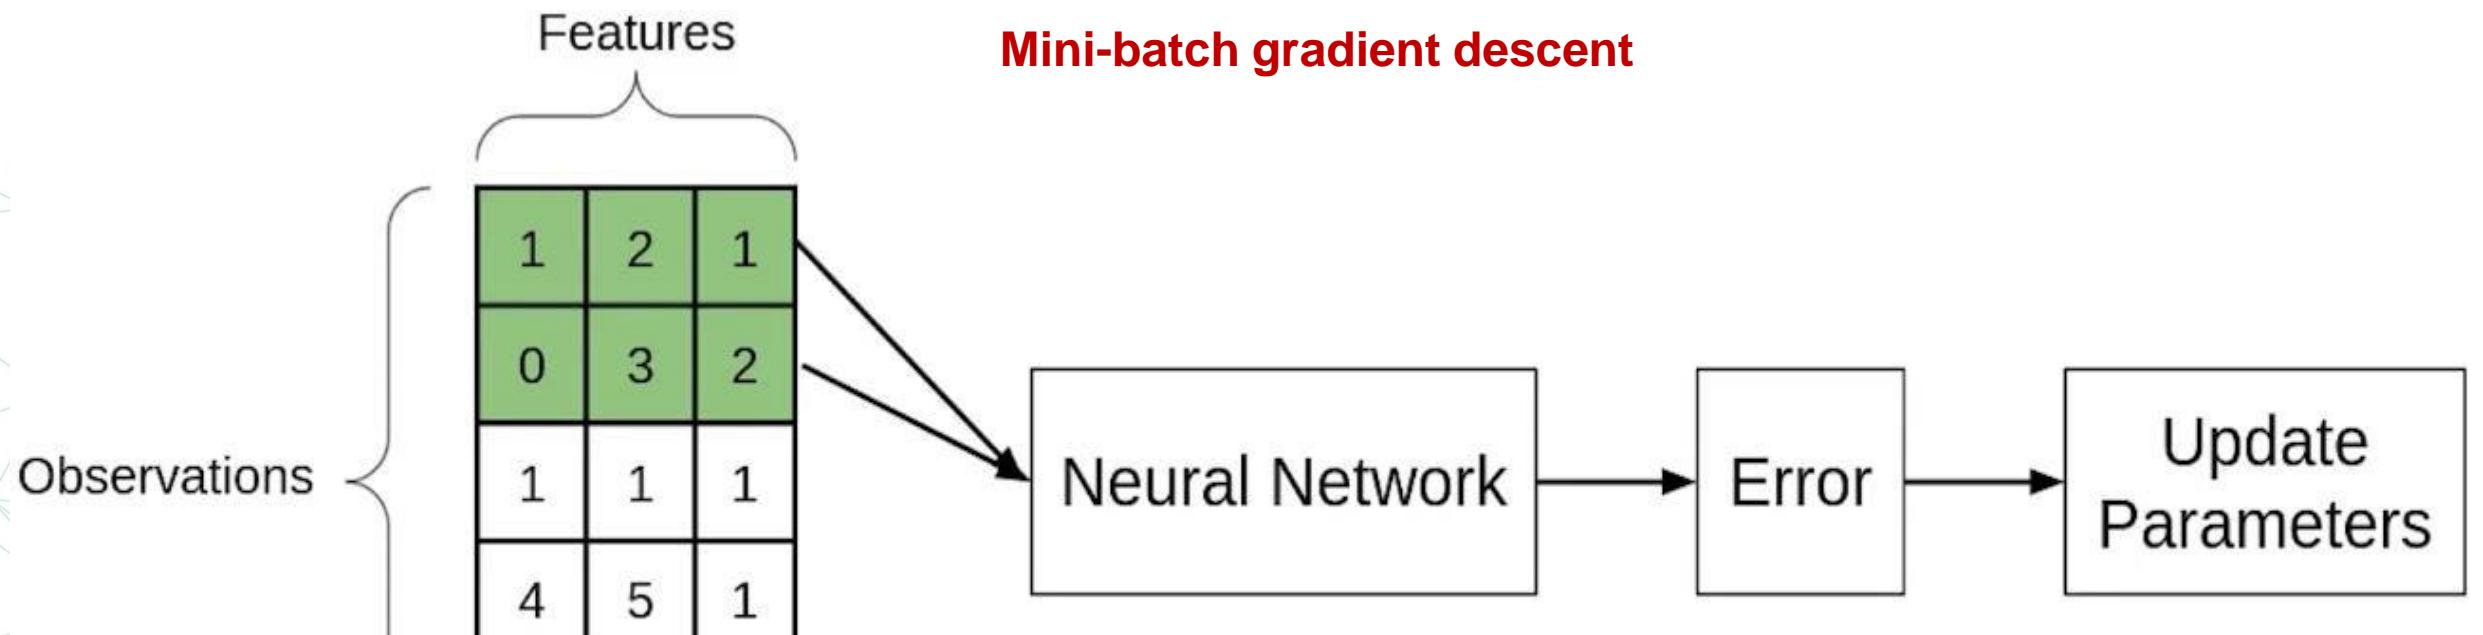


Variants of Gradient Descent Algorithm

Stochastic gradient descent



Variants of Gradient Descent Algorithm





Variants of Gradient Descent Algorithm

Batch Gradient Descent

- Entire dataset for updation
- Cost function reduces smoothly
- Computation cost is very high

Stochastic Gradient Descent (SGD)

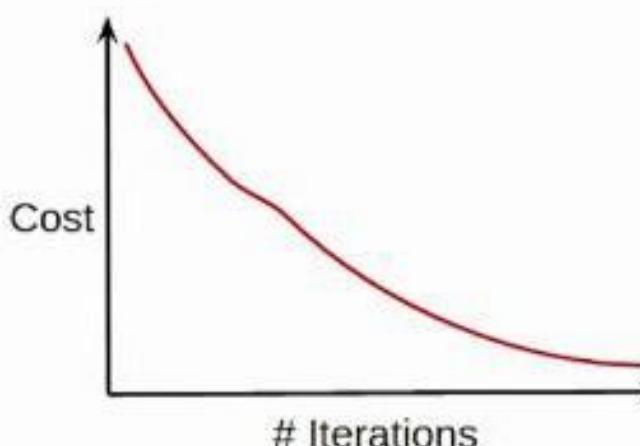
- Single observation for updation
- Lot of variations in cost function
- Computation time is more

Mini-Batch Gradient Descent

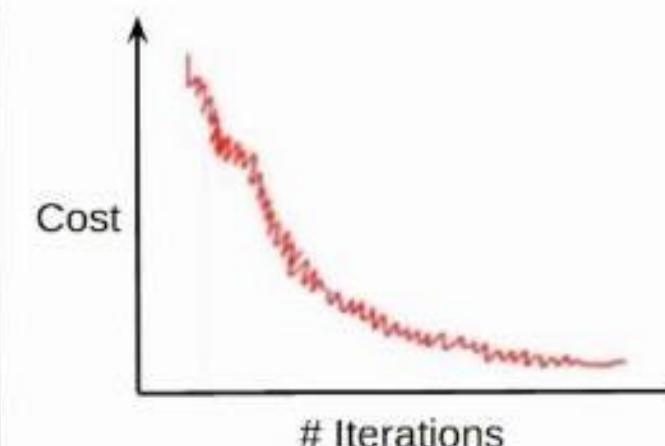
- Subset of data for updation
- Smoother cost function as compared to SGD
- Computation time is lesser than SGD
- Computation cost is lesser than Batch Gradient Descent

Variants of Gradient Descent Algorithm

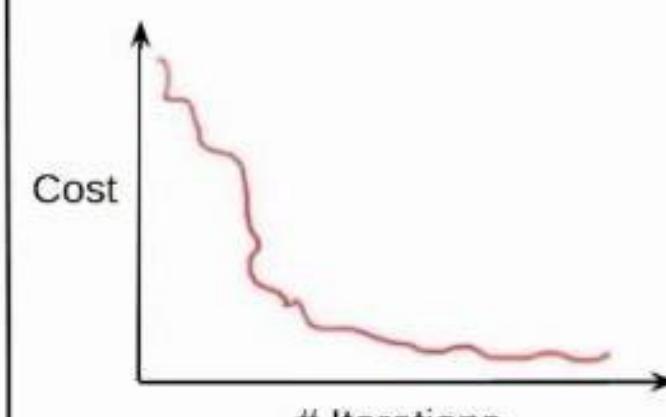
Batch Gradient Descent

- Cost function reduces smoothly
- 

Stochastic Gradient Descent (SGD)

- Lot of variations in cost function
- 

Mini-Batch Gradient Descent

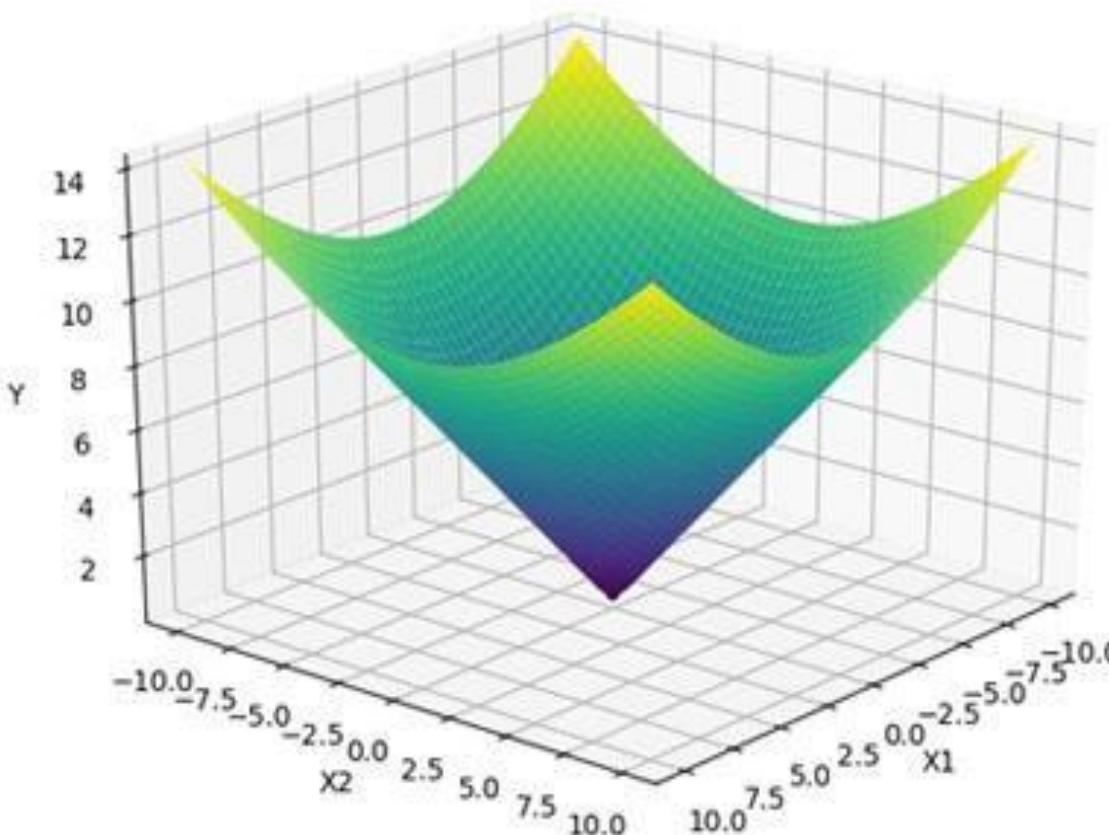
- Smoother cost function as compared to SGD
- 

Variants of Gradient Descent Algorithm

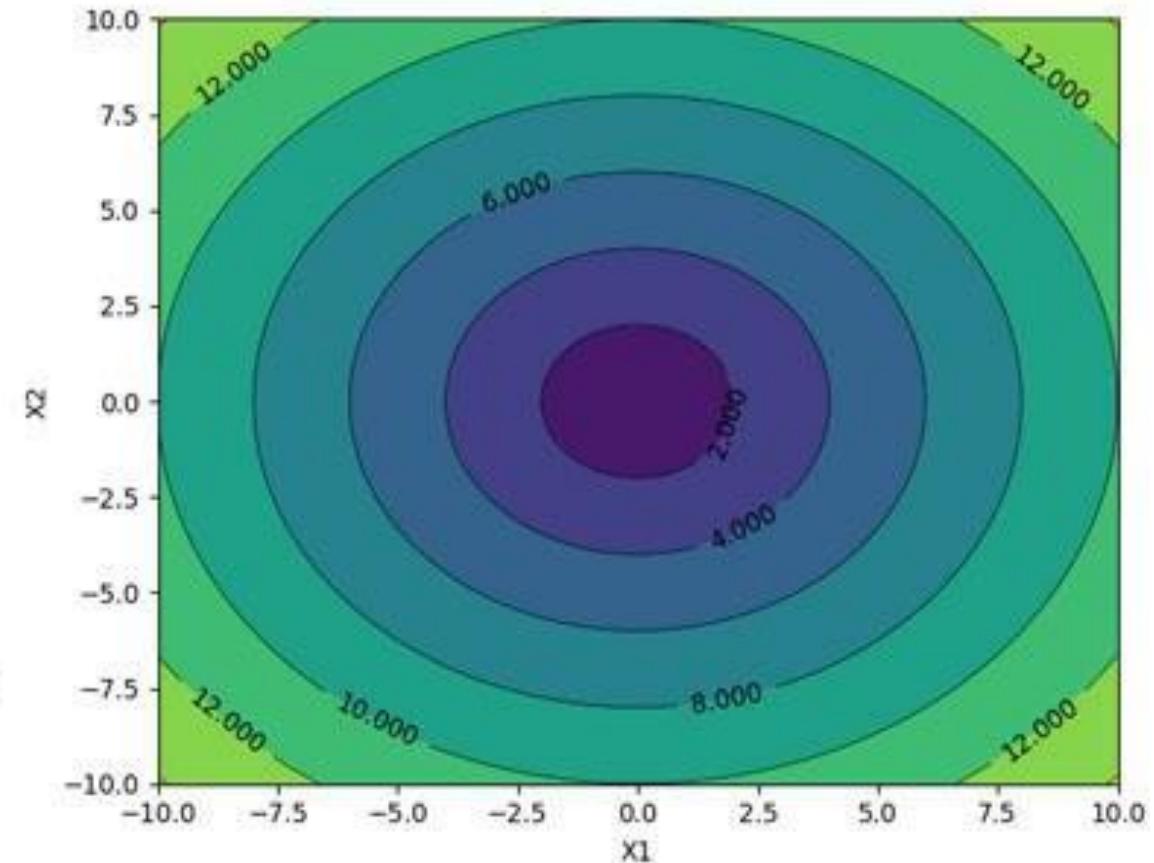




Contour Plot

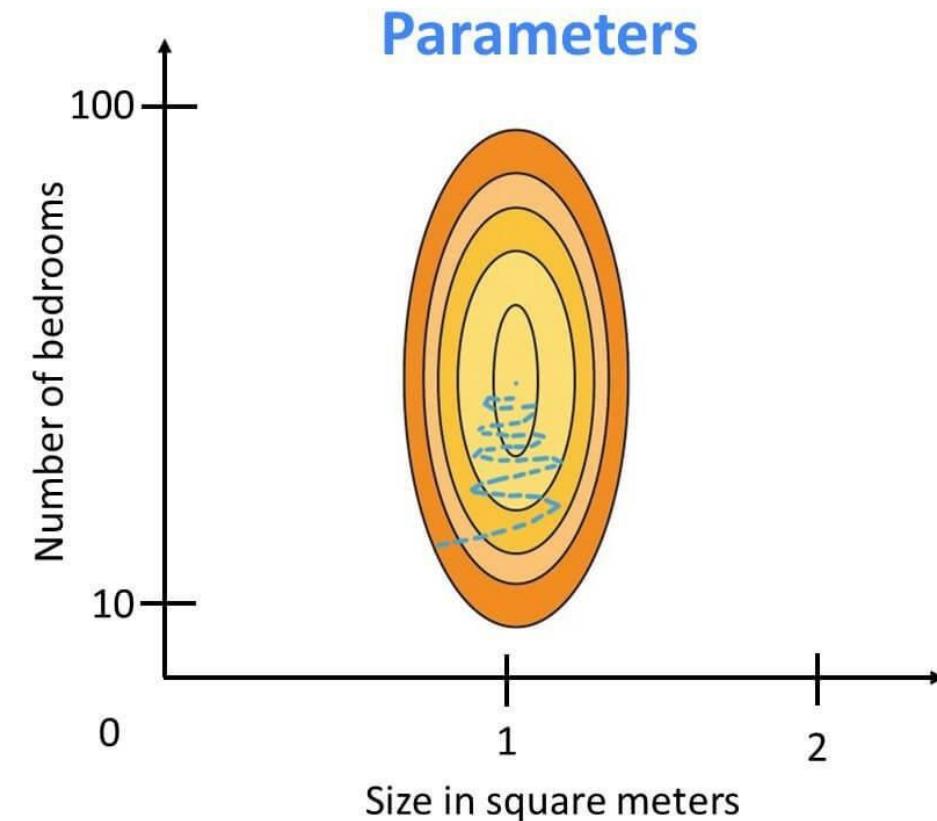
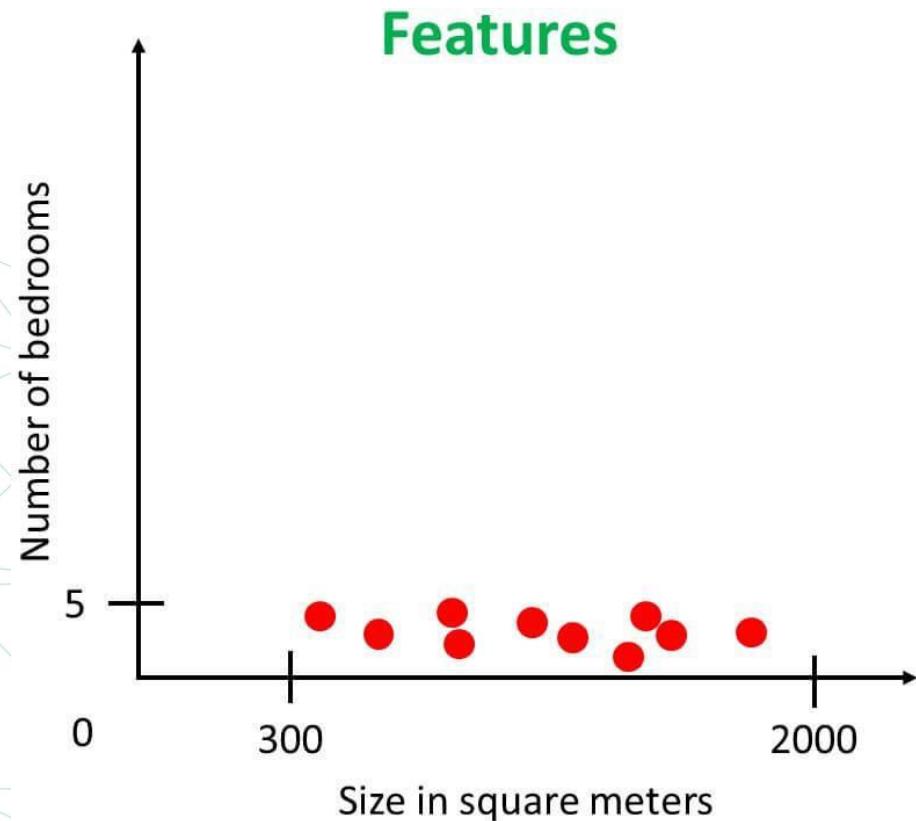


3D Plot

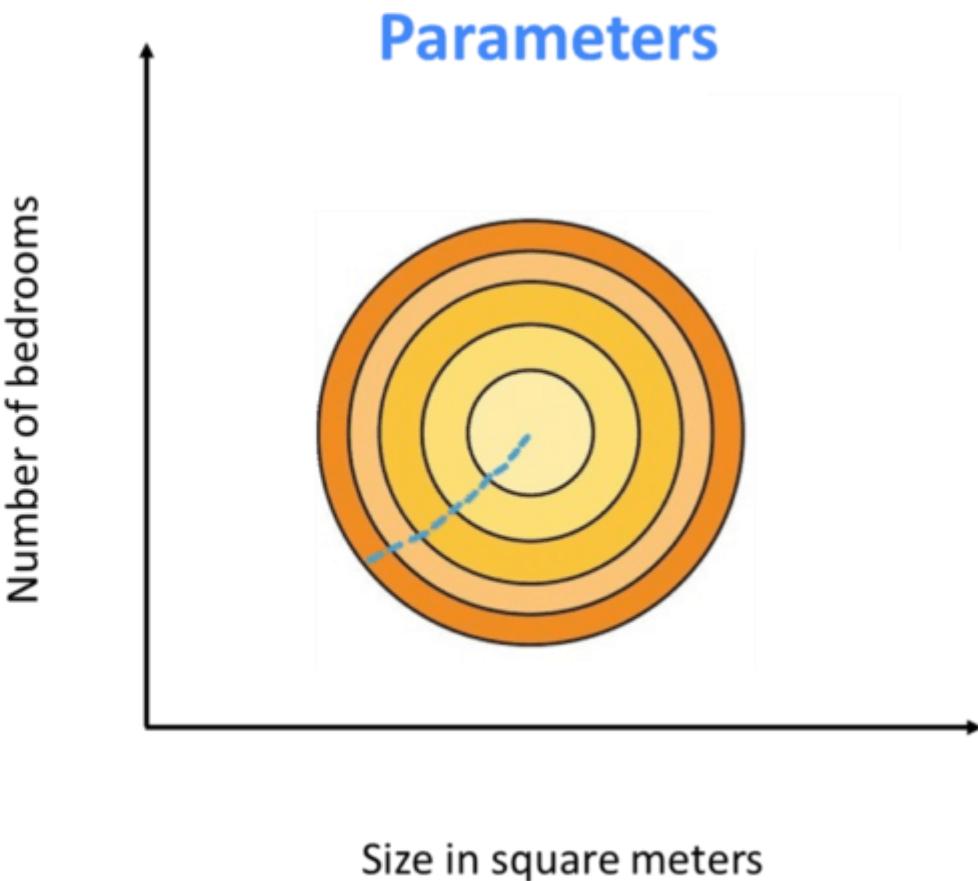
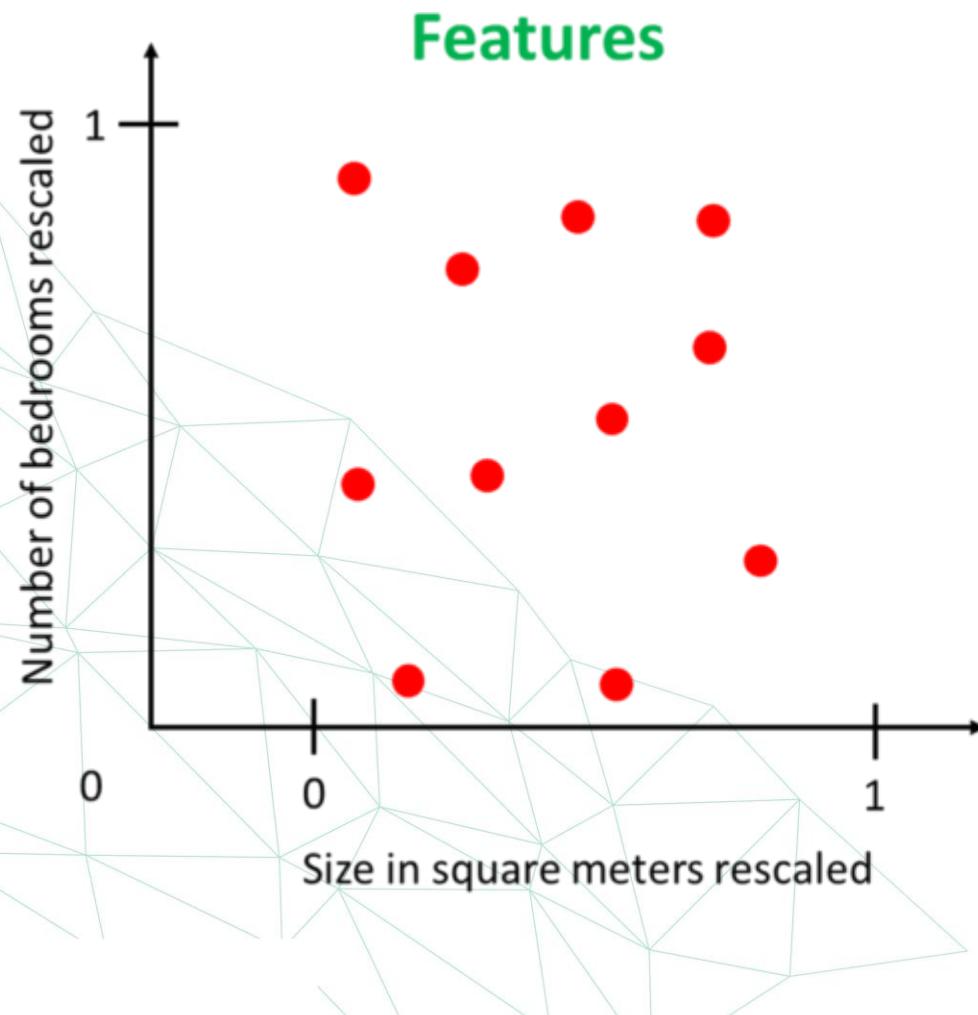


Contour Plot

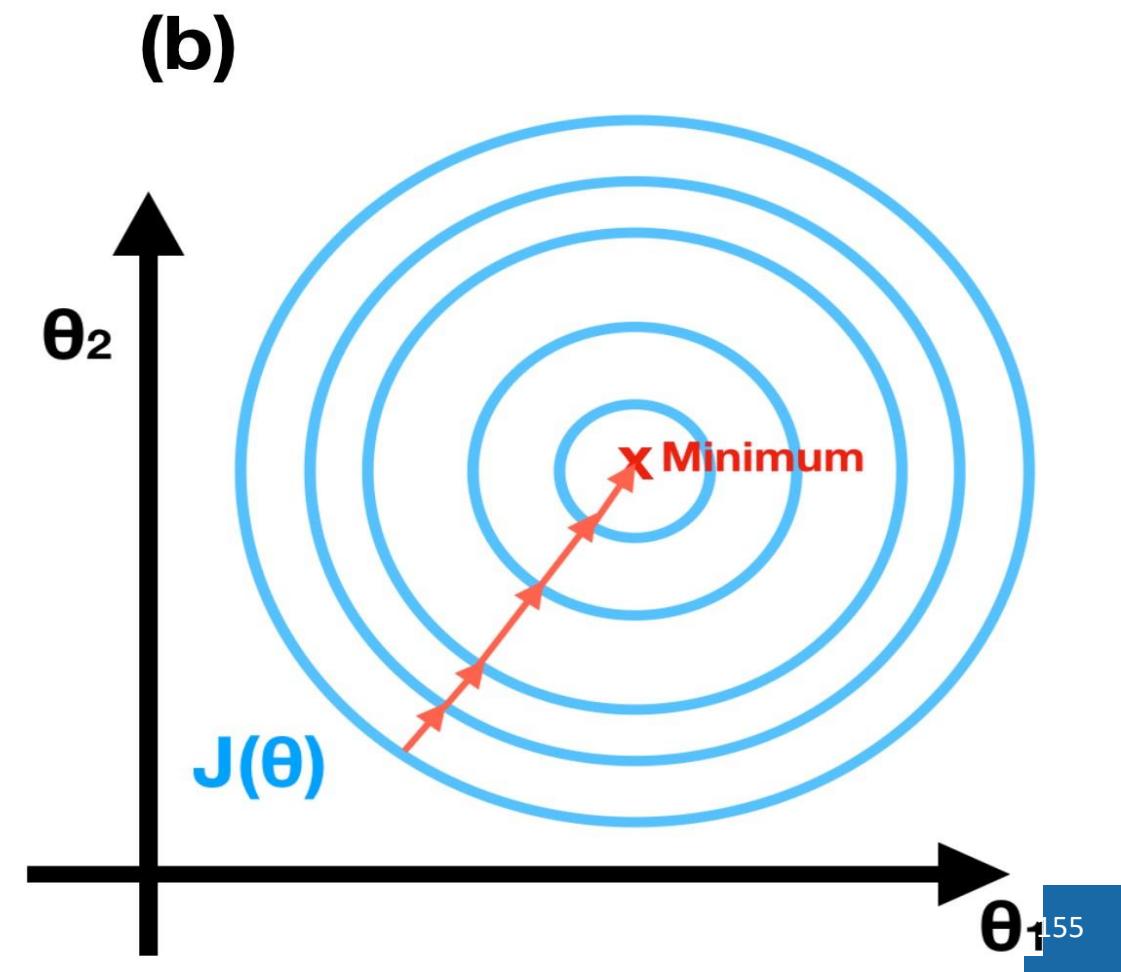
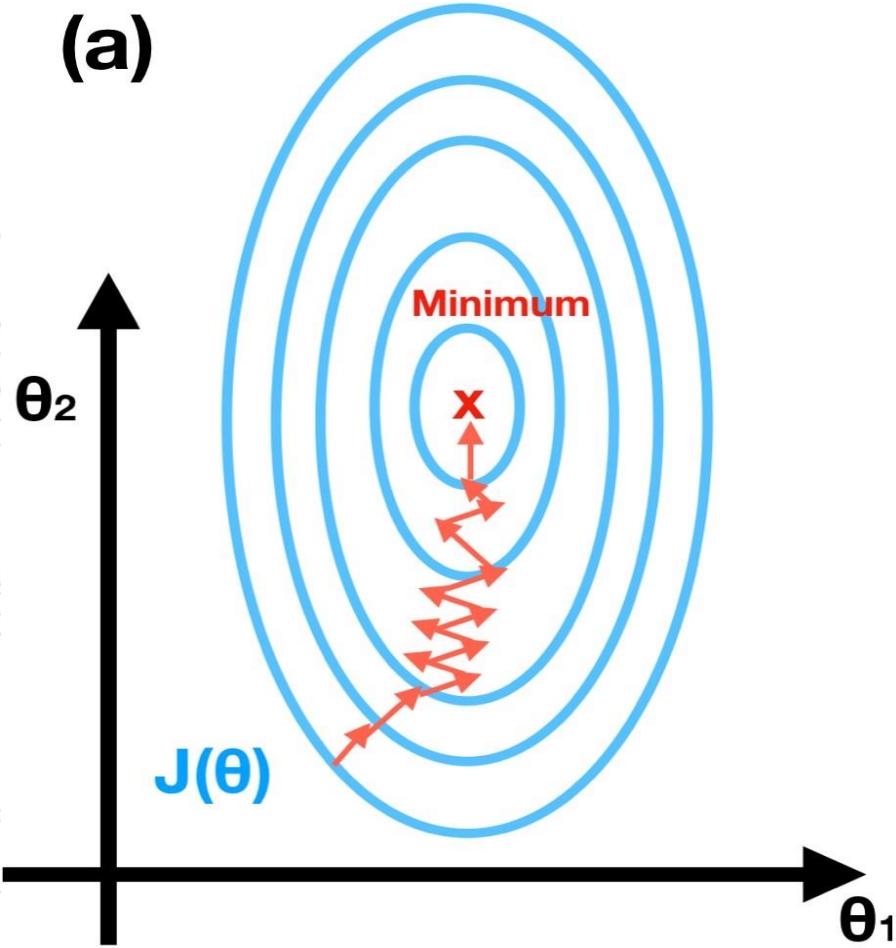
The Role of Feature Scaling in Gradient Descent



The Role of Feature Scaling in Gradient Descent



The Role of Feature Scaling in Gradient Descent



Simple
Linear
Regression

$$y = b_0 + b_1 x_1$$

Multiple
Linear
Regression

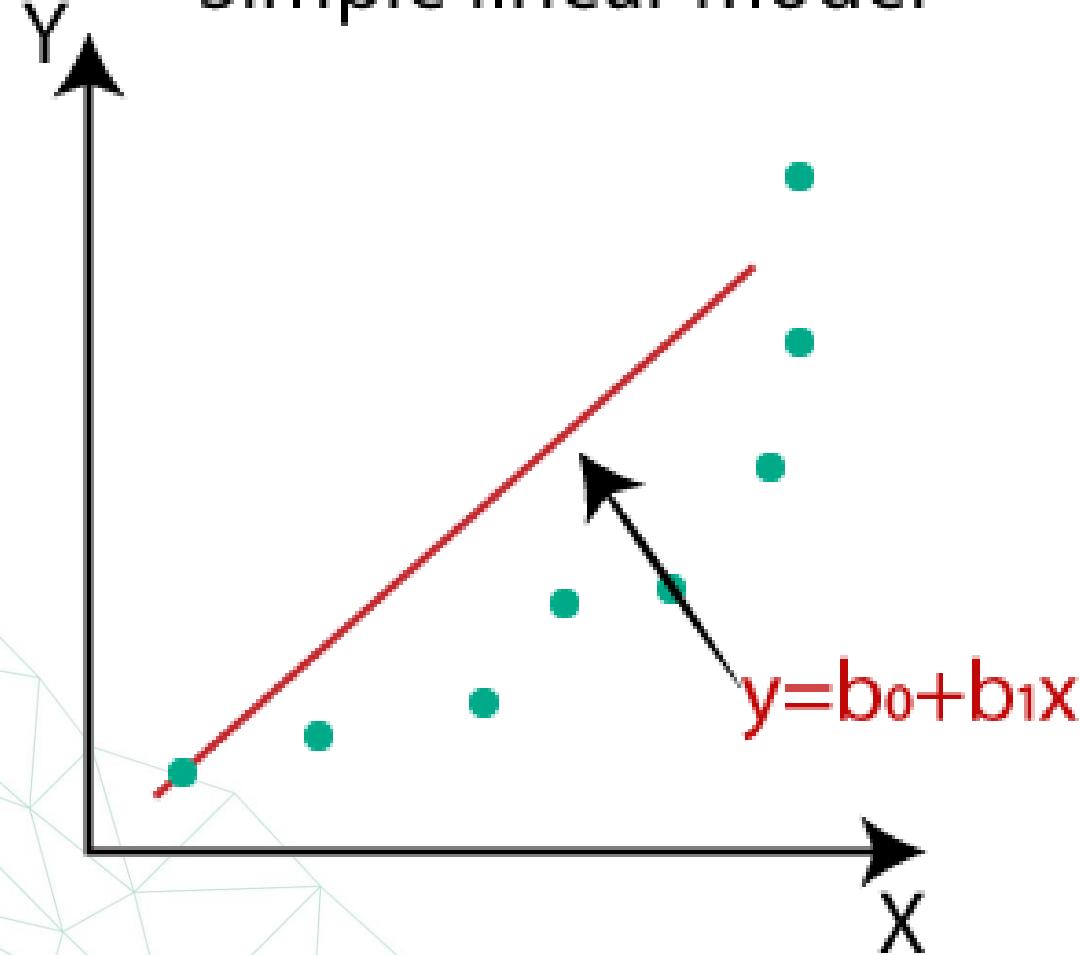
$$y = b_0 + b_1 x_1 + b_2 x_2 + \dots + b_n x_n$$

Polynomial
Linear
Regression

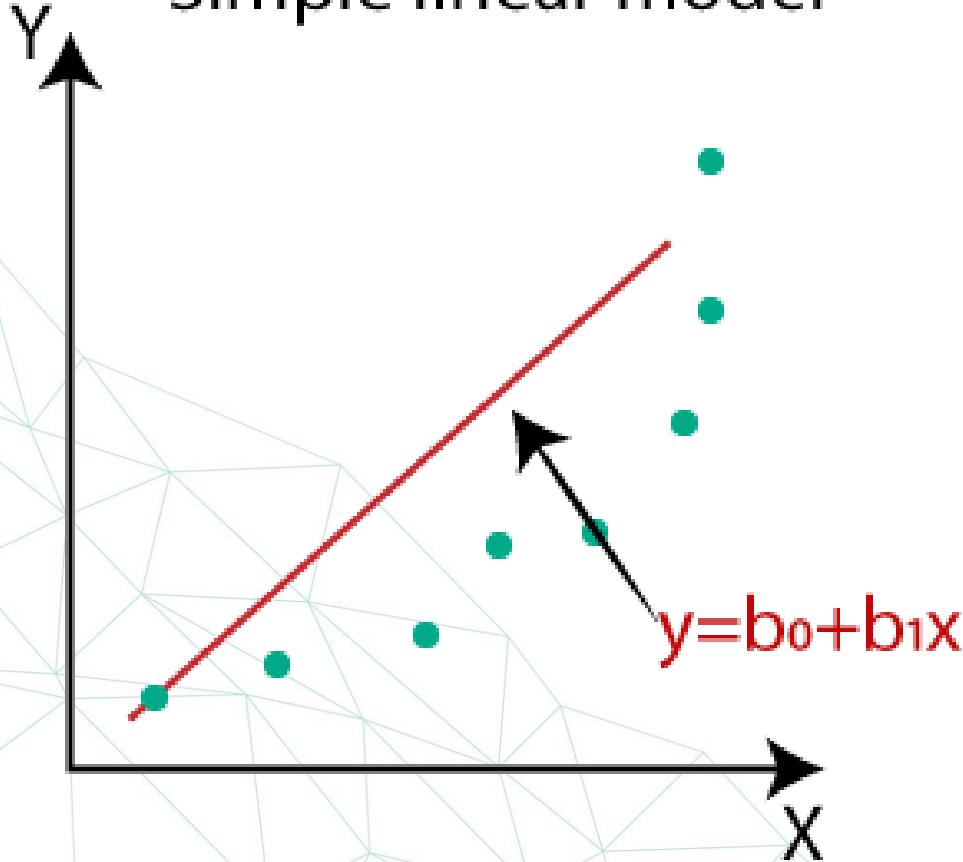
$$y = b_0 + b_1 x_1 + b_2 x_1^2 + \dots + b_n x_1^n$$

<https://data36.com/polynomial-regression-python-scikit-learn/>

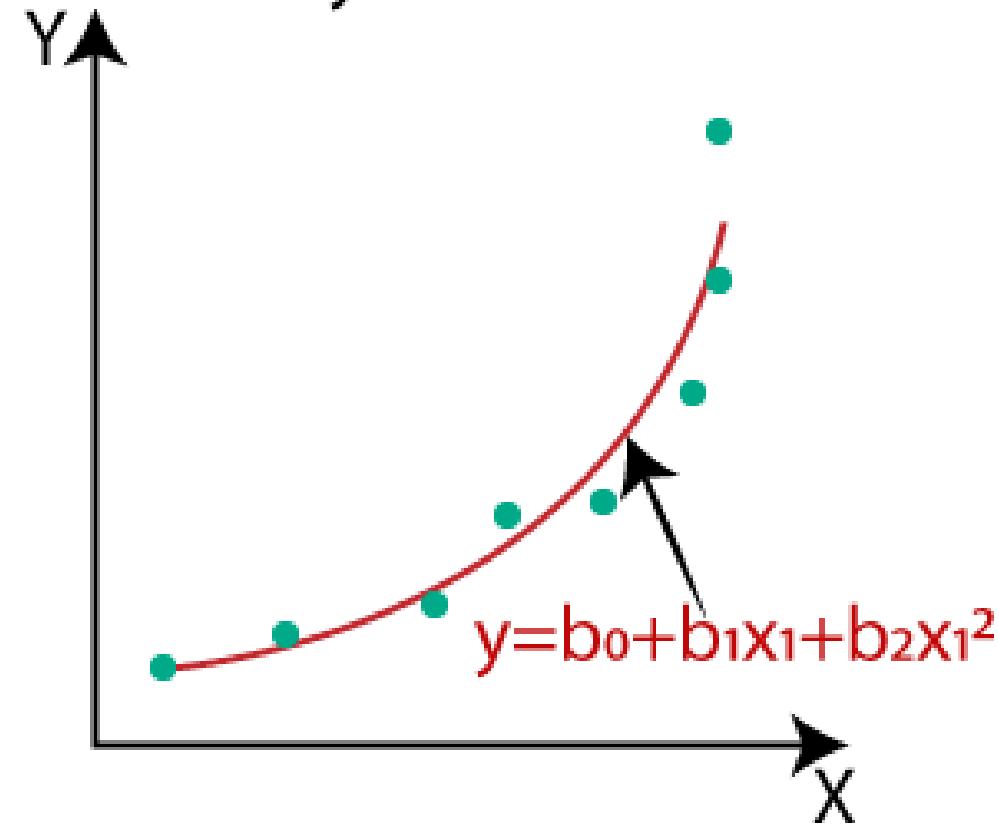
Simple linear model



Simple linear model



Polynomial model

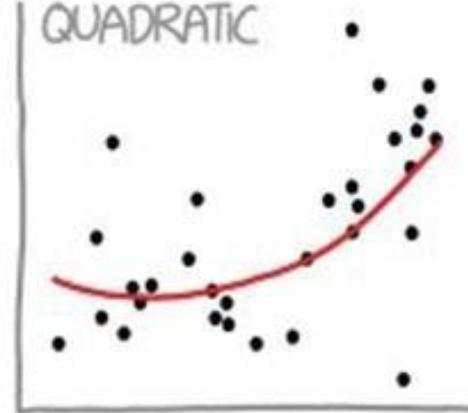




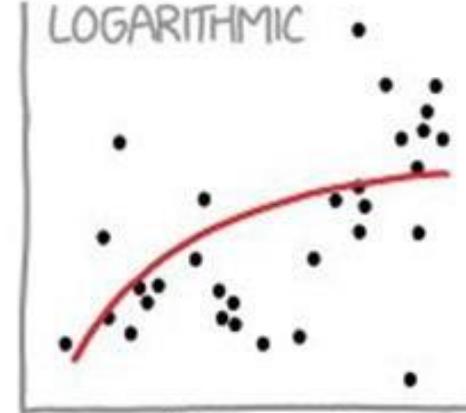
CURVE-FITTING METHODS AND THE MESSAGES THEY SEND



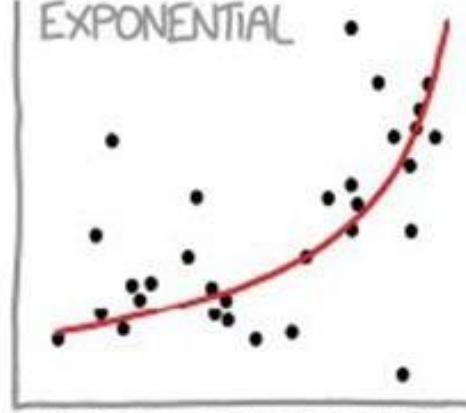
"HEY, I DID A
REGRESSION."



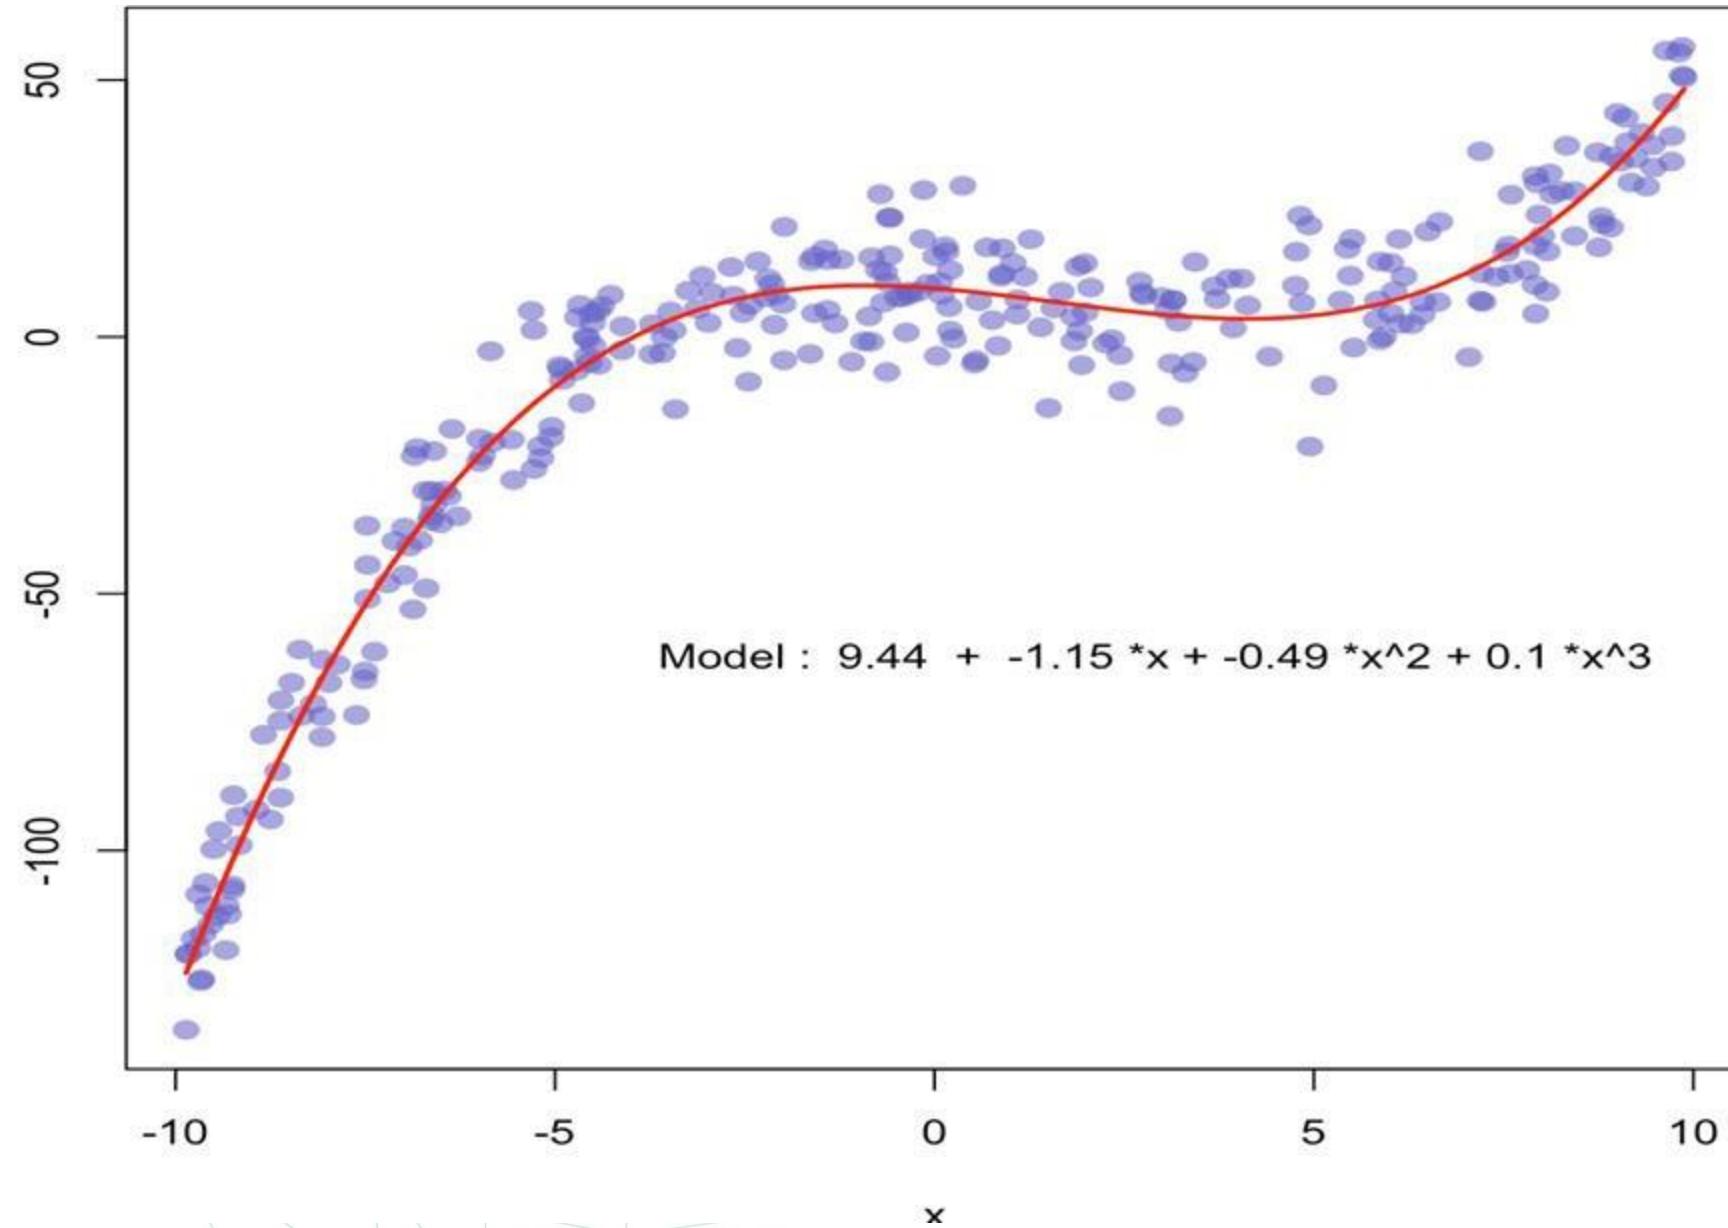
"I WANTED A CURVED
LINE, SO I MADE ONE
WITH MATH."



"LOOK, IT'S
TAPERING OFF!"



"LOOK, IT'S GROWING
UNCONTROLLABLY!"



1. What Are the Basic Assumption?

There are assumptions associated with a linear regression model:

1. Linearity: The relationship between X and the mean of Y is linear.
 2. Homoscedasticity: The variance of residual is the same for any value of X.
 3. Independence: Observations are independent of each other.
- Normality: For any fixed value of X, Y is normally distributed.
log normal transformation, Box-Cox transformation, etc

2. Advantages

1. Linear regression performs exceptionally well for linearly separable data
2. Easy to implement and train the model
3. It can handle overfitting using dimensionality reduction techniques and cross validation and regularization

3. Disadvantages

1. Sometimes Lot of Feature Engineering Is required
2. If the independent features are correlated it may affect performance
3. It is often quite prone to noise and overfitting

4. Whether Feature Scaling is required?

Yes

5. Impact of Missing Values?

It is sensitive to missing values

6. Impact of outliers?

linear regression needs the relationship between the independent and dependent variables to be linear. It is also important to check for outliers since linear regression is sensitive to outlier effects.

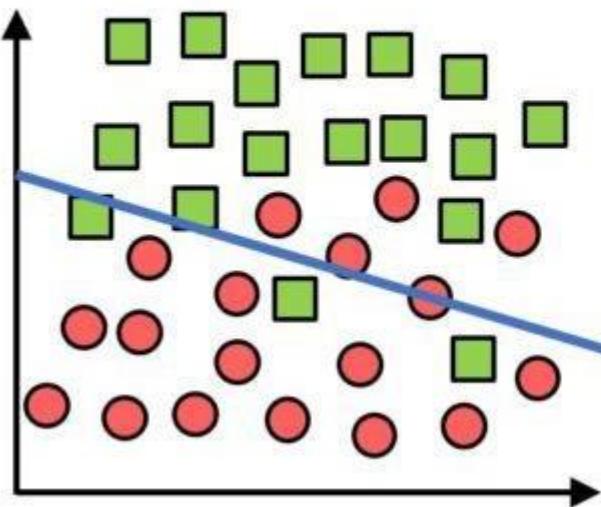
Practical Lab

House Price Prediction Dataset

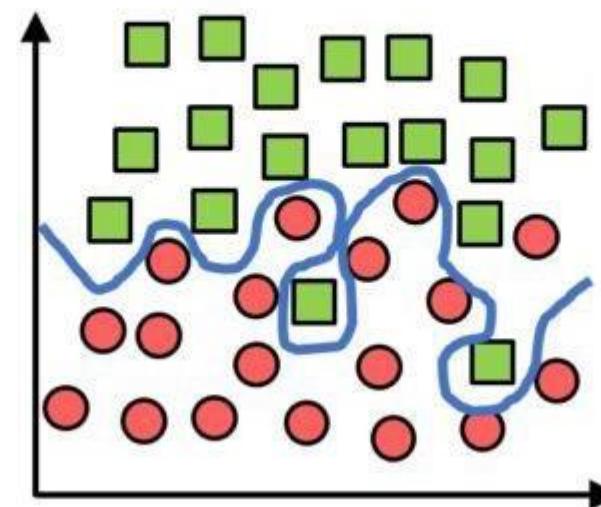
ML Concepts

Generalization, Overfitting, and Underfitting

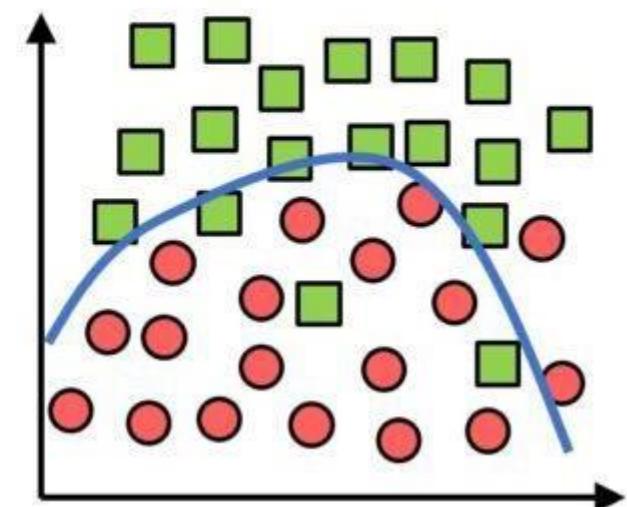
Underfitting



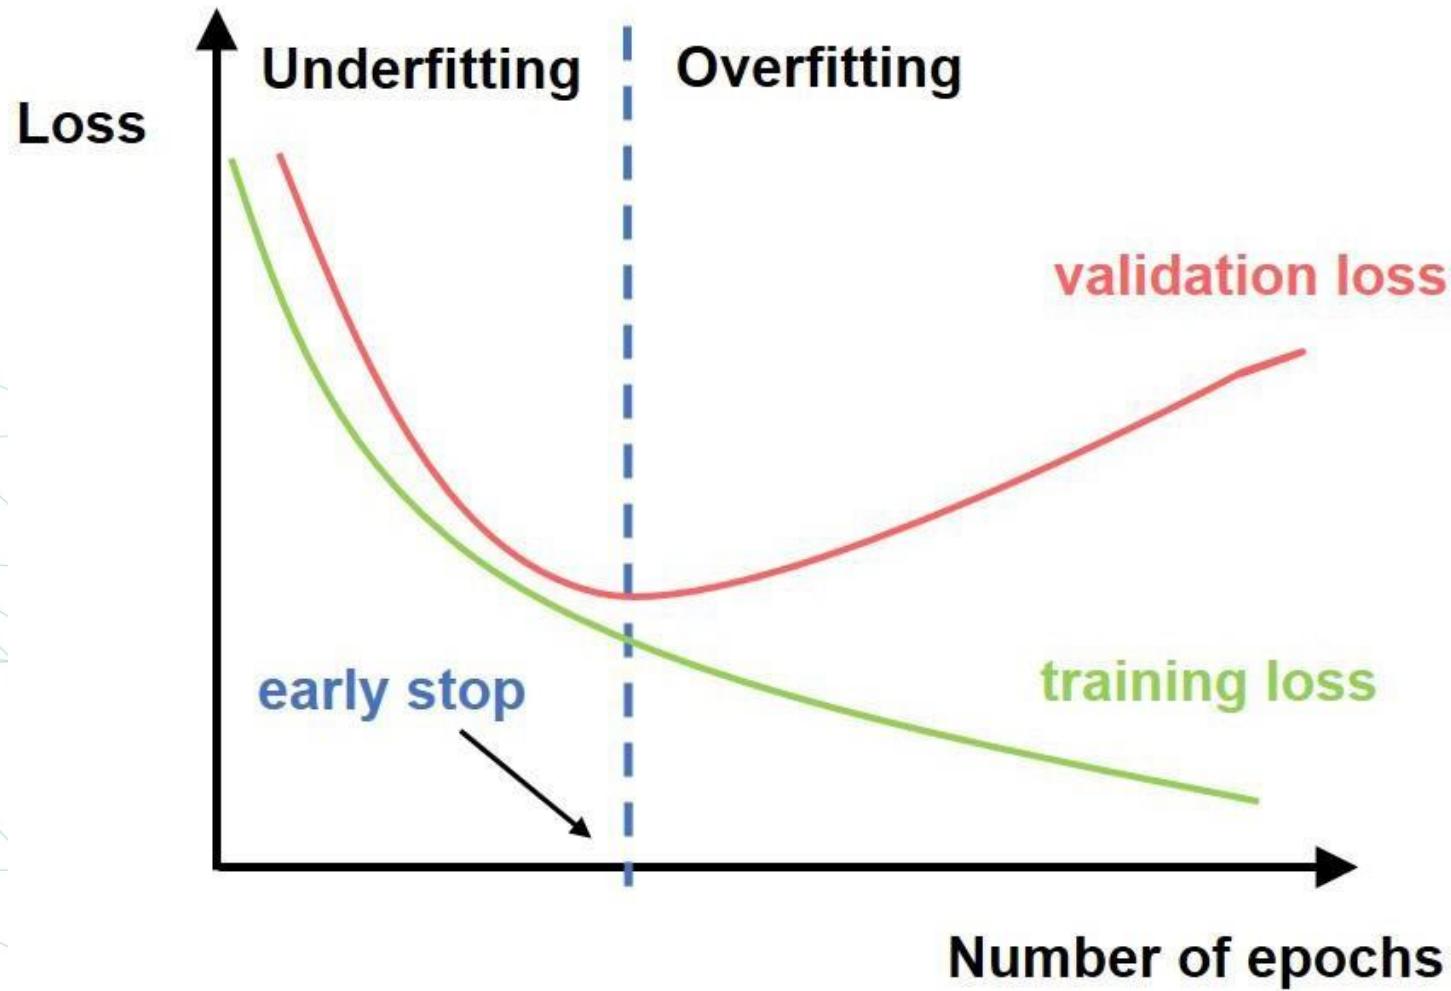
Overfitting

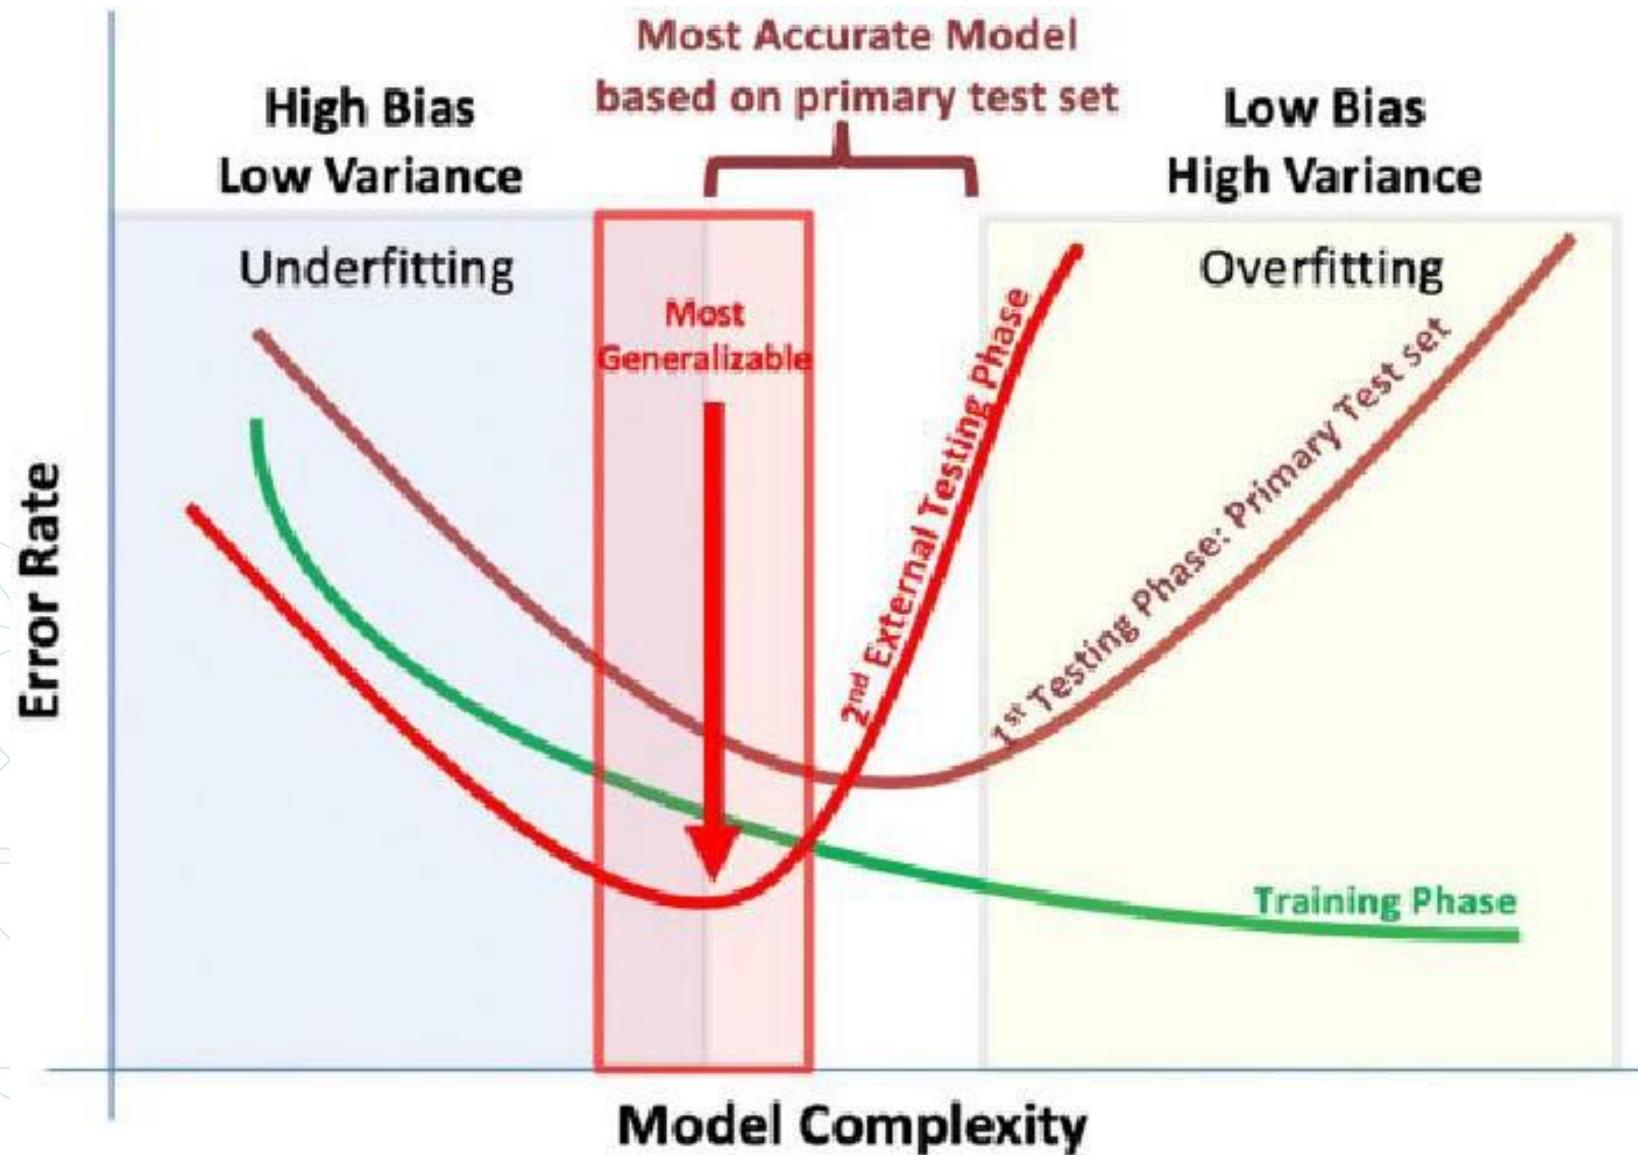


Optimal



Generalization, Overfitting, and Underfitting





How to avoid Overfitting

- Simplifying The Model
- Early Stopping
- Adding More Data To The Training Set
- Cross-validation
- Remove unnecessary features
- **Regularization**

Regularisation Techniques

**Handle Overfitting
With
Regularization**



L1 Regularization

- This adds a penalty equal to the **L1 norm** of the weights **vector(sum of the absolute value of the coefficients)**. It will shrink some parameters to **zero**.
- Hence some variables will not play any role in the model.
- L1 regression can be seen as a way to **select features** in a model.

$$\text{L1: } R(\theta) = \|\theta\|_1 = \sum_{i=1}^n |\theta_i|$$

$$\text{LossFunction} = \frac{1}{N} \sum_{i=1}^N (\hat{Y} - Y)^2 + \lambda \sum_{i=1}^N |\theta_i|$$



L2 Regularization

- This adds a **penalty** equal to the **L2 norm** of the weights vector(sum of the squared values of the coefficients).
- It will force the parameters to be **relatively small**.

$$L1: \quad R(\theta) = \|\theta\|_1 = \sum_{i=1}^n |\theta_i|$$

$$L2: \quad R(\theta) = \|\theta\|_2^2 = \sum_{i=1}^n \theta_i^2$$

$$LossFunction = \frac{1}{N} \sum_{i=1}^N (\hat{Y} - Y)^2 + \lambda \sum_{i=1}^N \theta_i^2$$

$$J_{\text{test}}(\theta) = \frac{1}{2m_{\text{test}}} \sum_{i=1}^{m_{\text{test}}} (\hat{y}_i - y_i)^2$$

where

- $h_{\theta}(x(i))$ is the predicted value of some datapoint $x(i)$
- $y(i)$ is original

The penalized cost function looks like this

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right)^2 + \lambda \sum_{j=1}^n \theta_j^2$$

Regularization Term

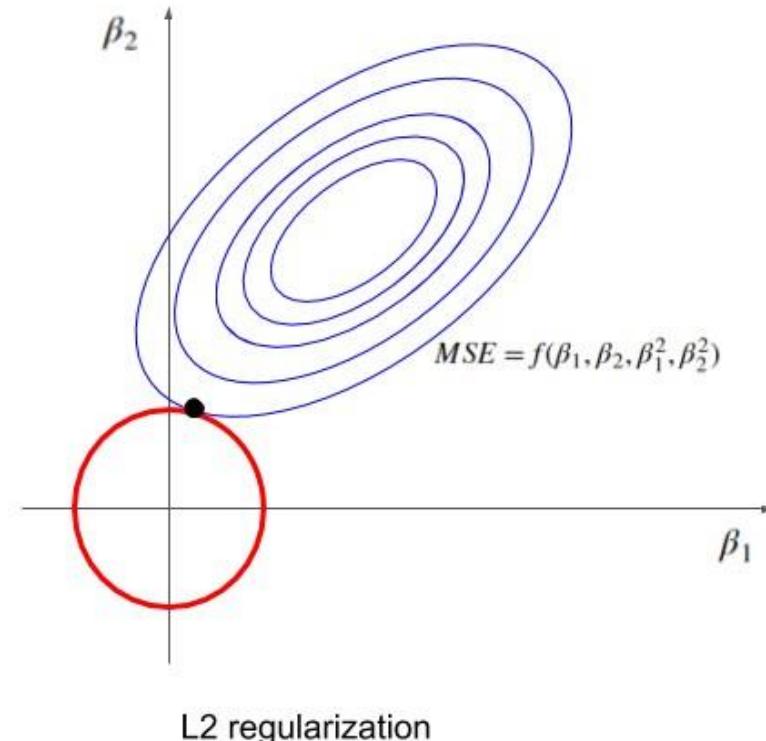
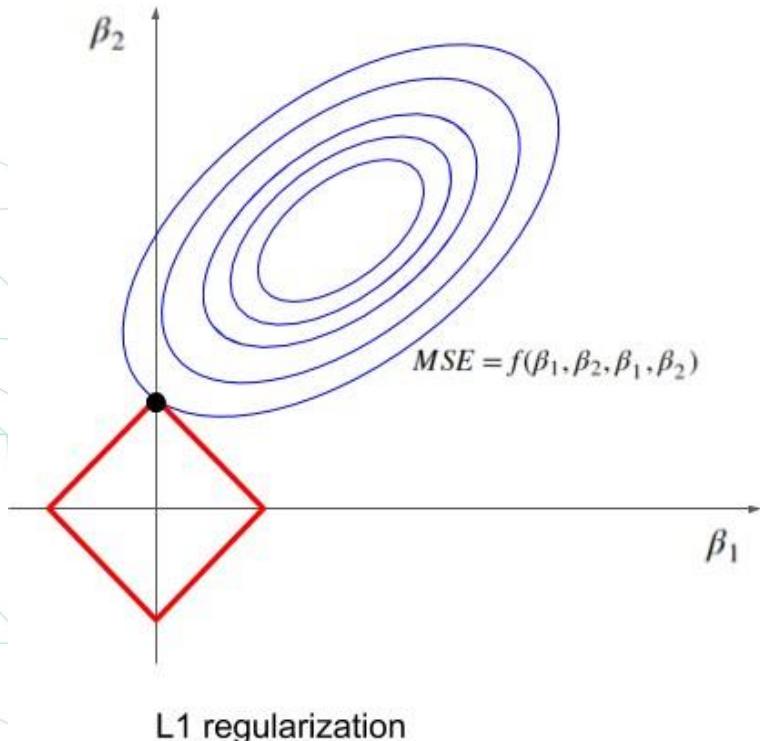
Regularization Parameter

where-

- λ is the tuning parameter that decides how much we want to penalize the flexibility of our model. It can be tuned using cross-validation.

The difference is that: while shrinking the quota, **L1 tends to cut off some factors by turning their coefficients to zero**, while **L2 tends to shrink these coefficients to a tiny number (none zero)**, keep some of their influence on Y.

Find β_1, β_2 to minimize MSE with restriction on $\beta_1\beta_2$





Comparison of L1 and L2 regularization

<i>L1 regularization</i>	<i>L2 regularization</i>
Sum of absolute value of weights	Sum of square of weights
Sparse solution	Non-sparse solution
Multiple solutions	One solution
Built-in feature selection	No feature selection
Robust to outliers	Not robust to outliers (due to the square term)

By this I mean the number of solutions to arrive at one point. L1 regularization uses Manhattan distances to arrive at a single point, so there are many routes that can be taken to arrive at a point. L2 regularization uses Euclidean distances, which will tell you the fastest way to get to a point. This means the L2 norm only has 1 possible solution.

Which solution is less Computationally expensive? L2

- Regularization Lab

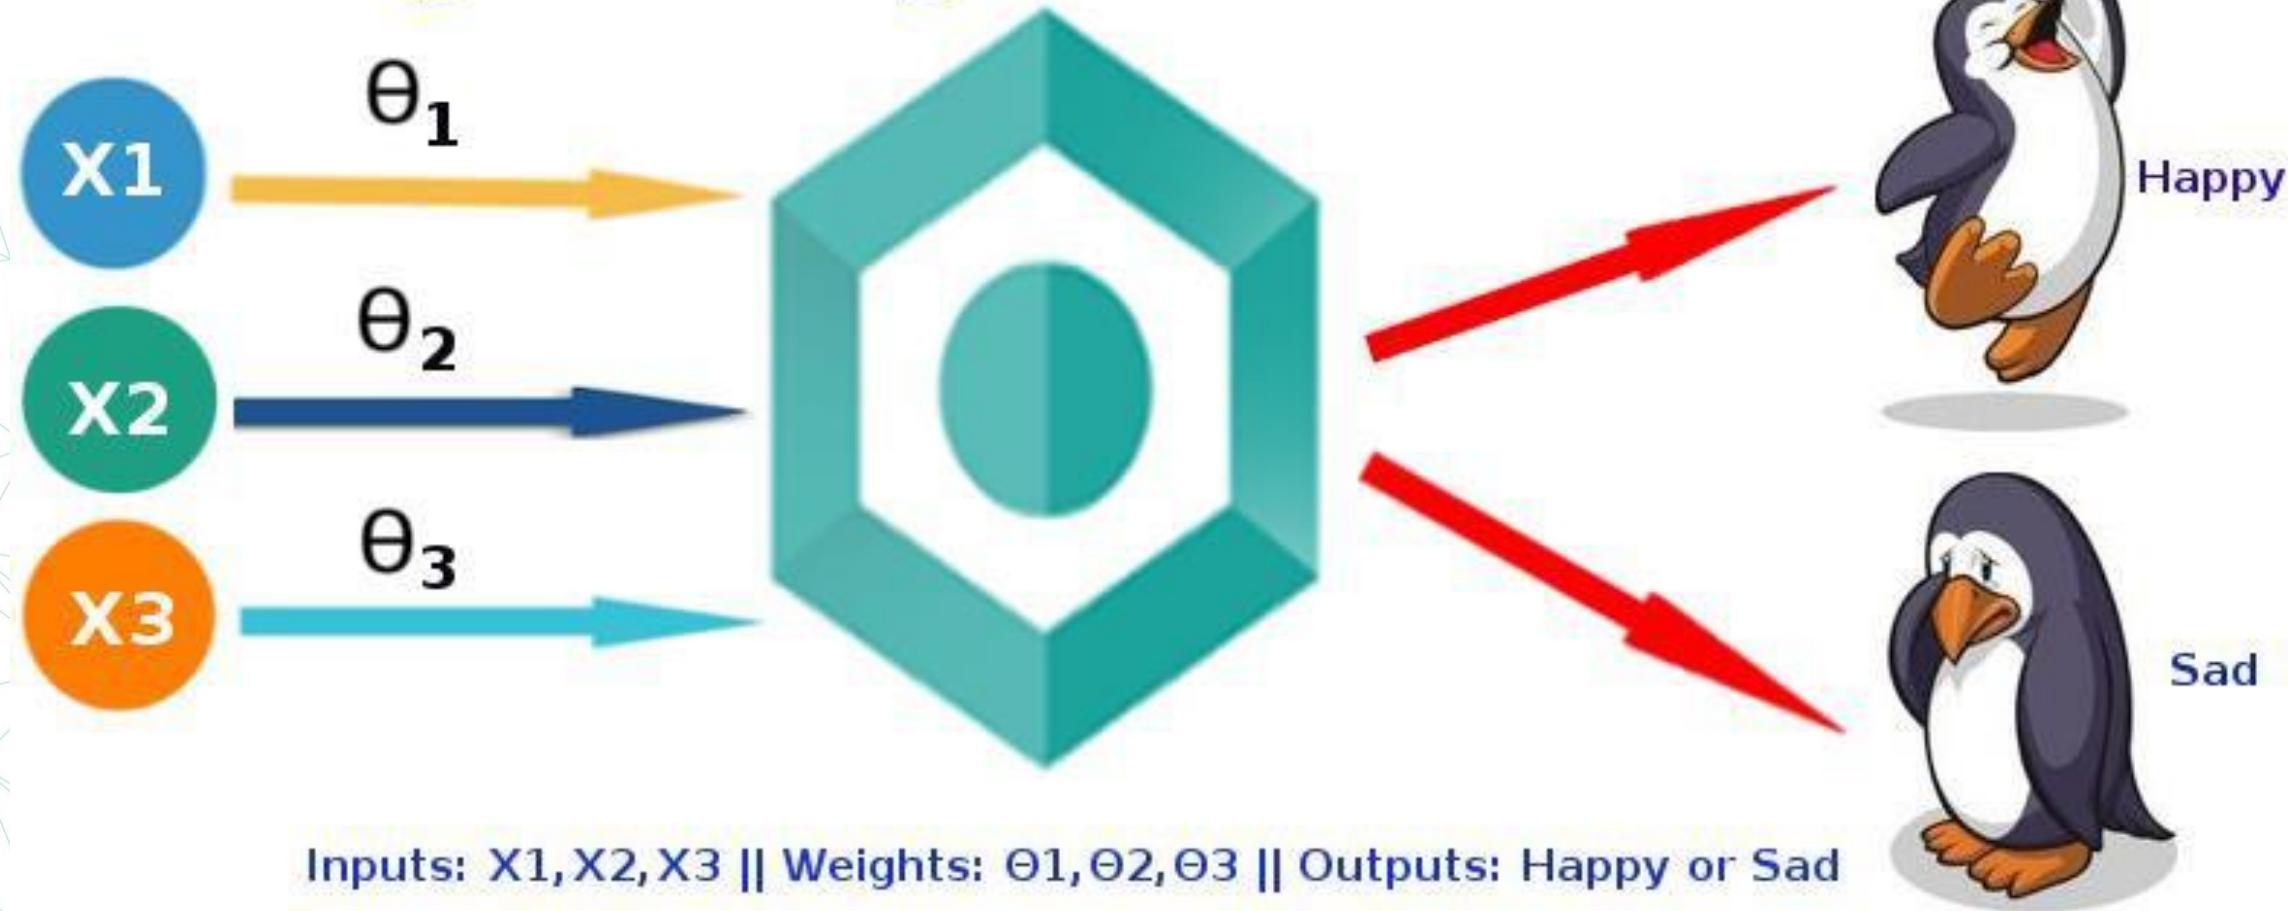
Q&A

Questions and answers

Logistic Regression

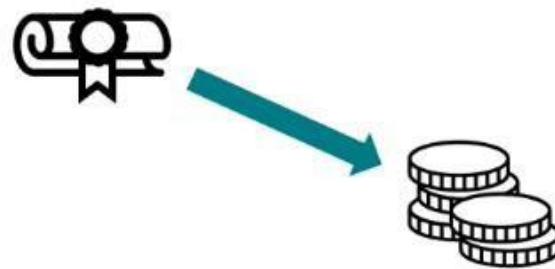


Logistic Regression Model

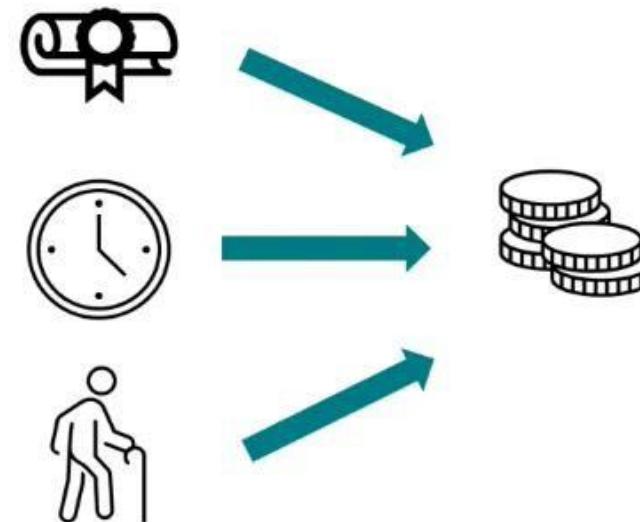


What is Logistic Regression?

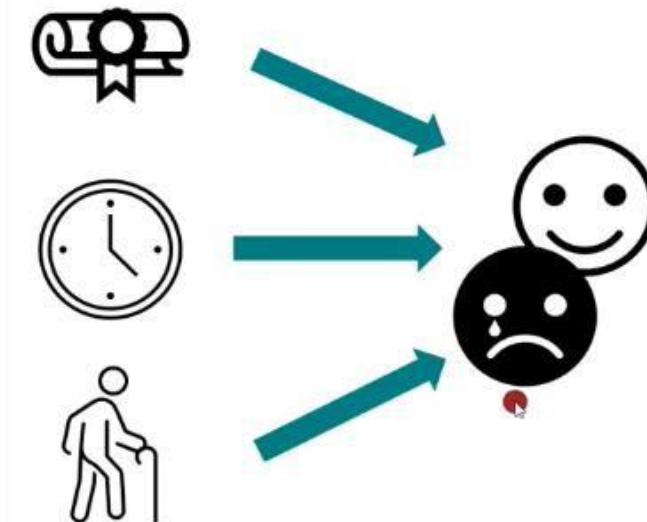
Simple linear regression



Multiple linear regression



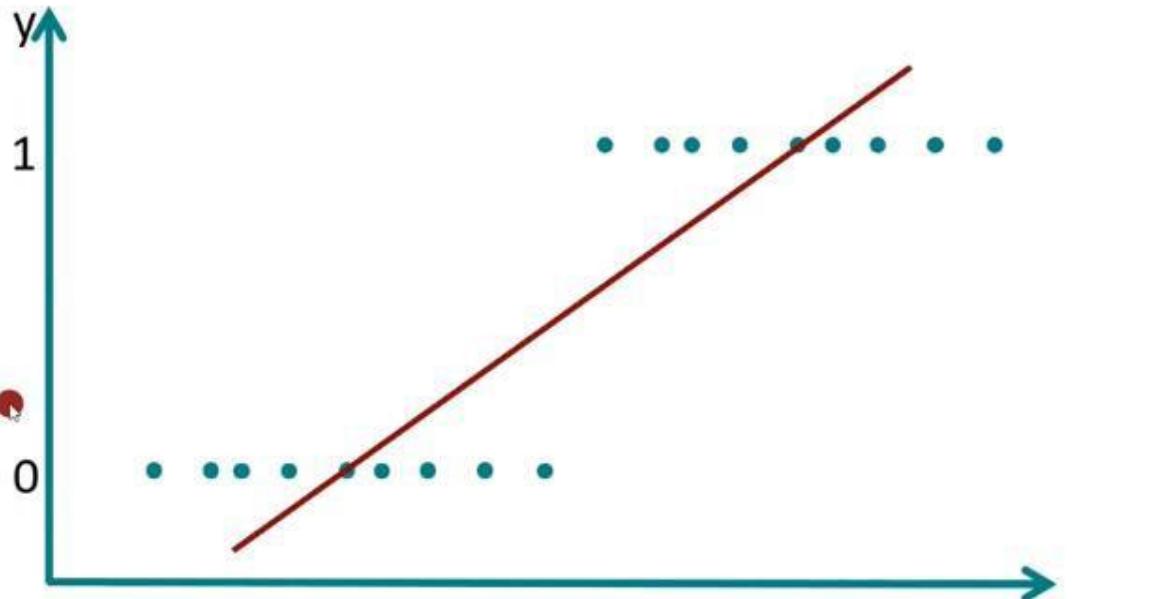
Logistic regression





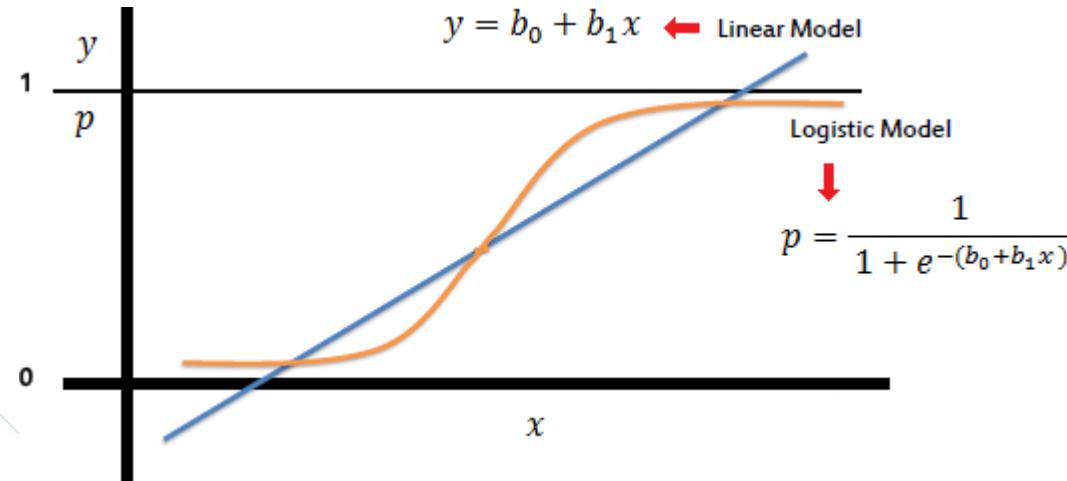
Why not just linear regression?

$$\hat{y} = b_1 \cdot x_1 + b_2 \cdot x_2 + \dots + b_k \cdot x_k + a$$



- The graph shows that values between **plus and minus infinity** can now occur.
- The goal of logistic regression is to estimate the **probability of occurrence**, not the value of the variable itself.
- The range of values for the prediction is restricted to the range between 0 and 1.
- To ensure that only values between 0 and 1 are possible, the logistic function f is used.

Logistic Regression



A linear regression will predict values outside the acceptable range (e.g. predicting probabilities outside the range 0 to 1)

Logistic Regression

In the logistic regression the constant (b_0) moves the curve left and right and the slope (b_1) defines the steepness of the curve. By simple transformation, the logistic regression equation can be written in terms of an odds ratio.

$$\frac{p}{1-p} = \exp(b_0 + b_1 x)$$

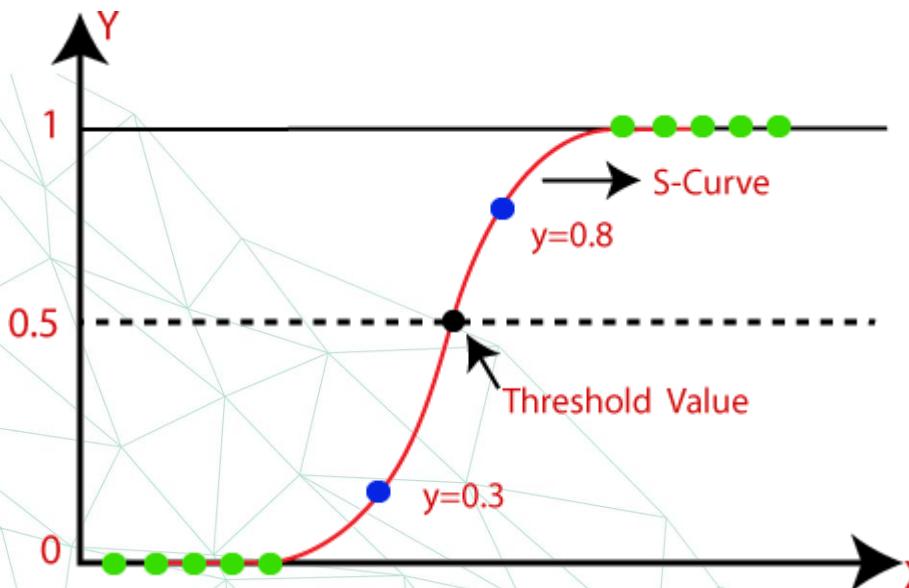
Finally, taking the natural log of both sides, we can write the equation in terms of log-odds (logit) which is a linear function of the predictors. The coefficient (b_1) is the amount the logit (log-odds) changes with a one unit change in x .

$$\ln\left(\frac{p}{1-p}\right) = b_0 + b_1 x$$

As mentioned before, logistic regression can handle any number of numerical and/or categorical variables.

$$p = \frac{1}{1 + e^{-(b_0 + b_1 x_1 + b_2 x_2 + \dots + b_p x_p)}}$$

Logistic Regression



$$y' = b_0 + b_1x_1 + b_2x_2$$

if p < 0.5 then N else Y

BUSAGE	DAYSDELQ	DEFAULT	y'	p	Prediction
87	2	N	-4.811	0.008	N
89	2	N	-4.795	0.008	N
100	26	Y	-2.261	0.094	N
275	54	Y	1.983	0.879	Y
42	57	Y	0.437	0.608	Y
88	53	N	0.395	0.597	Y

$$y' = -5.706 + 0.008 \times 87 + 0.102 \times 2 = -4.811$$

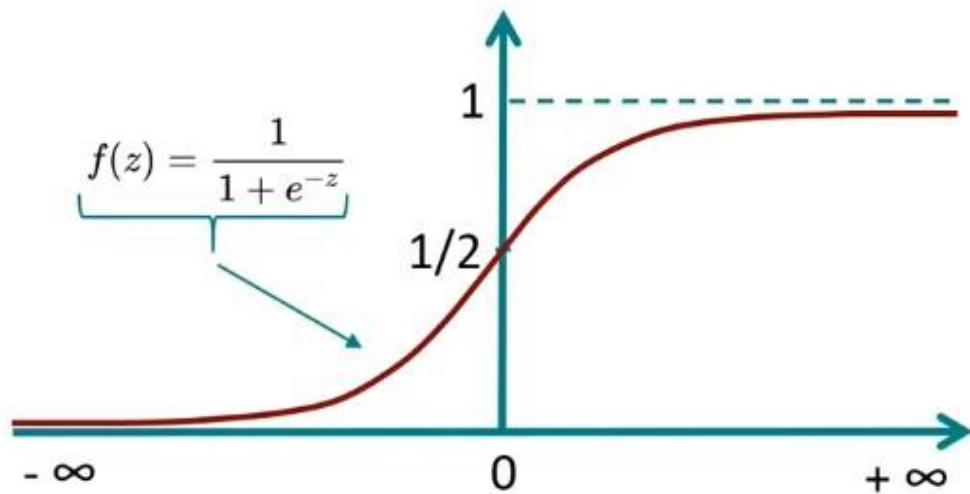
$$p = \frac{1}{1 + e^{-(b_0 + b_1x_1 + b_2x_2)}} = \frac{1}{1 + e^{-y'}}$$

$$p = \frac{1}{1 + e^{-4.811}} = 0.008 \quad \text{ $\leftarrow 0.5$ } \rightarrow \boxed{\text{N}}$$

What is Logistic Regression?

The logistic model is based on the logistic function.

The important thing about the logistic function is, that only values **between 0 and 1** are possible.

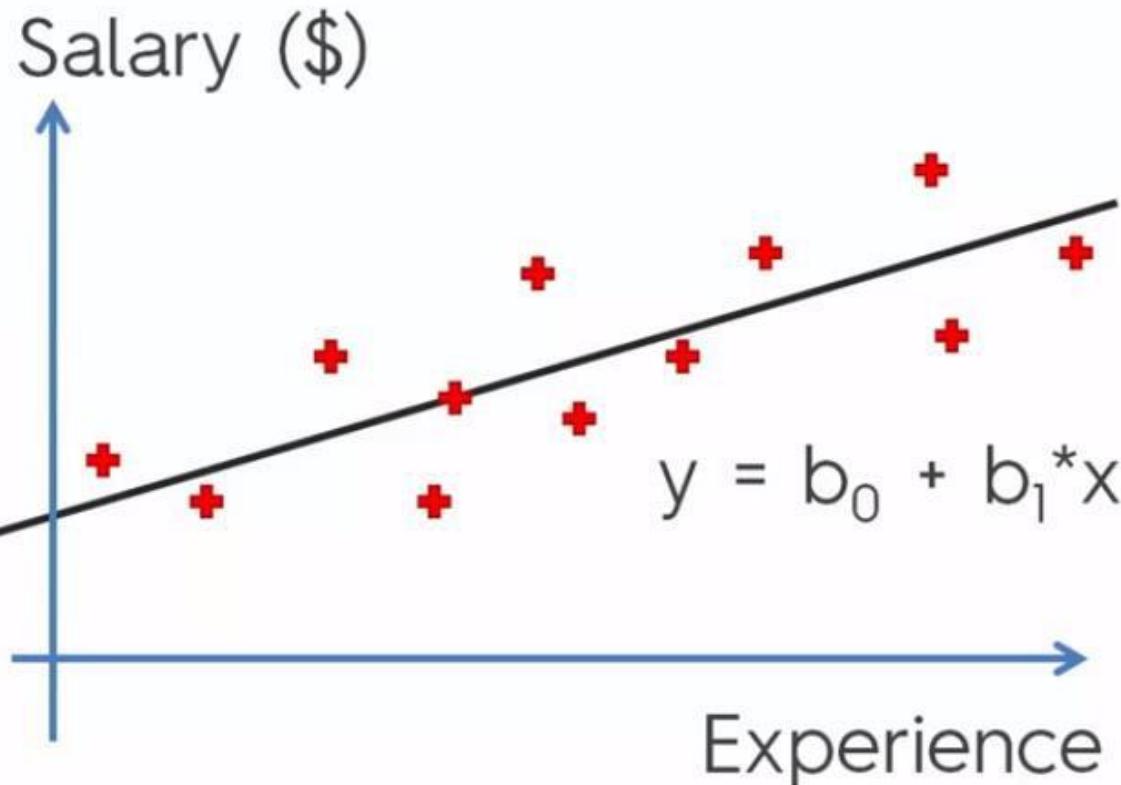


$$f(z) = \frac{1}{1 + e^{-z}} = \frac{1}{1 + e^{-(b_1 \cdot x_1 + \dots + b_k \cdot x_k + a)}}$$
$$\hat{y} = b_1 \cdot x_1 + b_2 \cdot x_2 + \dots + b_k \cdot x_k + a$$

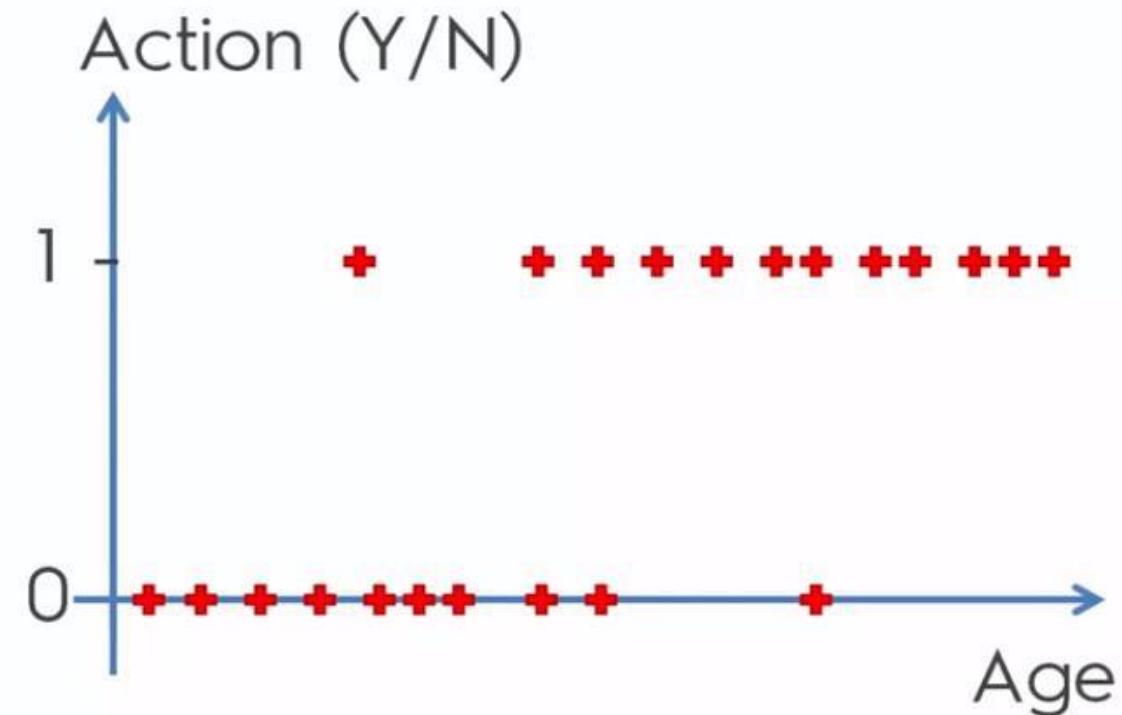


Logistic Regression

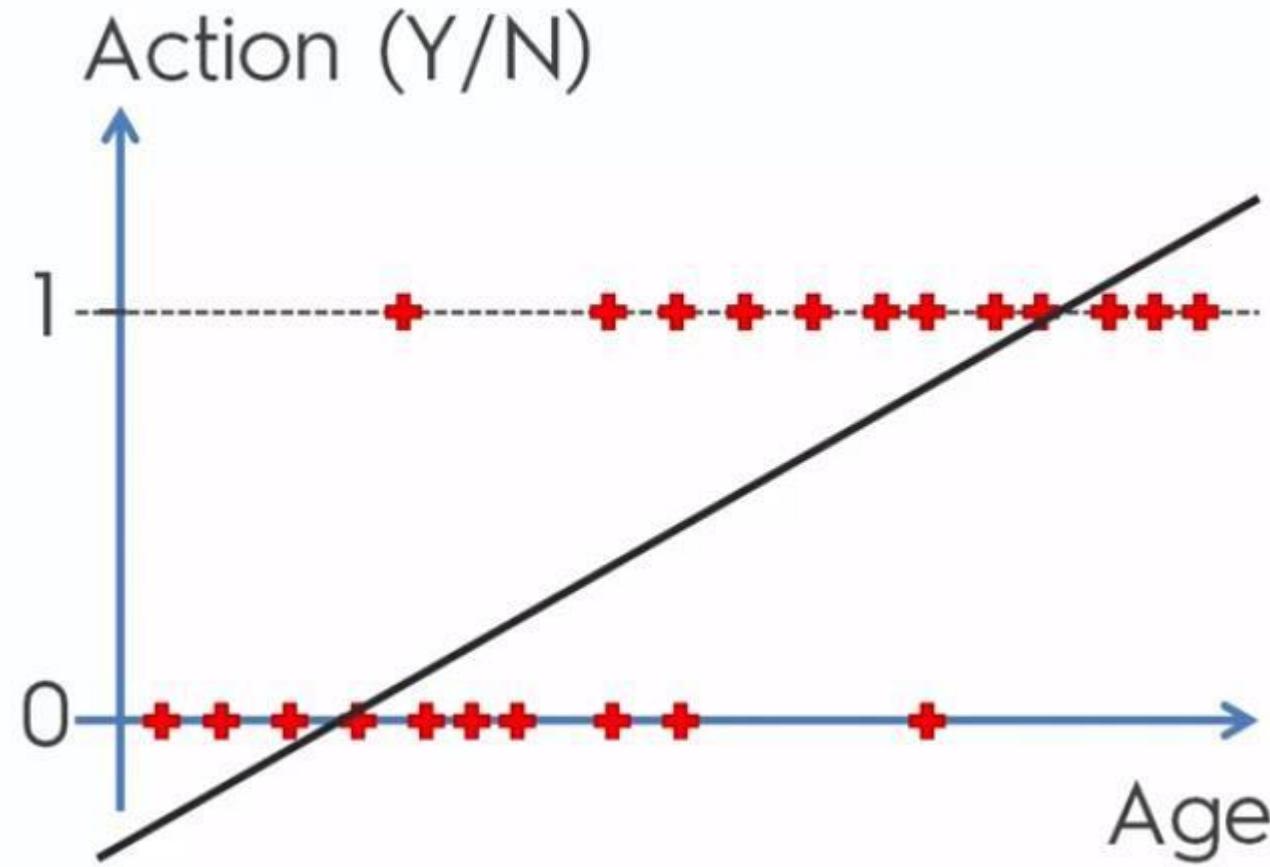
We know this:



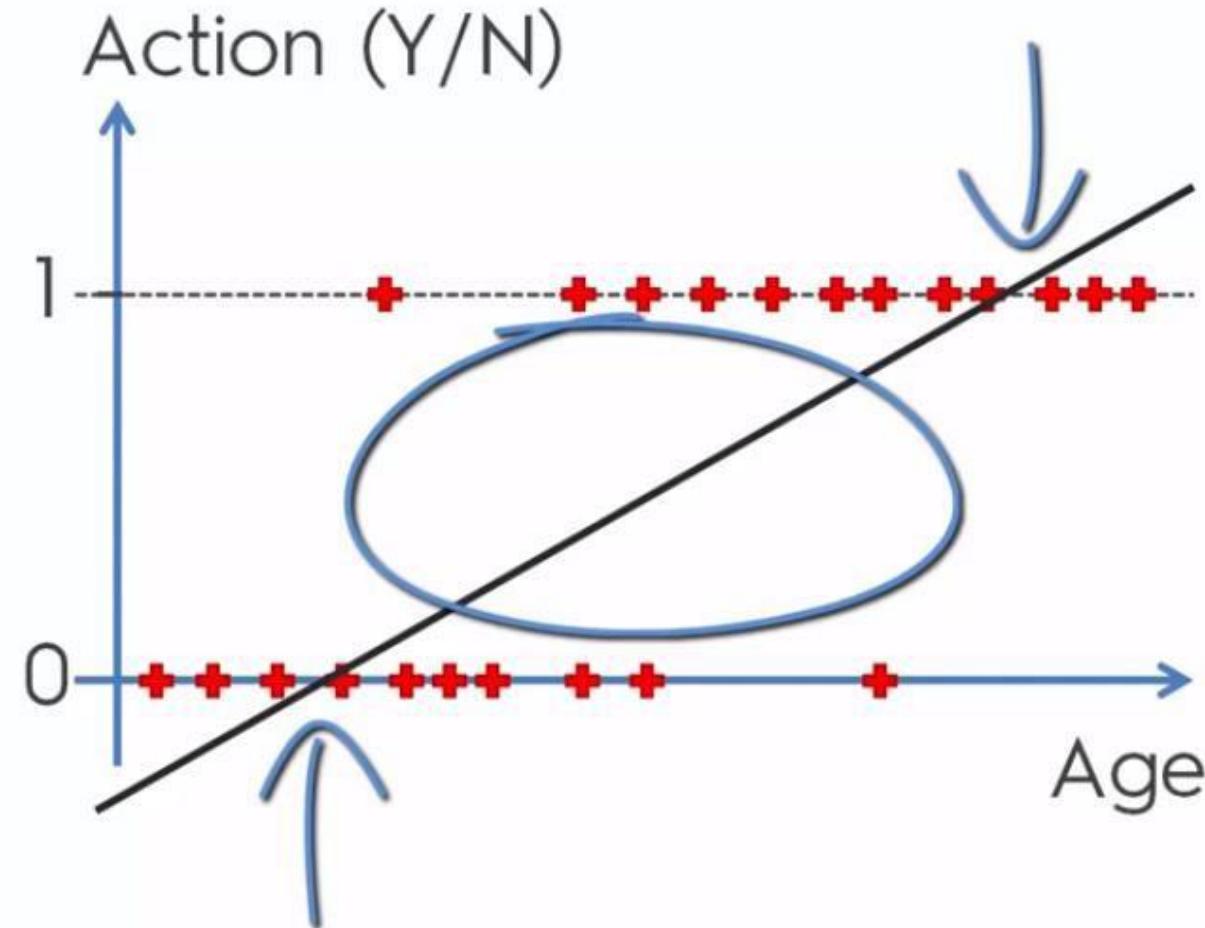
This is new:



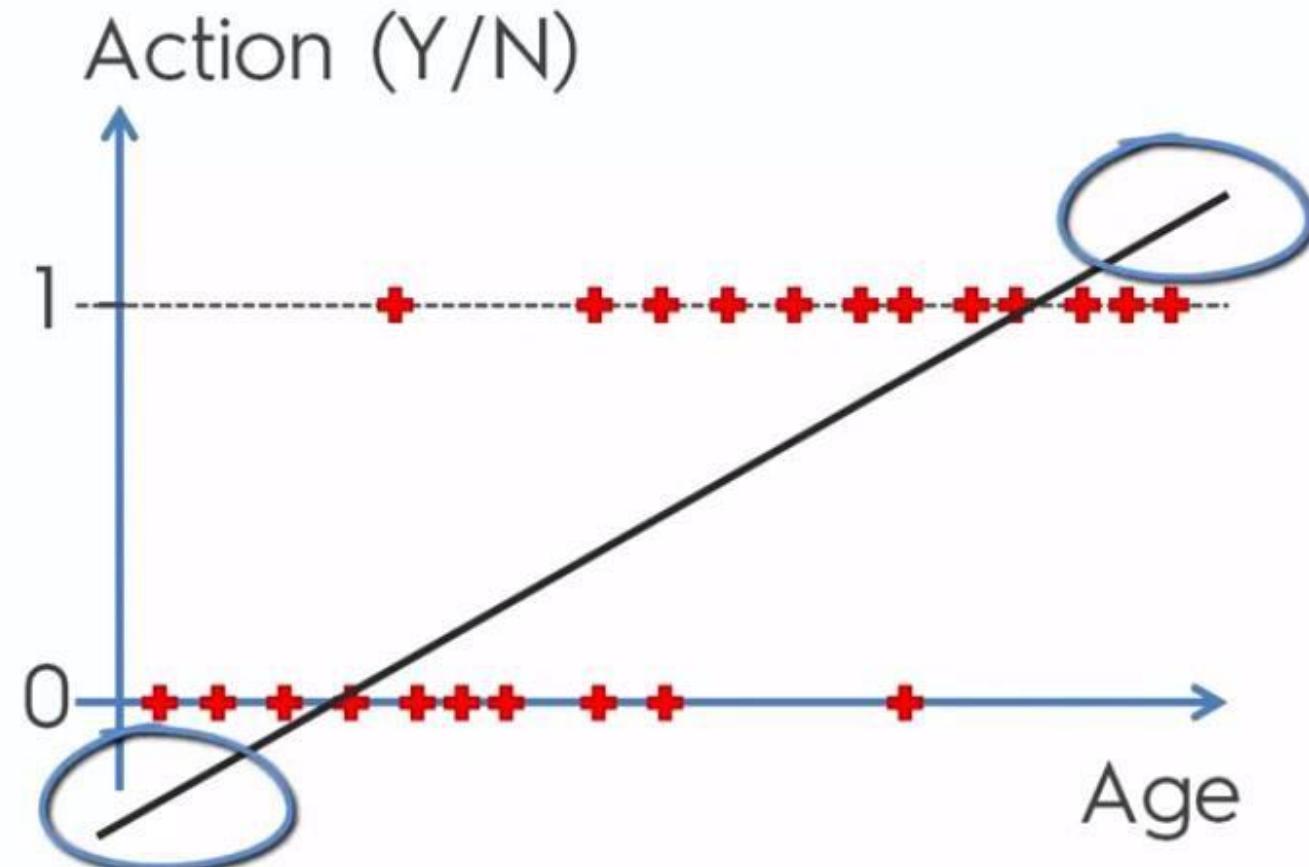
Logistic Regression



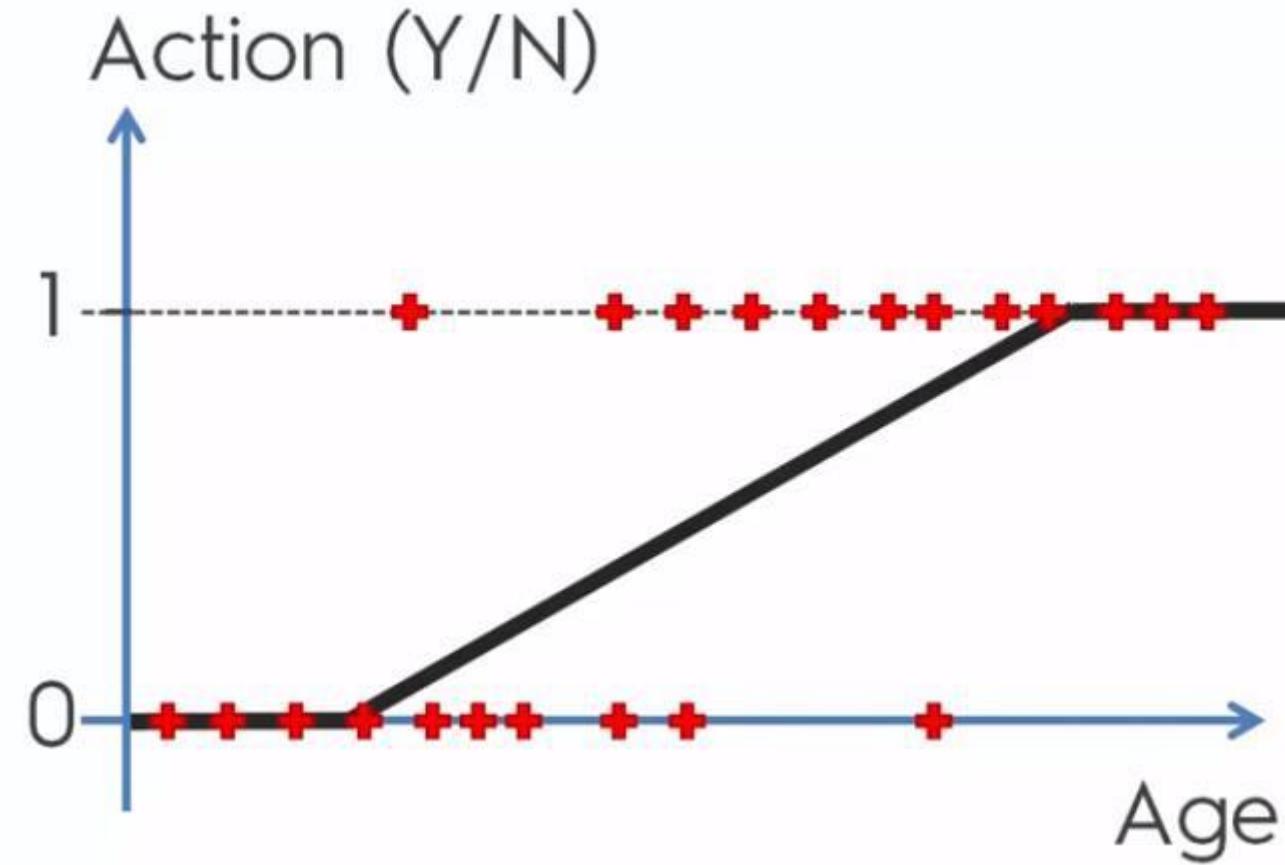
Logistic Regression



Logistic Regression



Logistic Regression



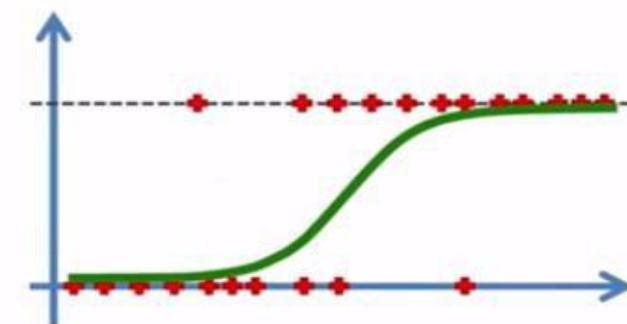
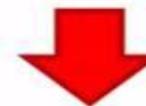
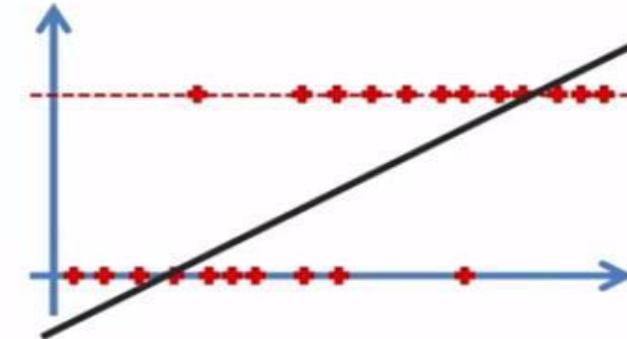
Logistic Regression

$$y = b_0 + b_1 * x$$

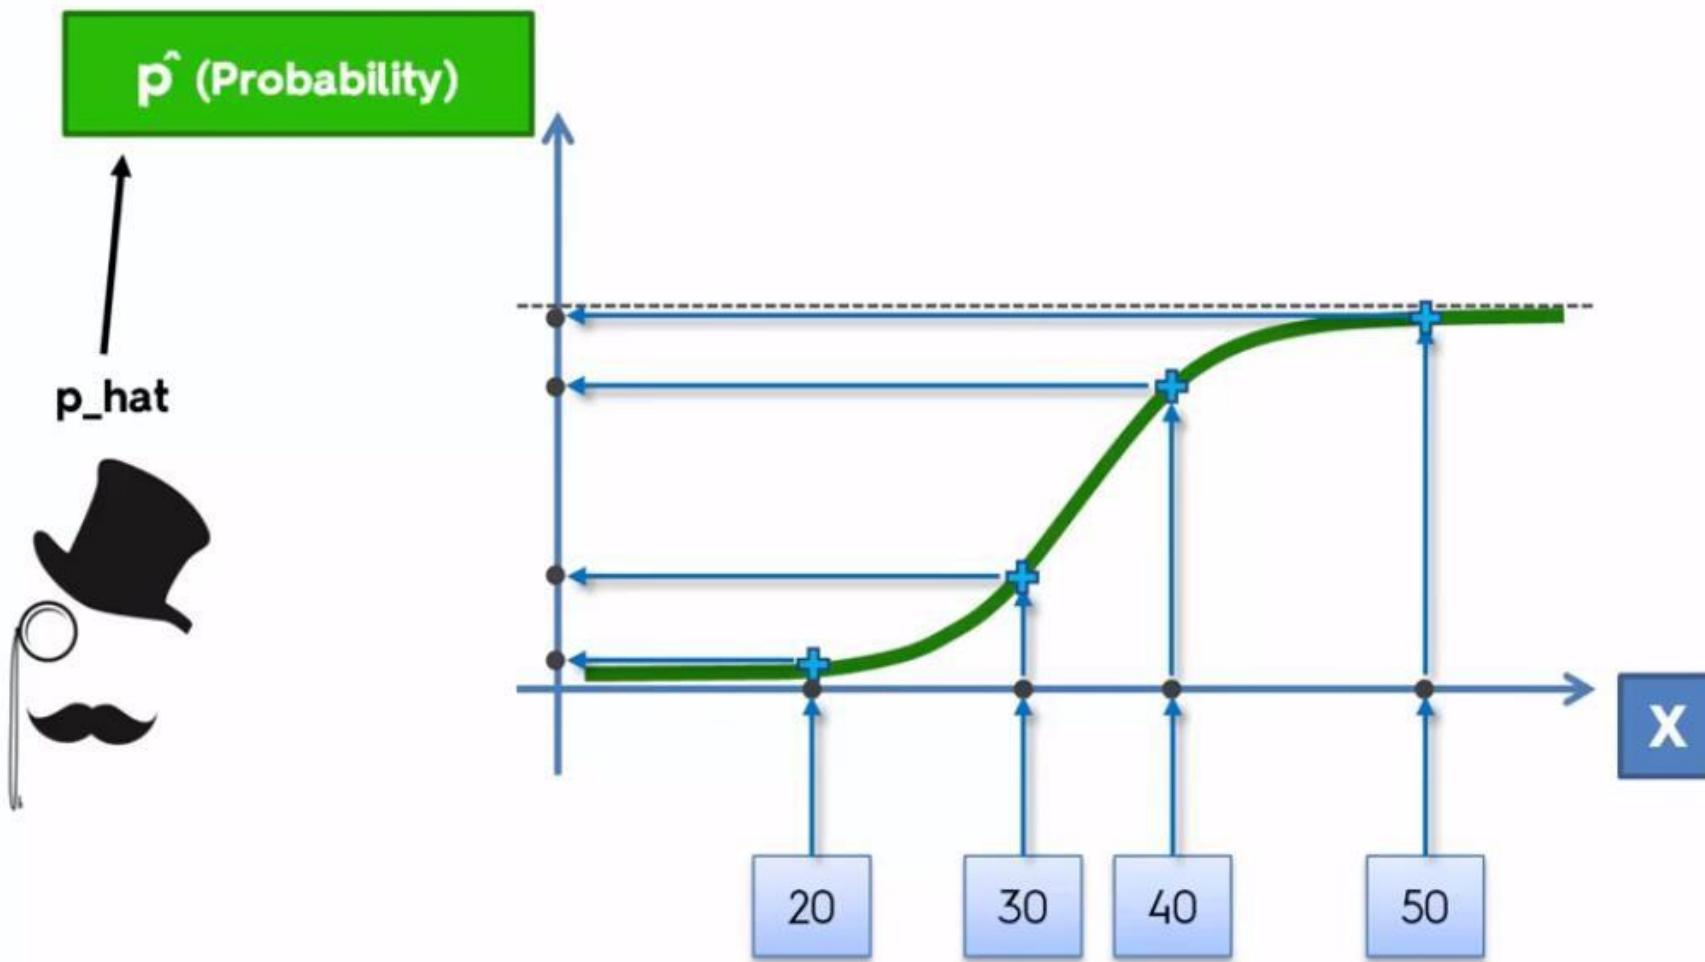


Sigmoid Function

$$P = \frac{1}{1 + e^{-y}}$$

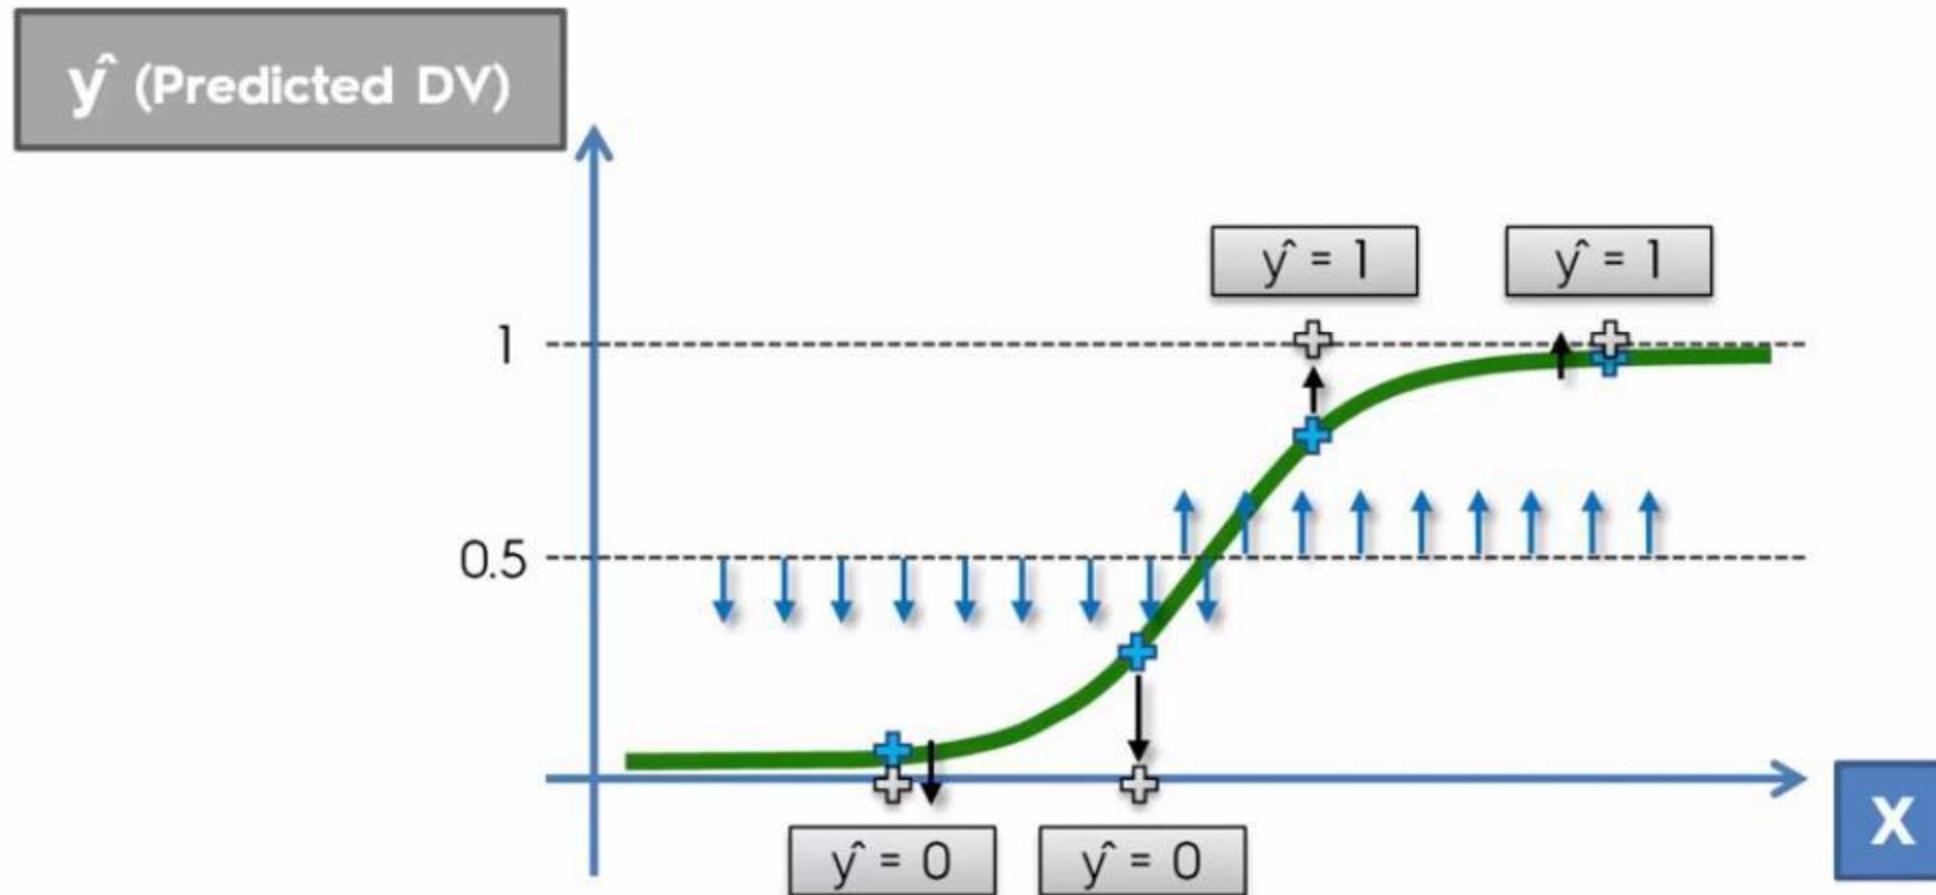


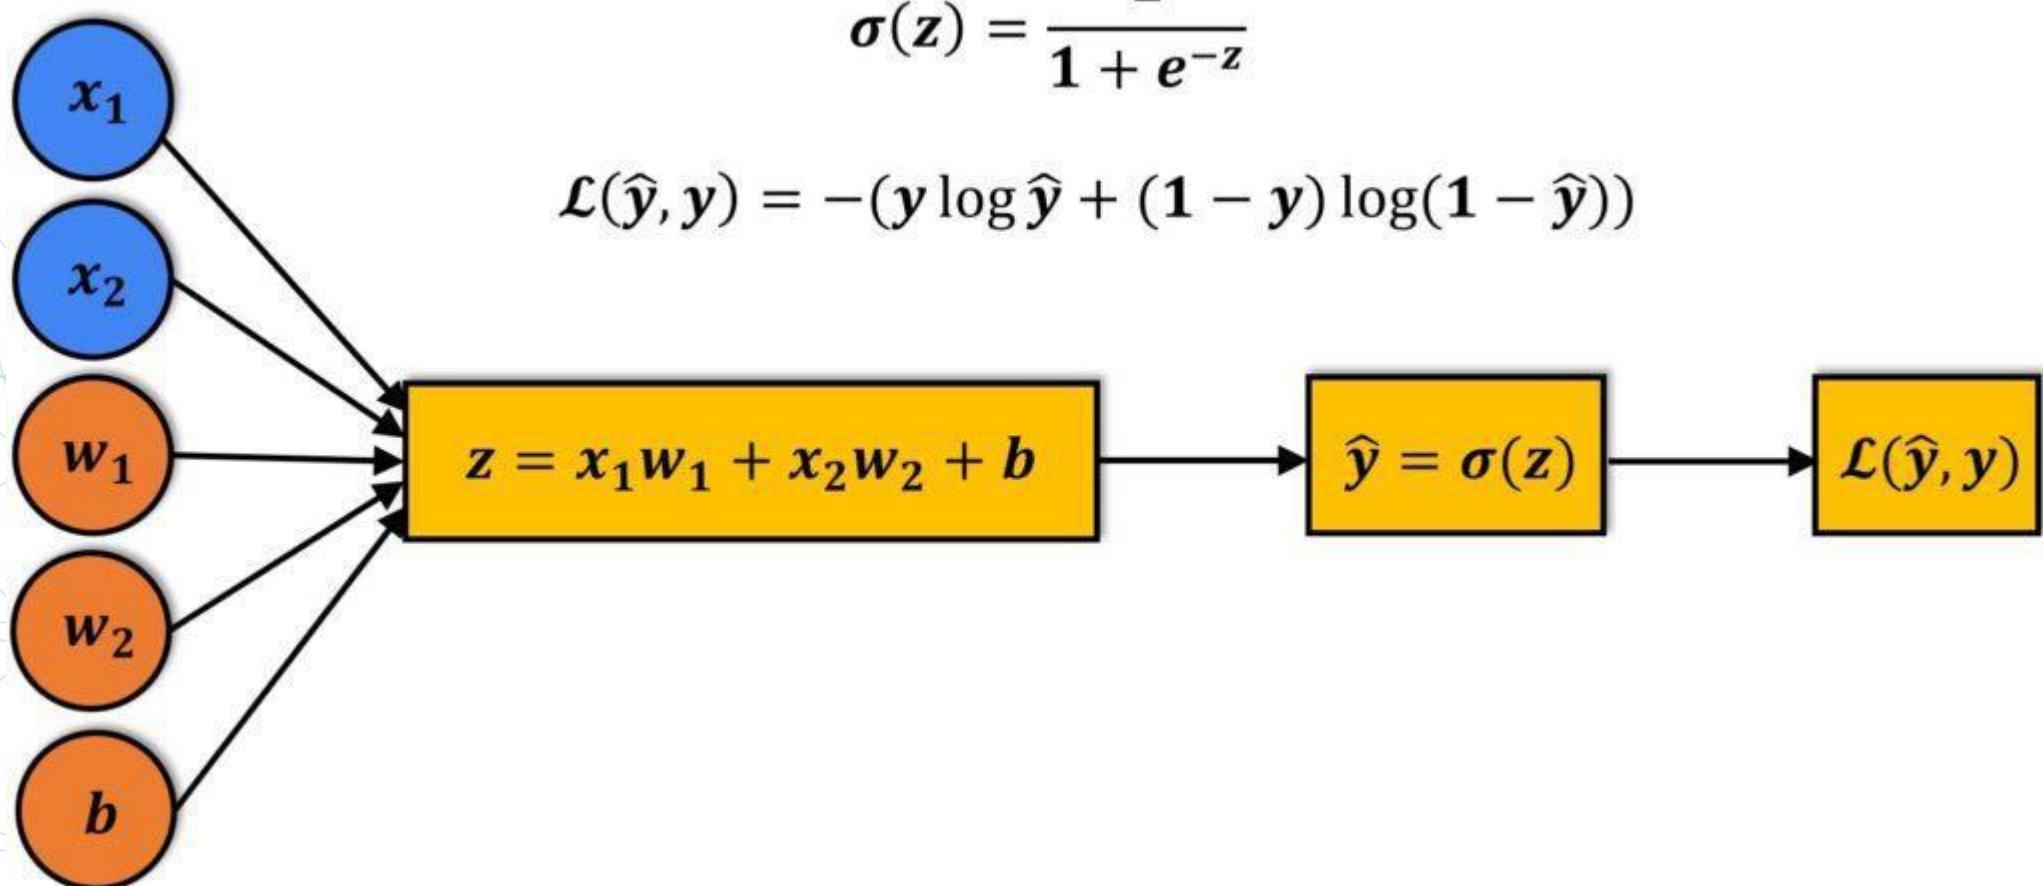
Logistic Regression





Logistic Regression

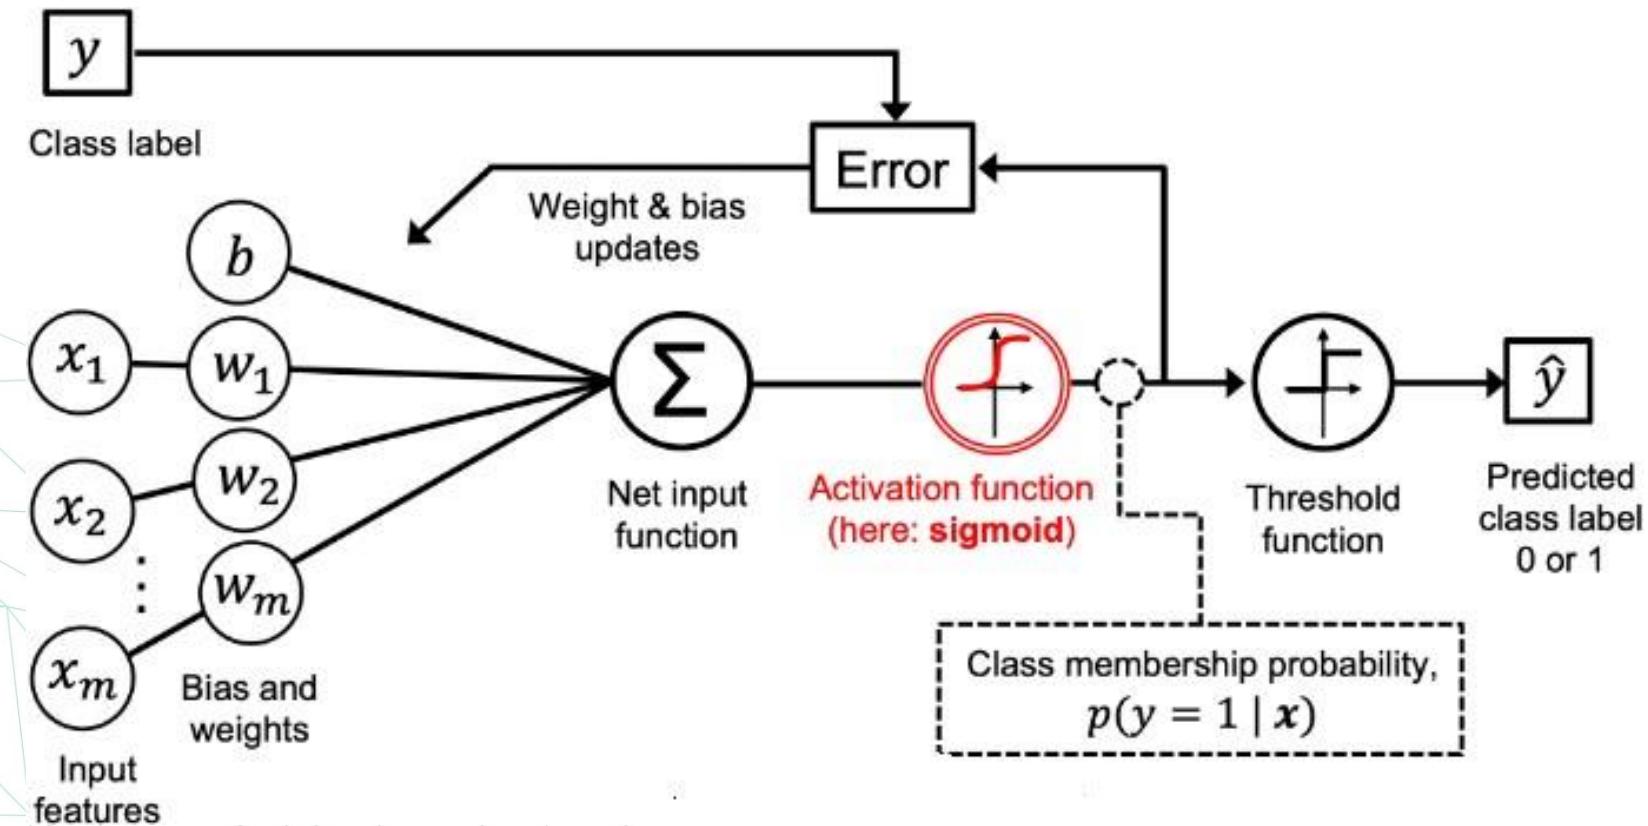




$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

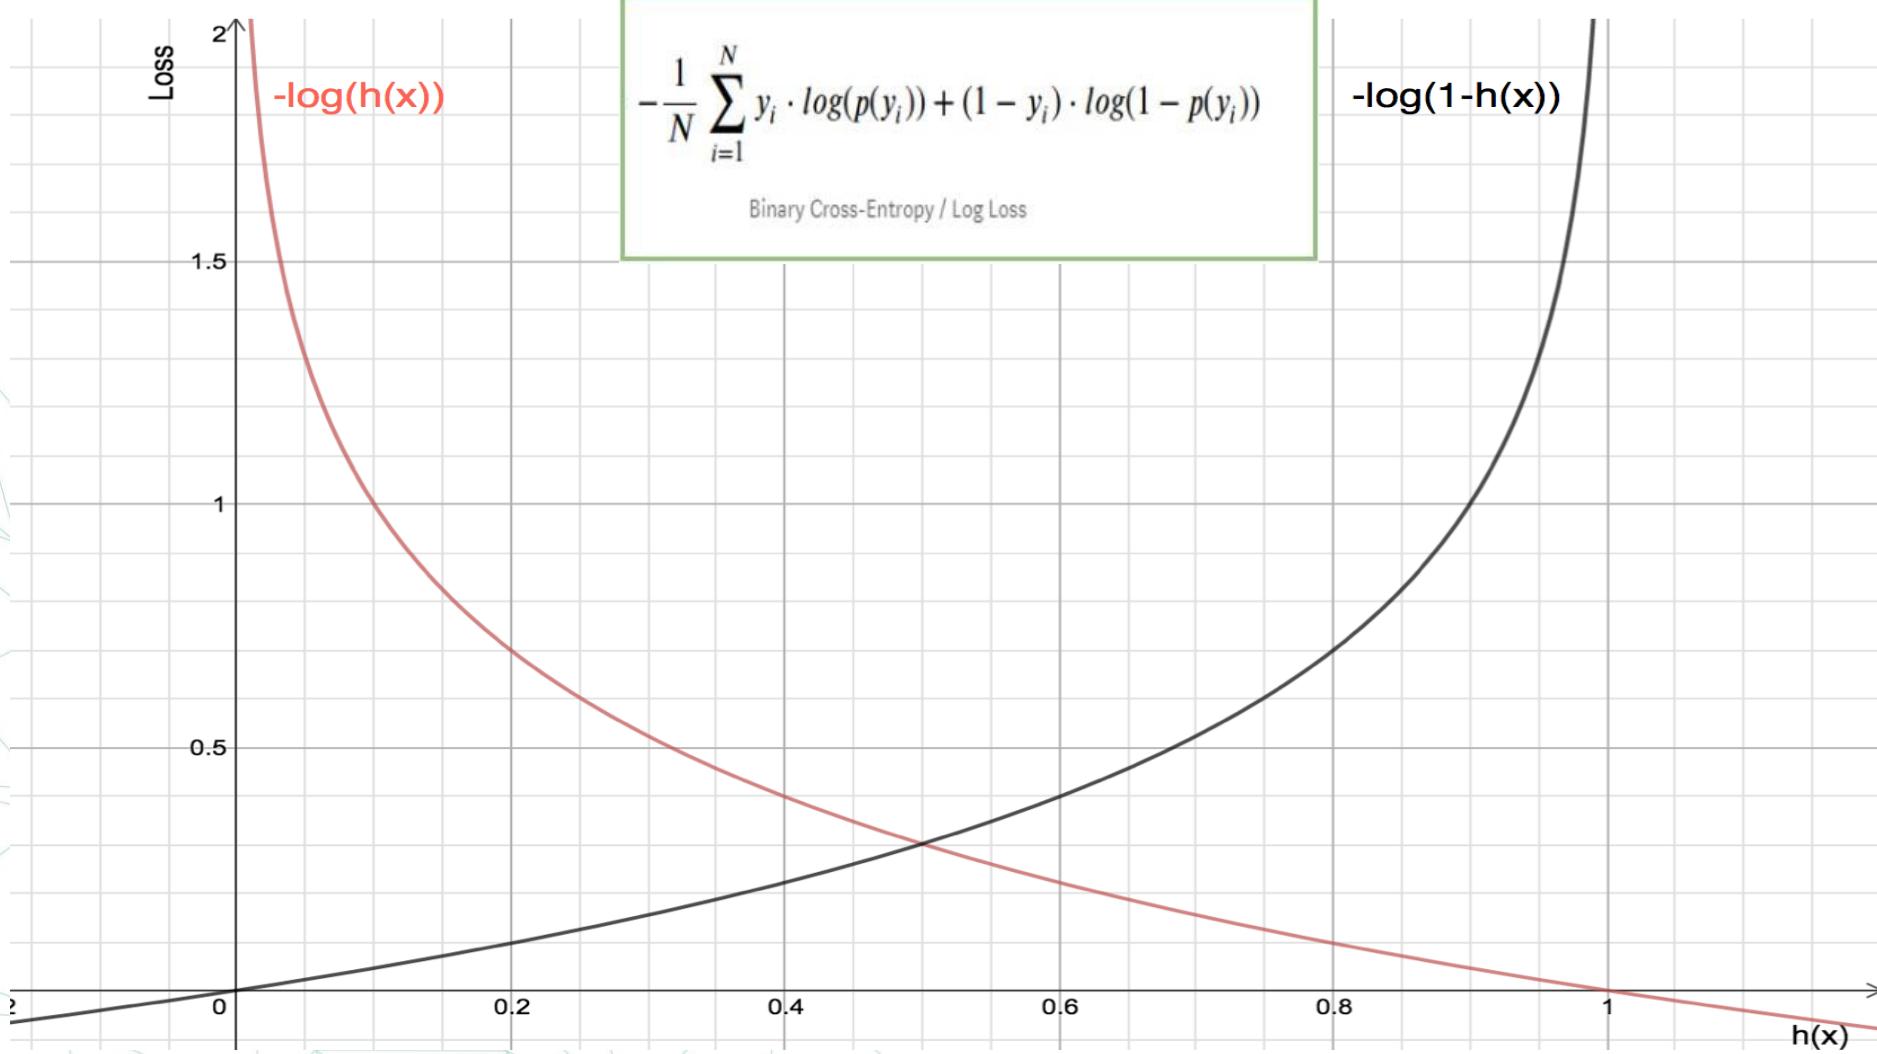
$$\mathcal{L}(\hat{y}, y) = -(y \log \hat{y} + (1 - y) \log(1 - \hat{y}))$$

Logistic Regression



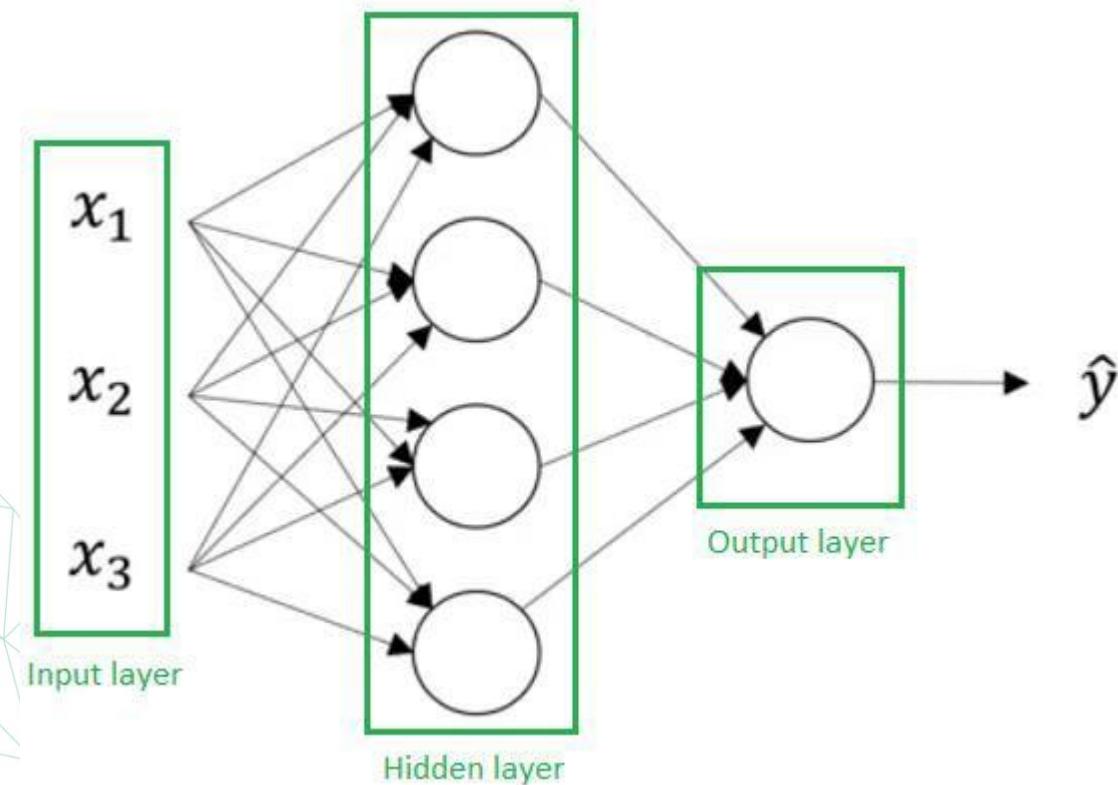
Cross Entropy Loss Function

$$H_p(q) = -\frac{1}{N} \sum_{i=1}^N y_i \cdot \log(p(y_i)) + (1 - y_i) \cdot \log(1 - p(y_i))$$





Neural Network





Classification

Pick one

Label 1	<input checked="" type="checkbox"/>
Label 2	<input type="checkbox"/>

Pick one

Label 1	<input type="checkbox"/>
Label 2	<input type="checkbox"/>
Label 3	<input type="checkbox"/>
Label 4	<input checked="" type="checkbox"/>
...	<input type="checkbox"/>
...	<input type="checkbox"/>
Label L	<input type="checkbox"/>

Pick all applicable

Label 1	<input type="checkbox"/>
Label 2	<input checked="" type="checkbox"/>
Label 3	<input type="checkbox"/>
Label 4	<input checked="" type="checkbox"/>
...	<input type="checkbox"/>
...	<input type="checkbox"/>
Label L	<input checked="" type="checkbox"/>

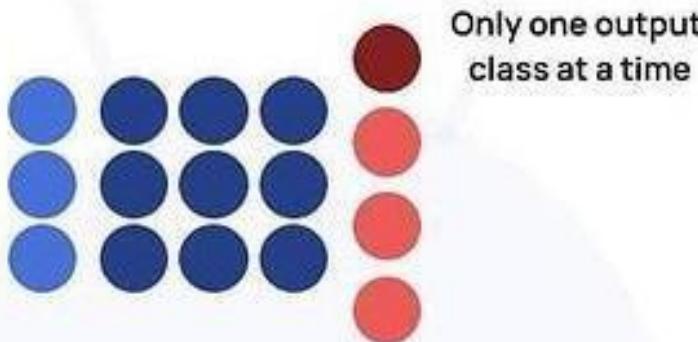
Binary

Multi-class

Multi-label

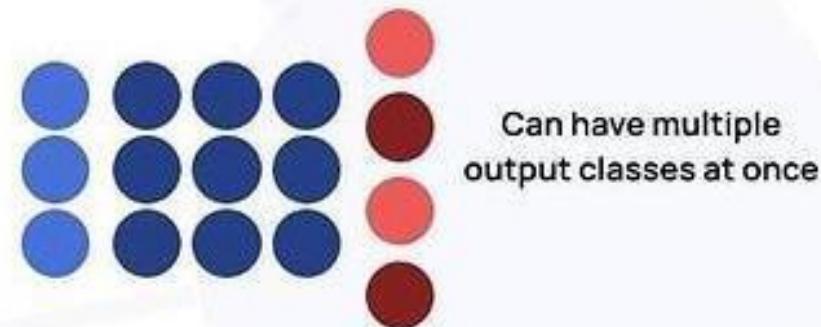
Multilabel classification

Multi-Class



Review	sentiment
Very good quality though	Positive
The design is very odd	Negative
I advise EVERYONE DO NOT BE FOOLED!	Negative
So Far So Good!	Positive
Works great!	positive

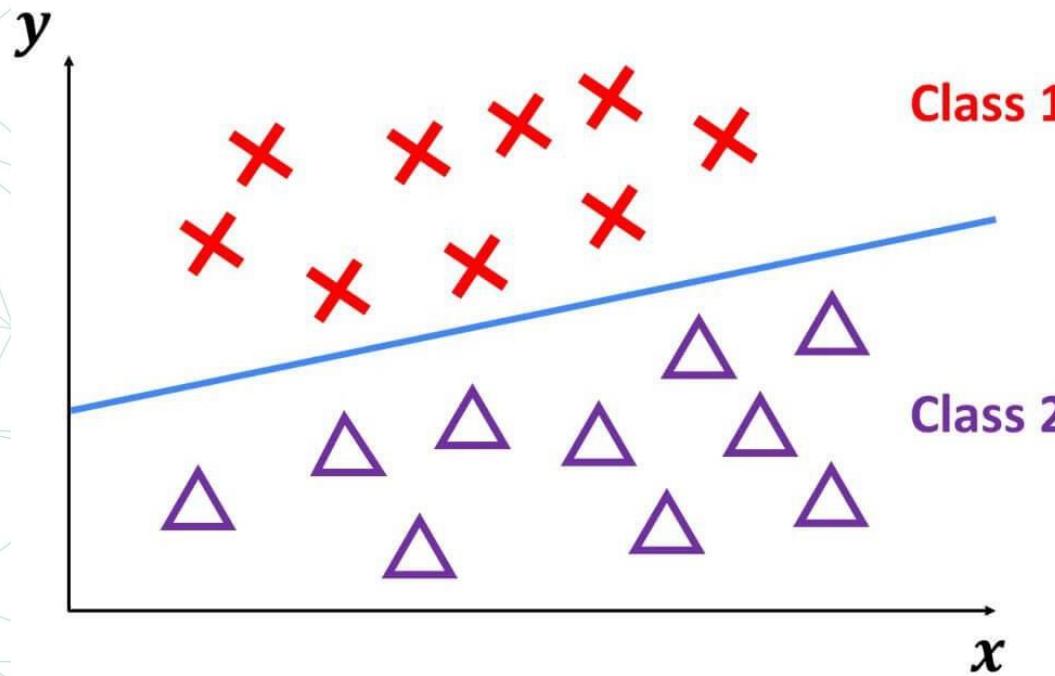
Multi-Label



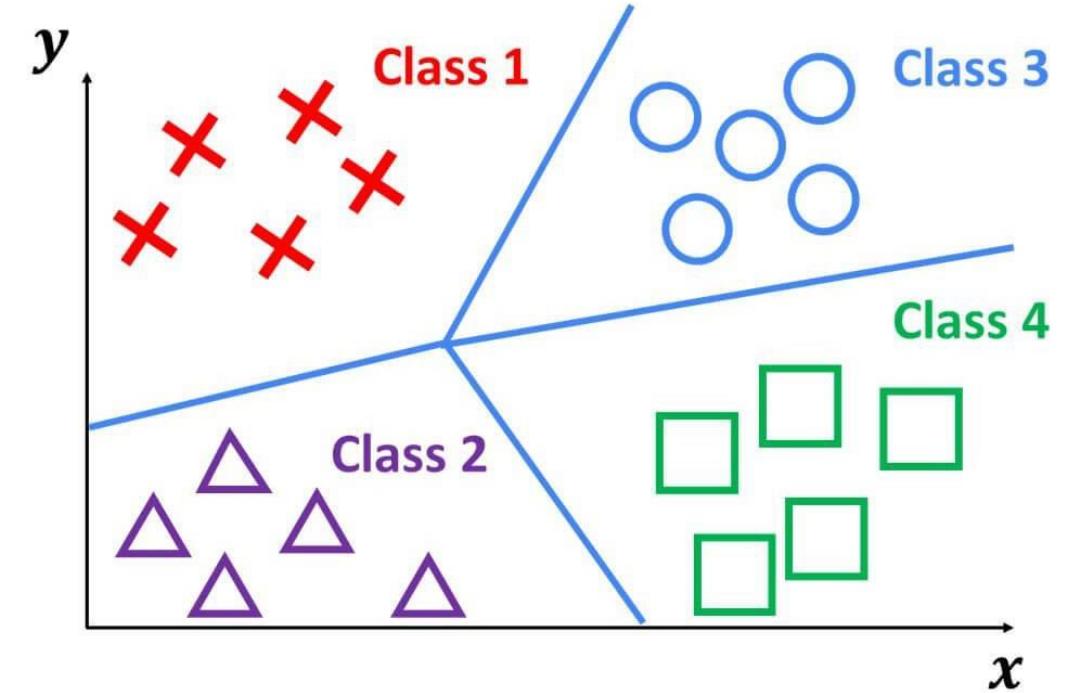
movie	Adventure	Comedy	Fantasy	Crime
Jumanji (1995)	1	0	1	0
Puccini for Beginners (2006)	0	1	0	0
How the Grinch Stole Christmas! (1966)	0	1	1	0

Multi-Class Classification

Binary Classification



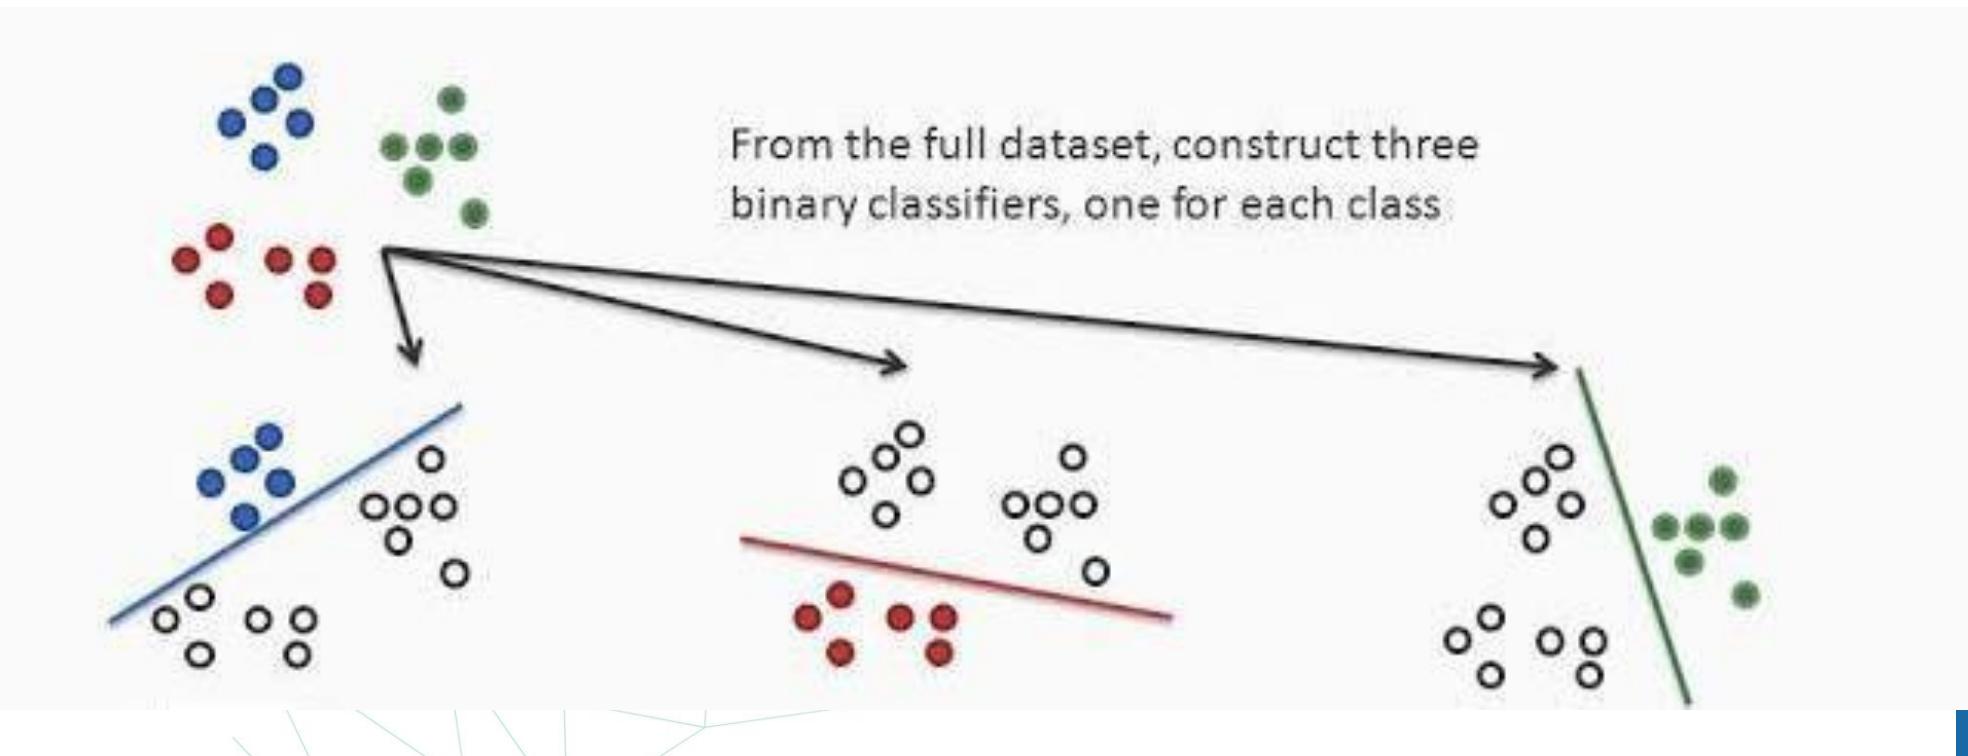
Multiclass Classification



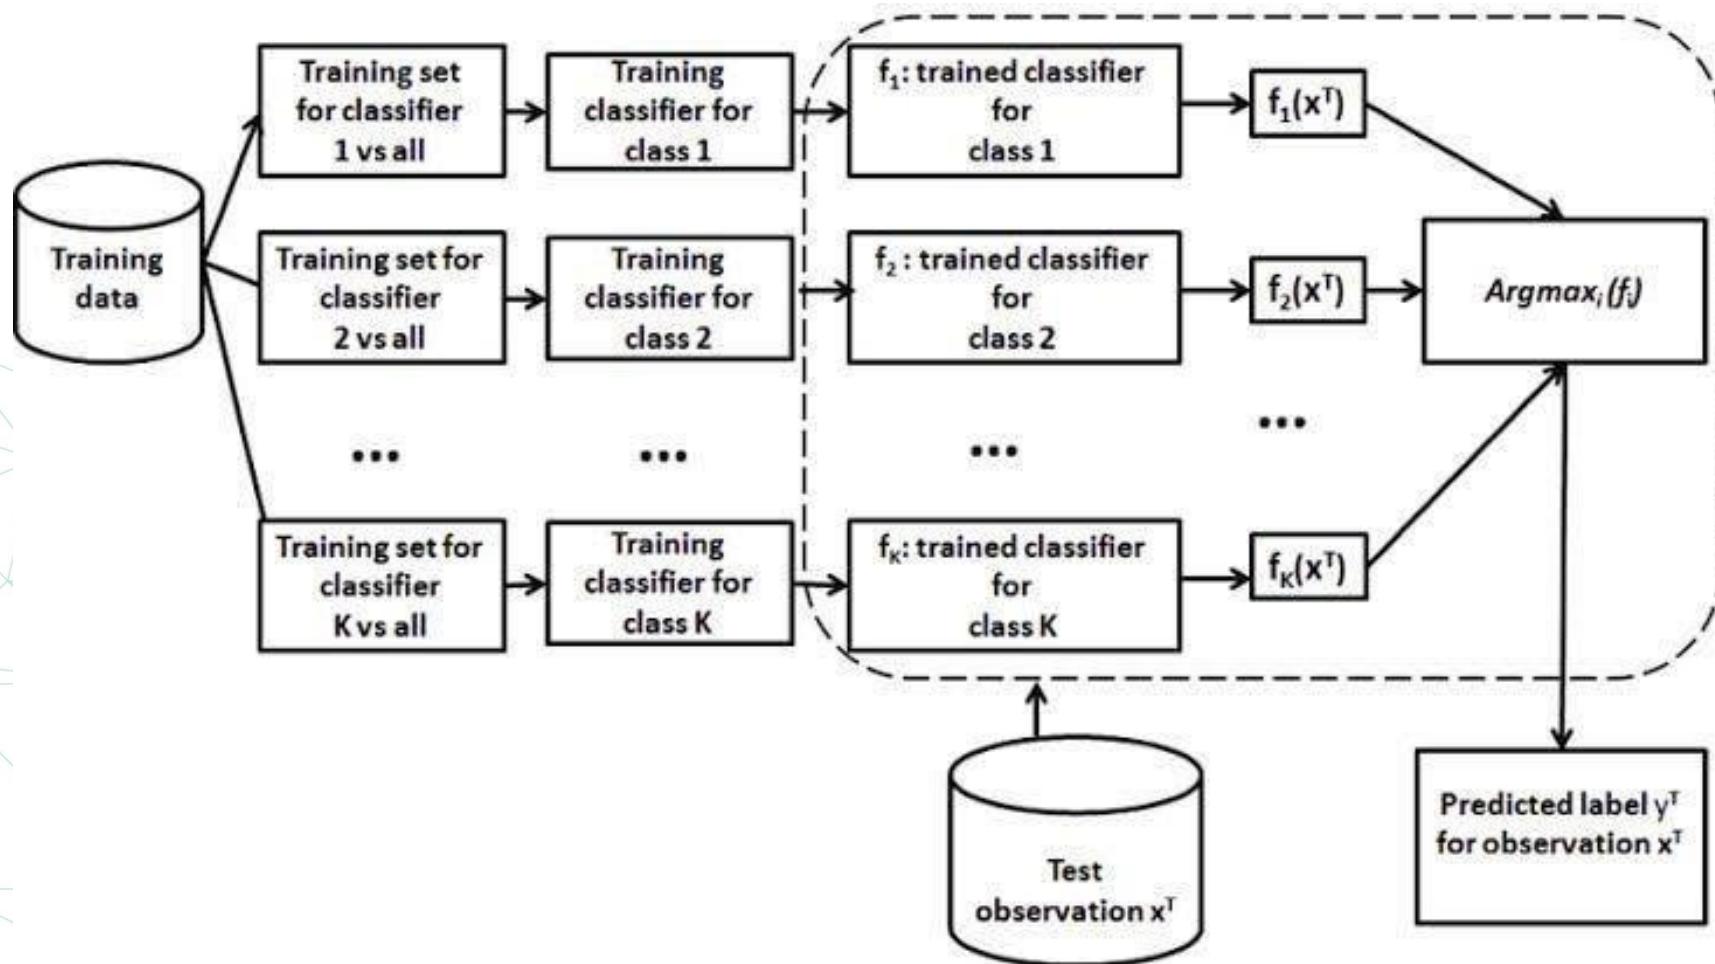
Multiclass Classification

- One-vs-All (One-vs-Rest):

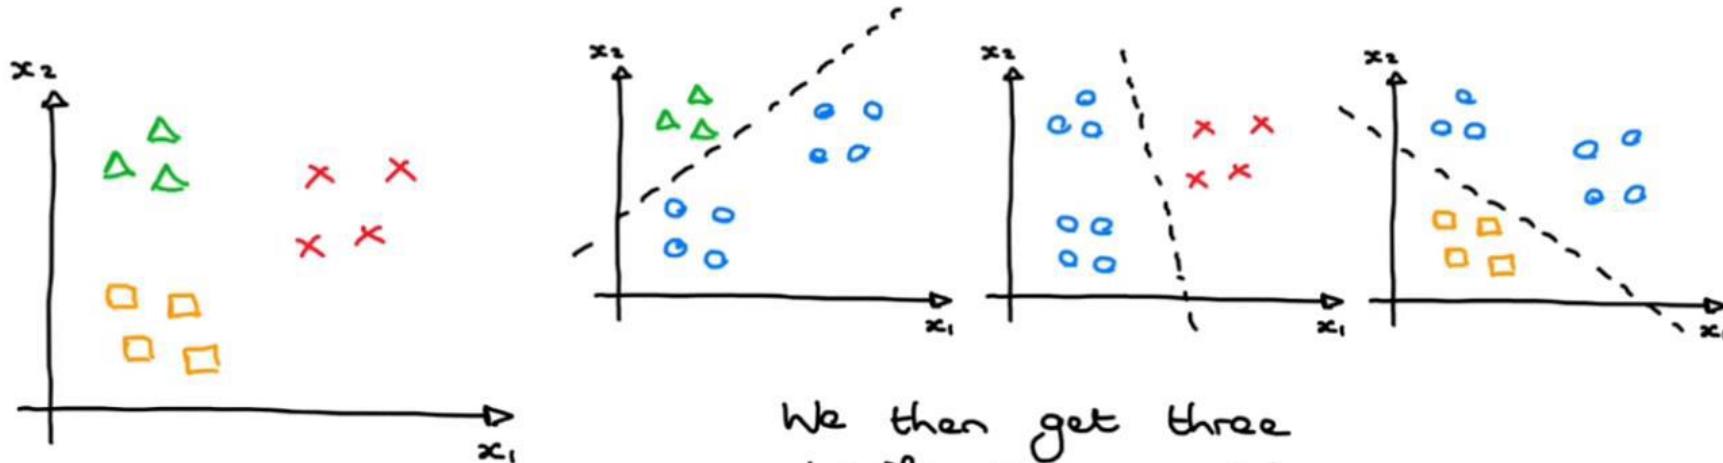
converts a multiclass problem into a series of binary tasks for each class in the target



One-vs-All



One-vs-rest classification



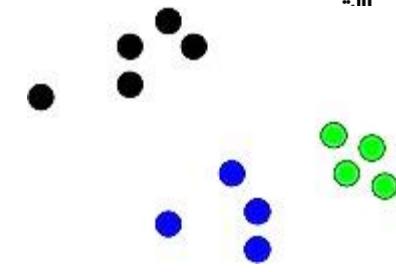
Strategy: Train three classifiers with $y \in \{0, 1\}$, where each classifier considers another class as the positive class.

We then get three classification models:

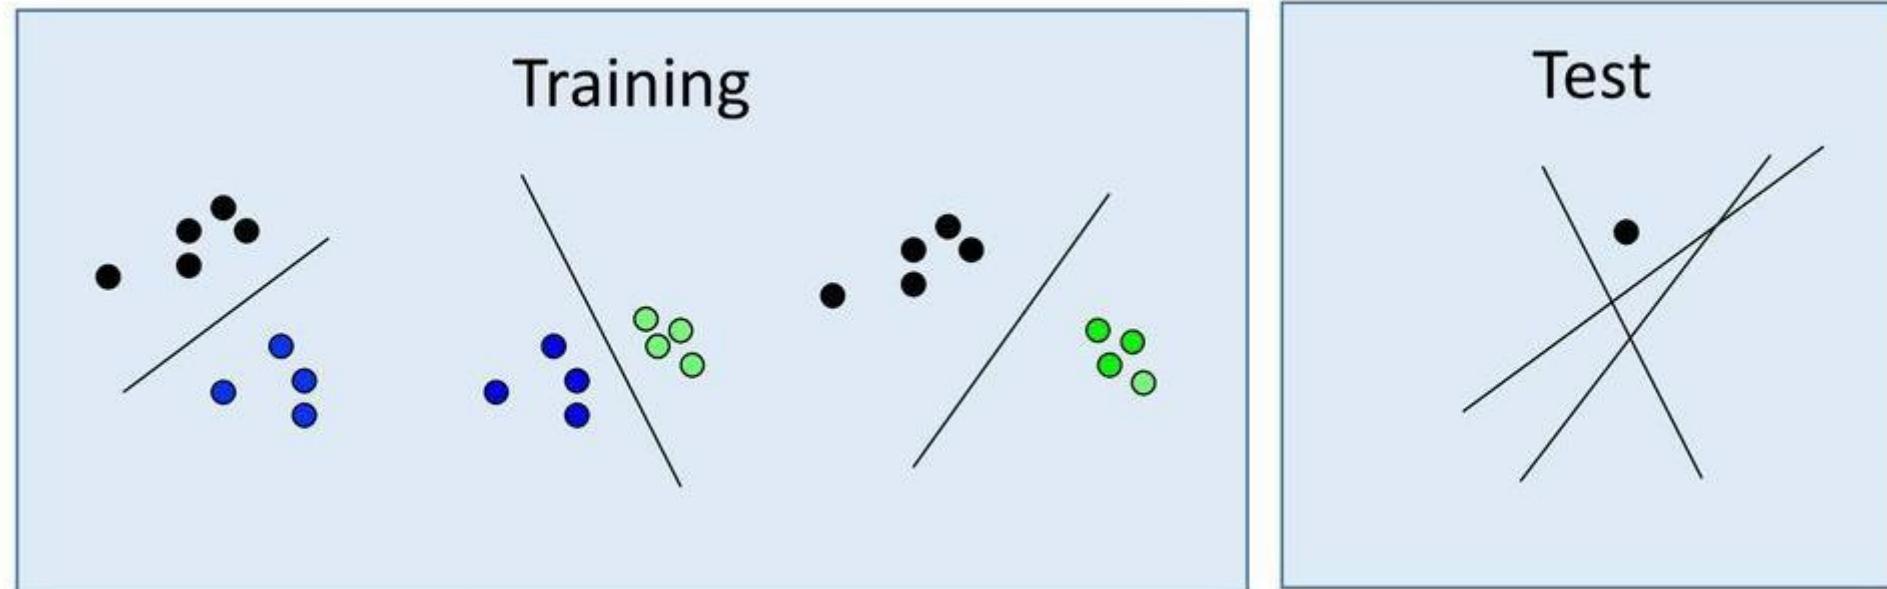
$$\begin{aligned} f_1(\underline{x}; \underline{w}_1) & \quad 1: \triangle \\ f_2(\underline{x}; \underline{w}_2) & \quad 2: \times \\ f_3(\underline{x}; \underline{w}_3) & \quad 3: \square \end{aligned}$$

Final predictions: $\arg \max_k f_k(\underline{x}; \underline{w}_k)$

Multiclass Classification



- All-vs-All (One-vs-One):
 - splits a multi-class problem into a single binary classification task for each pair of classes.



Practical Lab

Iris Dataset

Naive Bays



Naive Bayes

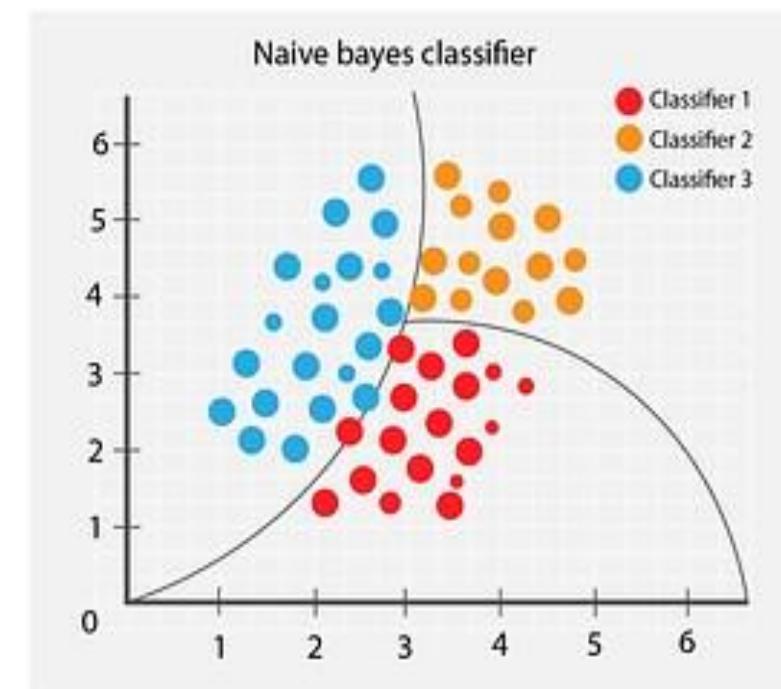
thatware.co

In machine learning, naive Bayes classifiers are a family of simple "probabilistic classifiers" based on applying Bayes' theorem with strong (naive) independence assumptions between the features.

$$P(A|B) = \frac{P(B|A) P(A)}{P(B)}$$

using Bayesian probability terminology, the above equation can be written as

$$\text{Posterior} = \frac{\text{prior} \times \text{likelihood}}{\text{evidence}}$$





Naïve Bayes

$$P(Y|X) = \frac{P(X|Y).P(Y)}{P(X)} \quad (1)$$

$$P(X) = P(X_1, X_2, \dots, \dots, X_n) = P(X_1).P(X_2) \dots \dots P(X_n)$$

because for independent events A and B, $P(AB) = P(A).P(B)$

Similarly because of conditional independence,

$$P(X|Y) = P(X_1, X_2, \dots, \dots, X_n|Y) = P(X_1|Y).P(X_2|Y) \dots \dots P(X_n|Y)$$

Substituting these in (1), we get

$$P(Y|X) = \frac{P(X_1|Y).P(X_2|Y) \dots \dots P(X_n|Y).P(Y)}{P(X_1).P(X_2) \dots \dots P(X_n)}$$

$$P(Y|X) \propto |P(X_1|Y).P(X_2|Y) \dots \dots P(X_n|Y).P(Y)|$$



Naïve Bayes

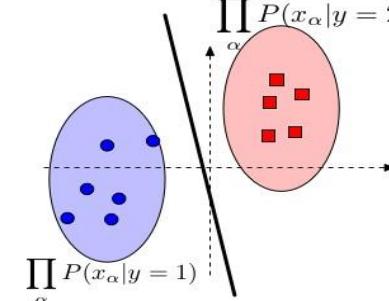
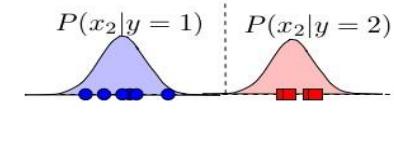
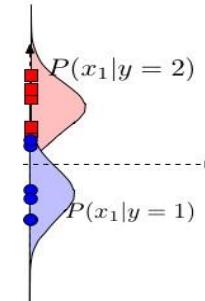
$$P(c | x) = \frac{P(x | c)P(c)}{P(x)}$$

Likelihood Class Prior Probability

Posterior Probability Predictor Prior Probability

$$P(c | X) = P(x_1 | c) \times P(x_2 | c) \times \dots \times P(x_n | c) \times P(c)$$

Naïve Bayes Classifier



Obs	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
X1	A	A	A	A	A	B	B	B	B	C	C	C	C	C	
X2	S	M	M	S	S	S	M	M	L	L	M	M	L	L	
Y	0	0	1	1	0	0	0	1	1	1	1	1	1	1	0

Prior Probability

$$P(Y=1) = 9/15 \quad P(Y=0) = 6/15$$

Conditional Probability

$$P(X1=A|Y=1)=2/9 \quad P(X1=B|Y=1)=3/9 \quad P(X1=C|Y=1)=4/9$$

$$P(X2=S|Y=1)=1/9 \quad P(X2=M|Y=1)=4/9 \quad P(X2=L|Y=1)=4/9$$

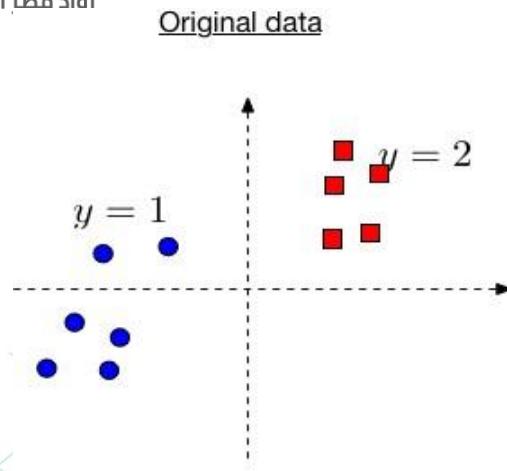
$$P(X1=A|Y=0)=3/6 \quad P(X1=B|Y=0)=2/6 \quad P(X1=C|Y=0)=1/6$$

$$P(X2=S|Y=0)=3/6 \quad P(X2=M|Y=0)=2/6 \quad P(X2=L|Y=0)=1/6$$

$$P(Y=1)P(X1=B|Y=1)P(X2=S|Y=1)=1/45$$

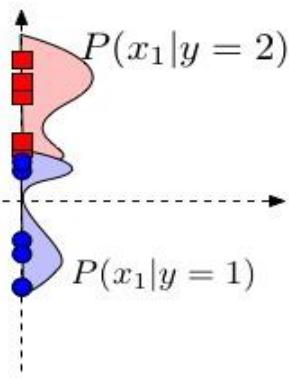
$$P(Y=0)P(X1=B|Y=0)P(X2=S|Y=0)=1/15$$

$P(Y=0)P(X1=B|Y=0)P(X2=S|Y=0) > P(Y=1)P(X1=B|Y=1)P(X2=S|Y=1)$, so $y=0$.

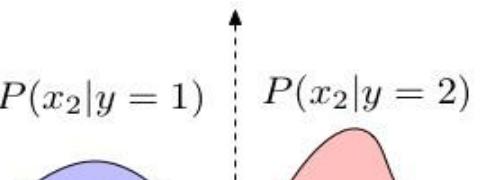


$$\begin{aligned}
 h(\vec{x}) &= \operatorname{argmax}_y P(y|\vec{x}) \\
 &= \operatorname{argmax}_y \frac{P(\vec{x}|y)P(y)}{P(\vec{x})} \\
 &= \operatorname{argmax}_y P(\vec{x}|y)P(y) \\
 &= \operatorname{argmax}_y \prod_{\alpha=1}^d P(x_\alpha|y)P(y)
 \end{aligned}$$

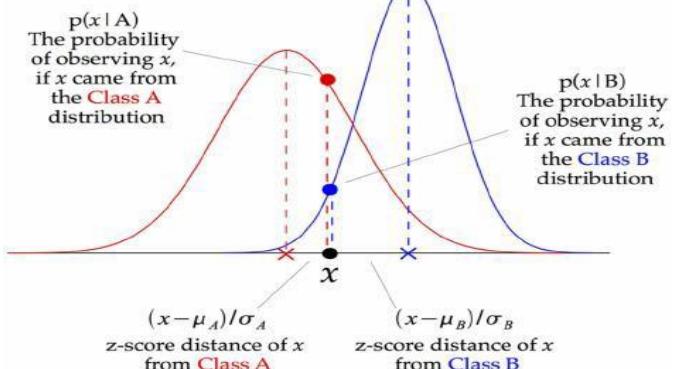
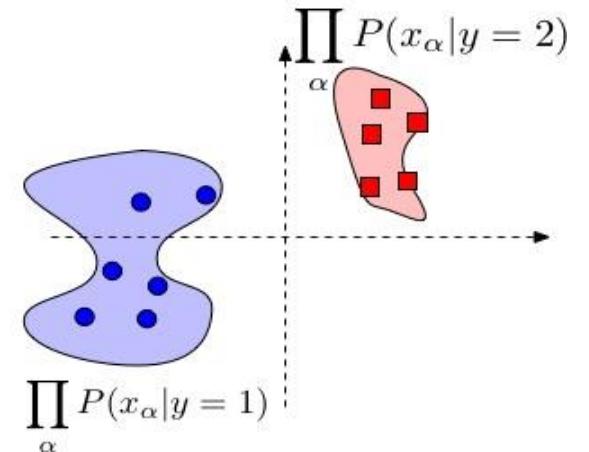
Estimation of first dimension



Estimation of second dimension



Resulting data distribution



$(P(\vec{x}) \text{ does not depend on } y)$

(by the naive assumption)

Naïve Bayes Classifier

- Naïve Bayes Classifier can be trained **easily** and **fast** and can be used as **benchmark model**.
- When **variable selection** is carried out properly, Naïve Bayes can **perform** as well as or even **better** than other **statistical** models such as logistic regression.
- Naive Bayes requires a **strong assumption of independent predictors**, so when the model has a **bad performance**, the reason leading to that may be the **dependence between predictors**.
- Doesn't require feature scaling

Naïve Bayes Example

Outlook	Temp	Humidity	Windy	Play Golf
Rainy	Hot	High	False	No
Rainy	Hot	High	True	No
Overcast	Hot	High	False	Yes
Sunny	Mild	High	False	Yes
Sunny	Cool	Normal	False	Yes
Sunny	Cool	Normal	True	No
Overcast	Cool	Normal	True	Yes
Rainy	Mild	High	False	No
Rainy	Cool	Normal	False	Yes
Sunny	Mild	Normal	False	Yes
Rainy	Mild	Normal	True	Yes
Overcast	Mild	High	True	Yes
Overcast	Hot	Normal	False	Yes
Sunny	Mild	High	True	No



		Play Golf	
		Yes	No
Outlook	Sunny	3	2
	Overcast	4	0
	Rainy	2	3

		Play Golf		
		Yes	No	
Outlook	Sunny	3/9	2/5	5/14
	Overcast	4/9	0/5	4/14
	Rainy	2/9	3/5	5/14
		9/14	5/14	

$$P(x | c) = P(\text{Sunny} | \text{Yes}) = 3 / 9 = 0.33$$

$$P(x) = P(\text{Sunny}) \\ = 5 / 14 = 0.36$$

$$P(c) = P(\text{Yes}) = 9 / 14 = 0.64$$

Posterior Probability:

$$P(c | x) = P(\text{Yes} | \text{Sunny}) = 0.33 \times 0.64 / 0.36 = 0.60$$

		Play Golf	
		Yes	No
Outlook	Sunny	3	2
	Overcast	4	0
	Rainy	2	3

		Play Golf		
		Yes	No	
Outlook	Sunny	3	2	5
	Overcast	4	0	4
	Rainy	2	3	5
		9	5	14

$$P(x | c) = P(\text{Sunny} | \text{No}) = 2 / 5 = 0.4$$

$$P(x) = P(\text{Sunny}) \\ = 5 / 14 = 0.36$$

$$P(c) = P(\text{No}) = 5 / 14 = 0.36$$

Posterior Probability:

$$P(c | x) = P(\text{No} | \text{Sunny}) = 0.40 \times 0.36 / 0.36 = 0.40$$



Frequency Table

		Play Golf	
		Yes	No
Outlook	Sunny	3	2
	Overcast	4	0
	Rainy	2	3

Likelihood Table

		Play Golf	
		Yes	No
Outlook	Sunny	3/9	2/5
	Overcast	4/9	0/5
	Rainy	2/9	3/5

		Play Golf	
		Yes	No
Humidity	High	3	4
	Normal	6	1

		Play Golf	
		Yes	No
Humidity	High	3/9	4/5
	Normal	6/9	1/5

		Play Golf	
		Yes	No
Temp.	Hot	2	2
	Mild	4	2
	Cool	3	1

		Play Golf	
		Yes	No
Temp.	Hot	2/9	2/5
	Mild	4/9	2/5
	Cool	3/9	1/5

		Play Golf	
		Yes	No
Windy	False	6	2
	True	3	3

		Play Golf	
		Yes	No
Windy	False	6/9	2/5
	True	3/9	3/5

Outlook	Temp	Humidity	Windy	Play
Rainy	Cool	High	True	?

$$P(Yes | X) = P(Rainy | Yes) \times P(Cool | Yes) \times P(High | Yes) \times P(True | Yes) \times P(Yes)$$

$$P(Yes | X) = 2/9 \times 3/9 \times 3/9 \times 3/9 \times 9/14 = 0.00529 \quad 0.2 = \frac{0.00529}{0.02057 + 0.00529}$$

$$P(No | X) = P(Rainy | No) \times P(Cool | No) \times P(High | No) \times P(True | No) \times P(No)$$

$$P(No | X) = 3/5 \times 1/5 \times 4/5 \times 3/5 \times 5/14 = 0.02057 \quad 0.8 = \frac{0.02057}{0.02057 + 0.00529}$$

Q&A

Questions and answers

Thanks